

Copyright	xix
Foreword	xxi
Preface	xxiii
1. Audience	xxiii
2. Book Layout	xxiii
3. Jenkins or Hudson?	xxiii
4. Font Conventions	xxiv
5. Command-Line Conventions	xxiv
6. Contributors	xxv
7. The Review Team	xxvi
8. Book Sponsors	xxvi
8.1. Wakaleo Consulting	xxvi
8.2. CloudBees	xxvii
8.3. Odd-e	xxvii
9. Using Code Examples	xxviii
10. Safari® Books Online	xxviii
11. How to Contact Us	xxix
12. Acknowledgments	xxix
1. Introducing Jenkins	1
1.1. Introduction	1
1.2. Continuous Integration Fundamentals	1
1.3. Introducing Jenkins (né Hudson)	2
1.4. From Hudson to Jenkins—A Short History	3
1.5. Should I Use Jenkins or Hudson?	4
1.6. Introducing Continuous Integration into Your Organization	5
1.6.1. Phase 1—No Build Server	5
1.6.2. Phase 2—Nightly Builds	5
1.6.3. Phase 3—Nightly Builds and Basic Automated Tests	5
1.6.4. Phase 4—Enter the Metrics	6
1.6.5. Phase 5—Getting More Serious About Testing	6
1.6.6. Phase 6—Automated Acceptance Tests and More Automated Deployment	6
1.6.7. Phase 7—Continuous Deployment	6
1.7. Where to Now?	7
2. Your First Steps with Jenkins	9
2.1. Introduction	9
2.2. Preparing Your Environment	9
2.2.1. Installing Java	10
2.2.2. Installing Git	10
2.2.3. Setting Up a GitHub Account	11
2.2.4. Configuring SSH Keys	11
2.2.5. Forking the Sample Repository	11
2.3. Starting Up Jenkins	12
2.4. Configuring the Tools	13

2.4.1. Configuring Your Maven Setup	14
2.4.2. Configuring the JDK	15
2.4.3. Notification	15
2.4.4. Setting Up Git	15
2.5. Your First Jenkins Build Job	16
2.6. Your First Build Job in Action	19
2.7. More Reporting—Displaying Javadocs	24
2.8. Adding Code Coverage and Other Metrics	25
2.9. Conclusion	28
3. Installing Jenkins	29
3.1. Introduction	29
3.2. Downloading and Installing Jenkins	29
3.3. Preparing a Build Server for Jenkins	30
3.4. The Jenkins Home Directory	32
3.5. Installing Jenkins on Debian or Ubuntu	33
3.6. Installing Jenkins on Red Hat, Fedora, or CentOS	34
3.7. Installing Jenkins on SUSE or OpenSUSE	35
3.8. Running Jenkins as a Stand-Alone Application	35
3.9. Running Jenkins Behind an Apache Server	38
3.10. Running Jenkins in an Application Server	39
3.11. Memory Considerations	40
3.12. Installing Jenkins as a Windows Service	40
3.13. What’s in the Jenkins Home Directory	42
3.14. Backing Up Your Jenkins Data	45
3.15. Upgrading Your Jenkins Installation	45
3.16. Conclusion	46
4. Configuring Your Jenkins Server	47
4.1. Introduction	47
4.2. The Configuration Dashboard—The Manage Jenkins Screen	47
4.3. Configuring the System Environment	49
4.4. Configuring Global Properties	50
4.5. Configuring Your JDKs	51
4.6. Configuring Your Build Tools	52
4.6.1. Maven	52
4.6.2. Ant	53
4.6.3. Shell-Scripting Language	54
4.7. Configuring Your Version Control Tools	54
4.7.1. Configuring Subversion	54
4.7.2. Configuring CVS	54
4.8. Configuring the Mail Server	54
4.9. Configuring a Proxy	55
4.10. Conclusion	56
5. Setting Up Your Build Jobs	57

5.1. Introduction	57
5.2. Jenkins Build Jobs	57
5.3. Creating a Freestyle Build Job	58
5.3.1. General Options	58
5.3.2. Advanced Project Options	59
5.4. Configuring Source Code Management	60
5.4.1. Working with Subversion	60
5.4.2. Working with Git	62
5.5. Build Triggers	69
5.5.1. Triggering a Build Job Once Another Build Job Has Finished	70
5.5.2. Scheduled Build Jobs	70
5.5.3. Polling the SCM	71
5.5.4. Triggering Builds Remotely	72
5.5.5. Manual Build Jobs	73
5.6. Build Steps	73
5.6.1. Maven Build Steps	74
5.6.2. Ant Build Steps	75
5.6.3. Executing a Shell or Windows Batch Command	76
5.6.4. Using Jenkins Environment Variables in Your Builds	77
5.6.5. Running Groovy Scripts	79
5.6.6. Building Projects in Other Languages	80
5.7. Post-Build Actions	80
5.7.1. Reporting on Test Results	81
5.7.2. Archiving Build Results	81
5.7.3. Notifications	83
5.7.4. Building Other Projects	84
5.8. Running Your New Build Job	84
5.9. Working with Maven Build Jobs	84
5.9.1. Building Whenever a SNAPSHOT Dependency Is Built	85
5.9.2. Configuring the Maven Build	85
5.9.3. Post-Build Actions	87
5.9.4. Deploying to an Enterprise Repository Manager	87
5.9.5. Deploying to Commercial Enterprise Repository Managers	89
5.9.6. Managing Modules	89
5.9.7. Extra Build Steps in Your Maven Build Jobs	90
5.10. Using Jenkins with Other Languages	90
5.10.1. Building Projects with Grails	90
5.10.2. Building Projects with Gradle	91
5.10.3. Building Projects with Visual Studio MSBuild	93
5.10.4. Building Projects with NAnt	94
5.10.5. Building Projects with Ruby and Ruby on Rails	94
5.11. Conclusion	95
6. Automated Testing	97

6.1. Introduction	97
6.2. Automating Your Unit and Integration Tests	98
6.3. Configuring Test Reports in Jenkins	98
6.4. Displaying Test Results	99
6.5. Ignoring Tests	101
6.6. Code Coverage	102
6.6.1. Measuring Code Coverage with Cobertura	103
6.6.2. Measuring Code Coverage with Clover	110
6.7. Automated Acceptance Tests	111
6.8. Automated Performance Tests with JMeter	113
6.9. Help! My Tests Are Too Slow!	119
6.9.1. Add More Hardware	119
6.9.2. Run Fewer Integration/Functional Tests	119
6.9.3. Run Your Tests in Parallel	120
6.10. Conclusion	120
7. Securing Jenkins	121
7.1. Introduction	121
7.2. Activating Security in Jenkins	121
7.3. Simple Security in Jenkins	121
7.4. Security Realms—Identifying Jenkins Users	122
7.4.1. Using Jenkins’s Built-in User Database	122
7.4.2. Using an LDAP Repository	124
7.4.3. Using Microsoft Active Directory	125
7.4.4. Using Unix Users and Groups	126
7.4.5. Delegating to the Servlet Container	126
7.4.6. Using Atlassian Crowd	126
7.4.7. Integrating with Other Systems	127
7.5. Authorization—Who Can Do What	128
7.5.1. Matrix-based Security	128
7.5.2. Project-based Security	132
7.5.3. Role-based Security	133
7.6. Auditing—Keeping Track of User Actions	134
7.7. Conclusion	137
8. Notification	139
8.1. Introduction	139
8.2. Email Notification	139
8.3. More Advanced Email Notification	140
8.4. Claiming Builds	143
8.5. RSS Feeds	143
8.6. Build Radiators	144
8.7. Instant Messaging	145
8.7.1. IM Notification with Jabber	145
8.7.2. IM Notification using IRC	148

8.8. IRC Notification	148
8.9. Desktop Notifiers	150
8.10. Notification via Notifo	152
8.11. Mobile Notification	153
8.12. SMS Notification	154
8.13. Making Noise	154
8.14. Extreme Feedback Devices	156
8.15. Conclusion	157
9. Code Quality	159
9.1. Introduction	159
9.2. Code Quality in Your Build Process	160
9.3. Popular Java and Groovy Code Quality Analysis Tools	161
9.3.1. Checkstyle	161
9.3.2. PMD/CPD	163
9.3.3. FindBugs	167
9.3.4. CodeNarc	169
9.4. Reporting on Code Quality Issues with the Violations Plugin	170
9.4.1. Working with Freestyle Build Jobs	171
9.4.2. Working with Maven Build Jobs	172
9.5. Using the Checkstyle, PMD, and FindBugs Reports	173
9.6. Reporting on Code Complexity	174
9.7. Reporting on Open Tasks	175
9.8. Integrating with Sonar	176
9.9. Conclusion	178
10. Advanced Builds	179
10.1. Introduction	179
10.2. Parameterized Build Jobs	179
10.2.1. Creating a Parameterized Build Job	179
10.2.2. Adapting Your Builds to Work with Parameterized Build Scripts	180
10.2.3. More Advanced Parameter Types	181
10.2.4. Building from a Subversion Tag	182
10.2.5. Building from a Git Tag	183
10.2.6. Starting a Parameterized Build Job Remotely	184
10.2.7. Parameterized Build Job History	184
10.3. Parameterized Triggers	184
10.4. Multiconfiguration Build Jobs	186
10.4.1. Setting Up a Multiconfiguration Build	186
10.4.2. Configuring a Slave Axis	186
10.4.3. Configuring a JDK Axis	187
10.4.4. Custom Axis	187
10.4.5. Running a Multiconfiguration Build	187
10.5. Generating Your Maven Build Jobs Automatically	189
10.5.1. Configuring a Job	189

10.5.2. Reusing Job Configuration with Inheritance	191
10.5.3. Plugin Support	192
10.5.4. Freestyle Jobs	194
10.6. Coordinating Your Builds	195
10.6.1. Parallel Builds in Jenkins	195
10.6.2. Dependency Graphs	196
10.6.3. Joins	196
10.6.4. Locks and Latches	197
10.7. Build Pipelines and Promotions	197
10.7.1. Managing Maven Releases with the M2Release Plugin	198
10.7.2. Copying Artifacts	200
10.7.3. Build Promotions	202
10.7.4. Aggregating Test Results	206
10.7.5. Build Pipelines	206
10.8. Conclusion	208
11. Distributed Builds	209
11.1. Introduction	209
11.2. The Jenkins Distributed Build Architecture	209
11.3. Master/Slave Strategies in Jenkins	210
11.3.1. The Master Starts the Slave Agent Using SSH	210
11.3.2. Starting the Slave Agent Manually Using Java Web Start	213
11.3.3. Installing a Jenkins Slave as a Windows Service	214
11.3.4. Starting the Slave Node in Headless Mode	215
11.3.5. Starting a Windows Slave as a Remote Service	215
11.4. Associating a Build Job with a Slave or Group of Slaves	216
11.5. Node Monitoring	217
11.6. Cloud Computing	218
11.6.1. Using Amazon EC2	218
11.7. Using the CloudBees DEV@cloud Service	221
11.8. Conclusion	222
12. Automated Deployment and Continuous Delivery	223
12.1. Introduction	223
12.2. Implementing Automated and Continuous Deployment	224
12.2.1. The Deployment Script	224
12.2.2. Database Updates	224
12.2.3. Smoke Tests	227
12.2.4. Rolling Back Changes	227
12.3. Deploying to an Application Server	228
12.3.1. Deploying a Java Application	228
12.3.2. Deploying Scripting-based Applications Like Ruby and PHP	235
12.4. Conclusion	236
13. Maintaining Jenkins	237
13.1. Introduction	237

13.2. Monitoring Disk Space	237
13.2.1. Using the Disk Usage Plugin	238
13.2.2. Disk Usage and the Jenkins Maven Project Type	239
13.3. Monitoring the Server Load	239
13.4. Backing Up Your Configuration	240
13.4.1. Fundamentals of Jenkins Backups	240
13.4.2. Using the Backup Plugin	241
13.4.3. More Lightweight Automated Backups	242
13.5. Archiving Build Jobs	243
13.6. Migrating Build Jobs	243
13.7. Conclusion	246
A. Automating Your Unit and Integration Tests	247
A.1. Automating Your Tests with Maven	247
A.2. Automating Your Tests with Ant	252
Index	255

List of Figures

2.1. Installing Java	10
2.2. Signing up for a GitHub account	11
2.3. Forking the sample code repository	12
2.4. Running Jenkins using Java Web Start from the book's website	12
2.5. Java Web Start will download and run the latest version of Jenkins	12
2.6. Java Web Start running Jenkins	13
2.7. The Jenkins start page	13
2.8. The Manage Jenkins screen	14
2.9. The Configure Jenkins screen	14
2.10. Configuring a Maven installation	15
2.11. Configuring a JDK installation	15
2.12. Managing plugins in Jenkins	16
2.13. Installing the Git plugin	16
2.14. Setting up your first build job in Jenkins	17
2.15. Telling Jenkins where to find the source code	17
2.16. Scheduling the build jobs	18
2.17. Adding a build step	18
2.18. Configuring JUnit test reports and artifact archiving	19
2.19. Your first build job running	19
2.20. The Jenkins dashboard	20
2.21. A failed build	22
2.22. The list of all the broken tests	23
2.23. Details about a failed test	23
2.24. Now the build is back to normal	24
2.25. Adding a new build step and report to generate Javadoc	24
2.26. Jenkins will add a Javadoc link to your build results	25
2.27. Jenkins has a large range of plugins available	26
2.28. Adding another Maven goal to generate test coverage metrics	26
2.29. Configuring the test coverage metrics in Jenkins	27
2.30. Jenkins displays code coverage metrics on the build home page	27
2.31. Jenkins lets you display code coverage metrics for packages and classes	27
2.32. Jenkins also displays a graph of code coverage over time	28
3.1. You can download the Jenkins binaries from the Jenkins website	29
3.2. Jenkins setup wizard in Windows	30
3.3. The Jenkins start page	30
3.4. Starting Jenkins using Java Web Start	41
3.5. Installing Jenkins as a Windows service	41
3.6. Configuring the Jenkins Windows Service	42
3.7. The Jenkins home directory	44
3.8. The Jenkins jobs directory	44

3.9. The builds directory	44
3.10. Upgrading Jenkins from the web interface	45
4.1. You configure your Jenkins installation in the Manage Jenkins screen	47
4.2. System configuration in Jenkins	49
4.3. Configuring environment variables in Jenkins	50
4.4. Using a configured environment variable	51
4.5. JDK configuration in Jenkins	51
4.6. Installing a JDK automatically	52
4.7. Configuring Maven in Jenkins	53
4.8. Configuring system-wide MVN_OPTS	53
4.9. Configuring Ant in Jenkins	54
4.10. Configuring an email server in Jenkins	54
4.11. Configuring an email server in Jenkins to use a Google Apps domain	55
4.12. Configuring Jenkins to use a proxy	55
5.1. Jenkins supports four main types of build jobs	58
5.2. Creating a new build job	58
5.3. Keeping a build job forever	59
5.4. To display the Advanced Options, you need to click on the Advanced button	59
5.5. The “Block build when upstream project is building” option is useful when a single commit can affect several related projects	60
5.6. Jenkins provides built-in support for Subversion	61
5.7. Source code browser showing what code changes caused a build	62
5.8. System-wide configuration of the Git plugin	63
5.9. Entering a Git repo URL	64
5.10. Advanced configuration of a Git repo URL	64
5.11. Advanced configuration of the Git branches to build	64
5.12. Branches and regions	65
5.13. Choosing strategy	66
5.14. Git executable global setup	67
5.15. Repository browser	67
5.16. Polling log	67
5.17. Results of Git polling	67
5.18. Gerrit Trigger	68
5.19. Git Publisher	68
5.20. Merge results	68
5.21. GitHub repository browser	69
5.22. GitHub repository browser	69
5.23. There are many ways that you can configure Jenkins to start a build job	70
5.24. Triggering another build job even if the current one is unstable	70
5.25. Triggering a build via a URL using a token	73
5.26. Adding a build step to a freestyle build job	74
5.27. Configuring an Ant build step	75
5.28. Configuring an Execute Shell step	76

5.29. Adding a Groovy installation to Jenkins	79
5.30. Running Groovy commands as part of a build job	80
5.31. Running Groovy scripts as part of a build job	80
5.32. Reporting on test results	81
5.33. Configuring build artifacts	81
5.34. Build artifacts are displayed on the build results page and on the build job home page	82
5.35. Archiving source code and a binary package	83
5.36. Email notification	84
5.37. Creating a new Maven build job	85
5.38. Specifying the Maven goals	86
5.39. Maven build jobs—advanced options	86
5.40. Deploying artifacts to a Maven repository	87
5.41. After deployment the artifact should be available on your Enterprise Repository Manager	88
5.42. Redeploying an artifact	88
5.43. Deploying to Artifactory from Jenkins	89
5.44. Jenkins displays a link to the corresponding Artifactory repository	89
5.45. Viewing the deployed artifact in Artifactory	89
5.46. Viewing the deployed artifact and the corresponding Jenkins build in Artifactory	89
5.47. Managing modules in a Maven build job	90
5.48. Configuring extra Maven build steps	90
5.49. Adding a Grails installation to Jenkins	90
5.50. Configuring a Grails build step	91
5.51. Configuring the Gradle plugin	92
5.52. Setting up a Gradle build job	93
5.53. Incremental Gradle job	93
5.54. Configuring .NET build tools in Jenkins	93
5.55. A build step using MSBuild	94
5.56. A build step using NAnt	94
5.57. A build step using Rake	95
5.58. Publishing code quality metrics for Ruby and Rails	95
6.1. You configure your Jenkins installation in the Manage Jenkins screen	99
6.2. Configuring Maven test reports in a freestyle project	99
6.3. Installing the xUnit plugin	99
6.4. Publishing xUnit test results	99
6.5. Jenkins displays test result trends on the project home page	100
6.6. Jenkins displays a summary of the test results	100
6.7. The details of a test failure	101
6.8. Build time trends can give you a good indicator of how fast your tests are running	101
6.9. Jenkins also lets you see how long your tests take to run	102
6.10. Jenkins displays skipped tests as yellow	102
6.11. Installing the Cobertura plugin	108
6.12. Your code coverage metrics build needs to generate the coverage data	108

6.13. Configuring the test coverage metrics in Jenkins	108
6.14. Test coverage results contribute to the project status on the dashboard	109
6.15. Configuring the test coverage metrics in Jenkins	110
6.16. Displaying code coverage metrics	110
6.17. Configuring Clover reporting in Jenkins	111
6.18. Clover code coverage trends	111
6.19. Using business-focused, behavior-driven naming conventions for JUnit tests	112
6.20. Installing the HTML Publisher plugin	112
6.21. Publishing HTML reports	113
6.22. Jenkins displays a special link on the build job home page for your report	113
6.23. The DocLinks plugin lets you archive both HTML and non-HTML artifacts	113
6.24. Preparing a performance test script in JMeter	114
6.25. Preparing a performance test script in JMeter	117
6.26. Setting up the performance build to run every night at midnight	117
6.27. Performance tests can require large amounts of memory	117
6.28. Configuring the Performance plugin in your build job	117
6.29. The Jenkins Performance plugin keeps track of response time and errors	118
6.30. You can also view performance results per request	119
7.1. Enabling security in Jenkins	121
7.2. The Jenkins Sign up page	122
7.3. The list of users known to Jenkins	123
7.4. Displaying the builds that a user participates in	123
7.5. Creating a new user account by signing up	123
7.6. Synchronizing email addresses	123
7.7. You can also manage Jenkins users from the Jenkins configuration page	123
7.8. The Jenkins user database	124
7.9. Configuring LDAP in Jenkins	124
7.10. Using LDAP Groups in Jenkins	125
7.11. Selecting the security realm	126
7.12. Using Atlassian Crowd as the Jenkins Security Realm	127
7.13. Using Atlassian Crowd as the Jenkins Security Realm	127
7.14. Using Atlassian Crowd groups in Jenkins	127
7.15. Using custom scripts to handle authentication	127
7.16. Matrix-based security configuration	129
7.17. Setting up an administrator	129
7.18. Setting up other users	129
7.19. Project-based security	132
7.20. Configuring project-based security	132
7.21. Viewing a project	133
7.22. Setting up Extended Read Permissions	133
7.23. Setting up Role-based security	133
7.24. The Manage Roles configuration menu	134
7.25. Managing global roles	134

7.26. Managing project roles	134
7.27. Assigning roles to users	134
7.28. Configuring the Audit Trail plugin	135
7.29. Setting up Job Configuration History	136
7.30. Viewing Job Configuration History	136
7.31. Viewing differences in Job Configuration History	136
8.1. Configuring email notification	139
8.2. Configuring advanced email notification	141
8.3. Configuring email notification triggers	142
8.4. Customized notification message	143
8.5. Claiming a failed build	143
8.6. RSS Feeds in Jenkins	143
8.7. Creating a build radiator view	144
8.8. Displaying a build radiator view	145
8.9. Installing the Jenkins IM plugins	145
8.10. Jenkins needs its own dedicated IM user account	146
8.11. Setting up basic Jabber notification in Jenkins	146
8.12. Advanced Jabber configuration	146
8.13. Jenkins Jabber messages in action	148
8.14. Install the Jenkins IRC plugins	149
8.15. Advanced IRC notification configuration	149
8.16. Advanced build job IRC notification configuration	149
8.17. IRC notification messages in action	150
8.18. Jenkins notifications in Eclipse	150
8.19. Jenkins connection in NetBeans	151
8.20. Launching the Jenkins Tray Application	152
8.21. Running the Jenkins Tray Application	152
8.22. Creating a Notifo service for your Jenkins instance	153
8.23. Configuring Notifo notifications in your Jenkins build job	153
8.24. Receiving a Notifo notification on an iPhone	153
8.25. Using the Hudson Helper iPhone app	153
8.26. Sending SMS notifications via an SMS Gateway Service	154
8.27. Receiving notification via SMS	154
8.28. Configuring Jenkins Sounds rules in a build job	155
8.29. Configuring Jenkins Sounds	155
8.30. Configuring Jenkins Speaks	155
8.31. A Nabaztag	156
8.32. Configuring your Nabaztag	157
9.1. It's easy to configure Checkstyle rules in Eclipse	161
9.2. Configuring PMD rules in Eclipse	164
9.3. Generating code quality reports in a Maven build	170
9.4. Configuring the violations plugin for a freestyle project	171
9.5. Violations over time	171

9.6. Violations for a given build	171
9.7. Configuring the violations plugin for a freestyle project	172
9.8. Configuring the violations plugin for a Maven project	172
9.9. Jenkins Maven build jobs understand Maven multimodule structures	173
9.10. Activating the Violations plugin for an individual module	173
9.11. Installing the Checkstyle and Static Analysis Utilities plugins	173
9.12. Configuring the Checkstyle plugin	174
9.13. Displaying Checkstyle trends	174
9.14. A coverage/complexity scatter plot	175
9.15. You can click on any point in the graph to investigate further	175
9.16. Configuring the Task Scanner plugin is straightforward	175
9.17. The Open Tasks Trend graph	176
9.18. Code quality reporting by Sonar	176
9.19. Jenkins and Sonar	176
9.20. Configuring Sonar in Jenkins	177
9.21. Configuring Sonar in a build job	177
9.22. Scheduling Sonar builds	177
10.1. Creating a parameterized build job	179
10.2. Adding a parameter to the build job	180
10.3. Adding a parameter to the build job	180
10.4. Demonstrating a build parameter	180
10.5. Adding a parameter to a Maven build job	181
10.6. Many different types of parameters are available	181
10.7. Configuring a Choice parameter	182
10.8. Configuring a Run parameter	182
10.9. Configuring a File parameter	182
10.10. Adding a parameter to build from a Subversion tag	183
10.11. Building from a Subversion tag	183
10.12. Configuring a parameter for a Git tag	183
10.13. Building from a Git tag	184
10.14. Jenkins stores what parameter values were used for each build	184
10.15. A parameterized build job with DATABASE parameter	185
10.16. Adding a parameterized trigger to a build job	185
10.17. The build job you trigger must also be a parameterized build job	185
10.18. Passing a predefined parameter to a parameterized build job	185
10.19. Creating a multiconfiguration build job	186
10.20. Adding an axis to a multiconfiguration build	186
10.21. Defining an axis of slave nodes	187
10.22. Defining an axis of JDK versions	187
10.23. Defining a user-defined axis	187
10.24. Multiconfiguration build results	188
10.25. Setting up a combination filter	188
10.26. Build results using a combination filter	188

10.27. A job generated by the Maven Jenkins plugin	190
10.28. jenkins-master job generated	191
10.29. Artifactory Jenkins plugin configuration	194
10.30. Triggering several other builds after a build job	195
10.31. A build job dependency graph	196
10.32. Configuring a join in the phoenix-web-tests build job	196
10.33. A more complicated build job dependency graph	197
10.34. Adding a new lock	197
10.35. Configuring a build job to use a lock	197
10.36. Configuring a Maven release using the M2Release plugin	199
10.37. The Perform Maven Release menu option	199
10.38. Performing a Maven release in Jenkins	199
10.39. Adding a “Copy artifacts from another project” build step	200
10.40. Running web tests against a copied WAR file	201
10.41. Copying from a multiconfiguration build	202
10.42. Build jobs in the promotion process	202
10.43. Configuring a build promotion process	203
10.44. Configuring a manual build promotion process	203
10.45. Viewing the details of a build promotion	203
10.46. Using fingerprints in the build promotion process	204
10.47. Fetching the WAR file from the upstream build job	204
10.48. Archiving the WAR file for use in the downstream job	205
10.49. Fetching the WAR file from the integration job	205
10.50. You need to determine the fingerprint of the WAR file we use	205
10.51. Fetching the latest promoted WAR file	205
10.52. Promoted builds are indicated by a star in the build history	206
10.53. Reporting on aggregate test results	206
10.54. Viewing aggregate test results	206
10.55. Configuring a manual step in the build pipeline	207
10.56. Creating a Build Pipeline view	207
10.57. Configuring a Build Pipeline view	207
10.58. A Build Pipeline in action	208
11.1. Managing slave nodes	210
11.2. Creating a new slave node	210
11.3. Creating a Unix slave node	211
11.4. Taking a slave off-line when idle	212
11.5. Configuring tool locations	212
11.6. Your new slave node in action	213
11.7. Creating a slave node for JNLP	213
11.8. Launching a slave via Java Web Start	213
11.9. The Jenkins slave agent in action	214
11.10. The Jenkins slave failing to connect to the master	214
11.11. Configuring the Jenkins slave port	214

11.12. Installing the Jenkins slave as a Windows service	215
11.13. Managing the Jenkins Windows service	215
11.14. Letting Jenkins control a Windows slave as a Windows service	216
11.15. Running a build job on a particular slave node	216
11.16. Jenkins proactively monitors your build agents	218
11.17. You manage your EC2 instances using the Amazon AWS Management Console	219
11.18. Configuring an Amazon EC2 slave	219
11.19. Configuring an Amazon EC2 slave	220
11.20. Creating a new Amazon EC2 image	221
11.21. Bringing an Amazon EC2 slave online manually	221
12.1. A simple automated deployment pipeline	229
12.2. Copying the binary artifact to be deployed	229
12.3. Deploying to Tomcat using the Deploy Plugin	229
12.4. Adding a “Build selector for Copy Artifact” parameter	230
12.5. Configuring a build selector parameter	231
12.6. Specify where to find the artifacts to be deployed	231
12.7. Choosing the build to redeploy	231
12.8. Using the “Specified by permalink” option	232
12.9. Using a specific build	232
12.10. Using a Maven Enterprise Repository	232
12.11. Deploying an artifact from a Maven repository	234
12.12. Preparing the WAR to be deployed	235
12.13. Configuring a remote host	235
12.14. Deploying files to a remote host in the build section	236
12.15. Deploying files to a remote host in the post-build actions	236
13.1. Discarding old builds	237
13.2. Discarding old builds—advanced options	238
13.3. Viewing disk usage	238
13.4. Displaying disk usage for a project	238
13.5. Displaying project disk usage over time	239
13.6. Maven build jobs—advanced options	239
13.7. Jenkins Load Statistics	240
13.8. The Jenkins Monitoring plugin	240
13.9. The builds directory	241
13.10. The Jenkins Backup Manager Plugin	242
13.11. Configuring the Jenkins Backup Manager	242
13.12. Configuring the Thin Backup plugin	242
13.13. Restoring a previous configuration	243
13.14. Reloading the configuration from disk	243
13.15. Jenkins will inform you if your data isn't compatible with the current version	244
13.16. Managing out-of-date build jobs data	245
A.1. A project containing freely-named test classes	249

Copyright

Copyright © 2010 John Ferguson Smart

Printed version published by O'Reilly Media, 1005 Gravenstein Highway North, Sebastopol, CA 95472.

Online version published by Wakaleo Consulting, Sydney Australia.

This work is licensed under a Creative Commons Attribution-Noncommercial-No Derivative Works 3.0 United States license. For more information about this license, see <http://creativecommons.org/licenses/by-nc-nd/3.0/us/>. You are free to share, copy, distribute, display, and perform the work under the following conditions:

- You must attribute the work to John Ferguson Smart
- You may not use this work for commercial purposes.
- You may not alter, transform, or build upon this work.

Java™ and all Java-based trademarks and logos are trademarks or registered trademarks of Sun Microsystems, Inc., in the United States and other countries.

Eclipse™ is a trademark of the Eclipse Foundation, Inc., in the United States and other countries.

Apache and the Apache feather logo are trademarks of The Apache Software Foundation.

Linux® is the registered trademark of Linus Torvalds in the U.S. and other countries.

Many of the designations used by manufacturers and sellers to distinguish their products are claimed as trademarks. Where those designations appear in this book, and Wakaleo Consulting was aware of a trademark claim, the designations have been printed in caps or initial caps.

While every precaution has been taken in the preparation of this book, the publisher and authors assume no responsibility for errors or omissions, or for damages resulting from the use of the information contained herein.

Foreword

Kohsuke Kawaguchi

Seven years ago, I wrote the first line of code that started this whole project that is now known as Jenkins, and was originally called Hudson. I used to be the guy who broke the build, so I needed a program to catch my mistakes before my colleagues did. It was just a simple tool that did a simple thing. But it rapidly evolved, and now I'd like to think that it's the most dominant CI server on the market bar none, encompassing a broad plugin ecosystem, commercial distributions, hosted Jenkins-as-a-Service, user groups, meet-ups, trainings, and so on.

As with most of my other projects, this project was open-sourced since its inception. Over its life it critically relied on the help and love of other people, without which the project wouldn't be in the current state. During this time I've also learned a thing or two about running open source projects. From that experience, I think people often overlook that there are many ways to help an open source project, of which writing code is just one of many. There's spreading words, helping other users, organizing meet-ups, and yes, there's writing documentation.

In this sense, John is an important part of the Jenkins community, even though he hasn't contributed code—instead, he makes Jenkins more approachable to new users. For example, he has a popular blog that's followed by many, where he regularly talks about continuous integration practices and other software development topics. He is good at explaining things so that people new to Jenkins can still understand them, which is something often hard for people like me who develop Jenkins day in day out. He is also well-known for his training courses, of which Jenkins is a part. This is another means by which he makes Jenkins accessible for more people. He clearly has a passion for evangelizing new ideas and teaching fellow developers to be more productive.

These days I spend my time at CloudBees where I focus my time on Open Source Jenkins, the CloudBees pro version of Jenkins where we build plugins on top of Jenkins, and taking Jenkins to the private and public cloud with CloudBees DEV@cloud service. In this role I now have more interaction with John than before, and my respect for his passion has only grown.

So I was truly delighted that he took on the daunting task of writing a book about Jenkins. It gives a great overview of the typical main ingredients of continuous integration. And for me personally, I always get asked if there's a book about Jenkins, and I can finally answer this question positively! But more importantly, this book reflects his passion, and his long experience in teaching people how to use Jenkins, in combination with other things. But don't take my words for it. You'll just need to read on to see it for yourself.

Preface

1. Audience

This book is aimed at relatively technical readers, though no prior experience with Continuous Integration is assumed. You may be new to Continuous Integration, and would like to learn about the benefits it can bring to your development team. Or, you might be using Jenkins or Hudson already, and want to discover how you can take your Continuous Integration infrastructure further.

Much of this book discusses Jenkins in the context of Java or JVM-based projects. Nevertheless, even if you're using another technology stack, this book should give you a good grounding in Continuous Integration with Jenkins. We discuss how to build projects using several non-Java technologies, including as Grails, Ruby on Rails, and .NET. In addition, many topics, such as general configuration, notification, distributed builds, and security are applicable no matter what language you're using.

2. Book Layout

Continuous Integration is like a lot of things: the more you put in, the more value you'll get out. While even a basic Continuous Integration setup will produce positive improvements in your team process, there are significant advantages to gradually assimilating and implementing some of the more advanced techniques as well. To this end, this book is organized as a progressive trek into the world of Continuous Integration with Jenkins, going from simple to more advanced. In the first chapter, we start off with a sweeping overview of what Jenkins is all about, in the form of a high-level guided tour. From there, we progress into how to install and configure your Jenkins server and how to set up basic build jobs. Once you've mastered the basics, we'll delve into more advanced topics, including automated testing practices, security, more advanced notification techniques, and measuring and reporting on code quality metrics. Next, we move on to more advanced build techniques such as matrix builds, distributed builds, and cloud-based CI, before discussing how to implement Continuous Deployment with Jenkins. Finally, we cover some tips on maintaining your Jenkins server.

3. Jenkins or Hudson?

As we mention in the introduction, Jenkins was originally, and up until recently, known as Hudson. In 2009, Oracle purchased Sun and inherited the Hudson code base. In early 2011, tensions between Oracle and the open source community reached rupture point and the project forked into two separate entities: Jenkins, run by most of the original Hudson developers, and Hudson, which remained under the control of Oracle.

As the title suggests, this book is primarily focused on Jenkins. However, much of the book was initially written before the fork, and the products remain very similar. So, although the examples and illustrations do usually refer to Jenkins, almost all of what's discussed will also apply to Hudson.

4. Font Conventions

This book follows certain conventions for font usage. Understanding these conventions up-front makes it easier to use this book.

Italic

Used for filenames, file extensions, URLs, application names, emphasis, and new terms when they're first introduced.

Constant width

Used for Java class names, methods, variables, properties, data types, database elements, and snippets of code that appear in text.

Constant width bold

Used for commands you enter at the command line and to highlight new code inserted in a running example.

Constant width italic

Used to annotate output.

5. Command-Line Conventions

From time to time, this book discusses command-line instructions. When we do, output produced by the console (e.g., command prompts or screen output) is displayed in normal characters, and commands (what you type) are written in **bold**. For example:

```
$ ls -al
total 168
drwxr-xr-x  16 johnsmart  staff    544 21 Jan 07:20 .
drwxr-xr-x+ 85 johnsmart  staff   2890 21 Jan 07:10 ..
-rw-r--r--   1 johnsmart  staff    30 26 May 2009 .owner
-rw-r--r--@   1 johnsmart  staff   1813 16 Apr 2009 config.xml
drwxr-xr-x 181 johnsmart  staff   6154 26 May 2009 fingerprints
drwxr-xr-x  17 johnsmart  staff    578 16 Apr 2009 jobs
drwxr-xr-x   3 johnsmart  staff    102 15 Apr 2009 log
drwxr-xr-x  63 johnsmart  staff   2142 26 May 2009 plugins
-rw-r--r--   1 johnsmart  staff     46 26 May 2009 queue.xml
-rw-r--r--@   1 johnsmart  staff     64 13 Nov 2008 secret.key
-rw-r--r--   1 johnsmart  staff  51568 26 May 2009 update-center.json
drwxr-xr-x   3 johnsmart  staff    102 26 May 2009 updates
drwxr-xr-x   3 johnsmart  staff    102 15 Apr 2009 userContent
drwxr-xr-x  12 johnsmart  staff    408 17 Feb 2009 users
drwxr-xr-x  28 johnsmart  staff    952 26 May 2009 war
```

Where necessary, the backslash character at the end of the line is used to indicate a line break: you can type this all on one line (without the backslash) if you prefer. Don't forget to ignore the ">" character at the start of the subsequent lines—it's a Unix prompt character:

```
$ wget -O - http://jenkins-ci.org/debian/jenkins-ci.org.key \
```



```
> | sudo apt-key add -
```

For consistency, unless we're discussing a Windows-specific issue, we'll use Unix-style command prompts (the dollar sign, "\$"), as shown here:

```
$ java -jar jenkins.war
```

or:

```
$ svn list svn://localhost
```

However, unless we say otherwise, Windows users can safely use these commands from the Windows command console:

```
C:\Documents and Settings\Owner> java -jar jenkins.war
```

or:

```
C:\Documents and Settings\Owner> svn list svn://localhost
```

6. Contributors

This book was not written alone. Rather, it's been a collaborative effort involving many people playing different roles. In particular, the following people generously contributed their time, knowledge, and writing skill to make this a better book:

- Evgeny Goldin is a Russian-born software engineer living in Israel. He's a lead developer at Thomson Reuters where he's responsible for a number of activities, some of which are directly related to Maven, Groovy, and build tools such as Artifactory and Jenkins. He has vast experience in a range of technologies, including Perl, Java, JavaScript, and Groovy. Build tools and dynamic languages are Evgeny's favorite subjects about which he often writes, presents, or blogs. These days he writes for GroovyMag, Methods & Tools, and runs two open source projects of his own: Maven-plugins¹ and GCommons². He blogs at <http://evgeny-goldin.com/blog> and can be found on Twitter as @evgeny_goldin.

Evgeny contributed a section on generating your Maven build jobs automatically in Chapter 10, Advanced Builds.

- Matthew McCullough is an energetic 15 year veteran of enterprise software development, open source education, and co-founder of Ambient Ideas, LLC, a Denver consultancy. Matthew currently is a trainer for GitHub.com, author of the Git Master Class series for O'Reilly, speaker at over 30 national and international conferences, author of 3 of the top 10 DZone RefCards, and President of the Denver Open Source Users Group. His current topics of research center around project automation: build tools (Maven, Leiningen, Gradle), distributed version control (Git), Continuous Integration (Jenkins), and Quality Metrics (Sonar). Matthew resides in Denver, Colorado with his

¹ <http://evgeny-goldin.com/wiki/Maven-plugins>

² <http://evgeny-goldin.com/wiki/GCommons>

beautiful wife and two young daughters, who are active in nearly every outdoor activity Colorado has to offer.

Matthew wrote the section on integrating Git with Jenkins in Chapter 5, *Setting Up Your Build Jobs*.

- Juven Xu is a software engineer from China who works for Sonatype. An active member of the open source community and recognized Maven expert, Juven was responsible for the Chinese translation of *Maven: The Definitive Guide* as well as an original Chinese reference book on Maven. He's also currently working on the Chinese translation of the present book.

Juven wrote the section on IRC notifications in Chapter 8, *Notification*.

- Rene Groeschke is a software engineer at Cassidian Systems, formerly known as EADS Deutschland GmbH, as well as an open source enthusiast. A certified ScrumMaster with about 7 years experience as a programmer in several enterprise Java projects, he's especially focused on Agile methodologies like Continuous Integration and Test-Driven Development. Besides his daily business, the University of Corporate Education in Friedrichshafen allows him to spread the word about scrum and scrum related topics by giving lectures to the bachelor-level students of information technology.

Rene contributed the section on building projects with Gradle in Chapter 5, *Setting Up Your Build Jobs*.

7. The Review Team

The technical review process for this book was a little different than the approach taken for most books. Rather than having one or two technical reviewers read the entire book near the end of the book writing process, a team of volunteers from the Jenkins community, including many key Jenkins developers, were able to read chapters as they were written. This review team was made up of the following people: Alan Harder, Andrew Bayer, Carlo Bonamico, Chris Graham, Eric Smalling, Gregory Boissinot, Harald Soevik, Julien Simpson, Juven Xu, Kohsuke Kawaguchi, Martijn Verburg, Ross Rowe, and Tyler Ballance.

8. Book Sponsors

This book would not have been possible without the help of several organizations who were willing to assist with and fund the book-writing process.

8.1. Wakaleo Consulting

Wakaleo Consulting³ is a consulting company that helps organizations optimize their software development process. Lead by John Ferguson Smart, author of this book and *Java Power Tools*⁴,

³ <http://www.wakaleo.com>

Wakaleo Consulting provides consulting, training and mentoring services in Agile Java Development and Testing Practices, Software Development Life Cycle optimization, and Agile Methodologies.

Wakaleo helps companies with training and assistance in areas such as Continuous Integration, Build Automation, Test-Driven Development, Automated Web Testing, and Clean Code, using open source tools such as Maven, Jenkins, Selenium 2, and Nexus. Wakaleo Consulting also runs public and on-site training around Continuous Integration and Continuous Deployment, Build Automation, Clean Code practices, Test-Driven Development, and Behavior-Driven Development, including Certified Scrum Developer (CSD) courses.

8.2. CloudBees

CloudBees⁵ is the only cloud company focused on servicing the complete develop-to-deploy life cycle of Java web applications in the cloud. The company is also the world's premier expert on the Jenkins/Hudson Continuous Integration tool.

Jenkins/Hudson creator Kohsuke Kawaguchi leads a CloudBees team of experts from around the world. They've created Nectar, a supported and enhanced version of Jenkins that's available on-premise by subscription. If you depend on Jenkins for mission-critical software processes, Nectar provides a highly-tested, stable, and fully-supported version of Jenkins. It also includes Nectar-only functionality such as automatic scaling to VMware virtual machines.

If you're ready to explore the power of Continuous Integration in the cloud, CloudBees makes Jenkins/Hudson available as part of its DEV@cloud build platform. You can get started with Jenkins instantly and can scale as needed—no big up-front investment in build servers, no more limited capacity for builds, and no maintenance hassles. Once an application is ready to go live, you can deploy on CloudBees's RUN@cloud Platform as a Service in just a few clicks.

With CloudBees's DEV@cloud and RUN@cloud services, you don't have to worry about servers, virtual machines, or IT staff. And with Nectar, you enjoy the most powerful, stable, supported Jenkins available.

8.3. Odd-e

Odd-e⁶ is an Asian-based company that builds products in innovative ways and helps others achieve the same. The team consists of experienced coaches and product developers who work according to the values of scrum, agile, lean, and craftsmanship, and the company is structured the same way. For example, Odd-e doesn't have an organizational hierarchy or managers making decisions for others. Instead, individuals self-organize and use all their skills to continuously improve their competence. The company provides training and follow-up coaching to help others collaboratively seek and develop a better way of working.

⁴ <http://oreilly.com/catalog/9780596527938>

⁵ <http://www.cloudbees.com>

⁶ <http://www.odd-e.com>

It's not the job but the values that bind Odd-e together. Its members love building software, value learning and contribution over maximizing profit, and are committed to supporting open source development in Asia.

9. Using Code Examples

This book is an open source book, published under the Creative Commons License. The book was written in DocBook, using XmlMind. The book's source code can be found on GitHub at <http://github.com/wakaleo/jenkins-the-definitive-guide-book>.

The sample Jenkins projects used in this book are open source and freely available online—see the book's web page at <http://www.wakaleo.com/books/jenkins-the-definitive-guide> for more details.

This book is here to help you get your job done. In general, you may use the code in this book in your programs and documentation. You don't need to contact us for permission unless you're reproducing a significant portion of the code. For example, writing a program that uses several chunks of code from this book doesn't require permission. Selling or distributing a CD-ROM of examples from O'Reilly books does require permission. Answering a question by citing this book and quoting example code doesn't require permission. Incorporating a significant amount of example code from this book into your product's documentation does require permission.

We appreciate, but do not require, attribution. An attribution usually includes the title, author, publisher, and ISBN. For example: "Jenkins: The Definitive Guide by John Ferguson Smart (O'Reilly). Copyright 2011 John Ferguson Smart, 978-1-449-30535-2."

If you feel your use of code examples falls outside fair use or the permission given above, feel free to contact us at [<permissions@oreilly.com>](mailto:permissions@oreilly.com).

10. Safari® Books Online

Note

Safari Books Online is an on-demand digital library that lets you easily search over 7,500 technology and creative reference books and videos to find the answers you need quickly.

With a subscription, you can read any page and watch any video from our library online. Read books on your cell phone and mobile devices. Access new titles before they're available for print, and get exclusive access to manuscripts in development and post feedback for the authors. Copy and paste code samples, organize your favorites, download chapters, bookmark key sections, create notes, print out pages, and benefit from tons of other time-saving features.

⁷ <http://my.safaribooksonline.com/?portal=oreilly>

O'Reilly Media has uploaded this book to the Safari Books Online service. To have full digital access to this book and others on similar topics from O'Reilly and other publishers, sign up for free at <http://my.safaribooksonline.com>⁷.

11. How to Contact Us

Please address comments and questions concerning this book to the publisher:

O'Reilly Media, Inc.
1005 Gravenstein Highway North
Sebastopol, CA 95472
800-998-9938 (in the United States or Canada)
707-829-0515 (international or local)
707-829-0104 (fax)

We have a web page for this book, where we list errata, examples, and any additional information. You can access this page at:

<http://www.oreilly.com/catalog/9781449305352>

To comment or ask technical questions about this book, send email to:

[<bookquestions@oreilly.com>](mailto:bookquestions@oreilly.com)

For more information about our books, courses, conferences, and news, see our website at <http://www.oreilly.com>.

Find us on Facebook: <http://facebook.com/oreilly>

Follow us on Twitter: <http://twitter.com/oreillymedia>

Watch us on YouTube: <http://www.youtube.com/oreillymedia>

12. Acknowledgments

First and foremost, my wonderful wife, Chantal, and boys, James and William, without whose love, support, and tolerance this book would not have been possible.

I'd like to thank Mike Loukides for working with me once again on this book project, and the whole O'Reilly team for their high standards of work.

Thank you to Kohsuke Kawaguchi for having created Jenkins, and for still being the driving force behind this brilliant product. Thanks also to Francois Dechery, Sacha Labourey, Harpreet Singh, and the rest of the CloudBees team for their help and support.

I'm also very grateful to those who took the time and energy to contribute work to the book: Evgeny Goldin, Matthew McCullough, Juven Xu, and Rene Groeschke.

A great thanks goes out to the following reviewers who provided valuable feedback throughout the whole writing process: Alan Harder, Andrew Bayer, Carlo Bonamico, Chris Graham, Eric Smalling, Gregory Boissinot, Harald Soevik, Julien Simpson, Juven Xu, Kohsuke Kawaguchi, Martijn Verburg, Ross Rowe, and Tyler Ballance.

Thank you to Andrew Bayer, Martijn Verburg, Matthew McCullough, Rob Purcell, Ray King, Andrew Walker, and many others, whose discussions and feedback provided me with inspiration and the ideas that made this book what it is.

Many other people have helped in various ways to make this book much richer and more complete than it would have been otherwise: Geoff and Alex Bullen, Pete Thomas, Gordon Weir, Jay Zimmerman, Tim O'Brien, Russ Miles, Richard Paul, Julien Simpson, John Stevenson, Michael Neale, Arnaud Héritier, and Manfred Moser.

And finally, a great thank you to the Hudson/Jenkins developer and user communities for their ongoing encouragement and support.

Chapter 1. Introducing Jenkins

1.1. Introduction

Continuous Integration, also known as CI, is a cornerstone of modern software development. In fact, it's a real game changer. Introducing CI into an organization radically alters the way teams think about the whole development process. CI has the potential to enable and trigger a series of incremental process improvements, going from a simple scheduled automated build right through to continuous delivery into production. A good CI infrastructure can streamline the development process, help detect and fix bugs faster, provide a useful project dashboard for both developers and non-developers, and, ultimately, help teams deliver more real value to the end user. Every professional development team, no matter how small, should be practicing CI.

1.2. Continuous Integration Fundamentals

Back in the days of waterfall projects and Gantt charts, before the introduction of CI practices, development team time and energy was regularly drained in the period leading up to a release by what was known as the Integration Phase. During this phase, the code changes made by individual developers or small teams were brought together piecemeal and forged into a working product. This was hard work, sometimes involving the integration of months of conflicting changes. It was very hard to anticipate the types of issues that would crop up, and even harder to fix them. Doing so could require reworking code that had been written weeks or months before. This painful process, fraught with risk and danger, often lead to significant delivery delays, unplanned costs, and, as a result, unhappy clients. CI was born to address these issues.

CI, in its simplest form, is a tool that monitors your version control system for changes. Whenever a change is detected, this tool automatically compiles and tests your application. If something goes wrong, the tool immediately notifies the developers so that they can fix the issue immediately.

But CI Integration can do much more than this. CI can also help you keep tabs on the health of your code base, automatically monitoring code quality and code coverage metrics, and helping keep technical conflicts and maintenance costs low. Publicly-visible code quality metrics can also encourage developers to take pride in the quality of their code and strive to improve it. Combined with automated end-to-end acceptance tests, CI can also act as a communication tool, publishing a clear picture of the current state of development efforts. Finally, it can simplify and accelerate delivery by helping automate the deployment process, letting you deploy the latest version of your application either automatically or as a one-click process.

In essence, CI is about reducing risk by providing faster feedback. First and foremost, it helps identify and fix integration and regression issues faster, resulting in smoother, quicker delivery, and fewer bugs. By providing better visibility for both technical and non-technical team members on the state of the project, CI can open and facilitate communication channels between team members and encourage

collaborative problem solving and process improvement. By automating the deployment process, CI helps get your software into the hands of the testers and the end users faster, more reliably, and with less effort.

This idea of automated deployment is important. Indeed, if you take automating the deployment process to its logical conclusion, you could push every build that passes the necessary automated tests into production. The practice of automatically deploying every successful build directly into production is generally known as Continuous Deployment.

However, a pure Continuous Deployment approach isn't always appropriate. For example, many users would not appreciate new versions falling into their laps several times a week, and prefer a more predictable (and transparent) release cycle. Commercial and marketing considerations might also play a role in deciding when a new release should actually be deployed.

The notion of Continuous Delivery is a slight variation on the idea of Continuous Deployment that takes these considerations into account. With Continuous Delivery, any and every successful build that has passed all the relevant automated tests and quality gates can potentially be deployed into production via a fully automated one-click process, and be in the hands of the end-user within minutes. However, the process isn't automatic: it's the business, rather than IT, that decides the best time to deliver the latest changes.

So, CI techniques, and in particular Continuous Deployment and Continuous Delivery, are very much about providing value to the end user faster. How long does it take your team to get a small code change out to production? How much of this process involves problems that could have been fixed earlier, had you known about the code changes that Joe down the corridor was making? How much is taken up by labor-intensive manual testing by QA teams? How much involves manual deployment steps, the secrets of which are known only to a select few? CI isn't a silver bullet by any means, but it can certainly help streamline many of these problems.

But CI is a mindset as much as a toolset. To get the most out of CI, a team needs to adopt a CI mentality. For example, your projects must have reliable, repeatable, and automated build and deployment processes, involving no human intervention. Fixing broken builds should take an absolute priority, and never be left to stagnate. And since the trust you place in your CI server depends to a great extent on the quality of your tests, the team needs to place a very strong emphasis on high quality tests and testing practices.

In this book we'll be looking at how to implement a robust and comprehensive CI solution using Jenkins or Hudson.

1.3. Introducing Jenkins (né Hudson)

Jenkins, originally called Hudson, is an open source CI tool written in Java. Boasting a dominant market share, Jenkins is used by teams of all sizes for projects in a wide variety of languages and technologies, including .NET, Ruby, Groovy, Grails, PHP, and more, as well as Java. So, what made Jenkins such a success and why use Jenkins for your CI infrastructure?

Firstly, Jenkins is easy to use. The user interface is simple, intuitive, and visually appealing, and Jenkins as a whole has a very low learning curve. As you'll see in the next chapter, you can get started with Jenkins in a matter of minutes.

However, Jenkins doesn't sacrifice power or extensibility: it's also extremely flexible and easy to adapt to your own purposes. Hundreds of open source plugins are available, with more coming out every week. These plugins cover everything from version control systems, build tools, code quality metrics, build notifiers, integration with external systems, UI customization, games, and much more. Installing them is quick and easy.

Last, but certainly not least, much of Jenkins's popularity comes from the size and vibrancy of its community. The Jenkins community is a large, dynamic, reactive, and welcoming bunch, with passionate champions, active mailing lists, IRC channels, and a very vocal blog and twitter account. The development pace is fast, with releases coming out weekly with the latest new features, bug fixes, and plugin updates.

However, Jenkins also caters to users who aren't comfortable with upgrading on a weekly basis. For those who prefer a less-hecktic release pace, there's also a Long-term Support, or LTS, release line that lags behind the latest release in favor of more stability and a slower rate of change. New LTS releases come out every three months or so, with important bug fixes being backported. This concept is similar to the Ubuntu LTS releases.

1.4. From Hudson to Jenkins—A Short History

Jenkins is the result of one visionary developer, Kohsuke Kawaguchi, who started the project as a hobby under the name of Hudson in late 2004 whilst working at Sun. As Hudson evolved over the years, it was adopted by more and more teams within Sun for their own projects. By early 2008, Sun recognized the quality and value of the tool, and asked Kohsuke to work on Hudson full-time to provide professional services and support. By 2010, Hudson had become the leading CI solution with a market share of over 70%.

In 2009, Oracle purchased Sun. Towards the end of 2010, tensions arose between the Hudson developer community and Oracle, initially triggered by problems with the Java.net infrastructure, and aggravated by issues related to Oracle's claim to the Hudson trademark. These tensions also reflected strong underlying disagreements about the way the project was being managed by Oracle. Indeed, Oracle wanted to move towards a more strictly controlled development process with a slower release schedule, whereas most of the core Hudson developers, led by Kohsuke, preferred to continue with the open, flexible, and fast-paced community-focused model that had worked so well for Hudson in the past.

In January 2011, the Hudson developer community decisively voted to rename the project to Jenkins. They subsequently migrated the original Hudson code base to a new GitHub project¹ and continued their work there. The vast majority of core and plugin developers broke away and followed Kohsuke

¹ <https://github.com/jenkinsci>

Kawaguchi and other core contributors to the Jenkins camp, where the bulk of the development activity can be seen today.

After the fork, a majority of users also followed the Jenkins developer community and switched to Jenkins. At the time of writing, polls show that some 75% of Hudson users had switched to Jenkins, while 13% were still using Hudson, and another 12% were using both Hudson and Jenkins or were in the process of migrating to Jenkins.

Nevertheless, Oracle and Sonatype (the company behind Maven and Nexus) have continued to work on the Hudson code base (now also hosted on GitHub at <https://github.com/hudson>), but with a very different focus. Indeed, the Sonatype developers have concentrated on major underlying infrastructure changes around, among other areas, Maven integration, the dependency injection framework, and the plugin architecture.

1.5. Should I Use Jenkins or Hudson?

So, should you use Jenkins or Hudson? Since this is a book on Jenkins, here are a few reasons why you might want to opt for Jenkins:

- Jenkins is the new Hudson. In fact, Jenkins is simply the old Hudson with a new name, so if you liked Hudson, you'll like Jenkins! Jenkins uses the Hudson code base, and the development team and project philosophy remain the same. In a nutshell, the original developers, who wrote the vast majority of the Hudson core, simply resumed business as usual after the fork working on the Jenkins project.
- The Jenkins community. Like many of the more successful Open Source projects, much of Hudson's strength came from its large and dynamic community, and its massive adoption. Bugs are identified (and generally fixed) much more rapidly, and, if you have a problem, chances are someone else will have had them too! If you run into trouble, post a question on the mailing list or IRC channel—there's sure to be someone who can help.
- The fast development pace. Jenkins continues the rapid release cycles that typified Hudson, which many developers love. New features, plugins, and bug fixes come out weekly, and the turn-around time for bug fixes can be very short indeed. And, if you prefer more stability, there are always the LTS releases

And, in the interest of balance, here are some reasons you might prefer to stick with Hudson:

- If it ain't broke, don't fix it. You already have a Hudson installation that you're happy with, and don't feel the need to upgrade to the latest version.
- Enterprise integration and Sonatype tools. Hudson is likely to place a strong emphasis on integration with enterprise tools such as LDAP/Active Directory, and the Sonatype products such as Maven 3, Nexus, and M2Eclipse, whereas Jenkins is more open to other competing tools such as Artifactory and Gradle.

- Plugin architecture. If you intend to write your own Jenkins/Hudson plugins, you should be aware that Sonatype is working on providing JSR-330 dependency injection for Hudson plugins. New developers may find this approach easier to use, though it does raise issues about future plugin compatibility between Jenkins and Hudson.

The good news is, no matter whether you're using Jenkins or Hudson, the products remain very similar, and the vast majority of techniques and tips discussed in this book apply equally well to both. Indeed, to illustrate this point, many screenshots in this book refer to Hudson rather than Jenkins.

1.6. Introducing Continuous Integration into Your Organization

CI isn't an all-or-nothing affair. In fact, introducing CI into an organization takes you on a path that progresses through several distinct phases. Each of these phases involves incremental improvements to the technical infrastructure as well as, perhaps more importantly, improvements in the practices and culture of the development team itself. In the following paragraphs, I've tried to paint an approximate picture of each phase.

1.6.1. Phase 1—No Build Server

Initially, the team has no central build server of any kind. Software is built manually on a developers' machines, though they may use an Ant script or similar to do so. Source code may be stored in a central source code repository, but developers don't necessarily commit their changes on a regular basis. Some time before a release is scheduled, a developer manually integrates the changes, a process which is generally associated with pain and suffering.

1.6.2. Phase 2—Nightly Builds

In this phase, the team has a build server, and automated builds are scheduled on a regular (typically nightly) basis. This build simply compiles the code, as there are no reliable or repeatable unit tests. Indeed, automated tests, if they're written, aren't a mandatory part of the build process, and may well not run correctly at all. However, developers now commit their changes regularly, at least at the end of every day. If a developer commits code changes that conflict with another developer's work, the build server alerts the team via email the following morning. Nevertheless, the team still tends to use the build server for information purposes only—they feel little obligation to fix a broken build immediately, and builds may stay broken on the build server for some time.

1.6.3. Phase 3—Nightly Builds and Basic Automated Tests

The team is now starting to take CI and automated testing more seriously. The build server is configured to kick off a build whenever new code is committed to the version control system, and team members are able to easily see what changes in the source code triggered a particular build, and what issues these changes address. In addition, the build script compiles the application and runs a set of automated unit and/or integration tests. In addition to email, the build server also alerts team members about integration

issues using more proactive channels such as Instant Messaging. Broken builds are now generally fixed quickly.

1.6.4. Phase 4—Enter the Metrics

Automated code quality and code coverage metrics are now run to help evaluate the quality of the code base and (to some extent, at least) the relevance and effectiveness of the tests. The code quality build also automatically generates API documentation for the application. All this helps teams keep the quality of the code base high, alerting team members if good testing practices are slipping. The team has also set up a “build radiator,” a dashboard view of the project status that's displayed on a prominent screen visible to all team members.

1.6.5. Phase 5—Getting More Serious About Testing

The benefits of CI are closely related to solid testing practices. Now, practices like Test-Driven Development are more widely practiced, resulting in a growing confidence in the results of the automated builds. The application is no longer simply compiled and tested, but if the tests pass, it's automatically deployed to a non-production application server for more comprehensive end-to-end tests and performance tests.

1.6.6. Phase 6—Automated Acceptance Tests and More Automated Deployment

Acceptance-Test Driven Development is practiced, guiding development efforts and providing high-level reporting on the state of the project. These automated tests use Behavior-Driven Development and Acceptance-Test Driven Development tools to act as communication and documentation tools as much as testing tools, publishing reports on test results in business terms that non-developers can understand. Since these high-level tests are automated at an early stage in the development process, they also provide a clear idea of what features have been implemented, and which remain to be done. The application is automatically deployed into test environments for testing by the QA team either when changes are committed or on a nightly basis; a version can be deployed (or “promoted”) to UAT and possibly production environments using a manually-triggered build when testers consider it ready. The team is also capable of using the build server to back out a release, rolling back to a previous version, if something goes horribly wrong.

1.6.7. Phase 7—Continuous Deployment

Confidence in the automated unit, integration, and acceptance tests is now such that teams can apply the automated deployment techniques developed in the previous phase to push out new changes directly into production.

The progression between levels here is of course somewhat approximate, and may not always match real-world situations. For example, you may well introduce automated web tests before integrating code quality and code coverage reporting. However, this section should give a general idea of how implementing a CI strategy in a real world organization generally works.

1.7. Where to Now?

Throughout the remainder of this book, as you study the various features Jenkins has to offer, as well as the practices required to make the most of these features, you'll see how you can progress through each of these levels with Jenkins. Remember, most of the examples used in the book are available online (see <http://www.wakaleo.com/books/jenkins-the-definitive-guide> for more details), so you can get your hands dirty too!

Chapter 2. Your First Steps with Jenkins

2.1. Introduction

In this chapter, you'll take a quick guided tour through some of Jenkins' key features. You'll see first-hand just how easy it is to install Jenkins and set up your first Jenkins automated build job. We won't dwell on the details too much—there are more details to come in the following chapters, as well as a detailed chapter on Jenkins Administration at the end of the book (Chapter 13, Maintaining Jenkins). This chapter is just an introduction. Still, by the end of the chapter, you'll also be keeping tabs on test results, generating javadoc, and publishing code coverage reports! You've got a lot of ground to cover, so let's get started!

2.2. Preparing Your Environment

There are two ways you can tackle this chapter. You can read through it without touching a keyboard, just to get an overview of what Jenkins is about. Or, you can get your hands dirty, and follow along on your own machine.

If you do want to follow along at home, you may need to set up some software on your machine. Remember, the most basic function of any CI tool is monitoring source code in a version control system, and fetching and building the latest version of source code whenever any changes are committed. So, you'll need a version control system. In your case, you'll be using Git¹. The central source code repository for our sample project is stored on GitHub². Don't worry about messing up this repository with your own changes, though: you'll be creating your own fork of the repository that you can use as you wish. If you haven't used Git and/or don't have an account on GitHub yet, don't worry. We'll walk through the basics, and the whole installation process is well documented on the GitHub website. We'll explain how to set it all up in great detail further on.

In this chapter, you'll be using Jenkins to build a Java application using Maven. Maven is a widely-used build tool in the Java world, with many powerful features, such as declarative dependency management, convention over configuration, and a large range of plugins. For your build, you'll also be using recent versions of the Java Development Kit (JDK) and Maven, but if you don't have these installed on your machine, don't fret! As you'll see, Jenkins installs them for you.

¹ <http://git-scm.com>

² <https://github.com>

2.2.1. Installing Java

The first thing you'll need to install on your machine is Java. Jenkins is a Java web application, so you'll need at least the Java Runtime Environment, or JRE, to run it. For the examples in this chapter, you'll need a recent version of Java (these examples were written with Java 6 update 17, and the latest release at the time of writing was Java 6 update 19). If you're not sure what version you're running, you can find out from the command line by running `java -version`. If Java is installed on your machine, you should see something like this:

```
$ java -version
java version "1.6.0_17"
Java(TM) SE Runtime Environment (build 1.6.0_17-b04-248-10M3025)
Java HotSpot(TM) 64-Bit Server VM (build 14.3-b01-101, mixed mode)
```

If you don't have Java already installed, or if your version is an older one, download and install the latest JRE installer from the Java website³, as shown in Figure 2.1, "Installing Java".

Figure 2.1. Installing Java

2.2.2. Installing Git

Since you'll be using Git, you'll need to install and configure Git on your machine. If you're new to Git, you might want to run through the basics on the Git Reference website⁴. If you get lost, the whole process is well documented on the GitHub help pages⁵.

First of all, you need to install Git on your machine. This involves downloading the appropriate installer for your operating system from the Git website⁶. There are packaged installers for both Windows and Mac OS X. If you're using Linux, you're in Git's home ground - most Linux distributions include Git. On Ubuntu or other Debian-based distributions, you could run something like:

```
$ sudo apt-get install git-core
```

On Fedora or another RPM-based distributions, you could use `yum` instead:

```
$ sudo yum install git-core
```

On Linux you also have the option of installing Git from source. There are instructions on how to do this on the Git website.

Once you're done, check that Git is installed and available by invoking it from the command line:

³ <http://java.sun.com/javase/downloads/index.jsp>

⁴ <http://gitref.org>

⁵ <http://help.github.com>

⁶ <http://git-scm.com>


```
$ git --version
git version 1.7.1
```

2.2.3. Setting Up a GitHub Account

Next, if you don't already have one, you'll need to create a GitHub account. This is easy and (for your purposes, at least) free of charge. All the cool kids have one. Go to the GitHub signup page⁷ and choose the "Create a free account" option. You'll just need to provide a username, a password, and your email address (see Figure 2.2, "Signing up for a GitHub account").

Figure 2.2. Signing up for a GitHub account

2.2.4. Configuring SSH Keys

GitHub uses SSH keys to establish a secure connection channel between your computer and the GitHub servers. Setting these up is not hard, but involves a bit of work. Fortunately, there are clear and detailed instructions on the GitHub website⁸.

2.2.5. Forking the Sample Repository

As we mentioned earlier, all the sample code for this book is stored on GitHub, at the following URL: <https://github.com/wakaleo/game-of-life>. This is a public repository, so you can freely view the source code online and check out your own working copy. However, if you want to make changes, you'll need to create your own fork. A fork is a personal copy of a repository that you can use as you wish. To create a fork, login to your GitHub account and navigate to the repository URL⁹. Then, click on the Fork button (see Figure 2.3, "Forking the sample code repository"). This creates your own personal copy of the repository.

Once you've forked the repository, you should clone a local copy on your machine to make sure everything is set up correctly. Go to the command line and run the following command (replacing `<username>` with your own GitHub username):

```
$ git clone git@github.com:<username>/game-of-life.git
```

This "clones" (or checks out, in Subversion terms) a copy of the project onto your local drive:

```
git clone git@github.com:john-smart/game-of-life.git
Initialized empty Git repository in /Users/johnsmart/.../game-of-life/.git/
remote: Counting objects: 1783, done.
remote: Compressing objects: 100% (589/589), done.
remote: Total 1783 (delta 1116), reused 1783 (delta 1116)
```

⁷ <https://github.com/plans>

⁸ <http://help.github.com/set-up-git-redirect>

⁹ <https://github.com/wakaleo/game-of-life>

```
Receiving objects: 100% (1783/1783), 14.83 MiB | 119 KiB/s, done.  
Resolving deltas: 100% (1116/1116), done.
```

You should now have a local copy of the project that you can build and execute. You'll use this project later on to trigger changes in the repository.

Figure 2.3. Forking the sample code repository

2.3. Starting Up Jenkins

There are several ways to run Jenkins. One of the easiest ways to run Jenkins for the first time is to use Java Web Start. This is a way to start a Java application on your local machine via a URL on a web page. In your case, this starts a Jenkins server running on your machine. All you need for this to work is a recent (Java 6 or later) version of the Java Runtime Environment (JRE), which you installed in the previous section.

For convenience, there's a link to the Jenkins Java Web Start instance on the book resources page¹⁰. Here you'll find a large orange Launch button in the Book Resources section (see Figure 2.4, “Running Jenkins using Java Web Start from the book’s website”). You can also find this link on the Meet Jenkins page on the Jenkins website¹¹, where, if you scroll down far enough, you should find a Test Drive section with an identical Launch button.

Figure 2.4. Running Jenkins using Java Web Start from the book’s website

When you click on the Launch button on either of these sites in Firefox, the browser asks if you want to open a file called `jenkins.jnlp` using Java Web Start. Click on OK—this downloads Jenkins and starts it up on your machine (see Figure 2.5, “Java Web Start will download and run the latest version of Jenkins”).

Figure 2.5. Java Web Start will download and run the latest version of Jenkins

In other browsers, clicking on this button may simply download the JNLP file. In Internet Explorer, you may even need to right click on the link and select “Save Target As” to save the JNLP file, and then run it from Windows Explorer. However, in both of these cases, when you open the JNLP file, Java Web Start downloads and starts Jenkins.

Java Web Start will only need to download a particular version of Jenkins once. From then on, when you click on the “Launch” button again, Java Web Start uses the copy of Jenkins it has already downloaded

¹⁰ <http://www.wakaleo.com/books/jenkins-the-definitive-guide>

¹¹ <http://wiki.jenkins-ci.org/display/JENKINS/Meet+Jenkins>

(that is, until the next version comes out). Ignore any messages your operating system or anti-virus software might produce—it's perfectly safe to run Jenkins on your local machine.

Once Jenkins has finished downloading, Java Web Start starts up Jenkins on your machine. You can see it running in a small window called “Jenkins Console” (see Figure 2.6, “Java Web Start running Jenkins”). To stop Jenkins at any time, just close this window.

Figure 2.6. Java Web Start running Jenkins

There are also installers available for the principal operating systems available on the Jenkins website¹². Or, if you're an experienced Java user versed in the ways of WAR files, you may prefer to simply download the latest version of Jenkins and run it from the command line. Jenkins comes in the form of an executable WAR file—you can download the most recent version from the Jenkins website home page¹³. For convenience, there's also a link to the latest version of Jenkins in the Resources section of this book's website¹⁴.

Once downloaded, you can start Jenkins from the command line as shown here:

```
$ java -jar jenkins.war
```

Whether you've started Jenkins using Java Web Start or from the command line, Jenkins should now be running on your machine. By default, Jenkins will be listening on port 8080, so you can access Jenkins from your web browser at <http://localhost:8080>.

Alternatively, if you're familiar with Java application servers such as Tomcat, you can simply deploy the Jenkins WAR file to your application server—with Tomcat, for example, you could simply place the `jenkins.war` file in Tomcat's `webapps` directory. If you're running Jenkins on an application server, the URL that you use to access Jenkins will be slightly different. On a default Tomcat installation, for example, you can access Jenkins from your web browser at <http://localhost:8080/jenkins>.

When you open Jenkins in your browser, you should see a screen like the one shown in Figure 2.7, “The Jenkins start page”. You're now ready to take your first steps with Jenkins!

Figure 2.7. The Jenkins start page

2.4. Configuring the Tools

Before you get started, you do need to do a little configuration. More precisely, you need to tell Jenkins about the build tools and JDK versions you'll be using for your builds.

¹² <http://jenkins-ci.org>

¹³ <http://http://jenkins-ci.org>

¹⁴ <http://www.wakaleo.com/books/jenkins-the-definitive-guide>

Click on the Manage Jenkins link on the home page (see Figure 2.7, “The Jenkins start page”). This takes you to the Manage Jenkins page, the central one-stop-shop for all your Jenkins configuration needs. From this screen, you can configure your Jenkins server, install and upgrade plugins, keep track of system load, manage distributed build servers, and more! For now, however, we’ll keep it simple. Just click on the Configure System link at the top of the list (see Figure 2.8, “The Manage Jenkins screen”).

Figure 2.8. The Manage Jenkins screen

This takes you to Jenkins’s main configuration screen (see Figure 2.9, “The Configure Jenkins screen”). From here you can configure everything from security configuration and build tools to email servers, version control systems, and integration with third-party software. The screen contains a lot of information, but most of the fields contain sensible default values, so you can safely ignore them for now.

Figure 2.9. The Configure Jenkins screen

For now, you just need to configure the tools required to build your sample project. The application you'll be building is a Java application, built using Maven. So, in this case, all you need to do is to set up a recent JDK and Maven installation.

However, before you start, take a look at the little blue question mark icons lined up on the right of the screen. These are Jenkins’s contextual help buttons. If you're curious about a particular field, click on the help icon next to it and Jenkins displays a very detailed description about what the field is and how it works.

2.4.1. Configuring Your Maven Setup

Our sample project uses Maven, so you need to install and configure Maven first. Jenkins provides great out-of-the-box support for Maven. Scroll down until you reach the Maven section in the Configure System screen (see Figure 2.10, “Configuring a Maven installation”).

Jenkins provides several options when it comes to configuring Maven. If you already have Maven installed on your machine, you can simply enter its path in the `MAVEN_HOME` field. Alternatively, you can install a Maven distribution by extracting a zip file located in a shared directory, or execute a home-rolled installation script. Or, you can let Jenkins do all the hard work and download Maven for you. To choose this option, just tick the Install automatically checkbox. Jenkins downloads and installs Maven from the Apache website the first time a build job needs it. Just choose the Maven version you want to install and Jenkins will do the rest. You'll also need to give a name to your Maven version (imaginatively called “Maven 2.2.1” in the example), so that you can refer to it in your build jobs.

For this to work, you need to have an Internet connection. If you're behind a proxy, you’ll need to provide your proxy information—we discuss how to set this up in Section 4.9, “Configuring a Proxy”.

Figure 2.10. Configuring a Maven installation

One of the nice things about the Jenkins Maven installation process is how well it works with remote build agents. Later on in the book, you'll see how Jenkins can also run builds on remote build servers. You can define a standard way of installing Maven for all of your build servers (downloading from the Internet, unzipping a distribution bundle on a shared server, etc.)—all of these options work when you add a new remote build agent or set up a new build server using this Jenkins configuration.

2.4.2. Configuring the JDK

Once you've configured your Maven installation, you'll also need to configure your Java environment (see Figure 2.11, “Configuring a JDK installation”). Again, if you have a Java Development Kit (as opposed to a Java Runtime Environment—the JDK contains extra development tools such as the Java compiler) already installed on your machine, you can simply provide the path to your JDK in the `JAVA_HOME` field. Otherwise, you can ask Jenkins to download the JDK from the Oracle website¹⁵ the first time a build job requires it. This is similar to the automatic Maven installation feature—just pick the JDK version you need and Jenkins takes care of all the logistics. However, for licensing reasons, you'll also need to tick a checkbox to indicate that you agree with the Java SDK License Agreement.

Figure 2.11. Configuring a JDK installation

Now, go to the bottom of the screen and click on the Save button.

2.4.3. Notification

Another important aspect you'd typically set up is notification. When a Jenkins build breaks, and when it works again, Jenkins can send out email messages to the team to spread the word. Using plugins, you can also get it to send instant messages or SMS messages, post entries on Twitter, or notify people in a few other ways. It all depends on what works best for your organizational culture. Email notification is easy enough to set up if you know your local SMTP server address—just provide this value in the Email Notification section towards the bottom of the main configuration page. However, to keep things simple, we're not going to worry about notifications just yet.

2.4.4. Setting Up Git

The last thing you need to configure for this demo is getting Jenkins working with Git. Jenkins comes with support for Subversion and CVS out of the box, but you'll need to install the Jenkins Git plugin to be able to complete the rest of this tutorial. Don't worry, the process is pretty simple. First of all, click on the Manage Jenkins link to the left of the screen to go back to the main configuration screen (see Figure 2.8, “The Manage Jenkins screen”). Then, click on Manage Plugins. This opens the plugin configuration

¹⁵ <http://www.oracle.com/technetwork/java/index.html>

screen, which is where you manage the extra features you want to install on your Jenkins server. You should see four tabs: Updates, Available, Installed, and Advanced (see Figure 2.12, “Managing plugins in Jenkins”).

Figure 2.12. Managing plugins in Jenkins

For now, just click on the Available tab. Here you'll see a very long list of available plugins. Find the Git Plugin entry in this list and tick the corresponding checkbox (see Figure 2.13, “Installing the Git plugin”), and then scroll down to the bottom of the screen and click on Install. This downloads and installs the Jenkins Git plugin into your local Jenkins instance.

Figure 2.13. Installing the Git plugin

Once it's done, restart Jenkins for the changes to take effect. To do this, you can simply click on the “Restart Jenkins when no jobs are running” button displayed on the installation screen, or alternatively, shut down and restart Jenkins by hand.

That's all you need to configure at this stage. You're now ready to set up your first Jenkins build job!

2.5. Your First Jenkins Build Job

Build jobs are at the heart of the Jenkins build process. Simply put, you can think of a Jenkins build job as a particular task or step in your build process. This may involve simply compiling your source code and running your unit tests. Or, you might want a build job to do other related tasks, such as running your integration tests, measuring code coverage or code quality metrics, generating technical documentation, or even deploying your application to a web server. A real project usually requires many separate but related build jobs.

Our sample application is a simple Java implementation of John Conway's “Game of Life.”¹⁶ The Game of Life is a mathematical game which takes place on a two dimensional grid of cells, which we'll refer to as the Universe. Each cell can be either alive or dead. Cells interact with their direct neighbors to determine whether they live or die in the next generation of cells. For each new generation of cells, the following rules apply:

- Any live cell with fewer than two live neighbors dies of underpopulation.
- Any live cell with more than three live neighbors dies of overcrowding.
- Any live cell with two or three live neighbors lives on to the next generation.
- Any dead cell with exactly three live neighbors becomes a live cell.

¹⁶See http://en.wikipedia.org/wiki/Conway%27s_Game_of_Life.

Our application is a Java module, built using Maven, that implements the core logic of the Game of Life. You'll worry about the user interfaces later on. For now, let's see how you can automate this build in Jenkins. If you're not familiar with Maven, or prefer Ant or another build framework—don't worry! The examples don't require much knowledge of Maven, and we'll be looking at plenty of examples of using other build tools later on in the book.

For your first build job, we'll keep it simple: you're just going to compile and test your sample application. Click on the New Job link. You should get to a screen similar to Figure 2.14, “Setting up your first build job in Jenkins”. Jenkins supports several different types of build jobs. The two most commonly-used are the freestyle builds and the Maven 2/3 builds. The freestyle projects allow you to configure just about any sort of build job: they're highly flexible and very configurable. The Maven 2/3 builds understand the Maven project structure, and can use this to let you set up Maven build jobs with less effort and a few extra features. There are also plugins that provide support for other types of build jobs. Nevertheless, although your project does use Maven, you're going to use a freestyle build job, just to keep things simple and general to start with. So, choose “Build a freestyle software project”, as shown in Figure 2.14, “Setting up your first build job in Jenkins”.

You'll also need to give your build job a sensible name. In this case, call it `game-of-life-default`, as it will be the default CI build for your Game of Life project.

Figure 2.14. Setting up your first build job in Jenkins

Once you click on OK, Jenkins displays the project configuration screen (see Figure 2.15, “Telling Jenkins where to find the source code”).

In a nutshell, Jenkins works by checking out the source code of your project and building it in its own workspace. So, the next thing you need to do is to tell Jenkins where it can find the source code for your project. You do this in the Source Code Management section (see Figure 2.15, “Telling Jenkins where to find the source code”). Jenkins provides support for CVS and Subversion out of the box, and many others such as Git, Mercurial, ClearCase, Perforce, and many more via plugins.

For this project, you'll be getting the source code from the GitHub repository you set up earlier. On the Jenkins screen, choose “Git” and enter the Repository URL you defined in Section 2.2.5, “Forking the Sample Repository” (see Figure 2.15, “Telling Jenkins where to find the source code”). Make sure this is the URL of your fork, and not of the original repository: it should have the form `git@github.com:<username>/game-of-life.git`, where `<username>` is the username for your own GitHub account. You can leave all of the other options up until here with their default values.

Figure 2.15. Telling Jenkins where to find the source code

Once you've told Jenkins where to find the source code for your application, you need to tell it how often it should check for updates. You want Jenkins to monitor the repository and start a build whenever any

changes have been committed. This is a common way to set up a build job in a CI context, as it provides fast feedback if the build fails. Other approaches include building on regular intervals (for example, once a day), requiring a user to kick off the build manually, or even triggering a build remotely using a “post-commit” hook in your SCM.

You configure all of this in the Build Triggers section (see Figure 2.16, “Scheduling the build jobs”). Pick the Poll SCM option and enter “* * * * *” (that’s five asterisks separated by spaces) in the Schedule box. Jenkins schedules are configured using `cron` syntax, well-known in the Unix world. The `cron` syntax consists of five fields separated by white space, indicating respectively the minute (0–59), hour (0–23), day of the month (1–31), month (1–12), and the day of the week (0–7, with 0 and 7 being Sunday). The star is a wildcard character which represents any valid value for that field. So, five stars means “every minute of every hour of every day”. You can also provide ranges of values: “* 9-17 * * *” would mean “every minute of every day, between 9am and 5pm”. You can also space out the schedule using intervals: for example “*/5 * * * *” means “every 5 minutes”. Finally, there are other convenient short-hands, such as “@daily” and “@hourly”.

Don’t worry if your Unix skills are a little rusty—if you click on the blue question mark icon on the right side of the schedule box, Jenkins brings up a very complete refresher.

Figure 2.16. Scheduling the build jobs

The next step is to configure the actual build itself. In a freestyle build job, you can break down your build job into a number of build steps. This makes it easier to organize builds in clean separate stages. For example, a build might run a suite of functional tests in one step, and then tag the build in a second step if all of the functional tests succeed. In technical terms, a build step might involve invoking an Ant task or a Maven target, or running a shell script. There are also Jenkins plugins that let you use additional types of build steps: Gant, Grails, Gradle, Rake, Ruby, MSBuild, and many other build tools are all supported.

For now, you just want to run a simple Maven build. Scroll down to the Build section and click on the “Add build step” and choose “Invoke top-level Maven targets” (see Figure 2.17, “Adding a build step”). Then, enter “clean package” in the Goals field. If you’re not familiar with Maven, this deletes any previous build artifacts, compiles your code, runs your unit tests, and generates a JAR file.

Figure 2.17. Adding a build step

By default, this build job fails if the code doesn’t compile or if any of the unit tests fail. That’s the most fundamental thing that you’d expect of any build server. But, Jenkins also does a great job of helping you display your test results and test result trends.

The de facto standard for test reporting in the Java world is the XML format used by JUnit. This format is also used by many other Java testing tools, such as TestNG, Spock, and Easyb. Jenkins understands

this format, so if your build produces JUnit XML test results, Jenkins can generate nice graphical test reports and statistics on test results over time, and also let you view the details of any test failures. Jenkins also keeps track of how long your tests take to run, both globally, and per test—this can come in handy if you need to track down performance issues.

So, the next thing you need to do is to get Jenkins to keep tabs on your unit tests.

Go to the Post-build Actions section (see Figure 2.18, “Configuring JUnit test reports and artifact archiving”) and tick the “Publish JUnit test result report” checkbox. When Maven runs unit tests in a project, it automatically generates the XML test reports in a directory called `surefire-reports` in the `target` directory. So, enter “`**/target/surefire-reports/*.xml`” in the “Test report XMLs” field. The two asterisks at the start of the path (“`**`”) are a best practice to make the configuration a bit more robust: they allow Jenkins to find the target directory no matter how you’ve configured Jenkins to check out the source code.

Another thing you often want to do is to archive your build results. Jenkins can store a copy of the binary artifacts generated by your build, allowing you to download the binaries produced by a build directly from the build results page. It also posts the latest binary artifacts on the project home page, which is a convenient way to distribute the latest and greatest version of your application. You can activate this option by ticking the “Archive the artifacts” checkbox and indicating which binary artifacts you want Jenkins to archive. In Figure 2.18, “Configuring JUnit test reports and artifact archiving”, for example, you’ve configured Jenkins to store all of the JAR files generated by this build job.

Figure 2.18. Configuring JUnit test reports and artifact archiving

Now you’re done—just click on the Save button at the bottom of the screen. Your build job should now be ready to run. So, let’s see it in action!

2.6. Your First Build Job in Action

Once you save your new build job, Jenkins displays the home page for this job (see Figure 2.19, “Your first build job running”). This is where Jenkins displays details about the latest build results and the build history.

If you wait a minute or so, the build should kick off automatically—you can see the stripy progress bar in the Build History section in the bottom left hand corner of Figure 2.19, “Your first build job running”. Or, if you’re impatient, you can also trigger the build manually using the Build Now button.

Figure 2.19. Your first build job running

The build will also now figure proudly on your Jenkins server’s home page (see Figure 2.20, “The Jenkins dashboard”). This page shows a summary of all of your build jobs, including the current build

status and general state of health of each of your builds. It tells you when each build ran successfully for the last time, when it last failed, and also the results of the last build.

One of Jenkins's specialities is the way it lets you get an idea of build behavior over time. For example, Jenkins uses a weather metaphor to help give you an idea of the stability of your builds. Essentially, the more your builds fail, the worse the weather gets. This helps you get an idea of whether a particular broken build is an isolated event, or if the build is breaking on a regular basis, in which case it might need some special attention.

You can also manually trigger a build job here, using the build schedule button (that's the one that looks a bit like a green play button on top of a clock).

Figure 2.20. The Jenkins dashboard

When the build finishes, the ball in the Build History box becomes solid blue. This means the build was a success. Build failures are generally indicated by a red ball. For some types of projects, you can also distinguish between a build error (such as a compiler error), indicated by a red ball, and other sorts of build failures, such as unit test failures or insufficient code coverage, which are indicated by a yellow ball. There are also some other details about the latest test results, when the last build was run, and so on. But, before you look at the details, let's get back to the core business model of a CI server—kicking off builds when someone changes the code!

You're going to commit a code change to GitHub and see what happens, using the source code you checked out in Section 2.2.5, "Forking the Sample Repository". You've now Jenkins configured to monitor your GitHub fork, so if you make any changes, Jenkins should be able to pick them up.

So, let's make a change. The idea is to introduce a code change that causes the unit tests to fail. If your Java is a bit rusty, don't worry. You won't need to know any Java to be able to break the build—just follow the instructions!

In normal development, you'd first modify the unit test that describes this behaviour. Then, you'd verify that the test fails with the existing code, and implement the code to ensure that the test passes. Then you'd commit your changes to your version control system, allowing Jenkins to build them. However, this would be a poor demonstration of how Jenkins handles unit test failures. So, in this example, you'll, against all best practices, simply modify the application code directly.

First of all, open the `Cell.java` file, which you'll find in the `gameoflife-core/src/main/java/com/wakaleo/gameoflife/domain` directory. Open this file in your favorite text editor. You should see something like this:

```
package com.wakaleo.gameoflife.domain;

public enum Cell {
    LIVE_CELL("*"), DEAD_CELL(".");
}
```

```

private String symbol;

private Cell(String symbol) {
    this.symbol = symbol;
}

@Override
public String toString() {
    return symbol;
}

static Cell fromSymbol(String symbol) {
    Cell cellRepresentedBySymbol = null;
    for (Cell cell : Cell.values()) {
        if (cell.symbol.equals(symbol)) {
            cellRepresentedBySymbol = cell;
            break;
        }
    }
    return cellRepresentedBySymbol;
}

public String getSymbol() {
    return symbol;
}
}

```

The application can print the state of the grid as a text array. Currently, the application prints your live cells as an asterisk (*), and dead cells appear as a minus character (-). So, a three-by-five grid containing a single living cell in the center would look like this:

```

-----
--*--
-----

```

Now, users have asked for a change to the application—they want pluses (+) instead of stars! So, you're going to make a slight change to the `Cell` class method, and rewrite it as follows (the modifications are in **bold**):

```

package com.wakaleo.gameoflife.domain;

public enum Cell {
    LIVE_CELL("+"), DEAD_CELL(".");

    private String symbol;

    private Cell(String symbol) {
        this.symbol = symbol;
    }

    @Override
    public String toString() {
        return symbol;
    }
}

```

```

    }

    static Cell fromSymbol(String symbol) {
        Cell cellRepresentedBySymbol = null;
        for (Cell cell : Cell.values()) {
            if (cell.symbol.equals(symbol)) {
                cellRepresentedBySymbol = cell;
                break;
            }
        }
        return cellRepresentedBySymbol;
    }

    public String getSymbol() {
        return symbol;
    }
}

```

Save these changes, and then commit them to the local Git repository by running `git commit`:

```

$ git commit -a -m "Changes stars to pluses"
[master 61ce946] Changes stars to pluses
1 files changed, 1 insertions(+), 1 deletions(-)

```

This commits the changes locally, but since Git is a distributed repository, you now have to push them through to your fork on GitHub. You do this by running `git push`:

```

$ git push
Counting objects: 21, done.
Delta compression using up to 4 threads.
Compressing objects: 100% (7/7), done.
Writing objects: 100% (11/11), 754 bytes, done.
Total 11 (delta 4), reused 0 (delta 0)
To git@github.com:john-smart/game-of-life.git
7882d5c..61ce946 master -> master

```

Now, go back to the Jenkins web page. After a minute or so, a new build should kick off, and fail. In fact, there are several other places which are affected by this change, and the regression tests related to these features are now failing. On the build job home page, you'll see a second build in the build history with an ominous red ball (see Figure 2.21, “A failed build”)—this tells you that the latest build has failed.

You might also notice some clouds next to the Build History title—this is the same “weather” icon that you saw on the home page, and serves the same purpose—to give you a general idea of how stable your build is over time.

Figure 2.21. A failed build

If you click on the new build history entry, Jenkins gives you some more details about what went wrong (see Figure 2.22, “The list of all the broken tests”). Jenkins tells you that there were 11 new test failures

in this build, something which can be seen at a glance in the Test Result Trend graph—red indicates test failures. You can even see which tests are failing, and how long they've been broken.

Figure 2.22. The list of all the broken tests

If you want to know exactly what went wrong, that's easy enough to figure out as well. If you click on the failed test classes, Jenkins brings up the actual details of the test failures (see Figure 2.23, “Details about a failed test”), which is a great help when it comes to reproducing and fixing the issue.

Figure 2.23. Details about a failed test

Jenkins displays a host of information about the failed test in a very readable form, including the error message the test produced, the stack trace, how long the test has been broken, and how long it took to run. Often, this in itself is enough to put a developer on the right track towards fixing the issue.

Now let's fix the build. To make things simple, just back out your changes and recommit the code in its original state (the end users just changed their mind about the asterisks, anyway). So, just undo the changes you made to the `Cell` class (again, the changes are highlighted in **bold**):

```
package com.wakaleo.gameoflife.domain;

public enum Cell {
    LIVE_CELL("*"), DEAD_CELL(".");

    private String symbol;

    private Cell(String symbol) {
        this.symbol = symbol;
    }

    @Override
    public String toString() {
        return symbol;
    }

    static Cell fromSymbol(String symbol) {
        Cell cellRepresentedBySymbol = null;
        for (Cell cell : Cell.values()) {
            if (cell.symbol.equals(symbol)) {
                cellRepresentedBySymbol = cell;
                break;
            }
        }
        return cellRepresentedBySymbol;
    }

    public String getSymbol() {
```

```
        return symbol;
    }
}
```

When you've done this, commit your changes again:

```
$ git commit -a -m "Restored the star"
[master bc924be] Restored the star
 1 files changed, 1 insertions(+), 1 deletions(-)
$ git push
Counting objects: 21, done.
Delta compression using up to 4 threads.
Compressing objects: 100% (7/7), done.
Writing objects: 100% (11/11), 752 bytes, done.
Total 11 (delta 4), reused 6 (delta 0)
To git@github.com:john-smart/game-of-life.git
 61ce946..bc924be  master -> master
```

Once you've committed these changes, Jenkins should pick them up and kick off a build. When this is done, you'll be able to see the fruit of your work on the build job home page (see Figure 2.24, “Now the build is back to normal”)—the build status is blue again and all is well. Also, notice the way you're building up a trend graph showing the number of succeeding unit tests over time—this sort of report really is one of Jenkins's strong points.

Figure 2.24. Now the build is back to normal

2.7. More Reporting—Displaying Javadocs

For many Java projects, Javadoc comments are an important source of low-level technical documentation. There are even tools, such as UmlGraph, that produce Javadoc with embedded UML diagrams to give you a better picture of how the classes fit together in the application. This sort of technical documentation has the advantage of being cheap to produce, accurate, and always up-to-date.

Jenkins can integrate Javadoc API documentation directly into the Jenkins website. This way, everyone can find the latest Javadoc easily, in a well known place. Often, this sort of task is performed in a separate build job, but, for simplicity, you're going to add another build step to the gameoflife-default build job to generate and display Javadoc documentation for the Game of Life API.

Start off by going into the “gameoflife-default” configuration screen again. Click on “Add build step”, and add a new build step to “Invoke top level Maven targets” (see Figure 2.25, “Adding a new build step and report to generate Javadoc”). In the Goals field, place `javadoc:javadoc`—this tells Maven to generate the Javadoc documentation.

Figure 2.25. Adding a new build step and report to generate Javadoc

Now, go to the “Post-build Action” and tick the “Publish Javadoc” checkbox. This project is a multimodule project, so a separate subdirectory is generated for each module (core, services, web, and so forth). For this example, you're interested in displaying the documentation for the core module. In the Javadoc directory field, enter `gameoflife-core/target/site/apidocs`—this is where Maven places the Javadocs it generates for the core module. Jenkins may display an error message saying that this directory doesn't exist at first. Jenkins is correct—this directory won't exist until you run the `javadoc:javadoc` goal, but since you haven't run this command yet you can safely ignore the message at this stage.

If you tick “Retain Javadoc for each successful build”, Jenkins also keeps track of the Javadocs for previous builds—not always useful, but it can come in handy at times.

Now, trigger a build manually. You can do this either from the build job's home page (using the Build Now link), or directly from the server home page. Once the build is finished, open the build job summary page. You should now see a Javadoc link featuring prominently on the screen—this link opens the latest version of the Javadoc documentation (see Figure 2.26, “Jenkins will add a Javadoc link to your build results”). You'll also see this link on the build details page, where it points to the Javadoc for that particular build, if you've asked Jenkins to store Javadoc for each build.

Figure 2.26. Jenkins will add a Javadoc link to your build results

2.8. Adding Code Coverage and Other Metrics

As we mentioned earlier, reporting is one of Jenkins's strong points. You've seen how easy it is to display test results and to publish Javadocs, but you can also publish a large number of other very useful reports using Jenkins plugins.

Plugins are another one of Jenkins's selling points—there are plugins for doing just about anything, from integrating new build tools or version control systems to notification mechanisms and reporting. In addition, Jenkins plugins are very easy to install and integrate into the existing Jenkins architecture.

To see how the plugins work, you're going to integrate code coverage metrics using the Cobertura plugin. Code coverage is an indication of how much of your application code is actually executed during your tests—it can be a useful tool in particular for finding areas of code that haven't been tested by your test suites. It can also give some indication as to how well a team is applying good testing practices such as Test-Driven Development or Behavior-Driven Development.

Cobertura¹⁷ is an open source code coverage tool that works well with both Maven and Jenkins. Your Maven demonstration project is already configured to record code coverage metrics, so all you need to do is to install the Jenkins Cobertura plugin and generate the code coverage metrics for Jenkins to record and display.

¹⁷ <http://cobertura.sourceforge.net>

Figure 2.27. Jenkins has a large range of plugins available

To install a new plugin, go to the Manage Jenkins page and click on the Manage Plugins entry. This displays a list of the available plugins as well as the plugins already installed on your server (see Figure 2.27, “Jenkins has a large range of plugins available”). If your build server doesn’t have an Internet connection, you can also manually install a plugin by downloading the plugin file elsewhere and uploading it to your Jenkins installation (just open the Advanced tab in Figure 2.27, “Jenkins has a large range of plugins available”), or by copying the plugin to the `$JENKINS_HOME/plugins` directory.

In your case, you're interested in the Cobertura plugin, so go to the Available tab and scroll down until you find the Cobertura Plugin entry in the Build Reports section. Click on the checkbox and then click on the Install button at the bottom of the screen.

This downloads and installs the plugin for you. Once it's done, you'll need to restart your Jenkins instance to see the fruits of your labor. When you've restarted Jenkins, go back to the Manage Plugins screen and click on the Installed tab—there should now be a Cobertura Plugin entry in the list of installed plugins on this page.

Once you've made sure the plugin was successfully installed, go to the configuration page for the `gameoflife-default` build job.

To set up code coverage metrics in your project, you need to do two things. First you need to generate the Cobertura coverage data in an XML form that Jenkins can use. Then, you need to configure Jenkins to display the coverage reports.

Our Game of Life project already has been configured to generate XML code coverage reports if you ask it. All you need to do is to run `mvn cobertura:cobertura` to generate the reports in XML form. Cobertura can also generate HTML reports, but in your case you'll be letting Jenkins take care of the reporting, so you can save on build time by not generating the report. For this example, for simplicity, you'll just add the `cobertura:cobertura` goal to the second build step (see Figure 2.28, “Adding another Maven goal to generate test coverage metrics”). You could also add a new build step just for the code coverage metrics. In a real-world project, code quality metrics like this are typically placed in a distinct build job, which is run less frequently than the default build.

Figure 2.28. Adding another Maven goal to generate test coverage metrics

Next, you need to tell Jenkins to keep track of your code coverage metrics. Scroll down to the “Post-build Actions” section. You should see a new checkbox labeled Publish Cobertura Reports. Jenkins often adds UI elements like this when you install a new plugin. When you tick this box, Jenkins displays the configuration options for the Cobertura plugin that you installed earlier (see Figure 2.29, “Configuring the test coverage metrics in Jenkins”).

Like most of the code-quality related plugins in Jenkins, the Cobertura plugin lets you fine-tune not only the way Jenkins displays the report data, but also how it interprets the data. In the Coverage Metrics Targets section, you can define what you consider to be the minimum acceptable levels of code coverage. In Figure 2.29, “Configuring the test coverage metrics in Jenkins”, you’ve configured Jenkins to list any builds with less than 50% test coverage as “unstable” (indicated by a yellow ball), and notify the team accordingly.

Figure 2.29. Configuring the test coverage metrics in Jenkins

This fine-tuning often comes in handy in real-world builds. For example, you may want to impose a special code coverage constraint in release builds, to ensure high code coverage in release versions. Another strategy that can be useful for legacy projects is to gradually increase the minimum tolerated code coverage level over time. This way you can avoid having to retro-fit unit tests on legacy code just to raise the code coverage, but you do encourage all new code and bug fixes to be well tested.

Now trigger a build manually. The first time you run the build job with Cobertura reporting activated, you’ll see coverage statistics for your build displayed on the build home page, along with a Coverage Report link where you can go for more details (see Figure 2.30, “Jenkins displays code coverage metrics on the build home page”). The Cobertura report shows different types of code coverage for the build you just ran. Since you’ve only run the test coverage metrics once, the coverage will be displayed as red and green bars.

Figure 2.30. Jenkins displays code coverage metrics on the build home page

If you click on the Coverage Report icon, you’ll see code coverage for each package in your application, and you can even drill down to see the code coverage (or lack thereof) for an individual class (see Figure 2.31, “Jenkins lets you display code coverage metrics for packages and classes”). When you get to this level, Jenkins displays both the overall coverage statistics for the class, and also highlights the lines that were executed in green, and those that weren’t in red.

This reporting gets better with time. Jenkins not only reports metrics data for the latest build, but also keeps track of metrics over time, so that you can see how they evolve throughout the life of the project.

For example, if you drill down into the coverage reports, you’ll notice that certain parts of this code are not tested (for example the `Cell.java` class in Figure 2.31, “Jenkins lets you display code coverage metrics for packages and classes”).

Figure 2.31. Jenkins lets you display code coverage metrics for packages and classes

Code coverage metrics are a great way to isolate code that hasn’t been tested, in order to add extra tests for corner cases that weren’t properly tested during the initial development, for example. The Jenkins

code coverage graphs are also a great way of keeping track of your code coverage metrics as the project grows. Indeed, as you add new tests, you'll notice that Jenkins displays a graph of code coverage over time, not just the latest results (see Figure 2.32, “Jenkins also displays a graph of code coverage over time”).

Figure 2.32. Jenkins also displays a graph of code coverage over time

Note that your objective here isn't to improve the code coverage just for the sake of improving code coverage—you're adding an extra test to verify some code that was not previously tested, and as a result the code coverage goes up. There's a subtle but important difference here—code coverage, as with any other metric, is very much a means to an end (high code quality and low maintenance costs), and not an end in itself.

Nevertheless, metrics like this can give you a great insight into the health of your project, and Jenkins presents them in a particularly accessible way.

This is just one of the code quality metrics plugins that have been written for Jenkins. There are many more (over fifty reporting plugins alone at the time of writing). We'll look at some more of them in Chapter 9, Code Quality.

2.9. Conclusion

In this chapter, you've gone through what you need to know to get started with Jenkins. You should be able to set up a new build job, report on JUnit test results, and create javadocs. You've also seen how to add a reporting plugin and keep tabs on code coverage. Well done! But there's still a lot more to learn about Jenkins. In the following chapters, you'll be looking at how Jenkins can help improve your build automation process in many other areas as well.

Chapter 3. Installing Jenkins

3.1. Introduction

One of the first things you'll probably notice about Jenkins is how easy it is to install. Indeed, in less than five minutes, you can have a Jenkins server up and running. However, as always, in the real world, things aren't always that simple, and there are a few details you should take into account when installing your Jenkins server for production use. In this chapter, we look at how to install Jenkins onto both your local machine and onto a fully fledged build server. We'll also look at how to take care of your Jenkins installation once it's up and running, and how to perform basic maintenance tasks such as backups and upgrades.

3.2. Downloading and Installing Jenkins

Jenkins is easy to install, and can run just about anywhere. You can run it either as a stand-alone application, or deployed on a conventional Java application server such as Tomcat or JBoss. This first option makes it easy to install and try out on your local machine. You can be up and running with a bare-bones installation in a matter of minutes.

Since Jenkins is a Java application, you'll need a recent version of Java on your machine. More precisely, you'll need at least Java 5. In fact, on your build server, you'll almost certainly need the full features of the Java Development Kit (JDK) 5.0 or better to execute your builds. If you're not sure, you can check the version of Java on your machine by executing the `java -version` command:

```
$ java -version
java version "1.6.0_17"
Java(TM) SE Runtime Environment (build 1.6.0_17-b04-248-10M3025)
Java HotSpot(TM) 64-Bit Server VM (build 14.3-b01-101, mixed mode)
```

Jenkins is distributed in the form of a bundled Java web application (a WAR file). You can download the latest version from the Jenkins website (<http://jenkins-ci.org>—see Figure 3.1, “You can download the Jenkins binaries from the Jenkins website”) or from the book website. Jenkins is a dynamic project, and new releases come out at a regular rate.

For Windows users, there's a graphical Windows installation package for Jenkins. The installer comes in the form of a ZIP file containing an MSI package for Jenkins, as well as a `setup.exe` file that can be used to install the .NET libraries if they haven't already been installed on your machine. In most cases, all you need to do is to unzip the ZIP file and run `jenkins-x.x.msi` inside (see Figure 3.2, “Jenkins setup wizard in Windows”). The MSI installer comes bundled with a bundled JRE, so no separate Java installation is required.

Figure 3.1. You can download the Jenkins binaries from the Jenkins website

Once you've run the installer, Jenkins will automatically start on port 8080 (see Figure 3.3, “The Jenkins start page”). The installer will have created a new Jenkins service that you can start and stop just like any other Windows service.

There are also excellent native packages for Mac OS X and most of the major Linux distributions, including Ubuntu, RedHat (including CentOS and Fedora), and OpenSolaris. We discuss how to install Jenkins on Ubuntu and Red Hat below.

If you aren't installing Jenkins using one of the native packages, you can simply download the latest binary distribution from the Jenkins website. Once you've downloaded the latest and greatest Jenkins release, place it in an appropriate directory on your build server. On a Windows environment, you might put it in a directory called `C:\Tools\Jenkins` (it's a good idea not to place Jenkins in a directory containing spaces in the path, such as `C:\Program Files`, as this can cause problems for Jenkins in some circumstances). On a Linux or Unix box, it might go in `/usr/local/jenkins`, `/opt/jenkins`, or in some other directory, depending on your local conventions and on the whim of your system administrator.

Figure 3.2. Jenkins setup wizard in Windows

Before we go any further, let's just start up Jenkins and take a look. If you didn't try this out in the previous chapter, now's the time to get your hands dirty. Open a console in the directory containing `jenkins.war` and run the following command:

```
$ java -jar jenkins.war
[Winstone 2008/07/01 20:54:53] - Beginning extraction from war file
...
INFO: Took 35 ms to load
...
[Winstone 2008/07/01 20:55:08] - HTTP Listener started: port=8080
[Winstone 2008/07/01 20:55:08] - Winstone Servlet Engine v0.9.10 running:
    controlPort=disabled
[Winstone 2008/07/01 20:55:08] - AJP13 Listener started: port=8009
```

Jenkins should now be running on port 8080. Open your browser on `http://localhost:8080` and take a look. (see Figure 3.3, “The Jenkins start page”).

Figure 3.3. The Jenkins start page

3.3. Preparing a Build Server for Jenkins

Installing Jenkins on your local development machine is one thing, but installing Jenkins on a proper build server requires a little more forethought and planning.

Before you start your installation, the first thing you'll need is a build server. To work well, Jenkins needs both CPU resources and memory. Jenkins itself is a relatively modest Java web application. However,

in most configurations, at least some of the builds will run on the principal build server. Builds tend to be both memory and CPU-intensive operations, and Jenkins can be configured to run several builds in parallel. Depending on the number of build jobs you're managing, Jenkins will also need memory of its own for its own internal use. The amount of memory required will depend largely on the nature of your builds, but memory is cheap these days (at least in non-hosted environments), and it's best not to be stingy.

A build server also needs CPU horsepower. As a rule of thumb, you'll need one CPU per parallel build, though, in practice, you can capitalize on I/O delays to do a little better than this. It's also in your best interest to dedicate your build server as much as possible to the task of running continuous builds. In particular, you should avoid running memory or CPU-intensive applications on your build server such as test servers, heavily-used enterprise applications, enterprise databases such as Oracle, enterprise mail servers, and so on.

One very practical option available in many organizations today is to use a virtual machine. This way, you can choose the amount of memory and number of CPUs you think appropriate for your initial installation, and easily add more memory and CPUs later on as required. However, if you're using a virtual machine, make sure that it has enough memory to support the maximum number of parallel builds you expect to be running. The memory usage of a CI server is best described as spiky—Jenkins will be creating additional JVMs as required for its build jobs, and these need memory.

Another useful approach is to set up multiple build machines. Jenkins makes it quite easy to set up “slaves” on other machines that can be used to run additional build jobs. The slaves remain inactive until a new build job is requested—then the main Jenkins installation dispatches the build job to the slave and reports on the results. This is a great way to absorb sudden spikes of build activity, for example just before a major release of your principal product. It's also a useful strategy if certain heavy-weight builds tend to “hog” the main build server—just put them on their own dedicated build agent! We'll look at how to do this in detail later on in the book.

If you're installing Jenkins on a Linux or Unix build server, it's a good idea to create a special user (and user group) for Jenkins. This makes it easier to monitor at a glance the system resources being used by the Jenkins builds, and to troubleshoot problematic builds in real conditions. The native binary installation packages discussed below do this for you. If you didn't use one of these, you can create a dedicated Jenkins user from the command line as shown here:

```
$ sudo groupadd build
$ sudo useradd --create-home --shell /bin/bash --groups build jenkins
```

The exact details may vary depending on your environment. For example, you may prefer to use a graphical administration console instead of the command line, or, on a Debian-based Linux server (such as Ubuntu), you might use the more user-friendly `adduser` and `addgroup` commands.

In most environments, you'll need to configure Java correctly for this user. For example, you can do this by defining the `JAVA_HOME` and `PATH` environment variables in `.bashrc`, as shown here:

```
export JAVA_HOME=/usr/local/java/jdk1.6.0
```

```
export PATH=$JAVA_HOME/bin:$PATH
```

This user will now be able to run Jenkins in an isolated environment.

3.4. The Jenkins Home Directory

Before you install Jenkins, however, there are some things you need to know about how Jenkins stores its data. Indeed, no matter where you store the Jenkins WAR file, Jenkins keeps all its important data in a special separate directory called the Jenkins home directory. Here, Jenkins stores information about your build server configuration, your build jobs, build artifacts, user accounts, and other useful information, as well as any plugins you may have installed. The Jenkins home directory format is backward compatible across versions, so you can freely update or reinstall your Jenkins executable without affecting your Jenkins home directory.

Needless to say, this directory will need a lot of disk space.

By default, the Jenkins home directory is called `.jenkins`, and is placed in your home directory. For example, if you're running a Windows 7 machine and your username is “john”, you'd find the Jenkins home directory under `C:\Users\john\.jenkins`. Under Windows XP, it would be `C:\Documents and Settings\john\.jenkins`. On a Linux machine, it would most likely be under `/home/john/.jenkins`. And so on.

You can force Jenkins to use a different directory as its home directory by defining the `JENKINS_HOME` environment variable. You may need to do this on a build server to conform to local directory conventions or to make your system administrator happy. For example, if your Jenkins WAR file is installed in `/usr/local/jenkins`, and the Jenkins home directory needs to be in the `/data/jenkins` directory, you might write a startup script containing the following lines:

```
export JENKINS_BASE=/usr/local/jenkins
export JENKINS_HOME=/data/jenkins
java -jar ${JENKINS_BASE}/jenkins.war
```

If you're running Jenkins in a Java EE container such as Tomcat or JBoss, you can configure the application server to expose its own environments variables. For example, if you're using Tomcat, you could create `jenkins.xml` in the `$CATALINA_BASE/conf/localhost` directory:

```
<Context docBase="../jenkins.war">
  <Environment name="JENKINS_HOME" type="java.lang.String"
    value="/data/jenkins" override="true"/>
</Context>
```

In a previous life, Jenkins was known as Hudson. Jenkins remains compatible with previous Hudson installations, and upgrading from Hudson to Jenkins can be as simple as replacing the old `hudson.war` file with `jenkins.war`. Jenkins will look for its home directory in the following places (by order of precedence):

1. A `JENKINS_HOME` JNDI environment entry

2. A HUDSON_HOME JNDI environment entry
3. A JENKINS_HOME system property
4. A HUDSON_HOME system property
5. A JENKINS_HOME environment variable
6. A HUDSON_HOME environment variable
7. The `.hudson` directory in the user's home directory, if it already exists
8. The `.jenkins` directory in the user's home directory

3.5. Installing Jenkins on Debian or Ubuntu

If you're installing Jenkins on Debian and Ubuntu, it's convenient to install the native binary package for these platforms. This is easy enough to do, though these binaries aren't provided in the standard repositories because of the high frequency of updates. First, you need to add the repository signing key to your system as shown here:

```
$ wget -q -O - http://pkg.jenkins-ci.org/debian/jenkins-ci.org.key \
| sudo apt-key add -
$ sudo echo "deb http://pkg.jenkins-ci.org/debian binary/" > \
/etc/apt/sources.list.d/jenkins.list
```

Now, update the Debian package repository:

```
$ sudo aptitude update
```

Once this is done, you can install Jenkins using the **aptitude** tool:

```
$ sudo aptitude install -y jenkins
```

This will install Jenkins as a service, with a correctly configured startup script in `/etc/init.d/jenkins` and a corresponding system user called “jenkins”. If you didn’t already have Java installed on your server, it will also install the OpenJDK version of Java. By default, you'll find the Jenkins WAR file in `/usr/share/jenkins`, and the Jenkins home directory in `/var/lib/jenkins`.

The installation process should have started Jenkins. In general, to start Jenkins, simply invoke this script:

```
$ sudo /etc/init.d/jenkins start
```

Jenkins will now be running on the default port of 8080 (`http://localhost:8080/`).

You can stop Jenkins as follows:

```
$ sudo /etc/inid.d/jenkins stop
```

Jenkins will write log files to `/var/log/jenkins/jenkins.log`. You can also fine-tune configuration parameters in `/etc/default/jenkins`. This is useful if you need to modify the Java startup arguments (`JAVA_ARGS`). You can also use this file to configure arguments that will be passed to Jenkins, such as the HTTP port or web application context (see Section 3.8, “Running Jenkins as a Stand-Alone Application”).

3.6. Installing Jenkins on Red Hat, Fedora, or CentOS

There are also native binary packages available for Red Hat, Fedora, and CentOS. First, you need to set up the repository as follows:

```
$ sudo wget -O /etc/yum.repos.d/jenkins.repo \  
http://jenkins-ci.org/redhat/jenkins.repo  
$ sudo rpm --import http://pkg.jenkins-ci.org/redhat/jenkins-ci.org.key
```

On a fresh installation, you may need to install the JDK:

```
$ sudo yum install java-1.6.0-openjdk
```

Next, you can install the package as shown here:

```
$ sudo yum install jenkins
```

This will install the latest version of Jenkins into `/usr/lib/jenkins`. The default Jenkins home directory will be `/var/lib/jenkins`.

Now you can start Jenkins using the `service` command:

```
$ sudo service jenkins start
```

Jenkins will now be running on the default port of 8080 (`http://localhost:8080/`).

Jenkins’s configuration parameters are placed in `/etc/sysconfig/jenkins`. However, at the time of writing the configuration options are more limited than those provided by the Ubuntu package: you can define the HTTP port using the `JENKINS_PORT` parameter, for example, but to specify an application context you need to modify the startup script by hand. The principal configuration options are listed here:

JENKINS_JAVA_CMD

The version of Java you want to use to run Jenkins

JENKINS_JAVA_OPTIONS

Command-line options to pass to Java, such as memory options

JENKINS_PORT

The port that Jenkins will run on

3.7. Installing Jenkins on SUSE or OpenSUSE

Binary packages are also available for SUSE and OpenSUSE, so the installation process on these platforms is straightforward. First, you need to add the Jenkins repository to the SUSE repository list:

```
$ sudo zypper addrepo http://pkg.jenkins-ci.org/opensuse/ jenkins
```

Finally, you simply install Jenkins using the `zypper` command:

```
$ sudo zypper install jenkins
```

As you can gather from the console output, this will install both Jenkins and the latest JDK from Oracle, if the latter isn't already installed. OpenSuse installations typically have the OpenJDK version of Java, but Jenkins prefers the Oracle variety. When it downloads the Oracle JDK, it will prompt you to validate the Oracle Java license before continuing with the installation.

This installation process will also create a `jenkins` user and install Jenkins as a service, so that it will start up automatically whenever the machine boots. To start Jenkins manually, you can invoke the `jenkins` startup script in the `/etc/init.d` directory:

```
$ sudo /etc/init.d/jenkins jenkins start
```

Jenkins will now be running on the default port of 8080 (`http://localhost:8080/`).

The configuration options are similar to the Red Hat installation (see Section 3.6, “Installing Jenkins on Red Hat, Fedora, or CentOS”). You can define a limited number of configuration variables in `/etc/sysconfig/jenkins`, but for any advanced configuration options, you need to modify the startup script in `/etc/init.d/jenkins`.

The `zypper` tool also makes it easy to update your Jenkins instance:

```
$ sudo zypper update jenkins
```

This will download and install the latest version of Jenkins from the Jenkins website.

3.8. Running Jenkins as a Stand-Alone Application

You can run the Jenkins server in one of two ways: either as a stand-alone application, or deployed as a standard web application in a Java Servlet container or application server such as Tomcat, JBoss, or GlassFish. Both approaches have their pros and cons, so we'll look at both here.

Jenkins comes bundled as a WAR file that you can run directly using an embedded servlet container. Jenkins uses the lightweight Winstone servlet engine to allow you to run the server out of the box, without having to configure an application server yourself. This is probably the easiest way to get started, getting you up and running with Jenkins in a matter of minutes. It's also a very flexible option, and provides some extra features unavailable if you deploy Jenkins in a conventional application server. In

particular, if you're running Jenkins as a stand-alone server, you'll be able to install plugins and upgrades on the fly, and restart Jenkins directly from the administration screens.

To run Jenkins using the embedded servlet container, just go to the command line and type the following:

```
C:\Program Files\Jenkins>
  java -jar jenkins.war
[Winstone 2011/07/01 20:54:53] - Beginning extraction from war file
[Winstone 2011/07/01 20:55:07] - No webapp classes folder found - C:\Users\john\
    .jenkins\war\WEB-INF\classes
jenkins home directory: C:\Users\john\.jenkins
...
INFO: Took 35 ms to load
...
[Winstone 2011/07/01 20:55:08] - HTTP Listener started: port=8080
[Winstone 2011/07/01 20:55:08] - Winstone Servlet Engine v0.9.10 running:
    controlPort=disabled
[Winstone 2011/07/01 20:55:08] - AJP13 Listener started: port=8009
```

In a Linux environment, the procedure is similar. Note how you start the Jenkins server from the “jenkins” user account you created earlier:

```
john@lambton:~$ sudo su - jenkins
jenkins@lambton:~$ java -jar /usr/local/jenkins/jenkins.war
[Winstone 2011/07/16 02:11:24] - Beginning extraction from war file
[Winstone 2011/07/16 02:11:27] - No webapp classes folder found - /home/jenkins/
    .jenkins/war/WEB-INF/classes
jenkins home directory: /home/jenkins/.jenkins
...
[Winstone 2011/07/16 02:11:31] - HTTP Listener started: port=8080
[Winstone 2011/07/16 02:11:31] - AJP13 Listener started: port=8009
[Winstone 2011/07/16 02:11:31] - Winstone Servlet Engine v0.9.10 running:
    controlPort=disabled
```

This will start the embedded servlet engine in the console window. The Jenkins web application will now be available on port 8080. When you run Jenkins using the embedded server, there's no web application context, so you access Jenkins directly using the server URL (e.g., <http://localhost:8080>).

To stop Jenkins, just press Ctrl-C.

By default, Jenkins will run on port 8080. If this doesn't suit your environment, you can specify the port manually, using the `--httpPort` option:

```
$ java -jar jenkins.war --httpPort=8081
```

In the real-world, Jenkins may not be the only web application running on your build server. Depending on the capacity of your server, Jenkins may have to cohabitate with other web applications or Maven repository managers, for example. If you're running Jenkins along side another application server, such as Tomcat, Jetty, or GlassFish, you'll also need to override the `ajp13` port, using the `--ajp13Port` option:

```
$ java -jar jenkins.war --httpPort=8081 --ajp13Port=8010
```

Some other useful options are:

--prefix

This option defines a context path for your Jenkins server. By default Jenkins will run on port 8080 with no context path (<http://localhost:8080>). However, if you use this option, you can force Jenkins to use whatever context path suits you, for example:

```
$ java -jar jenkins.war --prefix=jenkins
```

In this case, Jenkins will be accessible on <http://localhost:8080/jenkins>.

This option is often used when integrating a stand-alone instance of Jenkins with Apache.

--daemon

If you're running Jenkins on a Unix machine, you can use this option to start Jenkins as a background task, running as a Unix daemon.

--logfile

By default, Jenkins writes its logfile into the current directory. However, on a server, you often need to write your log files into a predetermined directory. You can use this option to redirect your messages into some other file:

```
$ java -jar jenkins.war --logfile=/var/log/jenkins.log
```

Stopping Jenkins using Ctrl-C is a little brutal, of course. In practice, you'd set up a script to start and stop your server automatically.

If you're running Jenkins using the embedded Winstone application server, you can also restart and shutdown Jenkins elegantly by calling the Winstone server directly. To do this, you need to specify the `controlPort` option when you start Jenkins, as shown here:

```
$ java -jar jenkins.war --controlPort=8001
```

A slightly more complete example in a Unix environment might look like this:

```
$ nohup java -jar jenkins.war --controlPort=8001 > /var/log/jenkins.log 2>&1 &
```

The key here is the `controlPort` option. This option lets you stop or restart Jenkins directly via the Winstone tools. The only problem is that you need a matching version of the Winstone JAR file. Fortunately, one comes bundled with your Jenkins installation, so you don't have to look far.

To restart the server, you can run the following command:

```
$ java -cp $JENKINS_HOME/war/winstone.jar winstone.tools.WinstoneControl reload: \
--host=localhost --port=8001
```

And to shut it down completely, you can use the following:

```
$ java -cp $JENKINS_HOME/war/winstone.jar winstone.tools.WinstoneControl shutdown \
--host=localhost --port=8001
```

Another way to shut down Jenkins cleanly is to invoke the special “/exit” URL, as shown here:

```
$ wget http://localhost:8080/exit
```

On a real server, only a system administrator should be able to access this URL. In this case, you'll need to provide a username and a password:

```
$ wget --user=admin --password=secret http://localhost:8080/exit
```

Note that you can actually do this from a different server, not just the local machine:

```
$ wget --user=admin --password=secret http://buildserver.acme.com:8080/exit
```

While both these methods will shut down Jenkins relatively cleanly (more so than killing the process directly, for example), they will interrupt any builds in progress. So, it's recommended practice to prepare the shutdown cleanly by using the Prepare for Shutdown button on the Manage Jenkins screen (see Section 4.2, “The Configuration Dashboard—The Manage Jenkins Screen”).

Running Jenkins as a stand-alone application may not be to everyone's taste. For a production server, you might want to take advantage of the more sophisticated monitoring and administration features of a full blown Java application server such as JBoss, GlassFish, or WebSphere. System administrators may be wary of the relatively little-known Winstone server, or may simply prefer Jenkins to fit into a known style of Java web application development. If this is the case, you may prefer to, or be obliged to, deploy Jenkins as a standard Java web application. We look at this option in the following section.

3.9. Running Jenkins Behind an Apache Server

If you're running Jenkins in a Unix environment, you may want to hide it behind an Apache HTTP server in order to harmonize the server URLs and simplify maintenance and access. This way, users can access Jenkins using a URL like `http://myserver.myorg.com/jenkins` rather than `http://myserver.myorg.com:8081`.

One way to do this is to use the Apache `mod_proxy` and `mod_proxy_ajp` modules. These modules implement proxying on your Apache server using the AJP13 (Apache JServer Protocol version 1.3). Using this module, Apache will transfer requests to particular URL patterns on your Apache server (running on port 80) directly to the Jenkins server running on a different port. So, when a user opens a URL like `http://www.myorg.com/jenkins`, Apache will transparently forward traffic to your Jenkins server running on `http://buildserver.myorg.com:8081/jenkins`. Technically, this is known as “Reverse Proxying,” as the client has no knowledge that the server is doing any proxying, or where the proxied server is located. So, you can safely tuck your Jenkins server away behind a firewall, while still providing broader access to your Jenkins instance via the public-facing URL.

The exact configuration of this module will vary depending on your Apache version and installation details, but one possible approach is shown here.

First of all, if you're running Jenkins as a stand-alone application, make sure you start up Jenkins using the `--prefix` option. The prefix you choose must match the suffix in the public-facing URL you want to use. So, if you want to access Jenkins via the URL `http://myserver.myorg.com/jenkins`, you'll need to provide `jenkins` as a prefix:

```
$ java -jar jenkins.war --httpPort=8081 --ajp13Port=8010 --prefix=jenkins
```

If you're running Jenkins in an application server such as Tomcat, it will already be running under a particular web context (`/jenkins` by default).

Next, make sure the `mod_proxy` and `mod_proxy_ajp` modules are activated. In `httpd.conf` (often in `/etc/httpd/conf`), you should have the following line:

```
LoadModule proxy_module modules/mod_proxy.so
```

The proxy is actually configured in `proxy_ajp.conf` often in `/etc/httpd/conf.d`. Note that the name of the proxy path (`/jenkins` in this example) must match the prefix or web context that Jenkins is using. An example of such a configuration file is given here:

```
LoadModule proxy_ajp_module modules/mod_proxy_ajp.so

ProxyPass          /jenkins http://localhost:8081/jenkins
ProxyPassReverse   /jenkins http://localhost:8081/jenkins
ProxyRequests      Off
```

Once this is done, you just need to restart your Apache server:

```
$ sudo /etc/init.d/httpd restart
Stopping httpd:          [ OK ]
Starting httpd:          [ OK ]
```

Now, you should be able to access your Jenkins server using a URL like `http://myserver.myorg.com/jenkins`.

3.10. Running Jenkins in an Application Server

Since Jenkins is distributed as an ordinary WAR file, it's easy to deploy it in any standard Java application server such as Tomcat, Jetty, or GlassFish. Running Jenkins in an application server is arguably more complicated to setup and to maintain. You also lose certain nice administration features such as the ability to upgrade Jenkins or restart the server directly from within Jenkins. On the other hand, your system administrators might be more familiar with maintaining an application running in Tomcat or GlassFish than in the more obscure Winstone server.

Let's look at how you'd typically deploy Jenkins in a Tomcat server. The easiest approach is undoubtedly to simply unzip the Tomcat binary distribution onto your disk (if it's not already installed) and copy the

`jenkins.war` file into the Tomcat `webapps` directory. You can download the Tomcat binaries from the Tomcat website¹.

You start Tomcat by running the `startup.bat` or `startup.sh` script in the Tomcat `bin` directory. Jenkins will be available when you start Tomcat. You should note that, in this case, Jenkins will execute in its own web application context (typically “jenkins”), so you'll need to include this in the URL you use to access your Jenkins server (e.g., `http://localhost:8080/jenkins`).

However, this approach isn't necessarily the most flexible or robust option. If your build server is a Windows box, for example, you probably should install Tomcat as a Windows service, so that you can ensure that it starts automatically whenever the server reboots. Similarly, if you're installing Tomcat in a Unix environment, it should be set up as a service.

3.11. Memory Considerations

CI servers use a lot of memory. This is the nature of the beast—builds will consume memory, and multiple builds running in parallel will consume still more memory. So, you should ensure that your build server has enough RAM to cope with however many builds you intend to run simultaneously.

Jenkins naturally needs RAM to run, but if you need to support a large number of build processes, it's not enough just to give Jenkins a lot of memory. In fact Jenkins creates a new Java process each time it kicks off a build, so during a large build, the build process needs the memory, not Jenkins.

You can define build-specific memory options for your Jenkins build jobs—we'll see how to do this later on in the book. However, if you have a lot of builds to maintain, you might want to define the `JAVA_OPTS`, `MAVEN_OPTS`, and `ANT_OPTS` environment variables to be used as default values for your builds. The `JAVA_OPTS` options will apply for the main Jenkins process, whereas the other two options will be used when Jenkins kicks off new JVM processes for Maven and Ant build jobs respectively.

Here's an example of how these variables might be configured on a Unix machine in the `.profile` file:

```
export JAVA_OPTS=-Djava.awt.headless=true -Xmx512m -DJENKINS_HOME=/data/jenkins
export MAVEN_OPTS="-Xmx512m -XX:MaxPermSize=256m"
export ANT_OPTS="-Xmx512m -XX:MaxPermSize=256m"
```

3.12. Installing Jenkins as a Windows Service

If you're running a production installation of Jenkins on a Windows box, it's essential to have it running as a Windows service. This way, Jenkins will automatically start whenever the server reboots, and can be managed using the standard Windows administration tools.

One of the advantages of running Jenkins on an application server such as Tomcat is that it's generally fairly easy to configure these servers to run as a Windows service. However, it's also fairly easy to install Jenkins as a service, without having to install Tomcat.

¹ <http://tomcat.apache.org>

Jenkins has a very convenient feature designed to make it easy to install Jenkins as a Windows service. There's currently no graphical installer that does this for you, but you get the next best thing—a web-based graphical installer.

First, you need to start the Jenkins server on your target machine. The simplest approach is to run Jenkins using Java Web Start (see Figure 3.4, “Starting Jenkins using Java Web Start”). Alternatively, you can do this by downloading Jenkins and running it from the command line, as we discussed earlier:

```
C:\jenkins> java -jar jenkins.war
```

This second option is useful if the default Jenkins port (8080) is already being used by another application. It doesn't actually matter which port you use—you can change this later.

Figure 3.4. Starting Jenkins using Java Web Start

Once you have Jenkins running, connect to it and go to the Manage Jenkins screen. Here you'll find an Install as Windows Service button. This will create a Jenkins service that will automatically start and stop Jenkins in an orderly manner (see Figure 3.5, “Installing Jenkins as a Windows service”).

Jenkins will prompt you for an installation directory. This will be the Jenkins home directory (`JENKINS_HOME`). The default value is the default `JENKINS_HOME` value: a directory called `.jenkins` in the current user's home directory. This is often not a good choice for a Windows installation. When running Jenkins on Windows XP, you should avoid installing your Jenkins home directory anywhere near your `C:\Documents And Settings` directory—not only is it a ridiculously long name, the spaces can wreak havoc with your Ant and Maven builds and any tests using classpath-based resources. It's much better to use a short and sensible name such as `C:\Jenkins`. The Vista and Windows 7 home directory paths like `C:\Users\john` also work fine.

Figure 3.5. Installing Jenkins as a Windows service

A short home directory path is sometimes required for other reasons, too. On many versions of Windows (Windows XP, Windows Server 2003, etc.), file path lengths are limited to around 260 characters. If you combine a nested Jenkins work directory and a deep classpath, you can often overrun this, which will result in very obscure build errors. To minimize the risks of over-running the Windows file path limits, you need to redefine the `JENKINS_HOME` environment variable to point to a shorter path, as we discussed above.

This approach won't always work with Windows Vista or Windows 7. An alternative strategy is to use the `jenkins.exe` program that the Web Start installation process will have installed in the directory you specified above. Open the command line prompt as an administrator (right-click, “Run as administrator”) and run the `jenkins.exe` executable with the `install` option:

```
C:\Jenkins> jenkins.exe install
```

This basic installation will work fine in a simple context, but you'll often need to fine-tune your service. For example, by default, the Jenkins service will be running under the local system account. However, if you're using Maven, Jenkins will need an `.m2` directory and a `settings.xml` file in the home directory. Similarly, if you're using Groovy, you might need a `.groovy/lib` directory. And so on. To allow this, and to make testing your Jenkins install easier, make sure you run this service under a real user account with the correct development environment set up (see Figure 3.6, “Configuring the Jenkins Windows Service”). Alternatively, run the application as the system user, but use the System Information page in Jenkins to check the `/home/my pc/Documents/jenkinsbook` directory, and place any files that must be placed in the user home directory here.

Figure 3.6. Configuring the Jenkins Windows Service

You configure the finer details of the Jenkins service in a file called `jenkins.xml`, in the same directory as your `jenkins.war` file. Here you can configure (or reconfigure) ports, JVM options, and the Jenkins work directory. In the following example, you give Jenkins a bit more memory and get it to run on port 8081:

```
<service>
  <id>jenkins</id>
  <name>Jenkins</name>
  <description>This service runs the Jenkins continuous integration system
  </description>
  <env name="JENKINS_HOME" value="D:\jenkins" />
  <executable>java</executable>
  <arguments>-Xrs -Xmx512m
  -Dhudson.lifecycle=hudson.lifecycle.WindowsServiceLifecycle
  -jar "%BASE%\jenkins.war" --httpPort=8081 --ajp13Port=8010</arguments>
</service>
```

Finally, if you need to uninstall the Jenkins service, you can do one of two things. The simplest is to run the Jenkins executable with the `uninstall` option:

```
C:\jenkins> jenkins.exe uninstall
```

The other option is to use the Windows service tool `sc`:

```
C:> sc delete jenkins
```

3.13. What's in the Jenkins Home Directory

The Jenkins home directory contains all the details of your Jenkins server configuration, details that you configure in the Manage Jenkins screen. These configuration details are stored in a set of XML files.

Much of the core configuration, for example, is stored in `config.xml`. Other tool-specific configuration is stored in other appropriately-named XML files. The details of your Maven installations, for example, are stored in `hudson.tasks.Maven.xml`. You rarely need to modify these files by hand, though occasionally doing so can come in handy.

The Jenkins home directory also contains a number of subdirectories (see Figure 3.7, “The Jenkins home directory”). Not all of the files and directories will be present after a fresh installation, as some are created when required by Jenkins. And if you look at an existing Jenkins installation, you'll see additional XML files relating to Jenkins configuration and plugins.

The main directories are described in more detail in Table 3.1, “The Jenkins home directory structure”.

Table 3.1. The Jenkins home directory structure

Directory	Description
<code>.jenkins</code>	The default Jenkins home directory (may be <code>.hudson</code> in older installations).
<code>fingerprints</code>	This directory is used by Jenkins to keep track of artifact fingerprints. We look at how to track artifacts later on in the book.
<code>jobs</code>	This directory contains configuration details about the build jobs that Jenkins manages, as well as the artifacts and data resulting from these builds. We look at this directory in detail below.
<code>plugins</code>	This directory contains any plugins that you've installed. Plugins allow you to extend Jenkins by adding extra features. Note that, with the exception of the Jenkins core plugins (subversion, cvs, ssh-slaves, maven, and scid-ad), plugins aren't stored with the <code>jenkins</code> executable, or in the expanded web application directory. This means that you can update your Jenkins executable and not have to reinstall all your plugins.
<code>updates</code>	This is an internal directory used by Jenkins to store information about available plugin updates.
<code>userContent</code>	You can use this directory for storing your own custom content on your Jenkins server. You can access files in this directory at <code>http://myserver/hudson/userContent</code> (if you are running Jenkins on an application server) or <code>http://myserver/userContent</code> (if you're running in stand-alone mode).
<code>users</code>	If you're using the native Jenkins user database, user accounts will be stored in this directory.
<code>war</code>	This directory contains the expanded web application. When you start Jenkins as a stand-alone application, it will extract the <code>WWAR</code> file into this directory.

Figure 3.7. The Jenkins home directory

The `jobs` directory is a crucial part of the Jenkins directory structure, and deserves a bit more attention. You can see an example of a real Jenkins jobs directory in Figure 3.8, “The Jenkins jobs directory”.

Figure 3.8. The Jenkins jobs directory

This directory contains a subdirectory for each Jenkins build job being managed by this instance of Jenkins. Each job directory in turn contains two subdirectories: `builds` and `workspace`, along with some other files. In particular, the job directory contains the build job `config.xml` file, which contains, as you might expect, the configuration details for this build job. There are also some other files used internally by Jenkins that you usually won't touch, such as `nextBuildNumber` (which contains the number that will be assigned to the next build in this build job), as well as symbolic links to the most recent successful build and the last stable one. A successful build is one that doesn't have any compilation errors. A stable build is a successful build that has passed whatever quality criteria you may have configured, such as unit tests, code coverage, and so forth.

Both the `build` and the `workspace` directories are important. The `workspace` directory is where Jenkins builds your project: it contains the source code Jenkins checks out, plus any files generated by the build itself. This workspace is reused for each successive build—there's only ever one `workspace` directory per project, and the disk space it requires tends to be relatively stable.

The `builds` directory contains a history of the builds executed for this job. You rarely need to directly manipulate these directories, but it can be useful to know what they contain. You can see a real example of the builds directory in Figure 3.9, “The builds directory”, where three builds have been performed. Jenkins stores build history and artifacts for each build it performs in a directory named as a timestamp (“2010-03-12_20-42-05” and so forth in Figure 3.9, “The builds directory”). It also contains symbolic links with the actual build numbers that point to the build history directories.

Figure 3.9. The builds directory

Each build directory contains information such as the build result log file, the Subversion revision number used for this build (if you're using Subversion), the changes that triggered this build, and any other data or metrics that you've asked Jenkins to keep track of. For example, if your build job keeps track of unit test results or test coverage metrics, this data will be stored here for each build. The build directory also contains any artifacts you're storing—binary artifacts, and also other generated files such as javadoc or code coverage metrics. Some types of build jobs, such as the Jenkins Maven build jobs, will also archive binary artifacts by default.

The size of the `build` directory will naturally grow over time, as the number of builds increases. You'll probably want to take this into account when designing your build server directory structure, especially

if your build server is running in a Unix-style environment with multiple disk partitions. A lot of this data take the form of text or XML files, which don't consume a large amount of extra space for each build. However, if your build archives some of the build artifacts, such as JAR or WAR files, they too will be stored here. The size of these artifacts should be factored into your disk space requirements. We'll see later on how to limit the number of builds stored for a particular build job if space is an issue. Limiting the number of build jobs that Jenkins stores is always a trade-off between disk space and keeping useful build statistics, as Jenkins does rely on this build history for its powerful reporting features.

Jenkins uses the files in this directory extensively to display build history and metrics data, so you should be particularly careful not to delete any of the build history directories without knowing exactly what you're doing.

3.14. Backing Up Your Jenkins Data

It's important to ensure that your Jenkins data is regularly backed up. This applies in particular to the Jenkins home directory, which contains your server configuration details as well as your build artifacts and build histories. This directory should be backed up frequently and automatically. The Jenkins executable itself is less critical, as it can easily be reinstalled without affecting your build environment.

3.15. Upgrading Your Jenkins Installation

Upgrading Jenkins is easy—you simply replace your local copy of `jenkins.war` and restart Jenkins. However, you should make sure there are no builds running when you restart your server. Since your build environment configuration details, plugins, and build history are stored in the Jenkins home directory, upgrading your Jenkins executable will have no impact on your installation. You can always check what version of Jenkins you're currently running by referring to the version number in the bottom right corner of every screen.

If you've installed Jenkins using one of the Linux packages, Jenkins can be upgraded using the same process as the other system packages on the server.

If you're running Jenkins as a stand-alone instance, you can also upgrade your Jenkins installation directly from the web interface, in the Manage Jenkins section. Jenkins will indicate if a more recent version is available, and give you the option to either download it manually or upgrade automatically (see Figure 3.10, “Upgrading Jenkins from the web interface”).

Figure 3.10. Upgrading Jenkins from the web interface

Once Jenkins has downloaded the upgrade, you can also tell it to restart when no jobs are running. This is probably the most convenient way to upgrade Jenkins, although it won't work in all environments. In particular, you need to be running Jenkins as a stand-alone application, and the user running Jenkins needs to have read-write access to `jenkins.war`.

If you're running Jenkins in an application server such as Tomcat or JBoss, you might need to do a bit more tidying up when you upgrade your Jenkins instance. Tomcat, for example, places compiled JSP pages in the `CATALINA_BASE/work` directory. When you upgrade your Jenkins version, these files need to be removed to prevent the possibility of any stale pages being served.

Any plugins you've installed will be unaffected by your Jenkins upgrades. However, plugins can also be upgraded independently of the main Jenkins executable. Upgrade your plugins directly in the Jenkins web application using the Jenkins Plugin Manager. We discuss plugins in more detail later on in this book.

3.16. Conclusion

In this chapter, you've seen how to install and run Jenkins in different environments, and you've learned a few basic tips on how to maintain your Jenkins installation once running. Jenkins is easy to install, both as a stand-alone application and as a WAR file deployed on an existing application server. The main things you need to consider when choosing a build server are CPU, memory, and disk space.

Chapter 4. Configuring Your Jenkins Server

4.1. Introduction

Before you can start creating your build jobs in Jenkins, you need to do a little configuration, to ensure that your Jenkins server works smoothly in your particular environment. Jenkins is highly configurable, and, although most options are provided with sensible default values, it's always a good idea to know exactly what your build server is doing.

Jenkins is globally very easy to configure. The administration screens are intuitive, and the contextual online help (the blue question mark icons next to each field) is detailed and precise. In this chapter, we'll look at how to configure your basic server setup in detail, including how to configure Jenkins to use different versions of Java, build tools such as Ant and Maven, and SCM tools such as CVS and Subversion. We'll look at more advanced server configuration, such as using other version control systems or notification tools, further on in the book.

4.2. The Configuration Dashboard—The Manage Jenkins Screen

In Jenkins, you manage virtually all features of system configuration in the Manage Jenkins screen (see Figure 4.1, “You configure your Jenkins installation in the Manage Jenkins screen”). You can also get to this screen directly from anywhere in the application by typing “manage” in the Jenkins search box. This screen changes depending on what plugins you install, so don't be surprised if you see more than what we show here.

Figure 4.1. You configure your Jenkins installation in the Manage Jenkins screen

This screen lets you configure different features of your Jenkins server. Each link on this page takes you to a dedicated configuration screen, where you can manage different parts of the Jenkins server. Some of the more interesting options are discussed here:

Configure System

This is where you manage paths to the various tools you use in your builds, such as JDKs, versions of Ant and Maven, as well as security options, email servers, and other system-wide configuration details. Many of the plugins that you install will also need to be configured here—Jenkins will add the fields dynamically when you install the plugins.

Reload Configuration from Disk

As we saw in the previous chapter, Jenkins stores all its system and build job configuration details as XML files stored in the Jenkins home directory (see Section 3.4, “The Jenkins Home Directory”). It also stores all of the build history in the same directory. If you're migrating build jobs from one Jenkins instance to another, or archiving old build jobs, you'll need to add or remove the corresponding build job directories in Jenkins's `builds` directory. You don't need to take Jenkins offline to do this—you can simply use the “Reload Configuration from Disk” option to reload the Jenkins system and build job configurations directly. This process can be a little slow if there's a lot of build history, as Jenkins loads not only the build configurations but also all of the historical data as well.

Manage Plugins

One of the best features of Jenkins is its extensible architecture. There's a large ecosystem of third-party open source plugins available, enabling you to add extra features to your build server, from support for different SCM tools such as Git, Mercurial, or ClearCase, to code quality and code coverage metrics reporting. We'll be looking at many of the more popular and useful plugins throughout this book. Plugins can be installed, updated, and removed through the Manage Plugins screen. Note that removing plugins needs to be done with some care, as it can sometimes affect the stability of your Jenkins instance—we'll look at this in more detail in Section 13.6, “Migrating Build Jobs”.

System Information

This screen displays a list of all the current Java system properties and system environment variables. Here, you can check exactly what version of Java Jenkins is running in, what user it's running under, and so forth. You can also check that Jenkins is using the correct environment variable settings. Its main use is for troubleshooting, so that you can make sure that your server is running with the system properties and variables you think it is.

System Log

The System Log screen is a convenient way to view the Jenkins log files in real time. Again, the main use of this screen is for troubleshooting.

You can also subscribe to RSS feeds for various levels of log messages. For example, as a Jenkins administrator, you might want to subscribe to all the ERROR and WARNING log messages.

Load Statistics

Jenkins keeps track of how busy your server is in terms of the number of concurrent builds and the length of the build queue (which gives an idea of how long your builds need to wait before being executed). These statistics can give you an idea of whether you need to add extra capacity or extra build nodes to your infrastructure.

Script Console

This screen lets you run Groovy scripts on the server. It's useful for advanced troubleshooting: since it requires a strong knowledge of the internal Jenkins architecture, it's mainly useful for plugin developers and the like.

Manage Nodes

Jenkins handles parallel and distributed builds well. In this screen, you can configure how many builds you want Jenkins to run simultaneously. If you're using distributed builds, you can set up build nodes. A build node is another machine that Jenkins can use to execute its builds. We'll look at how to configure distributed builds in detail in Chapter 11, *Distributed Builds*.

Prepare for Shutdown

If you need to shut down Jenkins, or the server Jenkins is running on, it's best not to do so when a build is being executed. To shut down Jenkins cleanly, you can use the Prepare for Shutdown link, which prevents any new builds from being started. Eventually, when all of the current builds have finished, you'll be able to shut down Jenkins cleanly.

We'll come back to some of these features in more detail later on in the book. In the following sections, we'll focus on how to configure the most important Jenkins system parameters.

4.3. Configuring the System Environment

The most important Jenkins administration page is the Configure System screen (Figure 4.2, “System configuration in Jenkins”). Here, you set up most of the fundamental tools that Jenkins needs to do its daily work. The default screen contains a number of sections, each relating to a different configuration area or external tool. In addition, when you install plugins, their system-wide configuration is also often done in this screen.

Figure 4.2. System configuration in Jenkins

The Configure System screen lets you define global parameters for your Jenkins installation, as well as external tools required for your build process. The first part of this screen lets you define some general system-wide parameters.

The Jenkins home directory is displayed, for reference. This way, you can check at a glance that you're working with the home directory that you expect. Remember, you can change this directory by setting the `JENKINS_HOME` environment variable in your environment (see Section 3.4, “The Jenkins Home Directory”).

The System Message field is useful for several purposes. This text is displayed at the top of your Jenkins home page. You can use HTML tags, so it's a simple way to customize your build server by including the name of your server and a short blurb describing its purpose. You can also use it to display messages for all users, such as to announce system outages and so on.

The Quiet Period is useful for SCM tools like CVS that commit file changes one by one, rather than grouped together in a single atomic transaction. Normally, Jenkins will trigger a build as soon as it detects a change in the source repository. However, this doesn't suit all environments. If you're using an SCM tool like CVS, you don't want Jenkins kicking off a build as soon as the first change comes in, as the repository will be in an inconsistent state until all of the changes have been committed. You

can use the Quiet Period field to avoid issues like this. If you set a value here, Jenkins will wait until no changes have been detected for the specified number of seconds before triggering the build. This helps to ensure that all of the changes have been committed and the repository is in a stable state before starting the build.

For most modern version control systems, such as Subversion, Git, or Mercurial, commits are atomic. This means that changes in multiple files are submitted to the repository as a single unit, and the source code on the repository is guaranteed to be in a stable state at all times. However, some teams still use an approach where one logical change set is delivered in several commit operations. In this case, you can use the Quiet Period field to ensure that the build always uses a stable source code version.

The Quiet Period value specified here is in fact the default system-wide value—if required, you can redefine this value individually for each project.

You also manage user accounts and user rights here. By default, Jenkins lets any user do anything. If you want a more restrictive approach, you'll need to activate Jenkins security here using the Enable Security field. There are many ways to do this, and we look at this aspect of Jenkins later on (see Chapter 7, *Securing Jenkins*).

4.4. Configuring Global Properties

The Global Properties (see Figure 4.3, “Configuring environment variables in Jenkins”) section lets you define variables that can be managed centrally but used in all of your build jobs. You can add as many properties as you want here, and use them in your build jobs. Jenkins will make them available within your build job environment, so you can freely use them within your Ant and Maven build scripts. Note that you shouldn't put periods (“.”) in the property names, as they won't be processed correctly. So, `ldapserver` or `ldap_server` is fine, but not `ldap.server`.

Figure 4.3. Configuring environment variables in Jenkins

There are two ways you typically use these variables. Firstly, you can use them directly in your build script, using the `${key}` or `$key` notation (so `${ldapserver}` or `$ldapserver` in the example given above). This is the simplest approach, but means that there's a tight coupling between your build job configuration and your build scripts.

If your script uses a different property name (one containing dots, for example), you can also pass the value to your build script in the build job configuration. In Figure 4.4, “Using a configured environment variable” you pass the `ldapserver` property value defined in Figure 4.3, “Configuring environment variables in Jenkins” to a Maven build job. Using the `-D` option means that this value will be accessible from within the build script. This is a flexible approach, as you can assign the global properties defined within Jenkins to script-specific variables in our build scripts. In Figure 4.4, “Using a configured environment variable”, for example, the `ldapserver` property will be available from within the Maven build via the internal `${ldapserver}` property.

Figure 4.4. Using a configured environment variable

4.5. Configuring Your JDKs

Historically, one of the most common uses of Jenkins has been to build Java applications. So, Jenkins naturally provides excellent built-in support for Java.

By default, Jenkins will build Java applications using whatever version of Java it finds on the system path, which is usually the version that Jenkins itself is running under. However, for a production build server, you'll probably want more control than this. For example, you may be running your Jenkins server under Java 6, for performance reasons. However, your production server might be running under Java 5 or even Java 1.4. Large organizations are often cautious when it comes to upgrading Java versions in their production environments, and some of the more heavyweight application servers on the market are notoriously slow to be certified with the latest JDKs.

In any case, it's always a wise practice to build your application using a version of Java that's close to the one running on your production server. While an application compiled with Java 1.4 will usually run fine under Java 6, the inverse isn't always true. Or, you may have different applications that need to be built using different versions of Java.

Jenkins provides good support for working with multiple JVMs. Indeed, Jenkins makes it very easy to configure and use as many versions of Java as you want. Like most system-level configuration, you do this in the Configure System screen (see Figure 4.2, “System configuration in Jenkins”). Here, you'll find a section called JDK which allows you to manage the JDK installations you need Jenkins to work with.

The simplest way to declare a JDK installation is simply to supply an appropriate name (which will be used to identify this Java installation later on when you configure your builds), along with the path to the Java installation directory (the same path you'd use for the `JAVA_HOME` variable), as shown in Figure 4.5, “JDK configuration in Jenkins”. Although you need to type the path manually, Jenkins will check in real time both that the directory exists and that it looks like a valid JDK directory.

Figure 4.5. JDK configuration in Jenkins

You can also ask Jenkins to install Java for you. In this case, Jenkins will download the JDK installation and install a copy on your machine (see Figure 4.6, “Installing a JDK automatically”). The first time a build needs to use this JDK, Jenkins will download and install the specified version of Java into the `tools` directory in the Jenkins home directory. If the build is running on a new build agent that doesn't have this JDK installed, it will download and install it onto the build agent machine as well.

This is also a great way to configure build agents. As we'll see later on in the book, Jenkins can delegate build jobs to other machines, or build agents. A build agent (or “slave”) is simply another computer that

Jenkins can use to run some of its builds. If you use Jenkins's Install automatically option, you don't need to manually install all the JDK versions you need on the build agent machines—Jenkins will do it for you the first time it needs to.

By default, Jenkins proposes to download the JDK from the Oracle website. If your Jenkins installation is behind a proxy server, you may need to configure your proxy settings to ensure that Jenkins can access the external download sites (see Section 4.9, “Configuring a Proxy”). Another option is to provide a URL pointing to your own internal copy of the JDK binaries (either in the form of a ZIP or a GZip-compressed TAR file), stored on a local server within your organization. This lets you provide standard installations on a local server and makes for faster automatic installations. When you use this option, Jenkins also lets you specify a label to restrict the use of this installation to the build nodes with this label. This is a useful technique if you need to install a specific version of a tool on certain build machines. The same approach can also be used for other build tools (such as Maven and Ant).

Figure 4.6. Installing a JDK automatically

The automatic installer won't work in all environments (if it can't find or identify your operating system to its satisfaction, for example, the installation will fail), but it's nevertheless a useful and convenient way to set up new build servers or distributed build agents in a consistent manner.

4.6. Configuring Your Build Tools

Build tools are the bread-and-butter of any build server, and Jenkins is no exception. Out of the box, Jenkins supports three principal build tools: Ant, Maven, and the basic shell-script (or batch script in Windows). Using Jenkins plugins, you can also add support for other build tools and other languages, such as Gant, Grails, MSBuild, and many more.

4.6.1. Maven

Maven is a high-level build scripting framework for Java that uses notions such as a standard directory structure and standard life cycles, Convention over Configuration, and Declarative Dependency Management to simplify a lot of the low-level scripting that you find in a typical Ant build script. In Maven, your project uses a standard, well-defined build life cycle—compile, test, package, deploy, and so forth. Each life cycle phase is associated with a Maven plugin. The various Maven plugins use the standard directory structure to carry out these tasks with a minimum of intervention on your part. You can also extend Maven by overriding the default plugin configurations or by invoking additional plugins.

Jenkins provides excellent support for Maven, and has a good understanding of Maven project structures and dependencies. You can either get Jenkins to install a specific version of Maven automatically (as we're doing with Maven 3 in the example), or provide a path to a local Maven installation (see Figure 4.7, “Configuring Maven in Jenkins”). You can configure as many versions of Maven for your build projects as you want, and use different versions of Maven for different projects.

Figure 4.7. Configuring Maven in Jenkins

If you tick the **Install automatically** checkbox, Jenkins will download and install the requested version of Maven for you. You can either ask Jenkins to download Maven directly from the Apache site, or from a (presumably local) URL of your choice. This is an excellent choice when you're using distributed builds, and, since Maven is cross-platform, it will work on any machine. You don't need to install Maven explicitly on each build machine—the first time a build machine needs to use Maven, it will download a copy and install it to the `tools` directory in the Jenkins home directory.

Sometimes you need to pass Java system options to your Maven build process. For instance it's often useful to give Maven a bit of extra memory for heavyweight tasks such as code coverage or site generation. Maven lets you do this by setting the `MAVEN_OPTS` variable. In Jenkins, you can set a system-wide default value, to be used across all projects (see Figure 4.8, “Configuring system-wide `MVN_OPTS`”). This comes in handy if you want to use certain standard memory options (for example) across all projects, without having to set it up in each project by hand.

Figure 4.8. Configuring system-wide `MVN_OPTS`

4.6.2. Ant

Ant is a widely-used and very well-known build scripting language for Java. It's a flexible, extensible, relatively low-level scripting language used in a large number of open source projects. An Ant build script (typically called `build.xml`) is made up of a number of targets. Each target performs a particular job in the build process, such as compiling your code or running your unit tests. It does so by executing tasks, which carry out a specific part of the build job, such as invoking `javac` to compile your code, or creating a new directory. Targets also have dependencies, indicating the order in which your build tasks need to be executed. For example, you need to compile your code before you can run your unit tests.

Jenkins provides excellent build-in support for Ant—you can invoke Ant targets from your build job, providing properties to customize the process as required. We look at how to do this in detail later on in this book.

If Ant is available on the system path, Jenkins will find it. However, if you want to know precisely what version of Ant you're using, or if you need to be able to use several different versions of Ant on different build jobs, you can configure as many version of Ant as required (see Figure 4.9, “Configuring Ant in Jenkins”). Just provide a name and installation directory for each version of Ant in the Ant section of the **Configure System** screen. You'll then be able to choose what version of Ant you want to use for each project.

If you're tick the **Install automatically** checkbox, Jenkins will download and install Ant into the `tools` directory of your Jenkins home directory, just like it does for Maven. It will download an Ant installation the first time a build job needs to use Ant, either from the Apache website or from a local URL. Again,

this is a great way to standardize build servers and make it easier to add new distributed build servers to an existing infrastructure.

Figure 4.9. Configuring Ant in Jenkins

4.6.3. Shell-Scripting Language

If you're running your build server on Unix or Linux, Jenkins lets you insert shell scripts into your build jobs. This is handy for performing low-level, OS-related tasks that you don't want to do in Ant or Maven. In the Shell section, you define the default shell that will be used when executing these shell scripts. By default, this is `/bin/sh`, but there are times you may want to modify this to another command interpreter such as `bash` or `Perl`.

In Windows, the Shell section doesn't apply—you use Windows batch scripting instead. So, on a Windows build server, you should leave this field blank.

4.7. Configuring Your Version Control Tools

Jenkins comes preinstalled with plugins for CVS and Subversion. Other version control systems are supported by plugins that you can download from the Manage Plugins screen.

4.7.1. Configuring Subversion

Subversion needs no special configuration, since Jenkins uses native Java libraries to interact with Subversion repositories. If you need to authenticate to connect to a repository, Jenkins will prompt you when you enter the Subversion URL in the build job configuration.

4.7.2. Configuring CVS

CVS needs little or no configuration. By default, Jenkins will look for tools like CVS on the system path, though you can provide the path explicitly if it isn't on the system path. CVS keeps login and password details in a file called `.cvspass`, which is usually in your home directory. If it isn't, you can provide a path where Jenkins can find this file.

4.8. Configuring the Mail Server

The last of the basic configuration options you need to set up is the email server configuration. Email is Jenkins's most fundamental notification technique—when a build fails, it will send an email message to the developer who committed the changes, and optionally to other team members as well. So, Jenkins needs to know about your email server (see Figure 4.10, “Configuring an email server in Jenkins”).

Figure 4.10. Configuring an email server in Jenkins

The System Admin email address is the address from which the notification messages are sent. You can also use this field to check the email setup—if you click on the Test configuration button, Jenkins will send a test email to this address.

In many organizations, you can derive a user's email address from their login by adding the organization domain name. For example, at ACME, user John Smith will have a login of "jsmith" and an email address of "jsmith@acme.com". If this extends to your version control system, Jenkins can save you a lot of configuration effort in this area. In the previous example, you could simply specify the default user email suffix of acme.com and Jenkins will figure out the rest.

You also need to provide a proper base URL for your Jenkins server (one that doesn't use localhost). Jenkins uses this URL in the email notifications so that users can go directly from the email to the build failure screen on Jenkins.

Jenkins also provides for more sophisticated email configuration, using more advanced features such as SMTP authentication and SSL. If you need these, click on the Advanced button to configure these options.

For example, many organizations use Google Apps for their email service. You can configure Jenkins to work with the Gmail service as shown in Figure 4.11, "Configuring an email server in Jenkins to use a Google Apps domain". All you need to do in this case is to use the Gmail SMTP server and provide your Gmail username and password in the SMTP Authentication (you also need to use SSL and the non-standard port of 465).

Figure 4.11. Configuring an email server in Jenkins to use a Google Apps domain

4.9. Configuring a Proxy

In most enterprise environments, your Jenkins server will be situated behind a firewall, and won't have direct access to the Internet. Jenkins needs Internet access to download plugins and updates, and also to install tools such as the JDK, Ant, and Maven from remote sites. If you need to go through an HTTP proxy server to get to the Internet, you can configure the connection details (the server and port, and, if required, the username and password) in the Advanced tab on the Plugin Manager screen (see Figure 4.12, "Configuring Jenkins to use a proxy").

If your proxy is using Microsoft's NTLM authentication scheme, then you'll need to provide a domain name as well as a username. You can place both in the User name field: just enter the domain name, followed by a back-slash (\), followed by the username, such as "MyDomain\Joe Bloggs".

Figure 4.12. Configuring Jenkins to use a proxy

Finally, if you're setting up Proxy access on your Jenkins build server, remember that all of the other tools running on this server will need to know about the proxy as well. In particular, this may include

tools such as Subversion (if you're accessing an external repository) and Maven (if you're not using an Enterprise Repository Manager).

4.10. Conclusion

You don't need a great deal of configuration to get started with Jenkins. The configuration that's required is fairly straightforward, and is centralised in the Configure System screen. Once this is done, you're ready to create your first Jenkins build job!

Chapter 5. Setting Up Your Build Jobs

5.1. Introduction

Build jobs are the basic currency of a CI server.

A build job is a particular way of compiling, testing, packaging, deploying, or otherwise doing something with your project. Build jobs come in a variety of forms; you may want to compile and unit test your application, report on code quality metrics related to the source code, generate documentation, bundle up an application for a release, deploy it to production, run an automated smoke test, or do any number of other similar tasks.

A software project will usually have several related build jobs. For example, you might choose to start off with a dedicated build job that runs all of your unit tests. If these pass, you might proceed to a build job that executes longer-running integration tests, runs code quality metrics, or generates technical documentation, before finally bundling up your web application and deploying it to a test server.

In Jenkins, build jobs are easy to set up. In this chapter, we'll look at the main types of build jobs and how to configure them. In later chapters, we'll take things further, looking at how to organize multiple build jobs, how to set up build promotion pipelines, and how to automate the deployment process. But, for now, let's start off with how to set up your basic build jobs in Jenkins.

5.2. Jenkins Build Jobs

Creating a new build job in Jenkins is simple: just click on the “New Job” menu item on the Jenkins dashboard. Jenkins supports several different types of build jobs which are presented to you when you choose to create a new job (see Figure 5.1, “Jenkins supports four main types of build jobs”).

Freestyle software project

Freestyle build jobs are general-purpose build jobs, which provides a maximum of flexibility.

Maven project

The “maven2/3 project” is a build job specially adapted to Maven projects. Jenkins understands Maven `POM` files and project structures, and can use the information gleaned from the `POM` file to reduce the work you need to do to set up your project.

Monitor an external job

The “Monitor an external job” build job lets you keep an eye on non-interactive processes, such as cron jobs.

Multiconfiguration job

The “multiconfiguration project” (also referred to as a “matrix project”) lets you run the same build job in many different configurations. This powerful feature can be useful for testing an

application in many different environments, with different databases, or even on different build machines. We'll be looking at how to configure multiconfiguration build jobs later on in the book.

Figure 5.1. Jenkins supports four main types of build jobs

You can also copy an existing job, which is a great way to create a new job that's very similar to an existing build job, except for a few configuration details.

In this chapter, we'll focus on the first two types of build jobs, which are the most commonly used. We'll discuss the others later on. Let's start with the most flexible option: the freestyle build job.

5.3. Creating a Freestyle Build Job

The freestyle build job is the most flexible and configurable option, and can be used for any type of project. It's relatively straightforward to set up, and many of the options you configure here also appear in other build jobs.

5.3.1. General Options

The first section you see when you create a new freestyle job contains general information about the project, such as a unique name and description, and other information about how and where the build job should be executed (see Figure 5.2, “Creating a new build job”).

Figure 5.2. Creating a new build job

The project name can be anything you like, but it's worth noting that it will be used for the project directory and the build job URL, so I generally avoid names with spaces. The project description will go on the project home page—use this to provide an overview of the build job's goals and context. HTML tags will work fine in this field.

The other options are more technical, and we'll be looking at some of them in detail later on in the book.

One important aspect that you should think about upfront is how you want to handle build history. Build jobs can consume a lot of disk space, especially if you store the build artifacts (the binary files, such as JARs, WARs, TARs, etc. generated by your build job). Even without artifacts, keeping a record of every build job consumes additional disk space and memory, which may or may not be justified, depending on the nature of your build job. For example, for a code quality metrics build that reports on static analysis and code coverage metrics over time, you might want to keep a record of the builds for the duration of the project, whereas, for a build job that automatically deploys an application to a test server, keeping the build history and artifacts for posterity might be less important.

The Discard Old Builds option lets you limit the number of builds you record in the build history. You can either tell Jenkins to only keep recent builds (Jenkins will delete builds after a certain number of

days), or to keep no more than a specified number of builds. If a certain build has particular sentimental value, you can always tell Jenkins to keep it forever by using the Keep forever button on the build details page (see Figure 5.3, “Keeping a build job forever”). Note that this button will only appear if you've asked Jenkins to discard old builds.

Figure 5.3. Keeping a build job forever

In addition, Jenkins will never delete the last stable and successful builds, no matter how old they are. For example, if you limit Jenkins to only keep the last twenty builds, and your last successful build was thirty builds ago, Jenkins will still keep the successful build job as well as the last twenty failing builds.

You also have the option to disable the build. A disabled build won't be executed until you enable it again. Using this option when you create a new build job is quite rare. On the other hand, this option often comes in handy to temporarily suspend a build during maintenance work or major refactoring, when notification of the build failures won't be useful for the team.

5.3.2. Advanced Project Options

The Advanced Project options contains, as the name suggests, configuration options that are less frequently required. You need to click on the Advanced button for them to appear (see Figure 5.4, “To display the Advanced Options, you need to click on the Advanced button”).

Figure 5.4. To display the Advanced Options, you need to click on the Advanced button

The Quiet Period option in the build job configuration simply lets you override the system-wide quiet period defined in the Jenkins System Configuration screen (see Section 4.3, “Configuring the System Environment”). This option is mainly used for version control systems that don't support atomic commits, such as CVS, but it's also sometimes used in teams where developers have the habit of committing their work in several small commits.

The “Block build when upstream project is building” option is useful when several related projects are affected by a single commit, but they must be built in a specific order. If you activate this option, Jenkins will wait until any upstream build jobs (see Section 5.5, “Build Triggers”) have finished before starting this build.

For instance, when you release a new version of a multimodule Maven project, version number updates will happen in many, if not all, of the project modules. Suppose, for example, that we've added a web application to the Game of Life project you used in Chapter 2, Your First Steps with Jenkins, setting it up as a separate Maven project. When you release a new version of this project, both the core and the web application version numbers will be updated (see Figure 5.5, “The “Block build when upstream project is building” option is useful when a single commit can affect several related projects”). Before you can build the web application, you need to build a new version of the original Game of Life core module. However, if you had a separate freestyle build job for each module, then the build jobs for both

the core and the web application would start simultaneously. The web application build job will fail if the core build job hasn't produced a new version of the core module, even if there are no test failures.

To avoid this issue, you could set up the web application build job to only start once the core build has successfully terminated. However, this would mean that the web application would never be built if changes were made that only affected it, and not the core module. A better approach is to use the “Block build when upstream project” option. In this case, when the version numbers are updated in version control, Jenkins will schedule both builds to be executed. However, it will wait until the core build has finished before starting the web application build.

Figure 5.5. The “Block build when upstream project is building” option is useful when a single commit can affect several related projects

You can also override the default workspace used by Jenkins to check out the source code and build your project. Normally, Jenkins will create a special workspace directory for your project, which can be found in the project's build job directory (see Section 3.13, “What's in the Jenkins Home Directory”). This works fine in almost all cases. However, there are times when you need to override this option and force Jenkins to use a special directory. One common example of this is if you want several build jobs to all work successively in the same directory. You can override the default directory by ticking the “Use custom workspace” option and providing the path yourself. The path can be either absolute or relative to Jenkins's home directory.

We'll look at some of the other more advanced options that appear in this section later on in the book.

5.4. Configuring Source Code Management

In its most basic role, a CI server monitors your version control system, and checks out the latest changes as they occur. The server then compiles and tests the most recent version of the code. Alternatively, it may simply check out and build the latest version of your source code on a regular basis. In either case, tight integration with your version control system is essential.

Because of its fundamental role, SCM configuration options in Jenkins are identical across all sorts of build jobs. Jenkins supports CVS and Subversion out of the box, with built-in support for Git, and also integrates with a large number of other version control systems via plugins. At the time of writing, SCM plugin support includes Accurev, Bazaar, BitKeeper, ClearCase, CMVC, Dimensions, Git, CA Harvest, Mercurial, Perforce, PVCS, StarTeam, CM/Synergy, Microsoft Team Foundation Server, and even Visual SourceSafe. In the rest of this section, we'll look at how to configure some of the more common SCM tools.

5.4.1. Working with Subversion

Subversion is one of the most widely used version control systems, and Jenkins comes bundled with full Subversion support (see Figure 5.6, “Jenkins provides built-in support for Subversion”). To use

source code from a Subversion repository, you simply provide the corresponding Subversion URL—it will work fine with any of the three Subversion protocols of (`http`, `svn`, or `file`). Jenkins will check that the URL is valid as soon as you enter it. If the repository requires authentication, Jenkins will prompt you for the corresponding credentials automatically, and store them for any other build jobs that access this repository.

Figure 5.6. Jenkins provides built-in support for Subversion

By default, Jenkins will check out the repository contents into a subdirectory of your workspace whose name will match the last element in the Subversion URL. So, if your Subversion URL is `svn://localhost/gameoflife/trunk`, Jenkins will check out the repository contents to a directory called `trunk` in the build job workspace. If you would prefer another directory name, just enter the directory name you want in the "Local module directory" field. Place a period (".") here if you want Jenkins to check the source code directly into the workspace.

Occasionally you may need to get source code from more than one Subversion URL. In this case, just use the "Add more locations..." button to add as many additional repository sources as you need.

A well-designed build process shouldn't modify the source code, or leave any extra files that might confuse your version control system or the build process. Both generated artifacts and temporary files (such as log files, reports, test data, or file-based databases) should go in a directory set aside for this purpose (such as the `target` directory in Maven builds), and/or be configured to be ignored by your version control repository. They should also be deleted as part of the build process, once the build has finished with them. This is also an important part of ensuring a clean and reproducible build process—for a given version of your source code, your build should behave in exactly the same way, no matter where or when it's run. Locally changed source code files, and the presence of temporary files, both have the potential of compromising this.

You can fine-tune the way Jenkins obtains the latest source code from your Subversion repository by selecting an appropriate value in the Check-out Strategy drop-down list. If your project is well-behaved, however, you may be able to speed things up substantially by selecting "Use 'svn update' as much as possible". This is the fastest option, but may leave artifacts and files from previous builds in your workspace. To be on the safe side, you may want to use the second option ("Use 'svn update' as much as possible, with 'svn revert' before update"), which will systematically run `svn revert` before running `svn update`. This will ensure that no local files have been modified, though it won't remove any new files that have been created during the build process. Alternatively, you can ask Jenkins to delete any unversioned or ignored files before performing an `svn update`, or play it safe by checking out a full clean copy for each build.

Another very useful feature is Jenkins's integration with source code browsers. A good source code browser is an important part of your CI setup. It lets you see at a glance what changes triggered a given build, which is very useful when it comes to troubleshooting broken builds (see Figure 5.7, "Source code browser showing what code changes caused a build"). Jenkins integrates with most of the major

source code browsers, including open source tools such as WebSVN and Sventon, and commercial ones like Atlassian's FishEye.

Figure 5.7. Source code browser showing what code changes caused a build

Jenkins also lets you refine the changes that will trigger a build. In the Advanced section, you can use the Excluded Regions field to tell Jenkins not to trigger a build if only certain files were changed. This field takes a list of regular expressions which identify files that should not trigger a build. For example, suppose you don't want Jenkins to start a new build if only images have been changed. To do this, you could use a set of regular expressions like the following:

```
/trunk/gameoflife/gameoflife-web/src/main/webapp/.*\.(jpg|gif|png)
```

Alternatively, you can specify the Included Regions, if you're only interested in changes in part of the source code directory structure. You can even combine the Excluded Regions and Included Regions fields—in this case a modified file will only trigger a build if it's in the Included Regions but not in the Excluded Regions.

You can also ignore changes coming from certain users (Excluded Users), or with certain commit messages (Excluded Commit Messages). For example, if your project uses Maven, you may want to use the Maven Release Plugin to promote your application from snapshot versions to official releases. This plugin will automatically bump up the version number of your application from a snapshot version used during development (such as 1.0.1-SNAPSHOT) to a release (1.0.1), bundle up and deploy a release of your application with this version number, and then move the version on to the next snapshot number (e.g., 1.0.2-SNAPSHOT) for ongoing development. During this process Maven takes care of many SCM bookkeeping tasks, such as committing the source code with the release version number and creating a tag for the released version of your application, and then committing the source code with the new snapshot version number.

Suppose you have a special build job for generating a new release using this process. The many commits generated by the Maven Release Plugin would normally trigger build jobs in Jenkins. However, since the release build job is already compiling and testing this version of your application, you don't need Jenkins to do it again in a separate build job. To ensure that Jenkins doesn't trigger a build for this case, you can use the Excluded Commit Messages field with the following value:

```
[maven-release-plugin] prepare release.*
```

This will ensure that Jenkins skips the changes corresponding to the new release version, but not those corresponding to the next snapshot version.

5.4.2. Working with Git

Contributed by Matthew McCullough

Git¹ is a popular distributed version control system that's a logical successor to Subversion² and a mind-share competitor to Mercurial³. Git support in Jenkins is both mature and full-featured. There are a number of plugins that can contribute to the overall success of Git in Jenkins. We'll begin by looking at the Git plugin, which provides core Git support in Jenkins. We'll discuss the supplemental plugins shortly.

5.4.2.1. Installing the plugin

The Git plugin is available in the Jenkins Plugin Manager and is documented on its own wiki page⁴. The plugin assumes that Git (version 1.3.3 or later) has already been installed on your build server, so you'll need to make sure that this is the case. You can do this by running the following command on your build server:

```
$ git --version
git version 1.7.1
```

Next, go back to Jenkins, check the corresponding check box in the Jenkins Plugin Manager page, and click the Install button.

5.4.2.1.1. System-wide configuration of the plugin

After installing the Git plugin, a small new set of configuration options will be available on the Manage Jenkins#Configure System page (see Figure 5.8, “System-wide configuration of the Git plugin”). In particular, you need to provide the path to your Git executable. If Git is already installed on the system path, just put “git” here.

Figure 5.8. System-wide configuration of the Git plugin

5.4.2.1.2. SSH key setup

If the Git repository you're accessing uses SSH passphrase-less authentication—for example, if the access address is similar to `git@github.com:matthewmccullough/some-repo.git`—you'll need to provide the private half of the key in `~/.ssh/id_rsa`, where `~` is the home directory of the user account under which Jenkins is running.

The fingerprint of the remote server will additionally need to be placed in `~/.ssh/known_hosts` to prevent Jenkins from invisibly prompting for authorization to access this Git server for the first time.

Alternatively, if the `jenkins` user is allowed to login, SSH into the Jenkins machine as `jenkins` and manually attempt to Git clone a remote repository. This will test your private key setup and establish the

¹ <http://git-scm.com/>

² <http://subversion.tigris.org/>

³ <http://mercurial.selenic.com/>

⁴ <http://wiki.hudson-ci.org/display/HUDSON/Git+Plugin>

`known_hosts` file in the `~/.ssh` directory. This is probably the simplest option for users unfamiliar with the intricacies of SSH configuration.

5.4.2.2. Using the plugin

On either an existing or a new Jenkins project, a new Source Code Management option for Git will be displayed. From here, you can configure one or more repository addresses (see Figure 5.9, “Entering a Git repo URL”). One repository is usually enough for most projects; adding a second repository can be useful in more complicated cases, and lets you specify distinct named locations for `pull` and `push` operations.

5.4.2.2.1. Advanced per-project source code management configuration

In most cases, the URL of the Git repository you're using should be enough. However, if you need more options, click on the Advanced button (see Figure 5.10, “Advanced configuration of a Git repo URL”). This provides more precise control of the `pull` behavior.

The Name of repository is a shorthand title (a.k.a. `remote` in Git parlance) for a given repository, that you can refer to later on in the merge action configuration.

The Refspec is a Git term⁵ for controlling precisely what's retrieved from remote servers and under what namespace it's stored locally.

5.4.2.2.2. Branches to build

The branch specifier (Figure 5.11, “Advanced configuration of the Git branches to build”) is the wildcard pattern or specific branch name that Jenkins should build. If left blank, all branches will be built. At the time writing, after the first time saving a job with a blank branches to build setting, it becomes populated with `**`, which means “build all branches.”

Figure 5.9. Entering a Git repo URL

Figure 5.10. Advanced configuration of a Git repo URL

Figure 5.11. Advanced configuration of the Git branches to build

5.4.2.2.3. Excluded regions

Regions (seen in Figure 5.12, “Branches and regions”) are named specific or wildcard paths in the codebase that, even when changed, shouldn't trigger a build. These are commonly noncompiled files

⁵ <http://progit.org/book/ch9-5.html>

such as localization bundles or images, which, understandably might not have an effect on unit or integration tests.

Figure 5.12. Branches and regions

5.4.2.2.4. Excluded users

The Git plugin also lets you ignore certain users, even if they make changes to the codebase that would typically trigger a build.

This isn't as spiteful as it sounds. Excluded users are typically automated users, not human developers, that happen to have distinct accounts with commit rights in the source code control system. These automated users often are performing small numeric changes such as bumping up version numbers in a `pom.xml` file, rather than making actual logic changes. If you want to exclude several users, just place each of them on separate lines.

5.4.2.2.5. Checkout/merge to local branch

There are times when you may want to create a local branch from the tree you've specified, rather than just using a direct detached HEAD checkout of the commit's hash. In this case, just specify your local branch in the "Checkout/merge to a local branch" field.

This is a little easier to illustrate with an example. Without specifying a local branch, the plugin would do something like this:

```
git checkout 73434e4a0af0f51c242f5ae8efc51a88383afc8a
```

On the other hand, if you use a local branch named `mylocalbranch`, Jenkins would do the following:

```
git branch -D mylocalbranch
git checkout -b mylocalbranch 73434e4a0af0f51c242f5ae8efc51a88383afc8a
```

5.4.2.2.6. Local subdirectory for repo

By default, Jenkins will clone the Git repository directly into the build job workspace. If you prefer to use a different directory, you can specify it here. Note that the directory you specify is relative to the build job workspace.

5.4.2.2.7. Merge before build

The typical recipe for using this option is to fold an integration branch into a branch more similar to `master`. Keep in mind that only conflict-less merges will happen automatically. More complex merges that require manual intervention will cause the build to fail.

The resultant merged branch won't automatically be pushed to another repository unless the later `push` post-build action is enabled.

5.4.2.2.8. Prune remote branches before build

Pruning removes local copies of remote branches that exist as a remnant of the previous clone but are no longer present on the remote. In short, this is cleaning the local clone to be in perfect sync with its remote siblings.

5.4.2.2.9. Clean after checkout

Activate Git's facilities for purging any untracked files or folders, returning your working copy to a pristine state.

5.4.2.2.10. Recursively update submodules

If you're using Git's submodule facilities in the project, this option lets you ensure that every submodule is up-to-date with an explicit call to `update`, even if submodules are nested within other submodules.

5.4.2.2.11. Use commit author in changelog

Jenkins tracks and displays the author of changed code in a summarized view. Git tracks both the committer and author of code distinctly, and this option lets you toggle which of those two usernames is displayed in the changelog.

5.4.2.2.12. Wipe out workspace

Typically Jenkins will reuse the workspace, merely freshening the checkout as necessary and, if you activated the “Clean after checkout” option, cleaning up untracked files. However, if you prefer to have a completely clean workspace, you can use the “Wipe out workspace” option to delete and rebuild the workspace from the ground up. Bear in mind that this may significantly lengthen the time it takes to initialize and build the project.

5.4.2.2.13. Choosing strategy

Jenkins decides which branches to build based on a strategy (see Figure 5.13, “Choosing strategy”). Users can influence this branch-search process. The default choice is to search for all branch HEADs. If the Gerrit plugin is installed, additional options for building all Gerrit-notified commits are displayed.

Figure 5.13. Choosing strategy

5.4.2.2.14. Git executable

In the Jenkins global options (see Figure 5.14, “Git executable global setup”), different Git executables can be set up and used on a per-build basis. This is infrequently used, and only when the clone or other

Git operations are highly sensitive to a particular version of Git. Git tends to be very version-flexible; slightly older repositories can easily be cloned with a newer version of Git and vice-versa.

Figure 5.14. Git executable global setup

5.4.2.2.15. Repository browser

Like Subversion, Git has several source code browsers that you can use. The most common ones are Gitorious, Git Web, or GitHub. If you provide the URL to the corresponding repository browser, Jenkins will be able to display a link to the source code changes that triggered a build (see Figure 5.15, “Repository browser”).

Figure 5.15. Repository browser

5.4.2.3. Build triggers

The basic Git plugin offers the ability to Poll SCM on a timed basis, looking for changes since the last inquiry. If changes are found, a build is started. The polling log (shown in Figure 5.16, “Polling log”) is accessible via a link on the left hand side of the page in the navigation bar when viewing a specific job. It offers information on the last time the repository was polled and if it replied with a list of changes (see Figure 5.17, “Results of Git polling”).

Figure 5.16. Polling log

The Git polling is distilled into a more developer-useful format that shows commit comments as well as hyperlinking usernames and changed files to more detailed views of each.

Figure 5.17. Results of Git polling

Installing the Gerrit Build Trigger adds a Gerrit event option that can be more efficient and precise than simply polling the repository.

5.4.2.3.1. Gerrit Trigger

Gerrit⁶ is an open source web application that facilitates code reviews⁷ for project sources hosted in a Git version control system. It reads a traditional Git repository, and provides a side-by-side comparison

⁶ <http://code.google.com/p/gerrit/>

⁷ <https://review.source.android.com/#q,status:open,n,z>

of changes. As the code is reviewed, Gerrit provides a location to comment and move the patch to an open, merged, or abandoned status.

The Gerrit Trigger⁸ is a Jenkins plugin that can trigger a Jenkins build of the code when any user-specified activity happens in a user-specified project in the Git repository (see Figure 5.18, “Gerrit Trigger”). It’s an alternative to the more typically-used Build periodically or Poll SCM.

Figure 5.18. Gerrit Trigger

The configuration for this plugin is minimal and focuses on the Project Type and Pattern and Branch Type and Pattern. In each pair, the type can be Plain, Path, or RegExp—pattern flavors of what to watch—and then the value (pattern) to evaluate using the type as the guide.

5.4.2.4. Post-build actions

The Git plugin for Jenkins adds Git-specific capabilities to the post-processing of the build artifacts. Specifically, the Git Publisher (shown in Figure 5.19, “Git Publisher”) offers merging and pushing actions. Check the Git Publisher checkbox to display four options.

Figure 5.19. Git Publisher

5.4.2.4.1. Push only if build succeeds

If a merge or other commit-creating action has been taken during the Jenkins build, it can be enabled to push to a remote.

5.4.2.4.2. Merge results

If prebuild merging is configured, push the merge-resultant branch to its origin (see Figure 5.20, “Merge results”).

Figure 5.20. Merge results

5.4.2.4.3. Tags

When pushing tags, each tag can be named and created if it doesn’t exist (which fails if it does already exist). Environment variables can be embedded in the tag name. Examples include the process ID,

⁸ <http://wiki.hudson-ci.org/display/HUDSON/Gerrit+Trigger>

such as `HUDSON_BUILD_$PPID`, or even a build number, if that's provided by a Jenkins plugin, such as `$HUDSON_AUTOTAG_$BUILDNUM`. Tags can be targeted to a specific remote such as `origin` or `integrationrepo`.

5.4.2.4.4. Branches

The current HEAD used in the Jenkins build of the application can be pushed to other remotes as an after-step of the build. You only need to provide the destination branch name and remote name.

Names of remotes are validated against the earlier configuration of the plugin. If the remote doesn't exist, a warning is displayed.

5.4.2.5. GitHub plugin

The GitHub plugin offers two integration points. First, it offers an optional link to the project's GitHub home page. Just enter the URL for the project (without the `tree/master` or `tree/branch` part). For example, `http://github.com/matthewmccullough/git-workshop`.

Secondly, the GitHub plugin offers per-file-changed links that are wired via the Repository browser section of a job's Source Code Management configuration (see Figure 5.21, "GitHub repository browser").

Figure 5.21. GitHub repository browser

With the `githubweb` repository browser chosen, all changed-detected files will be linked to the appropriate GitHub source-viewing web page (Figure 5.22, "GitHub repository browser").

Figure 5.22. GitHub repository browser

5.5. Build Triggers

Once you've configured your version control system, you need to tell Jenkins when to kick off a build. You set this up in the Build Triggers section.

In a freestyle build, there are three basic ways a build job can be triggered (see Figure 5.23, "There are many ways that you can configure Jenkins to start a build job"):

- Start a build job once another build job has completed
- Kick off builds at periodical intervals
- Poll the SCM for changes

Figure 5.23. There are many ways that you can configure Jenkins to start a build job

5.5.1. Triggering a Build Job Once Another Build Job Has Finished

The first option lets you set up a build that will be run whenever another build has finished. This is an easy way to set up a build pipeline. For example, you might set up an initial build job to run unit and integration tests, followed by another separate build job to run more CPU-intensive code quality metrics. You simply enter the name of the preceding build job in this field. If the build job can be triggered by several other build jobs, just list their names here, separated by commas. In this case, the build job will be triggered once any of the build jobs in the list finish.

There's a symmetrical field in the Post-build actions section of the preceding build job called (appropriately enough) “Build other projects”. This field will be automatically updated in the corresponding build jobs whenever you modify the “Build after other projects are built” field. However, unlike the “Build after other projects are built” field, this field gives you the option to trigger a build even if the build is unstable (see Figure 5.24, “Triggering another build job even if the current one is unstable”). This is useful, for example, if you want to run a code quality metrics build job even if there are unit test failures in the default build job.

Figure 5.24. Triggering another build job even if the current one is unstable

5.5.2. Scheduled Build Jobs

Another strategy is simply to trigger your build job at regular intervals. It's important to note that this isn't actually CI—it's simply scheduled builds, something you could also do, for example, with a Unix cron job. In the early days of automated builds, and even today in many shops, builds aren't run in response to changes committed to version control, but simply on a nightly basis. However, to be effective, a CI server should provide feedback much more quickly than once a day.

There are nevertheless a few cases where scheduled builds do make sense. This includes very long running build jobs, where quick feedback is less critical. For example, intensive load and performance tests which may take several hours to run, or . Sonar build jobs are an excellent way to keep tabs on code quality metrics across your projects and over time, but the Sonar server only stores one set of data per day, so running Sonar builds more frequently than this isn't useful.

For all scheduling tasks, Jenkins uses a cron-style syntax, consisting of five fields separated by white space in the following format:

MINUTE HOUR DOM MONTH DOW

with the following values possible for each field:

MINUTE

The minute within the hour (0–59)

HOURL

The hour of the day (0–23)

DOM

The day of the month (1–31)

MONTH

The month (1–12)

DOW

The day of the week (0–7), where 0 and 7 are Sunday.

There are also a few short-cuts:

- “*” represents all possible values for a field. For example, “* * * * *” means “once a minute.”
- You can define ranges using the “M–N” notation. For example “1–5” in the DOW field would mean “Monday to Friday.”
- You can use the slash notation to define skips through a range. For example, “*/5” in the MINUTE field would mean “every five minutes.”
- A comma-separated list indicates a list of valid values. For example, “15,45” in the MINUTE field would mean “at 15 and 45 minutes past every hour.”
- You can also use the shorthand values of “@yearly”, “@annually”, “@monthly”, “@weekly”, “@daily”, “@midnight”, and “@hourly”.

Typically, you'll only have one line in this field, but for more complicated scheduling setups, you may need multiple lines.

5.5.3. Polling the SCM

As you've seen, scheduled build jobs are usually not the best strategy for most CI build jobs. The value of any feedback is proportional to the speed in which you receive that feedback, and CI is no exception. That's why polling the SCM is generally a better option.

Polling involves asking the version control server at regular intervals if any source code changes have been committed. If so, Jenkins kicks off a build. Polling is usually a relatively cheap operation, so you can poll frequently to ensure that a build kicks off rapidly after changes have been committed. The more frequent the polling is, the faster the build jobs will start, and the more accurate the feedback about what change broke the build will be.

In Jenkins, SCM polling is easy to configure, and uses the same cron syntax we discussed previously.

The natural temptation for SCM polling is to poll as often as possible (for example, using “* * * * *”, or once every minute). Since Jenkins simply queries the version control system, and only kicks off a build if the source code has been modified, this approach is often reasonable for small projects. It shows its limits if there is a very large number of build jobs, as this may saturate the SCM server and the network with queries, many of them unnecessary. In this case, a more precise approach is better, where the Jenkins build job is triggered by the SCM when it receives a change. We discuss this option in Section 5.5.4, “Triggering Builds Remotely”.

If updates are frequently committed to the version control system, across many projects, this may cause many build jobs to be queued, which can in turn slow down feedback times further. You can reduce the build queue to some extent by polling less frequently, but at the cost of less precise feedback.

If you're using CVS, polling may not be a good option. When CVS checks for changes in a project, it checks each file one by one, which is a slow and tedious process. The best solution here is to migrate to a modern version control system such as Git or Subversion. The second-best solution is to use polling at very sparse intervals (for example, every 30 minutes).

5.5.4. Triggering Builds Remotely

Polling can be an effective strategy for smaller projects, but it doesn't scale particularly well—with large numbers of build jobs, it's wasteful of network resources, and there's always a small delay between the code change being committed and the build job starting. A more precise strategy is to get the SCM system to trigger the Jenkins build whenever a change is committed.

It's easy to start a Jenkins build job remotely. You simply invoke a URL of the following form:

`http://SERVER/jenkins/job/PROJECTNAME/build`

For example, if my Jenkins server was running on `http://myserver:8080/jenkins`, I could start the `gameoflife` build job by invoking the following URL using a tool like `wget` or `curl`:

```
$ wget http://myserver:8080/jenkins/job/gameoflife/build
```

The trick, then, is to get your version control server to do this whenever a change is committed. The details of how to do this are different for each version control system. In Subversion, for example, you'd need to write a post-commit hook script, which would trigger a build. You could, for example, write a Subversion hook script that parses the repository URL to extract the project name, and performs a `wget` operation on the URL of the corresponding build job:

```
JENKINS_SERVER=http://myserver:8080/jenkins
REPOS="$1"
PROJECT=<Regular Expression Processing Goes Here>❶
/usr/bin/wget $JENKINS_SERVER/job/${PROJECT}/build
```

- ❶ Use regular expression processing here to extract your project name from the Subversion repository URL.

However, this approach will only trigger one particular build, and relies on the convention that the default build job is based on the repository name in Subversion. A more flexible approach with Subversion is to use the Jenkins Subversion API directly, as shown here:

```
JENKINS_SERVER=http://myserver:8080/jenkins
REPOS="$1"
REV="$2"
UUID=`svnlook uuid $REPOS`
/usr/bin/wget \
  --header "Content-Type:text/plain;charset=UTF-8" \
  --post-data "`svnlook changed --revision $REV $REPOS`" \
  --output-document "-" \
  --timeout=2 \
  $JENKINS_SERVER/subversion/${UUID}/notifyCommit?rev=$REV
```

This would automatically start any Jenkins build jobs monitoring this Subversion repository.

If you've activated Jenkins security, things become a little more complicated. In the simplest case (where any user can do anything), you need to activate the “Trigger builds remotely” option (see Figure 5.25, “Triggering a build via a URL using a token”), and provide a special string that can be used in the URL:

`http://SERVER/jenkins/job/PROJECTNAME/build?token=DOIT`

Figure 5.25. Triggering a build via a URL using a token

This won't work if users need to be logged on to trigger a build (for example, if you're using matrix or project-based security). In this case, you'll need to provide a user name and password, as shown in the following example:

```
$ wget http://scott:tiger@myserver:8080/jenkins/job/gameoflife/build
```

or:

```
$ curl -u scott:tiger http://scott:tiger@myserver:8080/jenkins/job/gameoflife/build
```

5.5.5. Manual Build Jobs

A build doesn't have to be triggered automatically. Some build jobs should only be started manually, by human intervention. For example, you may want to set up an automated deployment to a UAT environment that should only be started on the request of the QA folks. In this case, you can simply leave the Build Triggers section empty.

5.6. Build Steps

Now, Jenkins should know where and how often to obtain the project source code. The next thing you need to explain to Jenkins is what to do with the source code. In a freestyle build, you do this by defining

build steps. Build steps are the basic building blocks for the Jenkins freestyle build process. They're how you tell Jenkins exactly how you want your project built.

A build job may have one or more steps. It may even occasionally have none. In a freestyle build, you can add as many build steps as you want to the Build section of your project configuration (see Figure 5.26, “Adding a build step to a freestyle build job”). In a basic Jenkins installation, you'll be able to add steps to invoke Maven and Ant, as well as running OS-specific shell or Windows batch commands. By installing additional plugins, you can also integrate other build tools, such as Groovy, Gradle, Grails, Jython, MSBuild, Phing, Python, Rake, and Ruby, just to name some of the more well-known tools.

In the remainder of this section, we'll delve into some of the more common types of build steps.

5.6.1. Maven Build Steps

Jenkins has excellent Maven support, and Maven build steps are easy to configure and very flexible. Just pick “Invoke top-level Maven targets” from the build step lists, pick a version of Maven to run (if you have multiple versions installed), and enter the Maven goals you want to run. Jenkins freestyle build jobs work fine with both Maven 2 and Maven 3.

Just like on the command line, you can specify as many individual goals as you want. You can also provide command-line options. A few useful Maven options in a CI context are:

`-B, --batch-mode`

This option tells Maven not to prompt for any input from the user, just use the default values, if any are required. If Maven does prompt for any input during the Jenkins build, the build will get stuck indefinitely.

`-U, --update-snapshots`

This option forces Maven to check for updated releases and snapshot dependencies on the remote repository. This makes sure you're building with the latest and greatest snapshot dependencies, and not just using older local copies which may not be in sync with the latest version of the source code.

`-Dsurefire.useFile=false`

This option forces Maven to write JUnit output to the console, rather than to text files in the target directory, as it normally would. This way, any test failure details are directly visible in the build job console output. The XML files that Jenkins needs for its test reporting will still be generated.

Figure 5.26. Adding a build step to a freestyle build job

The advanced options are also worth investigating (click on the Advanced button).

The optional **POM** field lets you override the default location of the Maven `pom.xml` file. This is the equivalent of running Maven from the command line with the `-f` or `--file` option. This is useful

for some multimodule Maven projects where the aggregate `pom.xml` file (the one containing the `<modules>` section) is located in a subdirectory rather than at the top level.

The Properties field lets you set property values that will be passed into the Maven build process, using the standard property file format illustrated here:

```
# Selenium test configuration
selenium.host=testserver.acme.com
selenium.port=8080
selenium.browser=firefox
```

These properties are passed to Maven as command-line options, as shown here:

```
$ mvn verify -Dselenium.host=testserver.acme.com ...
```

The JVM Options field lets you set any of the standard JVM options for your build job. So, if your build process is particularly memory intensive, you might add some extra heap space with the `-Xmx` option (for example, `-Xmx512m` would set the maximum heap size to 512 MB).

The final option lets you configure a private Maven repository for this build job. Normally, Maven will just use the default Maven repository (usually in the `.m2/repository` folder in the user's home directory). Occasionally, this can lead to build jobs interfering with each other, or use inconsistent snapshot versions from one build to another. To be sure that your build is run in clean laboratory conditions, you can use this option. Your build job will get its own private repository, reserved for its own exclusive use. On the downside, the first time the build job runs a build, it may take some time to download all of the Maven artifacts, and private repositories can take up a lot of space. However, it's the best way of guaranteeing that your build is run in a truly isolated environment.

5.6.2. Ant Build Steps

Freestyle build jobs work equally well with Ant. Apache Ant⁹ is a widely-used and very well-known Java build scripting tool. Indeed, a very large number of Java projects out there rely on Ant build scripts.

Ant isn't only used as a primary build scripting tool—even if your project uses Maven, you may resort to calling Ant scripts to do more specific tasks. There are Ant libraries available for many development tools and low-level tasks, such as using SSH, or working with proprietary application servers.

In its most basic form, configuring an Ant build step very is simple indeed—you just provide the version of Ant you want to use and the name of the target you want to invoke. In Figure 5.27, “Configuring an Ant build step”, for example, you're invoking an Ant script to run a JMeter test script.

Figure 5.27. Configuring an Ant build step

⁹ <http://ant.apache.org/>

As with the Maven build step, the “Advanced...” button provides you with more detailed options, such as specifying a different build script, or a build script in a different directory (the default will be `build.xml` in the root directory). You can also specify properties and JVM options, just as you can for Maven.

5.6.3. Executing a Shell or Windows Batch Command

Occasionally you may need to execute a command directly at the Operating System level. Some legacy build processes rely on OS-specific scripts, for example. In other cases, you may need to perform a low-level operation that's most easily done with an OS-level command.

You can do this in Jenkins with the `Execute Shell` (for Unix) or `Execute Windows Batch command` (for Windows). As an example, in Figure 5.28, “Configuring an Execute Shell step” we've added a step to execute the Unix `ls` command.

Figure 5.28. Configuring an Execute Shell step

The output from this build step is shown here:

```
[workspace] $ /bin/sh -xe /var/folders/.../jenkins2542160238803334344.s
+ ls -al
total 64
drwxr-xr-x  14 johnsmart  staff   476 30 Oct 15:21 .
drwxr-xr-x   9 johnsmart  staff   306 30 Oct 15:21 ..
-rw-r--r--@  1 johnsmart  staff   294 22 Sep 01:40 .checkstyle
-rw-r--r--@  1 johnsmart  staff   651 22 Sep 01:40 .classpath
-rw-r--r--@  1 johnsmart  staff   947 22 Sep 01:40 .project
drwxr-xr-x   5 johnsmart  staff   170 22 Sep 01:40 .settings
-rw-r--r--@  1 johnsmart  staff   437 22 Sep 01:40 .springBeans
drwxr-xr-x   9 johnsmart  staff   306 30 Oct 15:21 .svn
-rw-r--r--@  1 johnsmart  staff  1228 22 Sep 01:40 build.xml
-rw-r--r--@  1 johnsmart  staff    50 22 Sep 01:40 infinitest.filters
-rw-r--r--  1 johnsmart  staff  6112 30 Oct 15:21 pom.xml
drwxr-xr-x   5 johnsmart  staff   170 22 Sep 01:40 src
drwxr-xr-x   3 johnsmart  staff   102 22 Sep 01:40 target
drwxr-xr-x   5 johnsmart  staff   170 22 Sep 01:40 tools
```

You can either execute an OS-specific command (e.g., `ls`), or store a more complicated script as a file in your version control system, and execute this script. If you're executing a script, you just need to refer to the name of your script relative to the work directory.

Shell scripts are executed using the `-ex` option—the commands are printed to the console, as is the resulting output. If any of the executed commands return a nonzero value, the build will fail.

When Jenkins executes a script, it sets a number of environment variables that you can use within the script. We discuss these variable in more detail in the next section.

In fact, there are some very good reasons why you should avoid using OS-level scripts in your build jobs if you possibly can. In particular, OS scripts make your build job in the best of cases OS-specific,

and at worst dependant on the precise machine configuration. A more portable alternative to executing OS scripts include writing an equivalent script in a more portable scripting language, such as Groovy or Gant.

5.6.4. Using Jenkins Environment Variables in Your Builds

One useful trick that can be used in virtually any build step is to obtain information from Jenkins about the current build job. In fact, when Jenkins starts a build step, it makes the following environment variables available to the build script:

`BUILD_NUMBER`

The current build number, such as “153”.

`BUILD_ID`

A timestamp for the current build id, in the form YYYY-MM-DD_hh-mm-ss.

`JOB_NAME`

The name of the job, such as game-of-life.

`BUILD_TAG`

A convenient way to identify the current build job, in the form of `jenkins-${JOB_NAME}-${BUILD_NUMBER}` (e.g., `jenkins-game-of-life-2010-10-30_23-59-59`).

`EXECUTOR_NUMBER`

A number identifying the executor running this build among the executors of the same machine. This is the number you see in the “build executor status”, except that the number starts from 0, not 1.

`NODE_NAME`

The name of the slave if the build is running on a slave, or “” if the build is running on the master.

`NODE_LABELS`

The list of labels associated with the node that this build is running on.

`JAVA_HOME`

If your job is configured to use a specific JDK, this variable is set to the `JAVA_HOME` of the specified JDK. When this variable is set, `PATH` is also updated to contain `$JAVA_HOME/bin`.

`WORKSPACE`

The absolute path of the workspace.

`HUDSON_URL`

The full URL of the Jenkins server, for example `http://ci.acme.com:8080/jenkins/`.

`JOB_URL`

The full URL for this build job, for example `http://ci.acme.com:8080/jenkins/game-of-life`.

BUILD_URL

The full URL for this build, for example `http://ci.acme.com:8080/jenkins/game-of-life/20`.

SVN_REVISION

For Subversion-based projects, this variable contains the current revision number.

CVS_BRANCH

For CVS-based projects, this variable contains the branch of the module. If CVS is configured to check out the trunk, this environment variable won't be set.

These variables are easy to access. In an Ant script, you can access them using the `<property>` tag as shown here:

```
<target name="printinfo">
  <property environment="env" />
  <echo message="${env.BUILD_TAG}" />
</target>
```

In Maven, you can access the variables either in the same way (using the “env.” prefix), or directly using Jenkins environment variables. For example, in the following `pom.xml` file, the project URL will point to the Jenkins build job that ran the `mvn site` build:

```
<project...>
...
<groupId>com.wakaleo.gameoflife</groupId>
<artifactId>gameoflife-core</artifactId>
<version>0.0.55-SNAPSHOT</version>
<name>gameoflife-core</name>
<url>${JOB_URL}</url>
```

Alternatively, if you're building a web application, you can also use the `maven-war-plugin` to insert the build job number into the web application manifest, e.g.:

```
<project>
...
<build>
...
<plugins>
  <plugin>
    <artifactId>maven-war-plugin</artifactId>
    <configuration>
      <manifest>
        <addDefaultImplementationEntries>true</addDefaultImplementationEntries>
      </manifest>
      <archive>
        <manifestEntries>
          <Specification-Title>Continuous Integration with Hudson (Content)</Specification-Title>
          <Specification-Version>0.0.4-SNAPSHOT</Specification-Version>
          <Implementation-Version>${BUILD_TAG}</Implementation-Version>
        </manifestEntries>
```

```

        </archive>
    </configuration>
</plugin>
...
</plugins>
</build>
...
</project>

```

This will produce a `MANIFEST.MF` file along the following lines:

```

Manifest-Version: 1.0
Archiver-Version: Plexus Archiver
Created-By: Apache Maven
Built-By: johnsmart
Build-Jdk: 1.6.0_22
Jenkins-Build-Number: 63
Jenkins-Project: game-of-life
Jenkins-Version: 1.382
Implementation-Version: jenkins-game-of-life-63
Specification-Title: gameoflife-web
Specification-Version: 0.0.55-SNAPSHOT

```

And in a Groovy script, the environment variables can be obtained via the `System.getenv()` method:

```

def env = System.getenv()
env.each {
    println it
}

```

or:

```

def env = System.getenv()
println env['BUILD_NUMBER']

```

5.6.5. Running Groovy Scripts

Groovy isn't only a popular JVM dynamic language, it's also a convenient language for low-level scripting. The Jenkins Groovy Plugin¹⁰ lets you run arbitrary Groovy commands or invoke Groovy scripts as part of your build process.

Once you've installed the Groovy plugin in the usual way, you need to add a reference to your Groovy installation in the system configuration page (see Figure 5.29, “Adding a Groovy installation to Jenkins”).

Figure 5.29. Adding a Groovy installation to Jenkins

¹⁰ <http://wiki.jenkins-ci.org//display/HUDSON/Groovy+Plugin>

Now, you can add some Groovy scripting to your build job. When you click on “Add build step”, you’ll see two new entries in the drop-down menu: “Execute Groovy script” and “Execute system Groovy script”. The first option is generally what you want—this will simply execute a Groovy script in a separate JVM, as if you were invoking Groovy from the command line. The second option runs Groovy commands within Jenkins’s own JVM, with full access to Jenkins’s internals, and is mainly used to manipulate the Jenkins build jobs or build process itself. This is a more advanced topic that we’ll discuss later on in the book.

A Groovy build step can take one of two forms. For simple cases, you can just add a small snippet of Groovy, as shown in Figure 5.30, “Running Groovy commands as part of a build job”. For more involved or complicated cases, you’d probably write a Groovy script and place it under version control. Once your script is safely in your SCM, you can run it by selecting the “Groovy script file” option and providing the path to your script (relative to your build job workspace).

Figure 5.30. Running Groovy commands as part of a build job

In Figure 5.31, “Running Groovy scripts as part of a build job”, you can see a slightly more complicated example. Here we’re running a Groovy script called `run-fitness-tests.groovy`, which can be found in the `scripts` directory. This script takes the test suites to be executed as its parameters—you provide these in the Script parameters field. If you want to provide any options for Groovy itself, you can put these in the Groovy Parameters field. Alternatively, you can also provide command-line properties in the Properties field—this is simply a more convenient way of using the `-D` command-line option to pass property values to the Groovy script.

Figure 5.31. Running Groovy scripts as part of a build job

5.6.6. Building Projects in Other Languages

Jenkins is a flexible tool, and it can be used for much more than just Java and Groovy. For example, Jenkins also works well with Grails, .Net, Ruby, Python, and PHP, just to name a few. When using other languages, you generally need to install a plugin to support your favorite language, which will add a new build step type for this language. We’ll look at some examples in Section 5.10, “Using Jenkins with Other Languages”.

5.7. Post-Build Actions

Once the build is completed, there are still a few things you need to look after. You might want to archive some of the generated artifacts, to report on test results, and to notify people about the results. In this section, you look at some of the more common tasks you need to configure after the build is done.

5.7.1. Reporting on Test Results

One of the most obvious requirements of a build job is reporting test results such as whether there are any test failures, but also how many tests were executed, how long they took to execute, and so on. In the Java world, JUnit is the most commonly-used testing library around. The JUnit XML format for test results is widely used and understood by other tools as well.

Jenkins provides great support for test reporting. In a freestyle build job, you need to tick the “Publish JUnit test result report” option, and provide a path to your JUnit report files (see Figure 5.32, “Reporting on test results”). You can use a wildcard expression (such as `**/target/surefire-reports/*.xml` in a Maven project) to include JUnit reports from a number of different directories—Jenkins will aggregate the results into a single report.

Figure 5.32. Reporting on test results

We look at automated tests in much more detail in Chapter 6, Automated Testing.

5.7.2. Archiving Build Results

With a few exceptions, the principal goal of a build job is generally to build something. In Jenkins, we call this something an artifact. An artifact might be a binary executable (a JAR or WAR file for a Java project, for example), or some other related deliverable, such as documentation or source code. A build job can store any number of different artifacts, keeping only the latest copy or every artifact ever built.

Configuring Jenkins to store your artifacts is easy—just tick the “Archive the artifacts” checkbox in the Post-build Actions, and specify which artifacts you want to store (see Figure 5.33, “Configuring build artifacts”).

Figure 5.33. Configuring build artifacts

In the “Files to archive” field, you can provide the full paths of the files you want to archive (relative to the job workspace), or, use Ant-like wild cards (e.g., `**/*.jar`, for all the JAR files, anywhere in the workspace). One advantage of using wild cards is that it makes your build less dependent on your version control set up. For example, if you’re using Subversion (see Section 5.4, “Configuring Source Code Management”), Jenkins will check out your project either directly in your workspace, or into a subdirectory, depending on how you set it up. If you use a wild card expression like `**/target/*.war`, Jenkins will find the file no matter what directory the project is located in.

As usual, the Advanced button give access to a few extra options. If you’re using wild cards to find your artifacts, you might need to exclude certain directories from the search. You can do this by filling in the Excludes field. You enter a pattern to match any files that you don’t want to archive, even if they’d normally be included by the “Files to archive” field.

Archived artifacts can take a lot of disk space, especially if builds are frequent. For this reason, you may want to only keep the last successful one. To do this, just tick the “Discard all but the last successful/stable artifact” option. Jenkins will keep artifacts from the last stable build (if there were any). It will also keep the artifacts of the latest unstable build following the stable build (if any), and also from the last failed build that happened.

Archived build artifacts appear on the build results page (see Figure 5.34, “Build artifacts are displayed on the build results page and on the build job home page”). The most recent build artifacts are also displayed on the build job home page.

Figure 5.34. Build artifacts are displayed on the build results page and on the build job home page

You can also use permanent URLs to access the most recent build artifacts. This is a great way to reuse the latest artifacts from your builds, either in other Jenkins build jobs or in external scripts, for example. Three URLs are available: last stable build, last successful build, and last completed build.

Before looking at the URLs, we should discuss the concept of stable and successful builds.

A build is successful when the compilation reported no errors.

A build is considered stable if it was built successfully, and no publisher reports it as unstable. For example, depending on your project configuration, unit test failures, insufficient code coverage, or other code quality metrics issues could cause a build to be marked as unstable. So, a stable build is always successful, but the opposite isn't necessarily true—a build can be successful without being stable.

A completed build is simply a build that has finished, no matter what its result. Note that the archiving step will take place no matter what the outcome of the build was.

The format of the artifact URLs is intuitive, and takes the following form:

Latest stable build

```
<server-url>/job/<build-job>/lastStableBuild/artifact/<path-to-  
artifact>
```

Latest successful build

```
<server-url>/job/<build-job>/lastSuccessfulBuild/artifact/<path-to-  
artifact>
```

Latest completed build

```
<server-url>/job/<build-job>/lastCompletedBuild/artifact/<path-to-  
artifact>
```

This is best illustrated by some examples. Suppose your Jenkins server is running on `http://myserver:8080`, your build job is called `game-of-life`, and you're storing a file called

gameoflife.war, which is in the target directory of your workspace. The URLs for this artifact would be the following:

Latest stable build

```
http://myserver:8080/job/gameoflife/lastStableBuild/artifact/target/
gameoflife.war
```

Latest successful build

```
http://myserver:8080/job/gameoflife/lastSuccessfulBuild/artifact/
target/gameoflife.war
```

Latest completed build

```
http://myserver:8080/job/gameoflife/lastCompletedBuild/artifact/
target/gameoflife.war
```

Artifacts don't have to be just executable binaries. Imagine, for example, that your build process involves automatically deploying each build to a test server. For convenience, you want to keep a copy of the exact source code associated with each deployed WAR file. One way to do this would be to generate the source code associated with a build, and archive both this file and the WAR file. We could do this by generating a JAR file containing the application source code (for example, by using the Maven Source Plugin for a Maven project), and then including this in the list of artifacts to store (see Figure 5.35, “Archiving source code and a binary package”).

Figure 5.35. Archiving source code and a binary package

Of course, this example is a tad academic: it would probably be simpler just to use the revision number for this build (which is displayed on the build result page) to retrieve the source code from your version control system. But you get the idea.

Note that if you're using an Enterprise Repository Manager such as Nexus or Artifactory to store your binary artifacts, you may not need to keep them on the Jenkins server. You may prefer simply to automatically deploy your artifacts to your Enterprise Repository Manager as part of the build job, and retrieve them from there when required.

5.7.3. Notifications

The point of a CI server is to let people know when a build breaks. In Jenkins, this comes under the heading of Notification.

Out of the box, Jenkins provides support for email notification. You can activate this by ticking the “E-mail Notification” checkbox in the Post-build Actions (see Figure 5.36, “Email notification”). Then, enter the email addresses of the team members who will need to know when the build breaks. When the build does break, Jenkins will send a friendly email message to the users in this list containing a link to the broken build.

Figure 5.36. Email notification

You can also opt to send a separate email to the user whose commit (presumably) broke the build. For this to work, you need to have activated Security on your Jenkins server (see Chapter 7, *Securing Jenkins*).

Normally, Jenkins will send out an email notification whenever a build fails (for example, because of a compilation error). It will also send out a notification when the build becomes unstable for the first time (for example, if there are unit test failures). Unless you configure it to do so, Jenkins won't send emails for every unstable build, but only for the first one.

Finally, Jenkins will send a message when a previously failing or unstable build succeeds, to let everyone know that the problem has been resolved.

5.7.4. Building Other Projects

You can also start other build jobs in the Post-build Actions, using the “Build other projects” option. This is useful if you want to organize your build process in several, smaller steps, rather than one long build job. Just list the projects you want to start after this one. Normally, these projects will only be triggered if the build was stable, but you can optionally trigger another build job even if the current build is unstable. This might be useful, for example, if you wanted to run a code quality metrics reporting build job after a project’s main build job, even if there are test failures in the main build.

5.8. Running Your New Build Job

Now all you need to do is save your new build job. You can then trigger the first build manually, or just wait for it to kick off by itself. Once the build is finished, you can click on the build number to see the results of your work.

5.9. Working with Maven Build Jobs

In this section, we'll have a look at the other most commonly used build job: Maven 2/3 build jobs.

Maven build jobs are specifically adapted to Maven 2 and Maven 3 builds. Creating a Maven build job requires considerably less work than configuring the equivalent freestyle build job. Maven build jobs support advanced Maven-related features such as incremental builds on multimodule projects and triggering builds from changes in snapshot dependencies, making configuration and reporting much simpler.

However, there's a catch: Maven 2/3 build jobs are less flexible than freestyle build jobs, and don't support multiple build steps within the same build job. Some users also report that large Maven projects tend to run more slowly and use more memory when configured as Maven build jobs rather than as freestyle ones.

In this section, we'll investigate how to configure Maven 2/3 builds, when you can use them, as well as their advantages and limitations.

To create a new Maven build job, just choose the “Build a maven2/3 project” option in the New Job page (see Figure 5.37, “Creating a new Maven build job”).

Figure 5.37. Creating a new Maven build job

5.9.1. Building Whenever a SNAPSHOT Dependency Is Built

At first glance, the Maven 2/3 build job configuration screen is very similar to the one you saw for freestyle builds in the previous section. The first difference you may notice is in the Build Triggers section. In this section, an extra option is available: “Build whenever a SNAPSHOT dependency is built”. If you select this option, Jenkins will examine your `pom.xml` file (or files) to see if any SNAPSHOT dependencies are being built by other build jobs. If any other build jobs update a SNAPSHOT dependency that your project uses, Jenkins will build your project as well.

Typically in Maven, SNAPSHOT dependencies are used to share the latest bleeding-edge version of a library with other projects within the same team. Since they're by definition unstable, it isn't recommended practice to rely on SNAPSHOT dependencies from other teams or from external sources.

For example, imagine that you're working on a new game-of-life web application. You are using Maven for this project, so you can use a Maven build job in Jenkins. Your team is also working on a reusable library called cooltools. Since these two projects are being developed by the same team, you're using some of the latest cooltools features in the game-of-life web application. So, you have a SNAPSHOT dependency in the `<dependencies>` section of your game-of-life `pom.xml` file:

```
<dependencies>
  <dependency>
    <groupId>com.acme.common</groupId>
    <artifactId>cooltools</artifactId>
    <version>0.0.1-SNAPSHOT</version>
    <scope>test</scope>
  </dependency>
  ...
</dependencies>
```

On your Jenkins server, you've set up Maven build jobs for both the cooltools and the game-of-life applications. Since your game-of-life project needs the latest cooltools SNAPSHOT version, you tick the “Build whenever a SNAPSHOT dependency is built” option. This way, whenever the cooltools project is rebuilt, the game-of-life project will automatically be rebuilt as well.

5.9.2. Configuring the Maven Build

The next area where you'll notice a change is in the Build section. In a Maven build job, the build section is entirely devoted to running a single Maven goal (see Figure 5.38, “Specifying the Maven goals”). In

this section, you specify the version of Maven you want to execute, the location of the `pom.xml` file, and the Maven goal (or goals) to invoke. You can also add any command-line options you need here.

Figure 5.38. Specifying the Maven goals

In many cases, this is all you need to get your Maven build job configured. However, if you click on the “Advanced...” button, you can take your pick of some more advanced features (Figure 5.39, “Maven build jobs—advanced options”).

Figure 5.39. Maven build jobs—advanced options

The Incremental Build option comes in very handy for large, multimodule Maven builds. If you tick this option, when a change is made to one of the project modules, Jenkins will only rebuild that module and any modules that use the changed module. It performs this magic by using some new Maven features introduced in Maven 2.1 (so it won’t work if you’re using Maven 2.0.x). Jenkins detects which modules have been changed, and then uses the `-pl` (`--project-list`) option to build only the updated modules, and the `-amd` (`--also-make-dependents`) option to build the modules that use the updated modules. If nothing has been changed in the source code, all of the modules are built.

By default, Jenkins will archive all of the artifacts generated by a Maven build job. This can come in handy at times, but it can also be very expensive in disk storage. If you want to turn off this option, just tick the “Disable automatic artifact archiving” option. Alternatively, you can always limit the artifacts stored by using the “Discard Old Builds” option at the top of the configuration page.

The “Build modules in parallel” option tells Jenkins to run each individual module in parallel as a separate build. In theory, this could speed up your builds quite a bit. In practice, it will only really work if your modules are totally independent (that is, you aren’t using aggregation), which is rarely the case. If you think building your modules in parallel could really speed up your multimodule project, you may want to try a freestyle build with Maven 3 and its new parallel build feature.

Another useful option is “Use [a] private Maven repository”. Normally, when Jenkins runs Maven, it will behave in exactly the same way as Maven on the command line: it will store artifacts in, and retrieve artifacts from the local Maven repository (found in `~/.m2/repository` if you’ve## reconfigured it in the `settings.xml` file). This is efficient in terms of disk space, but not always ideal for CI builds. Indeed, if several build jobs are working on and with the same snapshot artifacts, the builds may end up interfering with each other.

When this option is checked, Jenkins will tell Maven to use `$WORKSPACE/.repository` as the local Maven repository. This means each job will get its own isolated Maven repository just for itself. It fixes the above problems, at the expense of additional disk space consumption.

Another way of addressing this problem is to override the default repository location by using the `maven.repo.local` property, as shown here:

```
$ mvn install -Dmaven.repo.local=~/.m2/staging-repository
```

This approach has the advantage of being able to share a repository across several build jobs, which is useful if you need to do a series of related builds. It will also work with freestyle jobs.

5.9.3. Post-Build Actions

The Post-Build actions in a Maven build job are considerably simpler to configure than in a freestyle job. This is simply because, since this is a Maven build, Jenkins knows where to look for a lot of the build output. Artifacts, test reports, Javadoc, and so forth are all generated in standard directories, which means you don't have to tell Jenkins where to find things. So, Jenkins will find, and report on, JUnit test results automatically, for example. Later on in the book, we'll see how the Maven projects also simplify the configuration of many code quality metrics tools and reports.

Most of the other Post-build Actions are similar to those we saw in the freestyle build job.

5.9.4. Deploying to an Enterprise Repository Manager

One extra option that does appear in the Maven build jobs is the ability to deploy your artifacts to a Maven repository (see Figure 5.40, “Deploying artifacts to a Maven repository”). An Enterprise Repository Manager is a server that acts as both a proxy/cache for public Maven artifacts, and as a central storage server for your own internal artifacts. Open Source Enterprise Repository Managers like Nexus (from Sonatype) and Artifactory (from JFrog) provide powerful maintenance and administration features that make configuring and maintaining your Maven repositories a lot simpler. Both these products have commercial versions, with additional features aimed at more sophisticated or high-end build infrastructures.

The advantage of getting Jenkins to deploy your artifacts (as opposed to simply running `mvn deploy`) is that, if you have a multimodule Maven build, the artifacts will only be deployed once the entire build has finished successfully. For example, suppose you have a multimodule Maven project with five modules. If you run `mvn deploy`, and the build fails after three modules, the first two modules will have been deployed to your repository, but not the last three, which leaves your repository in an instable state. Getting Jenkins to do the deploy ensures that the artifacts are only deployed as a group once the build has successfully finished.

Figure 5.40. Deploying artifacts to a Maven repository

To do this, just tick the “Deploy artifacts to Maven repository” option in the “Post-Build actions”. You'll need to specify the URL of the repository you want to deploy to. This needs to be the full URL to the repository (e.g., `http://nexus.acme.com/nexus/content/repositories/snapshots`, and not just `http://nexus.acme.com/nexus`)

Most repositories need you to authenticate before letting you deploy artifacts to them. The standard Maven way to do this is to place a `<server>` entry in your local `settings.xml` file, as shown here:

```
<settings...>
  <servers>
    <server>
      <id>nexus-snapshots</id>
      <username>scott</username>
      <password>tiger</password>
    </server>
    <server>
      <id>nexus-releases</id>
      <username>scott</username>
      <password>tiger</password>
    </server>
  </servers>
</settings>
```

For the more security-minded, you can also encrypt these passwords if required.

Then, enter the corresponding ID value in the Repository ID field in Jenkins. Jenkins will then be able to look up the right username and password, and deploy your artifacts. Once the build is finished, your artifacts should be available in your Maven Enterprise Repository (see Figure 5.41, “After deployment the artifact should be available on your Enterprise Repository Manager”).

Figure 5.41. After deployment the artifact should be available on your Enterprise Repository Manager

Using this option, you don’t always have to deploy straight away—you can always come back and deploy the artifacts from a previous build later. Just click on the “Redeploy Artifacts” menu on the left and specify the repository URL you want to deploy your artifact to (see Figure 5.42, “Redeploying an artifact”). As in the previous example, the Advanced button lets you provide the ID for the `<server>` entry in your local `settings.xml` file. As we’ll see later on in the book, you can also use this deployment as part of a build promotion process, configuring an automatic deployment to a different repository when certain quality metrics have been satisfied, for example.

Figure 5.42. Redeploying an artifact

This approach will work fine for any Enterprise Repository manager. However, if you’re using Artifactory, you may prefer to install the Jenkins Artifactory Plugin¹¹ which provides tighter two-way integration with the Artifactory Enterprise Repository Manager. It works by sending additional information to the Artifactory server during the deployment, allowing the server to refer back to the precise build that generated a given artifact. Once you’ve installed the plugin, you can activate it in your Maven build job by ticking the “Deploy artifacts to Artifactory” option in the Post-build Actions. Then, you choose what repositories your project should deploy to from a list of repositories on the server, along

¹¹ <http://wiki.jenkins-ci.org/display/JENKINS/Artifactory+Plugin>

with the username and password required to perform the deployment (see Figure 5.43, “Deploying to Artifactory from Jenkins”).

Figure 5.43. Deploying to Artifactory from Jenkins

Your build job will now automatically deploy to Artifactory. In addition, a link to the artifact on the server will now be displayed on the build job home and build results pages (see Figure 5.44, “Jenkins displays a link to the corresponding Artifactory repository”).

Figure 5.44. Jenkins displays a link to the corresponding Artifactory repository

This link takes you to a page on the Artifactory server containing the deployed artifact (see Figure 5.45, “Viewing the deployed artifact in Artifactory”). From this page, there's also a link that takes you back to the build that built the artifact.

Figure 5.45. Viewing the deployed artifact in Artifactory

5.9.5. Deploying to Commercial Enterprise Repository Managers

An Enterprise Repository Manager is an essential part of any Maven-based software development infrastructure. It also plays a key role for non-Maven projects using tools like Ivy and Gradle, both of which rely on standard Maven repositories.

Both of the principal Enterprise Repository Managers, Nexus and Artifactory, offer professional versions which come with extra integration features with Jenkins. Later on in the book, we'll discuss how you can use advanced features such as Nexus Pro's staging and release management to implement sophisticated build promotion strategies. On the deployment side of things, the commercial edition of Artifactory (Artifactory Pro Power Pack) extends the two-way integration you saw earlier. When you view an artifact in the repository browser, a “Builds” tab displays details about the Jenkins build that created the artifact, and a link to the Jenkins build page (see Figure 5.46, “Viewing the deployed artifact and the corresponding Jenkins build in Artifactory”). Artifactory also keeps track of the dependencies that were used in the Jenkins build, and will warn you if you try to delete them from the repository.

Figure 5.46. Viewing the deployed artifact and the corresponding Jenkins build in Artifactory

5.9.6. Managing Modules

When using Maven, it's common to split a project into several modules. Maven build jobs have an intrinsic understanding of multimodule projects, and add a Modules menu item that lets you display the structure of the project at a glance (see Figure 5.47, “Managing modules in a Maven build job”).

Figure 5.47. Managing modules in a Maven build job

Clicking on any of the modules will take you to the build page for that module. From here, you can view the detailed build results for each module, trigger a build of that module in isolation, and if necessary fine tune the configuration of individual module, overriding the configuration of the overall project.

5.9.7. Extra Build Steps in Your Maven Build Jobs

By default, the Maven build job only allows for a single Maven goal. There are times when this is a little limiting, and you would like to add some extra steps before or after the main build. You can do this with the Jenkins M2 Extra Steps Plugin. This plugin lets you add normal build steps before and after the main Maven goal, giving you the flexibility of a freestyle build while still having the convenience of the Maven build job configuration.

Install this plugin and go to the Build Environment section of your build job. Tick the “Configure Extra M2 Build Steps” option. You should now be able to add build steps that will be executed before and/or after your main Maven goal is executed (see Figure 5.48, “Configuring extra Maven build steps”).

Figure 5.48. Configuring extra Maven build steps

5.10. Using Jenkins with Other Languages

As we mentioned earlier, Jenkins provides excellent support for other languages. In this section, we'll look at how to use Jenkins with a few of the more common ones.

5.10.1. Building Projects with Grails

Grails is an open source dynamic web application framework built on Groovy and many well-established open source Java frameworks such as Spring and Hibernate.

Jenkins provides excellent support for Grails builds. First, you need to install the Jenkins Grails plugin¹². Once you've installed this and restarted Jenkins, you will need to provide at least one version of Grails for Jenkins to use in the Grails Builder section of the Configure System screen (see Figure 5.49, “Adding a Grails installation to Jenkins”).

Figure 5.49. Adding a Grails installation to Jenkins

Now, you can set up a freestyle build job to build your Grails project. The Grails plugin adds the “Build with Grails” build step, which you can use to build your Grails application (see Figure 5.50, “Configuring

¹² <http://wiki.jenkins-ci.org/display/HUDSON/Grails+Plugin>

a Grails build step”). Here, you provide the Grails target, or targets, you want to execute. Unlike the command line, you can execute several targets in the same command. However, if you need to pass any arguments to a particular target, you should enclose the target and its arguments in double quotes. In Figure 5.50, “Configuring a Grails build step”, for example, you run `grails clean`, followed by `grails test-app -unit -non-interactive`. To get this to work properly, we enclose the options of the second command in quotes, which gives us `grails clean "test-app -unit -non-interactive"`.

Figure 5.50. Configuring a Grails build step

The Grails build step takes many optional parameters. For example, Grails is finicky about versions—if your project was created by an older version, Grails will ask you to upgrade it. To be on the safe side, for example, you may want to tick the Force Upgrade checkbox, which makes sure that runs `grails upgrade --non-interactive` before it runs the main targets.

You can also specify the server port (useful if you're executing web tests), and any other properties you want to pass to the build.

5.10.2. Building Projects with Gradle

Contributed by Rene Groeschke

In comparison to the build tool veterans Ant and Maven, Gradle¹³ is a relatively new open source build tool for the JVM. Build scripts for Gradle are written in a Domain Specific Language (DSL) based on Groovy. Gradle implements convention over configuration, allows direct access to Ant tasks, and uses Maven-like declarative dependency management. The concise nature of Groovy scripting lets you write very expressive build scripts with very little code, albeit at the cost of losing the IDE support that exists for established tools like Ant and Maven.

There are two different ways to run your Gradle builds with Jenkins. You can either use the Gradle plugin for Jenkins or the Gradle wrapper functionality.

5.10.2.1. The Gradle plugin for Jenkins

You can install the Gradle plugin in the usual way—just go to the Manage Plugins screen and select the Jenkins Gradle plugin. Click Install and restart your Jenkins instance.

Once Jenkins has restarted, you'll need to configure your new Gradle plugin. You should now find a new Gradle section in your Configure System screen. Here you'll need to add the Gradle installation you want to use. The process is similar to that used for the other tool installations. First, click the Add Gradle button to add a new Gradle installation, and enter an appropriate name (see Figure 5.51, “Configuring the Gradle plugin”). If Gradle has already been installed on your build server, you can point to the local

¹³ <http://gradle.org>

Gradle home directory. Alternatively, you can use the “Install automatically” feature to download a Gradle installation, in the form of a ZIP or GZipped TAR file, directly from a URL. You can use a public URL (see <http://gradle.org/downloads.html>), or may prefer to make these installations available on a local server instead.

Figure 5.51. Configuring the Gradle plugin

You typically use freestyle build jobs to configure your Gradle builds. When you add a build step to a freestyle build job, you'll now have a new option called “Invoke Gradle script”, which lets you add Gradle specific settings to your build job.

As an example, here's a very simple Gradle build script. It's a simple Java project that uses a Maven directory structure and a Maven repository manager. There's a customizable task, called `uploadArchives`, to deploy the generated archive to the local Enterprise repository manager:

```
apply plugin: 'java'
apply plugin: 'maven'

version = '1.0-SNAPSHOT'
group = "org.acme"

repositories {
    mavenCentral()
    mavenRepo url: 'http://build.server/nexus/content/repositories/public'
}

dependencies {
    testCompile "junit:junit:4.8.2"
}

uploadArchives {
    repositories.mavenDeployer {
        configuration = configurations.archives
        repository(url: "http://build.server/nexus/content/repositories/snapshots") {
            authentication(userName: "admin", password: "password")
        }
    }
}
```

In Figure 5.52, “Setting up a Gradle build job”, you use the just configured “Gradle-0.9RC2” instance to run this Gradle build. In this case, you want to run the JUnit tests and upload the build artifacts to our local Maven repository. Furthermore, you configure your job to collect the test results from `**/build/test-results`, the default directory for storing test results in Gradle.

5.10.2.2. Incremental builds

While running a Gradle build job with unchanged sources, Gradle runs its builds incrementally. If the output of a Gradle task is still available and the sources haven't changed since the last build, Gradle

is able to skip the task execution and marks the according task as up-to-date. This incremental build feature can decrease the duration of a running build job considerably.

If Gradle evaluates the test task as up-to-date, even the execution of your unit tests is skipped. This can cause problems when running your Gradle build with Jenkins. In our sample build job above you configured a post build action to publish the JUnit reports of our build. If the test task is skipped by Gradle, the Jenkins job will be marked as failed with the following message:

Test reports were found but none of them are new. Did tests run?

You can easily fix this by invalidating the output and forcing a re-execution of your tests by adding the following snippet to your Gradle file:

```
test {  
    outputs.upToDateWhen { false }  
}
```

Figure 5.52. Setting up a Gradle build job

After adding the snippet above to your build file, your job console output should look like the one in Figure 5.53, “Incremental Gradle job”.

Figure 5.53. Incremental Gradle job

As you can see, all of the tasks except test and uploadArchives have been marked as up-to-date and not executed.

5.10.3. Building Projects with Visual Studio MSBuild

Jenkins is a Java application, but it also provides excellent support for .NET projects.

To build .NET projects in Jenkins, you need to install the MSBuild plugin¹⁴.

You may also want to install the MSTest plugin¹⁵ and the NUnit plugin¹⁶, to display your test results.

Once you've installed the .NET plugins and restarted Jenkins, you need to configure your .NET build tools. Go to the Configure System page and specify the path of the MSBuild executable (see Figure 5.54, “Configuring .NET build tools in Jenkins”).

Figure 5.54. Configuring .NET build tools in Jenkins

¹⁴ <http://wiki.jenkins-ci.org/display/HUDSON/MSBuild+Plugin>

¹⁵ <http://wiki.jenkins-ci.org/display/HUDSON/MSTest+Plugin>

¹⁶ <http://wiki.jenkins-ci.org/display/HUDSON/NUnit+Plugin>

Once you have this set up, you can return to your freestyle project and add your .NET build step configuration.

Go to the Build section and choose “Build a Visual project or solution using MSBuild” option in the Add Build Step menu. Then, enter the path to your MSBuild build script (a `.proj` or `.sln` file), along with any command-line options your build requires (see Figure 5.55, “A build step using MSBuild”).

Figure 5.55. A build step using MSBuild

5.10.4. Building Projects with NAnt

Another way to build your .NET projects is to use NAnt. NAnt is a .NET version of the Ant build scripting tool widely used in the Java world. NAnt build scripts are XML files (typically with a `.build` extension), with a very similar format to Ant build scripts.

To build with NAnt in Jenkins, you need to install the Jenkins NAnt plugin¹⁷. Once you've installed the plugin and restarted Jenkins, go to the Configure System page and specify the NAnt installation directory in the NAnt Builders section (see Figure 5.54, “Configuring .NET build tools in Jenkins”).

Now, go to the Build section of your freestyle project and choose “Execute NAnt build” (see Figure 5.56, “A build step using NAnt”). Here you specify your build script and the target you want to invoke. If you click on the “Advanced...” option, you can also set property values to be passed into the NAnt script.

Figure 5.56. A build step using NAnt

5.10.5. Building Projects with Ruby and Ruby on Rails

Jenkins makes an excellent choice when it comes to integrating CI into your Ruby and Ruby on Rails projects. The Rake Plugin lets you add Rake build steps to your build jobs. You can also use the Ruby Plugin which lets you run Ruby scripts directly in your build job. Finally, the Ruby Metrics Plugin provides support for Ruby code quality metrics tools such as RCov, Rails stats, and Flog.

Another invaluable tool in this area is `CI:Reporter`. This library is an add-on to `Test::Unit`, `RSpec`, and `Cucumber` that generates JUnit-compatible XML reports for your tests. As you'll see, JUnit-compatible test results can be used directly by Jenkins to report on your test results. You'd install `CI:Reporter` using Gem as illustrated here:

```
$ sudo gem install ci_reporter
Successfully installed ci_reporter-1.6.4
1 gem installed
```

Next, you'll need to set this up in your Rakefile by adding the following:

¹⁷ <http://wiki.jenkins-ci.org/display/HUDSON/NAnt+Plugin>

```
require 'rubygems'
gem 'ci_reporter'
require 'ci/reporter/rake/test_unit' # use this if you're using Test::Unit
```

In Chapter 9, Code Quality, we discuss integrating code quality metrics into your Jenkins builds. Jenkins also provides support for code coverage metrics in Ruby. The Ruby Metrics Plugin supports code coverage metrics using **rcov** as well as general code statistics with **Rails stats**. To install the **rcov-plugin**, you'll first need to run something along the following lines:

```
$ ./script/plugin install http://svn.codahale.com/rails_rcov
```

Once this is set up, you'll be able to display your test results and test result trends in Jenkins.

Finally, you can configure a Rake build simply by using a Rake build step, as illustrated in Figure 5.57, “A build step using Rake”.

Figure 5.57. A build step using Rake

You also need to configure Jenkins to report on the test and quality metrics results. You can do this by activating the “Publish JUnit test result report”, “Publish Rails stats report”, and “Public Rcov report” options (see Figure 5.58, “Publishing code quality metrics for Ruby and Rails”). The JUnit XML reports will be found in the `results` directory (enter `results/*.xml` in the “Test report XMLs” field), and the Rcov data in the `coverage/units` directory.

Figure 5.58. Publishing code quality metrics for Ruby and Rails

5.11. Conclusion

In this chapter we've covered the basics of creating new build jobs for the most common cases you're likely to encounter. Later on in the book, we'll build on these foundations to discuss more advanced options such as parameterized builds, matrix builds, and build promotion strategies.

Chapter 6. Automated Testing

6.1. Introduction

If you aren't using automated tests with your CI setup, you're really missing out on something big. Believe me—CI without automated tests is really just a small improvement over automatically scheduled builds. Now don't get me wrong, if you're coming from nothing, that's already a great step forward—but you can do much better. In short, if you're using Jenkins without any automated tests, you're not getting anywhere near as much value out of your CI infrastructure as you should.

One of the basic principles of CI is that a build should be verifiable. You have to be able to objectively determine whether a particular build is ready to proceed to the next stage of the build process. The most convenient way to do this is to use automated tests. Without proper automated testing, you find yourself having to retain many build artifacts and test them by hand, which is hardly in the spirit of CI.

There are many ways you can integrate automated tests into your application. One of the most efficient ways to write high quality tests is to write them first, using techniques such as Test-Driven Development (TDD) or Behavior-Driven Development (BDD). In these approaches, commonly used in many Agile projects, the aim of your unit tests is to both clarify your understanding of the code's behavior and to write an automated test showing the code does indeed implement this behavior. Focusing on testing the expected behavior, rather than the implementation, of your code also makes for more comprehensive and more accurate tests, and thus helps Jenkins to provide more relevant feedback.

Of course, more classical unit testing, done once the code has been implemented, is another commonly-used approach, and is certainly better than no tests at all.

Jenkins isn't limited to unit testing, though. There are many other types of automated testing that you should consider, depending on the nature of your application, including integration testing, web testing, functional testing, performance testing, load testing, and so on. All of these have their place in an automated build setup.

Jenkins can also be used, in conjunction with techniques like BDD and acceptance TDD, as a communications tool aimed at both developers and other project stakeholders. BDD frameworks such as easyb, fitness, jbehave, rspec, Cucumber, and many others, try to present acceptance tests in terms that testers, product owners, and end users can understand. With the use of such tools, Jenkins can report on project progress in business terms, and facilitate communication between developers and non-developers within a team.

For existing or legacy applications with little or no automated testing in place, it can be time-consuming and difficult to retro-fit comprehensive unit tests onto the code. In addition, the tests may not be very effective, as they'll tend to validate the existing implementation rather than verify the expected business behavior. One useful approach in these situations is to write automated functional tests ("regression")

tests that simulate the most common ways that users manipulate the application. For example, automated web testing tools such as Selenium and WebDriver can be effectively used to test web applications at a high level. While this approach isn't as comprehensive as a combination of good quality unit, integration, and acceptance tests, it's still an effective and relatively cost-efficient way to integrate automated regression testing into an existing application.

In this chapter, we'll see how Jenkins helps you keep track of automated test results, and how you can use this information to monitor and dissect your build process.

6.2. Automating Your Unit and Integration Tests

The first thing we'll look at is how to integrate your unit tests into Jenkins. Whether you're practicing TDD or writing unit tests using a more conventional approach, these are probably the first tests that you'll want to automate with Jenkins.

Jenkins does an excellent job of reporting on your test results. However, it's up to you to write the appropriate tests and to configure your build script to run them automatically. Fortunately, integrating unit tests into your automated builds is generally relatively easy.

There are many unit testing tools out there, with the xUnit family holding a predominant place. In the Java world, JUnit is the de facto standard, although TestNG is another popular Java unit testing framework with a number of innovative features. For C# applications, the NUnit testing framework provides similar functionalities as those provided by JUnit, as does `Test::Unit` for Ruby. For C/C+++, there's CppUnit, and PHP developers can use PHPUnit. This isn't even an exhaustive list!

These tools can also be used for integration tests, functional tests, web tests, and so forth. Many web testing tools, such as Selenium, WebDriver, and Watir, generate xUnit-compatible reports. BDD and automated Acceptance-Test tools such as easyb, Fittesse, and Concordion are also xUnit-friendly. In the following sections we make no distinction between these different types of tests since, from a configuration point of view, they're treated by Jenkins in exactly the same manner. However, you will almost certainly need to make the distinction in your build jobs. In order to get the fastest possible feedback loop, your tests should be grouped into well-defined categories, starting with the fast-running unit tests, and then proceeding to the integration tests, before finally running the slower functional and web tests.

A detailed discussion of how to automate your tests is beyond the scope of this book, but we do cover a few useful techniques for Maven and Ant in the Appendix A, Automating Your Unit and Integration Tests.

6.3. Configuring Test Reports in Jenkins

Once your build generates test results, you need to configure your Jenkins build job to display them. As mentioned above, Jenkins will work fine with any xUnit-compatible test reports, no matter what language they're written in.

For Maven build jobs, no special configuration is required—just make sure you invoke a goal that will run your tests, such as `mvn test` (for your unit tests) or `mvn verify` (for unit and integration tests). An example of a Maven build job configuration is shown in Figure 6.1, “You configure your Jenkins installation in the Manage Jenkins screen”.

Figure 6.1. You configure your Jenkins installation in the Manage Jenkins screen

For freestyle build jobs, you need to do a little more configuration work. In addition to ensuring that your build actually runs the tests, you need to tell Jenkins to publish the JUnit test report. You configure this in the “Post-build Actions” section (see Figure 6.2, “Configuring Maven test reports in a freestyle project”). Here, you provide a path to the JUnit or TestNG XML reports. Their exact location will depend on a project—for a Maven project, a path like `**/target/surefire-reports/*.xml` will find them for most projects. For an Ant-based project, it will depend on how you configured the Ant JUnit task, as we discussed above.

Figure 6.2. Configuring Maven test reports in a freestyle project

For Java projects, whether they're using JUnit or TestNG, Jenkins does an excellent job out of the box. If you're using Jenkins for non-Java projects, you might need the xUnit Plugin. This plugin lets Jenkins process test reports from non-Java tools in a consistent way. It provides support for MSUnit and NUnit (for C# and other .NET languages), UnitTest++ and Boost Test (for C++), PHPUnit (for PHP), as well as a few other xUnit libraries via additional plugins (see Figure 6.3, “Installing the xUnit plugin”).

Figure 6.3. Installing the xUnit plugin

Once you've installed the xUnit Plugin, you'll need to configure the reporting for your particular xUnit reports in the “Post-build Actions” section. Check the “Publish testing tools result report” checkbox, and enter the path to the XML reports generated by your testing library (see Figure 6.4, “Publishing xUnit test results”). When the build job runs, Jenkins will convert these reports to JUnit reports so that they can be displayed in Jenkins.

Figure 6.4. Publishing xUnit test results

6.4. Displaying Test Results

Once Jenkins knows where to find the test reports, it does a great job of reporting on them. Indeed, one of Jenkins's main jobs is to detect and to report on build failures. A failing unit test is one of the most obvious symptoms.

As we mentioned earlier, Jenkins makes the distinction between failed builds and unstable builds. A failed build (indicated by a red ball) indicates test failures, or a build job that's broken in some brutal manner, such as a compilation error. An unstable build, on the other hand, is a build that isn't considered of sufficient quality. This is intentionally a little vague: what defines “quality” in this sense is largely up to you, but it's typically related to code quality metrics such as code coverage or coding standards that we'll be discussing later on in the book. For now, let's focus on the failed builds.

In Figure 6.5, “Jenkins displays test result trends on the project home page” we can see how Jenkins displays a Maven build job containing test failures. This is the build job home page, which should be your first port of call when a build breaks. When a build results in failing tests, the Latest Test Result link will indicate the current number of test failures in this build job (“5 failures” in the illustration), and also the change the number of test failures since the last build (“+5” in the illustration—five new test failures). You can also see how the tests have been faring over time—test failures from previous builds will also appear in red in the Test Result Trend graph.

Figure 6.5. Jenkins displays test result trends on the project home page

If you click on the Latest Test Result link, Jenkins will give you a rundown of the current test results (see Figure 6.6, “Jenkins displays a summary of the test results”). Jenkins understands Maven multimodule project structures, and for a Maven build job, Jenkins will initially display a summary view of test results per module. For more details about the failing tests in a particular module, just click on the module you're interested in.

Figure 6.6. Jenkins displays a summary of the test results

For freestyle build jobs, Jenkins will directly give you a summary of your test results, but organized by high-level packages rather than modules.

In both cases, Jenkins starts off by presenting a summary of test results for each package. From here, you can drill down, seeing test results for each test class and then finally the tests within the test classes themselves. If there are any failed tests, these will be prominently displayed at the top of the page.

This full view gives you both a good overview of the current state of your tests and an indication of their history. The Age column tells you how for how long a test has been broken, with a hyperlink that takes you back to the first build in which this test failed.

You can also add a description to the test results, using the Edit Description link in the top right-hand corner of the screen. This is a great way to annotate a build failure with some additional details, in order to add extra information about the origin of test failures or some notes about how to fix them.

When a test fails, you generally want to know why. To see the details of a particular test failure, just click on the corresponding link on this screen. This will display all the gruesome details, including the error message and the stack trace, as well as a reminder of how long the test has been failing (see Figure 6.7,

“The details of a test failure”). You should be wary of tests that have been failing for more than just a couple of builds—this is an indicator of either a tricky technical problem that might need investigating, or a complacent attitude to failed builds (developers might just be ignoring build failures), which is more serious and definitely should be investigated.

Figure 6.7. The details of a test failure

Make sure you also keep an eye on how long your tests take to run, and not just whether they pass or fail. Unit tests should be designed to run fast, and overly long-running tests can be the sign of a performance issue. Slow unit tests also delay feedback. In CI, fast feedback is the name of the game. For example, running one thousand unit tests in five minutes is good—taking an hour to run them isn't. So, it's a good idea to regularly check how long your unit tests are taking to run, and, if necessary, investigate why they're taking so long.

Luckily, Jenkins can easily tell you how long your tests have been taking to run over time. On the build job home page, click on the “trend” link in the Build History box on the left of the screen. This will give you a graph along the lines of the one in Figure 6.8, “Build time trends can give you a good indicator of how fast your tests are running”, showing how long each of your builds took to run. Tests aren't the only thing that happens in a build job, but if you have enough tests to worry about, they'll probably take a large proportion of the time. So, this graph is a great way to see how well your tests are performing as well.

Figure 6.8. Build time trends can give you a good indicator of how fast your tests are running

When you're on the Test Results page (see Figure 6.6, “Jenkins displays a summary of the test results”), you can also drill down and see how long the tests in a particular module, package, or class are taking to run. Just click on the test duration in the test results page (“Took 31 ms” in Figure 6.6, “Jenkins displays a summary of the test results”) to view the test history for a package, class, or individual test (see Figure 6.9, “Jenkins also lets you see how long your tests take to run”). This makes it easy to isolate a test that's taking more time than it should, or even decide when a general optimization of your unit tests is required.

6.5. Ignoring Tests

Jenkins distinguishes between test failures and skipped tests. Skipped tests are ones that have been deactivated, for example, by using the `@Ignore` annotation in JUnit 4:

```
@Ignore("Pending more details from the BA")
@Test
public void cashWithdrawalShouldDeductSumFromBalance() throws Exception {
    Account account = new Account();
    account.makeDeposit(100);
    account.makeCashWithdraw(60);
    assertThat(account.getBalance(), is(40));
}
```

Figure 6.9. Jenkins also lets you see how long your tests take to run

Skipping some tests is perfectly legitimate in some circumstances, such as to place an automated acceptance test, or higher-level technical test, on hold while you implement the lower levels. In such cases, you don't want to be distracted by the failing acceptance test, but you don't want to forget that the test exists either. Using techniques such as the `@Ignore` annotation is better than simply commenting out the test or renaming it (in JUnit 3), as it lets Jenkins keep tabs on the ignored tests for you.

In TestNG, you can also skip tests, using the `enabled` property:

```
@Test(enabled=false)
public void cashWithdrawalShouldDeductSumFromBalance() throws Exception {
    Account account = new Account();
    account.makeDeposit(100);
    account.makeCashWithdraw(60);
    assertEquals(account.getBalance(), 40);
}
```

In TestNG, you can also define dependencies between tests, so that certain tests will only run after another test or group of tests has run, as illustrated here:

```
@Test
public void serverStartedOk() {...}

@Test(dependsOnMethods = { "serverStartedOk" })
public void whenAUserLogsOnWithACorrectUsernameAndPasswordTheHomePageIsDisplayed() {...}
```

Here, if the first test (`serverStartedOk()`) fails, the following test will be skipped.

In all of these cases, Jenkins will mark the tests that were not run in yellow, both in the overall test results trend, and in the test details (see Figure 6.10, “Jenkins displays skipped tests as yellow”). Skipped tests aren't as bad as test failures, but it's important not to get into the habit of neglecting them. Skipped tests are like branches in a version control system: a test should be skipped for a specific reason, with a clear idea as to when they'll be reactivated. A skipped test that remains skipped for too long leaves a bad smell.

Figure 6.10. Jenkins displays skipped tests as yellow

6.6. Code Coverage

Another very useful test-related metric is code coverage. Code coverage gives an indication of what parts of your application were executed during the tests. While this in itself isn't a sufficient indication of quality testing (it's easy to execute an entire application without actually testing anything, and code coverage metrics provide no indication of the quality or accuracy of your tests), it's a very good indication of code that has not been tested. And, if your team is introducing rigorous testing practices such as TDD, code coverage can be a good indicator of how well these practices are being applied.

Code coverage analysis is a CPU and memory-intensive process, and will slow down your build job significantly. For this reason, you'll typically run code coverage metrics in a separate Jenkins build job, run after your unit and integration tests are successful.

There are many code coverage tools available, and several are supported in Jenkins, all through dedicated plugins. Java developers can pick between Cobertura and Emma, two popular open source code coverage tools, or Clover, a powerful commercial code coverage tool from Atlassian. For .NET projects, you can use NCover.

The behavior and configuration of all of these tools is similar. In this section, we'll look at Cobertura.

6.6.1. Measuring Code Coverage with Cobertura

Cobertura¹ is an open source code coverage tool for Java and Groovy that's easy to use and integrates well with both Maven and Jenkins.

Like almost all of the Jenkins code quality metrics plugins,² the Cobertura plugin for Jenkins won't run any test coverage metrics for you. It's left up to you to generate the raw code coverage data as part of your automated build process. Jenkins, on the other hand, does an excellent job of reporting on the code coverage metrics, including keeping track of code coverage over time, and providing aggregate coverage across multiple application modules.

Code coverage can be a complicated business, and it helps to understand the basic process that Cobertura follows, especially when you need to set it up in more low-level build scripting tools like Ant. Code coverage analysis works in three steps. First, it modifies (or “instruments”) your application classes to make them keep a tally of the number of times each line of code has been executed.³ They store all this data in a special data file (Cobertura uses a file called `cobertura.ser`).

When the application code has been instrumented, you run your tests against this instrumented code. At the end of the tests, Cobertura will have generated a data file containing the number of times each line of code was executed during the tests.

Once this data file has been generated, Cobertura can use this data to generate a report in a more usable format, such as XML or HTML.

6.6.1.1. Integrating Cobertura with Maven

Producing code coverage metrics with Cobertura in Maven is relatively straightforward. If all you're interested in is producing code coverage data, you just need to add the `cobertura-maven-plugin` to the build section of your `pom.xml` file:

```
<project>
```

¹ <http://cobertura.sourceforge.net>

² With the notable exception of Sonar, which we'll look at later on in the book.

³ This is actually a slight over-simplification; in fact, Cobertura stores other data as well, such as how many times each possible outcome of a boolean test was executed. However, this does not alter the general approach.

```

...
<build>
  <plugins>
    <plugin>
      <groupId>org.codehaus.mojo</groupId>
      <artifactId>cobertura-maven-plugin</artifactId>
      <version>2.5.1</version>
      <configuration>
        <formats>
          <format>html</format>
          <format>xml</format>
        </formats>
      </configuration>
    </plugin>
    ...
  </plugins>
</build>
...
</project>

```

This will generate code coverage metrics when you invoke the Cobertura plugin directly:

```
$ mvn cobertura:cobertura
```

The code coverage data will be generated in the `target/site/cobertura` directory, in a file called `coverage.xml`.

This approach, however, will instrument your classes and produce code coverage data for every build, which is inefficient. A better approach is to place this configuration in a special profile, as shown here:

```

<project>
  ...
  <profiles>
    <profile>
      <id>metrics</id>
      <build>
        <plugins>
          <plugin>
            <groupId>org.codehaus.mojo</groupId>
            <artifactId>cobertura-maven-plugin</artifactId>
            <version>2.5.1</version>
            <configuration>
              <formats>
                <format>html</format>
                <format>xml</format>
              </formats>
            </configuration>
          </plugin>
        </plugins>
      </build>
    </profile>
    ...
  </profiles>

```

```
</project>
```

In this case, you'd invoke the Cobertura plugin using the metrics profile to generate the code coverage data:

```
$ mvn cobertura:cobertura -Pmetrics
```

Another approach is to include code coverage reporting in your Maven reports. This approach is considerably slower and more memory-hungry than just generating the coverage data, but it can make sense if you're also generating other code quality metrics and reports at the same time. If you want to do this using Maven 2, you need to also include the Maven Cobertura plugin in the reporting section, as shown here:

```
<project>
...
<reporting>
  <plugins>
    <plugin>
      <groupId>org.codehaus.mojo</groupId>
      <artifactId>cobertura-maven-plugin</artifactId>
      <version>2.5.1</version>
      <configuration>
        <formats>
          <format>html</format>
          <format>xml</format>
        </formats>
      </configuration>
    </plugin>
  </plugins>
</reporting>
</project>
```

Now, the coverage data will be generated when you generate the Maven site for this project:

```
$ mvn site
```

If your Maven project contains modules (as is common practice for larger Maven projects), you just need to set up the Cobertura configuration in a parent `pom.xml` file—test coverage metrics and reports will be generated separately for each module. If you use the `aggregate` configuration option, the Maven Cobertura plugin will also generate a high-level report combining coverage data from all of the modules. However, whether you use this option or not, the Jenkins Cobertura plugin will take coverage data from several files and combine them into a single aggregate report.

At the time of writing, there's a limitation with the Maven Cobertura plugin—code coverage will only be recorded for tests executed during the **test** life cycle phase, and not for tests executed during the **integration-test** phase. This can be an issue if you're using this phase to run integration or web tests that require a fully packaged and deployed application—in this case, coverage from tests that are only performed during the integration test phase won't be counted in the Cobertura code coverage metrics.

6.6.1.2. Integrating Cobertura with Ant

Integrating Cobertura into your Ant build is more complicated than doing so in Maven. However, it does give you finer control over what classes are instrumented, and when coverage is measured.

Cobertura comes bundled with an Ant task that you can use to integrate Cobertura into your Ant builds. You'll need to download the latest Cobertura distribution, and unzip it somewhere on your hard disk. To make your build more portable, and therefore easier to deploy into Jenkins, it's a good idea to place the Cobertura distribution you're using within your project directory, and to save it in your version control system. This way it's easier to ensure that the build will use the same version of Cobertura no matter where it's run.

Assuming you've downloaded the latest Cobertura installation and placed it within your project in a directory called `tools`, you could do something like this:

```
<property name="cobertura.dir" value="/home/myipc/Documents/jenkinsbook/hudsonbook-content/tools/cobertura-1.9.4" />

<path id="cobertura.classpath">②
  <fileset dir="${cobertura.dir}">
    <include name="cobertura.jar" />③
    <include name="lib/**/*.jar" />④
  </fileset>
</path>

<taskdef classpathref="cobertura.classpath" resource="tasks.properties" />
```

- ❶ Tell Ant where your Cobertura installation is.
- ❷ Set up a classpath that Cobertura can use to run.
- ❸ The path containing the Cobertura application itself.
- ❹ And all of its dependencies.

Next, you need to instrument your application classes. You have to be careful to place these instrumented classes in a separate directory, so that they don't get bundled up and deployed to production by accident:

```
<target name="instrument" depends="init,compile">❶
  <delete file="cobertura.ser"/>②
  <delete dir="${instrumented.dir}" />③
  <cobertura-instrument todir="${instrumented.dir}">④
    <fileset dir="${classes.dir}">
      <include name="**/*.class" />
      <exclude name="**/*Test.class" />
    </fileset>
  </cobertura-instrument>
</target>
```

- ❶ Only instrument the application classes once they have been compiled.
- ❷ Remove any coverage data generated by previous builds.

- ③ Remove any previously instrumented classes.
- ④ Instrument the application classes (but not the test classes) and place them in the `${instrumented.dir}` directory.

At this stage, the `${instrumented.dir}` directory contains an instrumented version of our application classes. Now, all you need to do to generate some useful code coverage data is to run our unit tests against the classes in this directory:

```
<target name="test-coverage" depends="instrument">
  <junit fork="yes" dir="/home/mypc/Documents/jenkinsbook/hudsonbook-content">❶
    <classpath location="${instrumented.dir}" />
    <classpath location="${classes.dir}" />
    <classpath refid="cobertura.classpath" />❷

    <formatter type="xml" />
    <test name="${testcase}" todir="${reports.xml.dir}" if="testcase" />
    <batchtest todir="${reports.xml.dir}" unless="testcase">
      <fileset dir="${src.dir}">
        <include name="**/*Test.java" />
      </fileset>
    </batchtest>
  </junit>
</target>
```

- ❶ Run the JUnit tests against the instrumented application classes.
- ❷ The instrumented classes use Cobertura classes, so the Cobertura libraries also need to be on the classpath.

This will produce the raw test coverage data you need to produce the XML test coverage reports that Jenkins can use. To actually produce these reports, you need to invoke another task, as shown here:

```
<target name="coverage-report" depends="test-coverage">
  <cobertura-report srcdir="${src.dir}" destdir="${coverage.xml.dir}"
    format="xml" />
</target>
```

Finally, don't forget to tidy up after you're done: the **clean** target should delete not only the generated classes, but also the generated instrumented classes, the Cobertura coverage data, and the Cobertura reports:

```
<target name="clean"
  description="Remove all files created by the build/test process.">
  <delete dir="${classes.dir}" />
  <delete dir="${instrumented.dir}" />
  <delete dir="${reports.dir}" />
  <delete file="cobertura.log" />
  <delete file="cobertura.ser" />
</target>
```

Once this is done, you're ready to integrate your coverage reports into Jenkins.

6.6.1.3. Installing the Cobertura code coverage plugin

Once code coverage data is being generated as part of your build process, you can configure Jenkins to report on it. This involves installing the Jenkins Cobertura plugin. We went through this process in Section 2.8, “Adding Code Coverage and Other Metrics”, but we’ll run through it again to refresh your memory. Go to the Manage Jenkins screen, and click on Manage Plugins. This will take you to the Plugin Manager screen. If Cobertura has not been installed, you’ll find the Cobertura Plugin in the Available tab, in the Build Reports section (see Figure 6.11, “Installing the Cobertura plugin”). To install it, just tick the checkbox and press enter (or scroll down to the bottom of the screen and click on the “Install” button). Jenkins will download and install the plugin for you. Once the downloading is done, you’ll need to restart your Jenkins server.

Figure 6.11. Installing the Cobertura plugin

6.6.1.4. Reporting on code coverage in your build

Once you’ve installed the plugin, you can set up code coverage reporting in your build jobs. Since code coverage can be slow and memory-hungry, you’d typically create a separate build job for this and other code quality metrics, to be run after the normal unit and integration tests. For very large projects, you may even want to set up a build that only runs on a nightly basis. Indeed, feedback on code coverage and other such metrics is usually not as time-critical as feedback on test results. This will leave build executors free for build jobs that can benefit from snappy feedback.

As we mentioned earlier, Jenkins doesn’t do any code coverage analysis itself—you need to configure your build to produce the Cobertura `coverage.xml` file (or files) before you can generate any nice graphs or reports, typically using one of the techniques we discussed previously (see Figure 6.12, “Your code coverage metrics build needs to generate the coverage data”).

Figure 6.12. Your code coverage metrics build needs to generate the coverage data

Once you’ve configured your build to produce some code coverage data, you can configure Cobertura in the “Post-build Actions” section of your build job. When you tick the “Publish Cobertura Coverage Report” checkbox, you should see something like Figure 6.13, “Configuring the test coverage metrics in Jenkins”.

Figure 6.13. Configuring the test coverage metrics in Jenkins

The first and most important field here is the path to the Cobertura XML data that you generated. Your project may include a single `coverage.xml` file, or several. If you have a multimodule Maven project, for example, the Maven Cobertura plugin will generate a separate `coverage.xml` file for each module.

The path accepts Ant-style wildcards, so it's easy to include code coverage data from several files. For any Maven project, a path like `**/target/site/cobertura/coverage.xml` will include all of the code coverage metrics for all of the modules in the project.

There are actually several types of code coverage, and it can sometimes be useful to distinguish between them. The most intuitive is Line Coverage, which counts the number of times any given line is executed during the automated tests. “Conditional Coverage” (also referred to as “Branch Coverage”) takes into account whether the boolean expressions in `if` statements and the like are tested in a way that checks all the possible outcomes of the conditional expression. For example, consider the following code snippet:

```
if (price > 10000) {  
    managerApprovalRequired = true;  
}
```

To obtain full Conditional Coverage for this code, you'd need to execute it twice: once with a value that's more than 10,000, and one with a value of 10,000 or less.

Other more basic code coverage metrics include methods (how many methods in the application were exercised by the tests), classes, and packages.

Jenkins lets you define which of these metrics you want to track. By default, the Cobertura plugin will record Conditional, Line, and Method coverage, which is usually plenty. However, it's easy to add other coverage metrics if you think this might be useful for your team.

Jenkins code quality metrics aren't simply a passive reporting process—Jenkins lets you define how these metrics affect the build outcome. You can define threshold values for the coverage metrics that affect both the build outcome and the weather reports on the Jenkins dashboard (see Figure 6.14, “Test coverage results contribute to the project status on the dashboard”). Each coverage metric that you track takes three threshold values.

Figure 6.14. Test coverage results contribute to the project status on the dashboard

The first (the one with the sunny icon) is the minimum value necessary for the build to have a sunny weather icon. The second indicates the value below which the build will be assigned a stormy weather icon. Jenkins will extrapolate between these values for the other more nuanced weather icons.

The last threshold value is simply the value below which a build will be marked as “unstable”—the yellow ball. While not quite as bad as the red ball (for a broken build), a yellow ball will still result in a notification message and will look bad on the dashboard.

This feature is far from simply a cosmetic detail—it provides a valuable way of setting objective code quality goals for your projects. Although it can't be interpreted alone, falling code coverage is generally not a good sign in a project. So, if you're serious about code coverage, use these threshold values to provide some hard feedback about when things aren't up to scratch.

6.6.1.5. Interpreting code coverage metrics

Jenkins displays your code coverage reports on the build job home page. The first time it runs, it produces a simple bar chart (see Figure 2.30, “Jenkins displays code coverage metrics on the build home page”). From the second build onwards, a graph is shown, indicating the various types of coverage that you're tracking over time (see Figure 6.15, “Configuring the test coverage metrics in Jenkins”). In both cases, the graph will also show the code coverage metrics for the latest build.

Figure 6.15. Configuring the test coverage metrics in Jenkins

Jenkins also does a great job letting you drill down into the coverage metrics, displaying coverage breakdowns for packages, classes within a package, and lines of code within a class (see Figure 6.16, “Displaying code coverage metrics”). No matter what level of detail you're viewing, Jenkins will display a graph at the top of the page showing the code coverage trend over time. Further down, you'll find the breakdown by package or class.

Figure 6.16. Displaying code coverage metrics

Once you get to the class details level, Jenkins will also display the source code of the class, with the lines color-coded according to their level of coverage. Lines that have been completely executed during the tests are green, and lines that were never executed are marked in red. A number in the margin indicates the number of times a given line was executed. Finally, yellow shading in the margin is used to indicate insufficient conditional coverage (for example, an `if` statement that was only tested with one outcome).

6.6.2. Measuring Code Coverage with Clover

Clover is an excellent commercial code coverage tool from Atlassian⁴. Clover works well for projects using Ant, Maven, and even Grails. The configuration and use of Clover is well documented on the Atlassian website, so we won't describe these aspects in detail. However, to give some context, here's what a typically Maven 2 configuration of Clover for use with Jenkins would look like:

```
<build>
...
<plugins>
...
<plugin>
  <groupId>com.atlassian.maven.plugins</groupId>
  <artifactId>maven-clover2-plugin</artifactId>
  <version>3.0.4</version>
  <configuration>
    <includesTestSourceRoots>>false</includesTestSourceRoots>
    <generateXml>true</generateXml>
```

⁴ <http://www.atlassian.com/software/clover>

```
        </configuration>
      </plugin>
    </plugins>
  </build>
  ...

```

This will generate both an HTML and XML coverage report, including aggregated data if the Maven project contains multiple modules.

To integrate Clover into Jenkins, you need to install the Jenkins Clover plugin in the usual manner using the Plugin Manager screen. Once you've restarted Jenkins, you'll be able to integrate Clover code coverage into your builds.

Running Clover on your project is a multistep project: you instrument your application code, run your tests, aggregate the test data (for multimodule Maven projects), and generate the HTML and XML reports. Since this can be a fairly slow operation, you typically run it as part of a separate build job, and not with your normal tests. You can do this as follows:

```
$ clover2:setup test clover2:aggregate clover2:clover
```

Next, you need to set up the Clover reporting in Jenkins. Tick the Publish Clover Coverage Report checkbox to set this up. The configuration is similar to that of Cobertura—you need to provide the path to the Clover HTML report directory, and to the XML report file. You can also define threshold values for sunny and stormy weather, and for unstable builds (see Figure 6.17, “Configuring Clover reporting in Jenkins”).

Figure 6.17. Configuring Clover reporting in Jenkins

Once you've done this, Jenkins will display the current level of code coverage, as well as a graph of the code coverage over time, on your project build job home page (see Figure 6.18, “Clover code coverage trends”).

Figure 6.18. Clover code coverage trends

6.7. Automated Acceptance Tests

Automated acceptance tests play an important part in many agile projects, both for verification and for communication. As a verification tool, acceptance tests perform a similar role as integration tests, and aim to demonstrate that the application effectively does what's expected of it. But this is almost a secondary aspect of automated acceptance tests. The primary focus is actually on communication—demonstrating to nondevelopers (business owners, business analysts, testers, and so forth) precisely where the project is at.

Acceptance tests shouldn't be mixed with developer-focused tests, as both their aim and their audience is very different. Acceptance tests should be working examples of how the system works, with an emphasis on demonstration rather than exhaustive proof. The exhaustive tests should be done at the unit-testing level.

Acceptance Tests can be automated using conventional tools such as JUnit, but there's a growing tendency to use BDD frameworks for this purpose, as they tend to be a better fit for the public-facing nature of acceptance tests. BDD tools used for automated acceptance tests typically generate HTML reports with a specific layout that's well-suited to nondevelopers. They often also produce JUnit-compatible reports that can be understood directly by Jenkins.

BDD frameworks also have the notion of “Pending tests”, tests that are automated, but haven't yet been implemented by the development team. This distinction plays an important role in communication with other non-developer stakeholders. If you can automate these tests early on in the process, they can give an excellent indicator of which features have been implemented, which work, and which haven't been started yet.

As a rule, your acceptance tests should be displayed separately from the other more conventional automated tests. If they use the same testing framework as your normal tests (e.g., JUnit), make sure they're executed in a dedicated build job, so that non-developers can view them and concentrate on the business-focused tests without being distracted by low-level or technical ones. It can also help to adopt business-focused and behavioral naming conventions for your tests and test classes, to make them more accessible to non-developers (see Figure 6.19, “Using business-focused, behavior-driven naming conventions for JUnit tests”). The way you name your tests and test classes can make a huge difference when it comes to reading the test reports and understanding the actual business features and behavior that's being tested.

Figure 6.19. Using business-focused, behavior-driven naming conventions for JUnit tests

If you're using a tool that generates HTML reports, you can display them in the same build as your conventional tests, as long as they appear in a separate report. Jenkins provides a very convenient plugin for this sort of HTML report, called the HTML Publisher plugin (see Figure 6.20, “Installing the HTML Publisher plugin”). While it's still your job to ensure that your build produces the right reports, Jenkins can display the reports on your build job page, making them easily accessible to all team members.

Figure 6.20. Installing the HTML Publisher plugin

This plugin is easy to configure. Just go to the “Post-build Actions” section and tick the “Publish HTML reports” checkbox (see Figure 6.21, “Publishing HTML reports”). Next, give Jenkins the directory your HTML reports were generated in, an index page, and a title for your report. You can also ask Jenkins to store the reports generated for each build, or only keep the latest one.

Figure 6.21. Publishing HTML reports

Once this is done, Jenkins will display a special icon on your build job home page, with a link to your HTML report. In Figure 6.22, “Jenkins displays a special link on the build job home page for your report”, you can see the easyb reports we configured previously in action.

Figure 6.22. Jenkins displays a special link on the build job home page for your report

The HTML Publisher plugin works perfectly for HTML reports. If, on the other hand, you want to (also) publish non-HTML documents, such as text files, PDFs, and so forth, then the DocLinks plugin is for you. This plugin is similar to the HTML Publisher plugin, but lets you archive both HTML reports as well as documents in other formats. For example, in Figure 6.23, “The DocLinks plugin lets you archive both HTML and non-HTML artifacts”, we’ve configured a build job to archive both a PDF document and an HTML report. Both these documents will now be listed on the build home page.

Figure 6.23. The DocLinks plugin lets you archive both HTML and non-HTML artifacts

6.8. Automated Performance Tests with JMeter

Application performance is another important area of testing. Performance testing can be used to verify many things, such as how quickly an application responds to requests with a given number of simultaneous users, or how well the application copes with an increasing number of users. Many applications have Service Level Agreements, or SLAs, which define contractually how well they must perform.

Performance testing is often a one-off, ad-hoc activity, only undertaken right at the end of the project or when things start to go wrong. Nevertheless, performance issues are like any other sort of bug—the later in the process they’re detected, the more costly they are to fix. It therefore makes good sense to automate these performance and load tests so that you can spot any areas of degrading performance before they get out into the wild.

JMeter⁵ is a popular open source performance and load testing tool. It works by simulating load on your application, and measuring the response time as the number of simulated users and requests increase. It effectively simulates the actions of a browser or client application, sending requests of various sorts (HTTP, SOAP, JDBC, JMS and so on) to your server. You configure a set of requests to be sent to your application, as well as random pauses, conditions and loops, and other variations designed to better imitate real user actions.

⁵ <http://jakarta.apache.org/jmeter/>

JMeter runs as a Swing application, in which you can configure your test scripts (see Figure 6.24, “Preparing a performance test script in JMeter”). You can even run JMeter as a proxy, and then manipulate your application in an ordinary browser to prepare an initial version of your test script.

A full tutorial on using JMeter is beyond the scope of this book. However, it's fairly easy to learn, and you can find ample details about how to use it on the JMeter website. With a little work, you can have a very respectable test script up and running in a matter of hours.

What we're interested in here is the process of automating these performance tests. There are several ways to integrate JMeter tests into your Jenkins build process. Although at the time of writing, there was no official JMeter plugin for Maven available in the Maven repositories, there's an Ant plugin. So, the simplest approach is to write an Ant script to run your performance tests, and then either call this Ant script directly, or (if you're using a Maven project, and want to run JMeter through Maven) use the Maven Ant integration to invoke the Ant script from within Maven. A simple Ant script running some JMeter tests is illustrated here:

```
<project default="jmeter">
  <path id="jmeter.lib.path">
    <pathelement location="/home/mypc/Documents/jenkinsbook/hudsonbook-content/tools/jmeter/extras">
    </pathelement>
  </path>

  <taskdef name="jmeter"
    classname="org.programmerplanet.ant.taskdefs.jmeter.JMeterTask"
    classpathref="jmeter.lib.path" />

  <target name="jmeter">
    <jmeter jmeterhome="/home/mypc/Documents/jenkinsbook/hudsonbook-content/tools/jmeter"
      testplan="/home/mypc/Documents/jenkinsbook/hudsonbook-content/src/test/jmeter/gameoflif
      resultlog="/home/mypc/Documents/jenkinsbook/hudsonbook-content/target/jmeter-results.j
      <jvmarg value="-Xmx512m" />
    </jmeter>
  </target>
</project>
```

This assumes that the JMeter installation is available in the `tools` directory of your project. Placing tools such as JMeter within your project structure is a good habit, as it makes your build scripts more portable and easier to run on any machine, which is precisely what you need to run them on Jenkins.

Figure 6.24. Preparing a performance test script in JMeter

Note that we're also using the optional `<jvmarg>` tag to provide JMeter with an ample amount of memory—performance testing is a memory-hungry activity.

The script shown here will execute the JMeter performance tests against a running application. So, you need to ensure that the application you want to test is up and running before you start the tests. There are

several ways to do this. For more heavy-weight performance tests, you'll usually want to deploy your application to a test server before running the tests. For most applications this isn't usually too difficult—the Maven Cargo plugin, for example, lets you automate the deployment process to a variety of local and remote servers. We'll also see how to do this in Jenkins later on in the book.

Alternatively, if you're using Maven for a web application, you can use the Jetty or Cargo plugin to ensure that the application is deployed before the integration tests start, and then call the JMeter Ant script from within Maven during the integration test phase. Using Jetty, for example, you could do something like this:

```
<project...>
  <build>
    <plugins>
      <plugin>
        <groupId>org.mortbay.jetty</groupId>
        <artifactId>jetty-maven-plugin</artifactId>
        <version>7.1.0.v20100505</version>
        <configuration>
          <scanIntervalSeconds>10</scanIntervalSeconds>
          <connectors>
            <connector
              implementation="org.eclipse.jetty.server.nio.SelectChannelConnector">
                <port>${jetty.port}</port>
                <maxIdleTime>60000</maxIdleTime>
            </connector>
          </connectors>
          <stopKey>foo</stopKey>
          <stopPort>9999</stopPort>
        </configuration>
      <executions>
        <execution>
          <id>start-jetty</id>
          <phase>pre-integration-test</phase>
          <goals>
            <goal>run</goal>
          </goals>
          <configuration>
            <scanIntervalSeconds>0</scanIntervalSeconds>
            <daemon>true</daemon>
          </configuration>
        </execution>
        <execution>
          <id>stop-jetty</id>
          <phase>post-integration-test</phase>
          <goals>
            <goal>stop</goal>
          </goals>
        </execution>
      </executions>
    </plugin>
    ...
  </plugins>
</build>
```

```
</project>
```

This will start up an instance of Jetty, deploy your web application to it just before the integration tests, and shut it down afterwards.

Finally, you need to run the JMeter performance tests during this phase. You can do this by using the `maven-antrun-plugin` to invoke the Ant script you wrote earlier on during the integration test phase:

```
<project...>
...
<profiles>
  <profile>
    <id>performance</id>
    <build>
      <plugins>
        <plugin>
          <artifactId>maven-antrun-plugin</artifactId>
          <version>1.4</version>
          <executions>
            <execution>
              <id>run-jmeter</id>
              <phase>integration-test</phase>
              <goals>
                <goal>run</goal>
              </goals>
              <configuration>
                <tasks>
                  <ant antfile="build.xml" target="jmeter" >
                </tasks>
              </configuration>
            </execution>
          </executions>
        </plugin>
      </plugins>
    </build>
  </profile>
</profiles>
...
</project>
```

Now, all you need to do is to run the integration tests with the performance profile to get Maven to run the JMeter test suite. You can do this by invoking the **integration-test** or **verify** Maven life cycle phase:

```
$ mvn verify -Pperformance
```

Once you've configured your build script to handle JMeter, you can set up a performance test build in Jenkins. For this, we'll use the Performance Test Jenkins plugin, which understands JMeter logs and can generate nice statistics and graphs using this data. So, go to the Plugin Manager screen on your Jenkins server and install this plugin (see Figure 6.25, "Preparing a performance test script in JMeter"). When you've installed the plugin, you'll need to restart Jenkins.

Figure 6.25. Preparing a performance test script in JMeter

Once you've installed the plugin, you can set up a performance build job in Jenkins. This build job will typically be fairly separate from your other builds. In Figure 6.26, “Setting up the performance build to run every night at midnight”, we've set up the performance build to run on a nightly basis, which is probably enough for a long-running load or performance test.

Figure 6.26. Setting up the performance build to run every night at midnight

All that remains is to configure the build job to run your performance tests. In Figure 6.27, “Performance tests can require large amounts of memory”, we're running the Maven build you configured earlier on. Note that we're using the MAVEN_OPTS field (accessible by clicking on the Advanced button) to provide plenty of memory for the build job.

Figure 6.27. Performance tests can require large amounts of memory

To set up performance reporting, just tick the “Publish Performance test result report” option in the Post-build Actions section (see Figure 6.28, “Configuring the Performance plugin in your build job”). You'll need to tell Jenkins where to find your JMeter test results (the output files, not the test scripts). The Performance plugin is happy to process multiple JMeter results, so you can put wildcards in the path to make sure all of your JMeter reports are displayed.

If you take your performance metrics seriously, then the build should fail if the required SLA isn't met. In a CI environment, any sort of metrics build that doesn't fail if minimum quality criteria aren't met will tend to be ignored.

You can configure the Performance plugin to mark a build as unstable or failing if a certain percentage of requests result in errors. By default, these values will only be raised in the event of real application errors (i.e., bugs) or server crashes. However, you really should configure your JMeter test scripts to place a ceiling on the maximum acceptable response time for your requests. This is particularly important if your application has contractual obligations in this regard. One way to do this in JMeter is by adding a Duration Assertion element to your script. This will cause an error if any request takes longer than a certain fixed time to execute.

Figure 6.28. Configuring the Performance plugin in your build job

Now, when the build job runs, the Performance plugin will produce graphs keeping track of overall response times and of the number of errors (see Figure 6.29, “The Jenkins Performance plugin keeps track of response time and errors”). There will be a separate graph for each JMeter report you've generated. If there's only one graph, it will appear on the build home page; otherwise you can view them on a dedicated page that you can access via the Performance Trend menu item.

Figure 6.29. The Jenkins Performance plugin keeps track of response time and errors

This graph gives you an overview of performance over time. You'd typically use this graph to ensure that your average response times are within the expected limits, and also spot any unusually high variations in the average or maximum response times. However, if you need to track down and isolate performance issues, the Performance Breakdown screen can be more useful. From within the Performance Trend report, click on the Last Report link at the top of the screen. This will display a breakdown of response times and errors per request (see Figure 6.30, “You can also view performance results per request”). You can do the same thing for previous builds, by clicking on the Performance Report link in the build details page.

With some minor variations, a JMeter test script basically works by simulating a given number of simultaneous users. Typically, however, you will want to see how your application performs for different numbers of users. The Jenkins Performance plugin handles this quite well, and can process graphs for multiple JMeter reports. Just make sure you use a wildcard expression when you tell Jenkins where to find the reports.

Of course, it would be nice to be able to reuse the same JMeter test script for each test run. JMeter supports parameters, so you can easily reuse the same JMeter script with different numbers of simulated users. You just use a property expression in your JMeter script, and then pass the property to JMeter when you run the script. If your property is called `request.threads`, then the property expression in your JMeter script would be `${__property(request.threads)}`. Then, you can use the `<property>` element in the `<jmeter>` Ant task to pass the property when you run the script. The following Ant target, for example, runs JMeter three times, for 200, 500, and 1000 simultaneous users:

```
<target name="jmeter">
  <jmeter jmeterhome="/home/my pc/Documents/jenkinsbook/hudsonbook-content/tools/jmeter"
    testplan="/home/my pc/Documents/jenkinsbook/hudsonbook-content/src/test/jmeter/gameoflif
    resultlog="/home/my pc/Documents/jenkinsbook/hudsonbook-content/target/jmeter-results-20
    <jvmarg value="-Xmx512m" />
    <property name="request.threads" value="200"/>
    <property name="request.loop" value="20"/>
  </jmeter>
  <jmeter jmeterhome="/home/my pc/Documents/jenkinsbook/hudsonbook-content/tools/jmeter"
    testplan="/home/my pc/Documents/jenkinsbook/hudsonbook-content/src/test/jmeter/gameoflif
    resultlog="/home/my pc/Documents/jenkinsbook/hudsonbook-content/target/jmeter-results-50
    <jvmarg value="-Xmx512m" />
    <property name="request.threads" value="500"/>
    <property name="request.loop" value="20"/>
  </jmeter>
  <jmeter jmeterhome="/home/my pc/Documents/jenkinsbook/hudsonbook-content/tools/jmeter"
    testplan="/home/my pc/Documents/jenkinsbook/hudsonbook-content/src/test/jmeter/gameoflif
    resultlog="/home/my pc/Documents/jenkinsbook/hudsonbook-content/target/jmeter-results-10
    <jvmarg value="-Xmx512m" />
    <property name="request.threads" value="1000"/>
    <property name="request.loop" value="20"/>
  </jmeter>
</target>
```

Figure 6.30. You can also view performance results per request

6.9. Help! My Tests Are Too Slow!

One of the underlying principles of designing your CI builds is that the value of information about a build failure diminishes rapidly with time. In other words, the longer the news of a build failure takes to get to you, the less it's worth, and the harder it is to fix.

Indeed, if your functional or integration tests are taking several hours to run, chances are they won't be run for every change. They're more likely to be scheduled as a nightly build. The problem with this is that a lot can happen in twenty-four hours, and, if the nightly build fails, it will be difficult to figure out which of the many changes committed to version control during the day was responsible. This is a serious issue, and penalizes your CI server's ability to provide fast useful feedback.

Of course some builds are slow, by their very nature. Performance or load tests fall into this category, as do some more heavyweight code quality metrics builds for large projects. However, integration and functional tests most definitely do not fall into this category. You should do all you can to make these tests as fast as possible. Under ten minutes is probably acceptable for a full integration/functional test suite. Two hours isn't.

So, if you find yourself needing to speed up your tests, here are a few strategies that might help, in approximate order of difficulty.

6.9.1. Add More Hardware

Sometimes the easiest way to speed up your builds is to throw more hardware into the mix. This could be as simple as upgrading your build server. Compared to the time and effort saved in identifying and fixing integration-related bugs, the cost of buying a shiny new build server is relatively modest.

Another option is to consider using a virtual or cloud-based approach. Later on in the book, you'll see how you can use VMware virtual machines, or cloud-based infrastructure such as Amazon Web Services (EC2) or CloudBees to increase your build capacity on an "as-needed" basis, without having to invest in permanent new machines.

This approach can also involve distributing your builds across several servers. While this won't in itself speed up your tests, it may result in faster feedback if your build server is under heavy demand, and if build jobs are constantly being queued.

6.9.2. Run Fewer Integration/Functional Tests

In many applications, integration or functional tests are used by default as the standard way to test almost all aspects of the system. However, these tests aren't the best way to detect and identify bugs. Because

of the large number of components involved in a typical end-to-end test, it can be very hard to know where something has gone wrong. In addition, with so many moving parts, it's extremely difficult, if not completely unfeasible, to cover all of the possible paths through the application.

For this reason, wherever possible, you should prefer quick-running unit tests to much slower integration and functional tests. When you're confident that the individual components work well, you can complete the picture by a few end-to-end tests that step through common use cases for the system, or use cases that have caused problems in the past. This will help ensure that the components do fit together correctly, which is, after all, what integration tests are supposed to do. But, leave the more comprehensive tests where possible to unit tests. This strategy is probably the most sustainable approach for keeping your feedback loop short, but it does require some discipline and effort.

6.9.3. Run Your Tests in Parallel

If your functional tests take two hours to run, it's unlikely that they all need to be run back-to-back. It's also unlikely that they will be consuming all of the available CPU on your build machine. So breaking your integration tests into smaller batches and running them in parallel makes a lot of sense.

There are several strategies you can try, and your mileage will probably vary depending on the nature of your application. One approach, for example, is to set up several build jobs to run different subsets of your functional tests, and to run these jobs in parallel. Jenkins lets you aggregate test results. This is a good way to take advantage of a distributed build architecture to speed up your builds even further. Essential to this strategy is the ability to run subsets of your tests in isolation, which may require some refactoring.

At a lower level, you can also run your tests in parallel at the build scripting level. As you saw earlier, both TestNG and the more recent versions of JUnit support running tests in parallel. Nevertheless, you'll need to ensure that your tests can be run concurrently, which may take some refactoring. For example, common files or shared instance variables within test cases will cause problems here.

In general, you need to be careful of interactions between your tests. If your web tests start up an embedded web server such as Jetty, for example, you need to make sure the port used is different for each set of concurrent tests.

Nevertheless, if you can get it to work for your application, running your tests in parallel is one of the more effective way to speed up your tests.

6.10. Conclusion

Automated testing is a critical part of any CI environment, and should be taken very seriously. As in other areas on CI, and perhaps even more so, feedback is king, so it's important to ensure that your tests run fast, even the integration and functional ones.

Chapter 7. Securing Jenkins

7.1. Introduction

Jenkins supports several security models, and can integrate with several user repositories. In smaller organizations, where developers work in close proximity, security on your Jenkins machine may not be a large concern—you may simply want to prevent unidentified users from tampering with your build job configurations. For larger organizations, with multiple teams, a stricter approach might be required, where only team members and system administrators are allowed to modify build job configurations. In situations where the Jenkins server may be exposed to a broader audience, such as on an internal corporate website, or even on the Internet, certain build jobs may be visible to all users whereas others will need to be hidden from unauthorized users.

In this chapter, we'll look at how to configure different security configurations in Jenkins for different environments and circumstances.

7.2. Activating Security in Jenkins

Setting up basic security in Jenkins is easy enough. Go to the main configuration page and check the Enable security checkbox (see Figure 7.1, “Enabling security in Jenkins”). This will display a number of options, that we'll investigate in detail in this chapter. The first section, Security Realms, determines where Jenkins will look for users during authentication, and includes options such as using users stored in an LDAP server, using the underlying Unix user accounts (assuming, of course, that Jenkins is running on a Unix machine), or using a simple built-in user database managed by Jenkins.

The second section, Authorization, determines what users can do once they're logged in. This ranges from simple options like “Anyone can do anything” or “Logged-in users can do anything,” to more sophisticated role and project-based authorization policies.

Figure 7.1. Enabling security in Jenkins

In the remainder of this chapter, we'll look at how to configure Jenkins security for a number of common scenarios.

7.3. Simple Security in Jenkins

The most simple usable security model in Jenkins allows authenticated users to do anything, whereas non-authenticated users will just have a read-only view of the build jobs. This is great for small teams—developers can manage the build jobs, whereas other users (testers, BAs, project managers, and so on) can view the build jobs as required to view the status of the project. Indeed, certain build jobs may be

set up just for this purpose, displaying the results of automated acceptance tests or code quality metrics, for example.

You can set up this sort of configuration to choose “Logged-in users can do anything” in the Authorization section. There are several ways that Jenkins can authenticate users (see Section 7.4, “Security Realms—Identifying Jenkins Users”), but for this example, we’ll be using the simplest option, which is to use Jenkins’s own built in database (see Section 7.4.1, “Using Jenkins’s Built-in User Database”). This is the configuration illustrated in Figure 7.1, “Enabling security in Jenkins”.

Make sure you tick the “Allow users to sign up” option. This option will display a Sign up link at the top of the screen to let users create their own user account as required (see Figure 7.2, “The Jenkins Sign up page”). It’s a good idea for developers to use their SCM username here: in this case, Jenkins will be able to work out what users contributed to the SCM changes that triggered a particular build.

Figure 7.2. The Jenkins Sign up page

This approach is obviously a little too simple for many situations—it’s useful for small teams working in close proximity, where the aim is to know whose changes caused (or broke) a particular build, rather than to manage access in any more restrictive way. In the following sections, we’ll discuss the two orthogonal aspects of Jenkins security: identifying your users (Security Realms) and determining what they’re allowed to (Authorization).

7.4. Security Realms—Identifying Jenkins Users

Jenkins lets you identify and manage users in a number of ways, ranging from a simple, built-in user database suitable for small teams to integration with enterprise directories, with many other options in between.

7.4.1. Using Jenkins’s Built-in User Database

The easiest way to manage user accounts is to use Jenkins’s internal user database. This is a good option if you want to keep things simple, as very little setup or configuration is required. Users who need to log in to the Jenkins server can sign up and create an account for themselves, and, depending on the security model chosen, an administrator can then decide what these users are allowed to do.

Jenkins automatically adds all SCM users to this database whenever a change is committed to source code monitored by Jenkins. These user names are used mainly to record who is responsible for each build job. You can view the list of currently known users by clicking on the People menu entry (see Figure 7.3, “The list of users known to Jenkins”). Here, you can visualize the users that Jenkins currently knows about, and also see the last project they committed changes to. Note that this list contains all of the users who have ever committed changes to the projects that Jenkins monitors—they may not be (and usually aren’t) all active Jenkins users who are able to log in to the Jenkins server.

Figure 7.3. The list of users known to Jenkins

If you click on a user in this list, Jenkins takes you to a page displaying various details about this user, including the user's full name and the build jobs they've contributed to (see Figure 7.4, "Displaying the builds that a user participates in"). From here, you can also modify or complete the details about this user, such as their password or email address.

Figure 7.4. Displaying the builds that a user participates in

A user appearing in this list can't necessarily log in to Jenkins. To be able to log in to Jenkins, the user account needs to be set up with a password. There are essentially two ways to do this. If you've activated the "Allow users to sign up" option, users can simply sign up with their SCM user name and provide their email address and a password (see Section 7.3, "Simple Security in Jenkins"). Alternatively, you can activate a user by clicking on the Configure menu option in the user details screen, and provide an email address and password yourself (see Figure 7.5, "Creating a new user account by signing up").

Figure 7.5. Creating a new user account by signing up

It's worth noting that if your email addresses are synchronized with your version control user names (for example, if you work at acme.com, and user "joe" in your version control system has an email address of joe@acme.com), you can get Jenkins to derive the email address from a user name by adding a suffix that you configure in the Email Notification section (see Figure 7.6, "Synchronizing email addresses"). If you've set up this sort of configuration, you don't need to specify the email address for new users unless it doesn't follow this convention.

Figure 7.6. Synchronizing email addresses

Another way to manage the current active users (those who can actually log in to Jenkins) is by clicking on the Manage Users link in the main Jenkins configuration page (see Figure 7.7, "You can also manage Jenkins users from the Jenkins configuration page").

Figure 7.7. You can also manage Jenkins users from the Jenkins configuration page

From here, you can view and edit the users who can log in to Jenkins (see Figure 7.8, "The Jenkins user database"). This includes both users that have signed up manually (if this option has been activated) and SCM users that you've activated by configuring them with a password. You can also edit a user's details

(for example modifying their email address or resetting their password), or even remove them from the list of active users. Doing this won't remove them from the overall user list (their name will still appear in the build history, for example), but they'll no longer be able to log in to the Jenkins server.

Figure 7.8. The Jenkins user database

The internal Jenkins database is sufficient for many teams and organizations. However, for larger organizations, it may become tedious and repetitive to manage large numbers of user accounts by hand, especially if this information already exists elsewhere. In the following sections, we'll look at how to hook Jenkins up to other user management systems, such as LDAP repositories and Unix users and groups.

7.4.2. Using an LDAP Repository

Many organizations use LDAP directories to store user accounts and passwords across applications. Jenkins integrates well with LDAP, with no special plugins required. It can authenticate users using the LDAP repository, check group membership, and retrieve the email address of authenticated users.

To integrate Jenkins with your LDAP repository, Just select “LDAP” in the Security Realm section, and fill in the appropriate details about your LDAP server (see Figure 7.9, “Configuring LDAP in Jenkins”). The most important field is the repository server. If you're using a non-standard port, you'll need to provide this as well (for example, `ldap.acme.org:1389`). Or, if you're using LDAPS, you'll need to specify this as well (for example, `ldaps://ldap.acme.org`)

If your server supports anonymous binding, this will probably be enough to get you started. If not, you can use the Advanced options to fine-tune your configuration.

Most of the Advanced fields can safely be left blank unless you have a good reason to change them. If your repository is extremely large, you may want to specify a root DN value (e.g., `dc=acme, dc=com`) and/or a User and Group search base (e.g., `ou=people`) to narrow down the scope of user queries. This isn't usually required unless you notice performance issues. Or, if your server doesn't support anonymous binding, you'll need to provide a Manager DN and a Manager DN password so that Jenkins can connect to the server to perform its queries.

Figure 7.9. Configuring LDAP in Jenkins

Once you've set up LDAP as your Security Realm, you can configure your favorite security model as described previously. When users log in to Jenkins, they'll be authenticated against the LDAP repository.

You can also use LDAP groups, though the configuration isn't immediately obvious. Suppose you've defined a group called JenkinsAdmin in your LDAP repository, with a DN of `cn=JenkinsAdmin, ou=Groups, dc=acme, dc=com`. To refer to this group in Jenkins, you need to take the common name

(cn) in uppercase, and prefix it with `ROLE_`. So `cn=JenkinsAdmin` becomes `ROLE_JENKINSADMIN`. You can see an example of LDAP groups used in this way in Figure 7.10, “Using LDAP Groups in Jenkins”.

Figure 7.10. Using LDAP Groups in Jenkins

7.4.3. Using Microsoft Active Directory

Microsoft Active Directory is a directory service product widely used in Microsoft environments. Although Active Directory does provide an LDAP service, it can be a little tricky to set up, and it's simpler to get Jenkins to talk directly to the Active Directory server. Fortunately, there's a plugin for that.

The Jenkins Active Directory plugin lets you configure Jenkins to authenticate against a Microsoft Active Directory server. You can both authenticate users and retrieve their groups for Matrix and Project-based authorization. Note that, unlike the conventional LDAP integration (see Section 7.4.2, “Using an LDAP Repository”), there's no need to prefix group names with `ROLE_`—you can use Active Directory groups (such as “Domain Admins”) directory.

To configure the plugin, you need to provide the full domain name of your Active Directory server. If you have more than one domain, you can provide a comma-separated list. If you provide the forest name (say “acme.com” instead of “europe.acme.com”), then the search will be done against the global catalog. Note that if you do this without specifying the bind DN (see below), the user would have to login as “europe\joe” or “joe@europe”.

The advanced options let you specify a site name (to improve performance by restricting the domain controllers that Jenkins queries), and a Binding DN and password, which come in handy if you're connecting to a multidomain forest. You need to provide a valid Binding DN and password values, that Jenkins can use to connect to your server so that it can establish the full identity of the user being authenticated. This way, the user can simply type in “jack” or “jill”, and have the system automatically figure out that they're `jack@europe.acme.com` or `jack@asia.acme.com`. You need to provide the full user principal name with domain name, like `admin@europe.acme.com`, or an LDAP-style distinguished name, such as `CN=Administrator,OU=europe,DC=acme,DC=com`.

Another nice thing about this plugin is that it works both in a Windows environment and in a Unix environment. So, if Jenkins is running on a Unix server, it can still authenticate against a Microsoft Active Directory service running on another machine.

More precisely, if Jenkins is running on a Windows machine and you don't specify a domain, that machine must be a member of the domain you wish to authenticate against. Jenkins will use ADSI to figure out all the details, so no additional configuration is required.

On a non-Windows machine (or if you specify one or more domains), you need to tell Jenkins the name of Active Directory domain(s) to authenticate with. Jenkins then uses DNS SRV records and LDAP service of Active Directory to authenticate users.

Jenkins can determine which groups in Active Directory the user belongs to, so you can use these as part of your authorisation strategy. For example, you can use these groups in matrix-based security, or allow “Domain Admins” to administer Jenkins.

7.4.4. Using Unix Users and Groups

If you're running Jenkins on a Unix machine, you can also ask Jenkins to use the local user and group accounts. In this case, users will log into Jenkins using their Unix account logins and passwords. This uses Pluggable Authentication Modules (PAM), and also works fine with NIS.

In its most basic form, this is somewhat cumbersome, as it requires new user accounts to be set up and configured for each new Jenkins user. It's only really useful if these accounts need to be set up for other purposes.

7.4.5. Delegating to the Servlet Container

Another way to identify Jenkins users is to let your Servlet container do it for you. This approach is useful if you're running Jenkins on a Servlet container such as Tomcat or GlassFish, and you already have an established way to integrate the Servlet container with your local enterprise user directory. Tomcat, for example, allows you to authenticate users against a relational database (using direct JDBC or a DataSource), JNDI, JAAS, or an XML configuration file. You can also use the roles defined in the Servlet container's user directory for use with Matrix and Project-based authorization strategies.

In Jenkins, this is easy to configure—just select this option in the Security Realm section (see Figure 7.11, “Selecting the security realm”). Once you've done this, Jenkins will let the server take care of everything.

Figure 7.11. Selecting the security realm

7.4.6. Using Atlassian Crowd

If your organization is using Atlassian products such as JIRA and Confluence, you may also be using Crowd. Crowd is a commercial Identity Management and Single-Sign On (SSO) application from Atlassian that lets you manage single user accounts across multiple products. It lets you manage both an internal database of users, groups, and roles, and integrates with external directories such as LDAP directories or custom user stores.

Using the Jenkins Crowd plugin, you can use Atlassian Crowd as the source of your Jenkins users and groups. Before you start, you need to set up a new application in Crowd (see Figure 7.12, “Using Atlassian Crowd as the Jenkins Security Realm”). Just set up a new Generic Application called “hudson” (or something similar), and step through the tabs. In the Connections tab, you need to provide the IP address of your Jenkins server. Then, map the Crowd directories that you'll be using to retrieve Jenkins user accounts and group information. Finally, you'll need to tell Crowd which users from these

directories can connect to Jenkins. One option is to allow all users to authenticate, and let Jenkins sort out the details. Alternatively, you can list the Crowd user groups who are allowed to connect to Jenkins.

Figure 7.12. Using Atlassian Crowd as the Jenkins Security Realm

Once you've set this up, you need to install the Jenkins Crowd plugin, which you do as usual via the Jenkins Plugin Manager. Once you've installed the plugin and restarted Jenkins, you can define Crowd as your Security Realm in the main Jenkins configuration screen (see Figure 7.13, “Using Atlassian Crowd as the Jenkins Security Realm”).

Figure 7.13. Using Atlassian Crowd as the Jenkins Security Realm

With this plugin installed and configured, you can use users and groups from Crowd for any of the Jenkins Authorization strategies we discussed earlier on in the chapter. For example, in Figure 7.14, “Using Atlassian Crowd groups in Jenkins”, we're using user groups defined in Crowd to set up Matrix-based security in the main configuration screen.

Figure 7.14. Using Atlassian Crowd groups in Jenkins

7.4.7. Integrating with Other Systems

In addition to the authentication strategies discussed here, there are a number of other plugins that allow Jenkins to authenticate against other systems. At the time of writing, these include Central Authentication Service (CAS)—an open source single sign-on tool—and the Collabnet Source Forge Enterprise Edition (SFEE) server.

If no plugin is available, you can also write your own custom authentication script. To do this, you need to install the Script Security Realm plugin. Once you've installed the script and restarted Jenkins, you can write two scripts in your favorite scripting language. One script authenticates users, whereas the other determines the groups of a given user (see Figure 7.15, “Using custom scripts to handle authentication”).

Figure 7.15. Using custom scripts to handle authentication

Before invoking the authentication script, Jenkins sets two environment variables: `U`, containing the username, and `P`, containing the password. This script uses these environment variables to authenticate using the specified username and password, returning 0 if the authentication is successful, and some other value otherwise. If authentication fails, the output from the process will be reported in the error message displayed to the user. Here is a simple Groovy authentication script:

```
def env = System.getenv()
```

```

def username = env['U']
def password = env['P']

println "Authenticating user $username"

if (authenticate(username, password)) {
    System.exit 0
} else {
    System.exit 1
}

def authenticate(def username, def password) {
    def userIsAuthenticated = true
    // Authentication logic goes here
    return userIsAuthenticated
}

```

This script is sufficient if all you have to deal with is basic authentication without groups. If you want to use groups from your custom authentication source in your Matrix-based or Project-based authorizations (see Section 7.5, “Authorization—Who Can Do What”), you can write a second script, which determines the groups for a given user. This script uses the U environment variable to determine which user is trying to log in, and prints a comma-separated list of groups for this user to the standard output. If you don’t like commas, you can override the separating character in the configuration. A simple Groovy script to do this job is shown here:

```

def env = System.getenv()
def username = env['U']

println findGroupsFor(username)

System.exit 0

def findGroupsFor(def username) {
    return "admin,game-of-life-developer"
}

```

Both these scripts must return 0 when called for a user to be authenticated.

7.5. Authorization—Who Can Do What

Once you've defined how to identify your users, you need to decide what they're allowed to do. Jenkins supports a variety of strategies in this area, ranging from a simple approach where a logged-in user can do anything to more involved roles and project-based authentication strategies.

7.5.1. Matrix-based Security

Letting signed-in users do anything is certainly flexible, and may be all you need for a small team. For larger or multiple teams, or cases where Jenkins is being used outside the development environment, a more sophisticated approach is generally required.

Matrix-based security is a more sophisticated approach, where different users are assigned different rights, using a role-based approach.

7.5.1.1. Setting up matrix-based security

The first step in setting up matrix-based security in Jenkins is to create an administrator. This is an essential step, and must be done before all others. Your administrator can be an existing user, or one created specially for the purpose. If you want to create a dedicated administrator user, simply create one by signing up in the usual way (see Figure 7.2, “The Jenkins Sign up page”). It doesn’t have to be associated with an SCM user.

Once your admin user is ready, you can activate matrix-based security by selecting “Matrix-based security” in the Authorization section of the main configuration page. Jenkins will display a table containing authorized users, and checkboxes corresponding to the various permissions that you can assign to these users (see Figure 7.16, “Matrix-based security configuration”).

Figure 7.16. Matrix-based security configuration

The special “anonymous” user is always present in the table. This user represents unauthenticated users. Typically, you only grant very limited rights to unauthenticated users, such as read-only access, or no access at all (as shown in Figure 7.16, “Matrix-based security configuration”).

The first thing you need to do now is to grant administration rights to your administrator. Enter your administration user in the “User/group to add” field and click on Add. Your administrator will now appear in the permissions matrix. Now, make sure you grant this user every permission (see Figure 7.17, “Setting up an administrator”), and save your configuration. You should now be able to log in with your administrator account (if you aren’t already logged in with this account) and continue to set up your other users.

Figure 7.17. Setting up an administrator

7.5.1.2. Fine-tuning user permissions

Once you’ve set up your administrator account, you can add any other users that need to access your Jenkins instance. Simply add the user names and tick the permissions you want to grant them (see Figure 7.18, “Setting up other users”). If you’re using an LDAP server or Unix users and groups as the underlying authentication schema (see Section 7.4.2, “Using an LDAP Repository”), you can also configure permissions for groups of users.

Figure 7.18. Setting up other users

You can grant a range of permissions, which are organized into several groups: Overall, Slave, Job, Run, View, and SCM. Most of the permissions are fairly obvious, but some need a little more explanation. The individual permissions are as follows:

Overall

This group covers basic system-wide permissions:

Administer

Lets a user make system-wide configuration changes and other sensitive operations, for example, in the main Jenkins configuration pages. This should be reserved for the Jenkins administrator.

Read

Provides read-only access to virtually all of the pages in Jenkins. If you want anonymous users to be able to view build jobs freely, but not to be able to modify or start them, grant the Read role to the special “anonymous” user. If not, simply revoke this permission for the Anonymous user. If you want all authenticated users to be able to see build jobs, then add a special user called “authenticated”, and grant this user Overall/Read permission.

Slave

This group covers permissions about remote build nodes, or slaves:

Configure

Create and configure new build nodes.

Delete

Delete build nodes.

Job

This group covers job-related permissions:

Create

Create a new build job.

Delete

Delete an existing build job.

Configure

Update the configuration of an existing build job.

Read

View build jobs.

Build

Start a build job.

Workspace

View and download the workspace contents for a build job. Remember, the workspace contains source code and artifacts, so if you want to protect these from general access, you should revoke this permission.

Release

Start a Maven release for a project configured with the M2Release plugin.

Run

This group covers rights related to particular builds in the build history:

Delete

Delete a build from the build history.

Update

Update the description and other properties of a build in the build history. This can be useful if a user wants to leave a note about the cause of a build failure, for example.

View

This group covers managing views:

Create

Create a new view.

Delete

Delete an existing view.

Configure

Configure an existing view.

SCM

Permissions related to your version control system:

Tag

Create a new tag in the source code repository for a given build.

Others

There can also be other permissions available, depending on the plugins installed. One useful one is:

Promote

If the Promoted Builds plugin is installed, allow users to manually promote a build.

7.5.1.3. Help! I've locked myself out!

It may happen that, during this process, you might lock yourself out of Jenkins. This can happen if, for example, you save the matrix configuration without having correctly set up your administrator. If this

happens, don't panic—there's an easy fix, as long as you have access to Jenkins's home directory. Simply open up the `config.xml` file at the root of the Jenkins home directory. This will contain something like this:

```
<hudson>
  <version>1.391</version>
  <numExecutors>2</numExecutors>
  <mode>NORMAL</mode>
  <useSecurity>true</useSecurity>
  ...
```

The thing to look for is the `<useSecurity>` element. To restore your access to Jenkins, change this value to `false`, and restart your server. You'll now be able to access Jenkins again, and set up your security configuration correctly.

7.5.2. Project-based Security

Project-based security lets you build on the matrix-based security model we just discussed, and apply it to individual projects. Not only can you assign system-wide roles for your users, you can also configure more specific rights for certain individual projects.

To activate project-level security, select “Project-based Matrix Authorization Strategy” in the Authorization section of the main configuration screen (see Figure 7.19, “Project-based security”). Here, you set up the default rights for users and groups, as you saw with Matrix-based security (see Section 7.5.1, “Matrix-based Security”).

Figure 7.19. Project-based security

These are the default permissions that apply to all projects that haven't been specially configured. However, when you use project-based security, you can also set up special project-specific permissions. You do this by selecting “Enable project-based security” in the project configuration screen (see Figure 7.20, “Configuring project-based security”). Jenkins will display a table of project-specific permissions. You can configure these permissions for different users and groups, just like on the system-wide configuration page. These permissions will be added to the system-wide permissions to produce a project-specific set of permissions applicable for this project.

Figure 7.20. Configuring project-based security

The way this works is easiest to understand with a few practical examples. In Figure 7.19, “Project-based security”, for instance, no permissions have been granted to the anonymous user, so by default all build jobs will remain invisible until a user signs on. However, we're using project-based security, so you can override this on a project-by-project basis. In Figure 7.20, “Configuring project-based security”, for example, we've set up the game-of-life project to have read-only access for the special “anonymous” user.

When you save this configuration, unauthenticated users will be able to see the game-of-life project in read-only mode (see Figure 7.21, “Viewing a project”). This same principle applies with all of the project-specific permissions.

Figure 7.21. Viewing a project

Note that Jenkins permissions are cumulative—at the time of writing, there's no way to revoke a system-wide permission for a particular project. For example, if the anonymous user has read-access to build jobs at the system level, you can't revoke read-only access for an individual project. So, when using project-based security, use the system level matrix to define minimum default permissions applicable across all of your projects, and set up projects with additional project-specific authorizations.

There are many approaches to managing project permissions, and they depend as much on organizational culture as on technical considerations. One common strategy approach is to allow team members to have full access to their own projects, and read-only access to other projects. The Extended Read Permission plugin is a useful extension to have for this scenario. This plugin lets you let users from other teams see a read-only view of your project configuration, without being able to modify anything (see Figure 7.22, “Setting up Extended Read Permissions”). This is a great way to share build configuration practices and tips with other teams without letting them tamper with your builds.

Figure 7.22. Setting up Extended Read Permissions

It's worth noting that, whenever large and/or multiple teams are involved, the internal Jenkins database reaches its limits quite quickly. If this happens, it's worth considering integrating with a more specialized directory service such as an LDAP server, Active Directory, or Atlassian Crowd, or possibly a more sophisticated permission system such as role-based security, discussed in the following section.

7.5.3. Role-based Security

Sometimes managing user permissions individually can be cumbersome, and you may not want to integrate with an LDAP server to set up groups that way. A more recent alternative option is to use the Role Strategy plugin, which allows you to define global and project-level roles, and assign these roles to users.

You install the plugin in the usual way, via the Plugin Manager. Once installed, you can activate this authorization strategy in the main configuration page (see Figure 7.23, “Setting up Role-based security”).

Figure 7.23. Setting up Role-based security

Once you've set this up, you can define roles that regroup sets of related permissions. You set up and configure your roles, and assign these roles to your users, in the Manage Roles screen, which you can access in the Manage Jenkins screen (see Figure 7.24, “The Manage Roles configuration menu”).

Figure 7.24. The Manage Roles configuration menu

In the Manage Roles screen, you can set up global and project-level permissions. Global permissions apply across all projects, and are typically system-wide administration or general access permissions (see Figure 7.25, “Managing global roles”). Setting these roles up is intuitive, and is similar to setting up user permissions in the other security models we've seen.

Figure 7.25. Managing global roles

Project roles are slightly more complicated. A project role regroups a set of permissions that are applicable to one or more (presumably related) projects. You define the relevant projects using a regular expression, so it helps to have a clear and consistent set of naming conventions in place for your project names (see Figure 7.26, “Managing project roles”). For example, you may wish to create roles distinguishing developers with full configuration rights for their own project from users who can simply trigger a build and view the build results. Or, create roles where developers can configure certain automated deployment build jobs, but only production teams are allowed to execute these jobs.

Figure 7.26. Managing project roles

Once you've defined these roles, you can go to the Assign Roles screen to set up individual users or groups with these roles (see Figure 7.27, “Assigning roles to users”).

Figure 7.27. Assigning roles to users

Role-based strategy is relatively new in Jenkins, but it's an excellent way to simplify the task of managing permissions in large, multiteam, and multiproject organizations.

7.6. Auditing—Keeping Track of User Actions

In addition to configuring user accounts and access rights, it can also be useful to keep track of the individual user actions: in other words, who did what to your server configuration. This sort of audit trail facility is even required in many organizations.

There are two Jenkins plugins that can help you do this. The Audit Trail plugin keeps a record of user changes in a special log file. The JobConfigHistory plugin lets you keep a copy of previous versions of the various system and job configuration files that Jenkins uses.

The Audit Trail Plugin keeps track of the main user actions in a set of rolling log files. To set this up, go to the Plugin Manager page and select the Audit Trail plugin in the list of available plugins. Then, as usual, click on Install and restart Jenkins once the plugin has been downloaded.

You can set up the audit trail configuration in the Audit Trail section of the main Jenkins configuration page (see Figure 7.28, “Configuring the Audit Trail plugin”). The most important field is the Log Location, which is where you specify the directory in which the log files are to be written. The audit trail is designed to produce system-style log files, which are often placed in a special system directory such as `/var/log`. You can also configure the number of log files to be maintained, and the (approximate) maximum size of each file. The simplest option is to provide an absolute path (such as `/var/log/hudson.log`), in which case Jenkins will write to log files with names like `/var/log/hudson.log.1`, `/var/log/hudson.log.2`, and so forth. Of course, you need to ensure that the user running your Jenkins instance is allowed to write to this directory.

Figure 7.28. Configuring the Audit Trail plugin

You can also use the format defined in the Java logging `FileHandler`¹ class for more control over the generated log files. In this format, you can insert variables such as `%h`, for the current user’s home directory, and `%t`, for the system temporary directory, to build a more dynamic file path.

By default, the details recorded in the audit logs are fairly sparse—they effectively record key actions performed, such as creating, modifying, or deleting job configurations or views, and the user who performed the actions. The log also shows how individual build jobs started. An extract from the default log is shown here:

```
Dec 27, 2010 9:16:08 AM /job/game-of-life/configSubmit by johnsmart
Dec 27, 2010 9:16:42 AM /view/All/createItem by johnsmart
Dec 27, 2010 9:16:57 AM /job/game-of-life-prod-deployment/doDelete by johnsmart
Dec 27, 2010 9:24:38 AM job/game-of-life/ #177 Started by user johnsmart
Dec 27, 2010 9:25:57 AM job/game-of-life-acceptance-tests/ #107 Started by upstream
    project "game-of-life" build number 177
Dec 27, 2010 9:25:58 AM job/game-of-life-functional-tests/ #7 Started by upstream
    project "game-of-life" build number 177
Dec 27, 2010 9:28:15 AM /configSubmit by johnsmart
```

This audit trail is certainly useful, especially from a system administration perspective. However, it doesn’t provide any information about the exact changes that were made to the Jenkins configuration. Nevertheless, one of the most important reasons to keep track of user actions in Jenkins is to keep tabs on exactly what changes were made to build job configurations. When something goes wrong, it can be useful to know what changes were done and so be able to undo them. The `JobConfigHistory` plugin lets you do just this.

The `JobConfigHistory` plugin is a powerful tool that lets you keep a full history of changes made to both job and system configuration files. You install it from the Plugin Manager in the usual way. Once

¹ <http://download.oracle.com/javase/1.5.0/docs/api/java/util/logging/FileHandler.html>

installed, you can fine-tune the job history configuration in the Manage Jenkins screen (see Figure 7.29, “Setting up Job Configuration History”).

Figure 7.29. Setting up Job Configuration History

Here, you can configure a number of useful nonstandard options. In particular, you should specify a directory where Jenkins can store configuration history in the “Root history folder” field. This is the directory where Jenkins will store a record of both system-related and job-related configuration changes. It can be either an absolute directory (such as `/var/hudson/history`), or a relative directory, calculated from the Jenkins home directory (see Section 3.4, “The Jenkins Home Directory”). If you don’t do this, job configuration history will be stored with the jobs, and will be lost if you delete a job.

There are a few other useful options in the Advanced section. The “Save system configuration changes” checkbox lets you keep track of system-wide configuration updates, and not just job-specific ones. In addition, the “Don’t save duplicate history” checkbox allows you to avoid recording configuration updates if no actual changes have been made. If not, a new version of the configuration will be recorded, even if you’ve only pressed the Save button without making any changes. Jenkins can also cause this to happen internally—for example, system configuration settings are all saved whenever the main configuration page is saved, even if no changes have been made.

Once you’ve set up this plugin, you can access the configuration history both for the whole server, including system configuration updates, as well as the changes made to the configuration of each project. In both cases, you can view these changes by clicking on the Job Config History icon to the right of the screen. Clicking on this icon from the Jenkins dashboard will display a view of all of your configuration history, including job changes and system-wide changes (see Figure 7.30, “Viewing Job Configuration History”).

Figure 7.30. Viewing Job Configuration History

If you click on a system-wide change (indicated by the “(system)” suffix in the list), Jenkins takes you to a screen that lists all of the versions of that file, and allows you to view the differences between the different versions (see Figure 7.31, “Viewing differences in Job Configuration History”). The differences are displayed as diff files, which isn’t particularly readable in itself. However, for small changes, the readable XML format of most of the Jenkins configuration files makes this sufficient to understand what changes were made.

Figure 7.31. Viewing differences in Job Configuration History

The JobConfigHistory plugin is a powerful tool. However, at the time of writing, it does have its limits. As mentioned, the plugin only displays the differences in raw `diff` format, and you can’t

restore a previous version of a configuration file (doing this out of context could be dangerous in some circumstances, particularly for system-wide configuration files). Nevertheless, it gives a very clear picture of the changes that have been made, both to your build jobs and to your system configuration.

7.7. Conclusion

In this chapter we've looked at a variety of ways to configure security in Jenkins. The Jenkins security model, with the two orthogonal concepts of Authentication and Authorization, is flexible and extensible. For a Jenkins installation of any size, you should try to integrate your Jenkins security strategy with the organization as a whole. This can go from simply integrating with your local LDAP repository to setting up or using a full-blown SSO solution such as Crown or CAS. In either case, it will make the system considerably easier to administrate in the long run.

Chapter 8. Notification

8.1. Introduction

While it's important to get your build server building your software, it's even more important to get your build server to let people know when it can't do so. A crucial part of the value proposition of any CI environment is to improve the flow of information about the health of your project, be it failing unit tests or regressions in the integration test suite, or other quality related issues such as a drop in code coverage or code quality metrics. In all cases, a CI server must let the right people know about any new issues, and it must be able to do so fast. This is what we call Notification.

There are two main classes of notification strategies, which we call passive and active (or pull/push). Passive notification (pull) requires the developers to consciously consult the latest build status, including RSS feeds, build radiators, and (to a certain extent) email. Active notification (push) pro-actively alerts the developers when a build fails, and includes methods such as desktop notifiers, chat, and SMS. Both approaches have their good and bad points. Passive notification strategies, such as build radiators, can raise general awareness about failed builds, and help install a team culture where fixing broken builds takes a high priority. More direct forms of notification can actively encourage developers to take matters into their own hands and fix broken builds more quickly.

8.2. Email Notification

Email notification is the most obvious and most common form of CI notification. Email is well-known, ubiquitous, easy to use, and easy to configure (see Section 4.8, “Configuring the Mail Server”). So, when teams set up their first CI environment, it's usually the most common initial notification strategy they try.

You activate email notification in Jenkins by ticking the E-mail Notification checkbox and providing the list of email addresses of the people who need to be notified (see Figure 8.1, “Configuring email notification”). By default, Jenkins will send an email message for every failed or unstable build. Remember, it will also send a new email for the first successful build after a series of failed or unstable builds, to indicate that the issue has been fixed.

Figure 8.1. Configuring email notification

Normally, a build shouldn't take too many tries to get working again—developers should diagnose and reproduce the issue locally, fix it locally, and only then commit their fix to version control. Repeated build failures usually indicate either a chronic configuration issue or poor developer practices (for example, developers committing changes without checking that they work locally first).

You can also opt to send a separate email message to any developers who have committed changes to the broken build. This is generally a good idea, as developers who have committed changes since the

last build are naturally the people who should be the most interested in the build results. Jenkins will get the email address of the user from the currently-configured security realm (see Section 7.4, “Security Realms—Identifying Jenkins Users”), or by deriving the email address from the SCM username if you’ve set this up (see Section 4.8, “Configuring the Mail Server”).

If you use this option, it may be less useful to include the entire team in the main distribution list. You may want to simply include people who will be interested in monitoring the result of every build (such as technical leads), and let Jenkins inform contributing developers directly.

This assumes, of course, that the changes caused the build failure, which is generally (but not always) the case. However, if the builds are infrequent (for example, nightly builds, or if a build is queued for several hours before finally kicking off), many changes may have been committed, and it’s hard to know which one was actually responsible for the build failure.

Not all builds are alike when it comes to email notification. Developers committing changes are particularly interested in the results of the unit and integration test builds (especially those triggered by their own changes), whereas BAs and testers might be more interested in keeping tabs on the status of the automated acceptance tests. So, the exact email notification setup for each build job will be different. In fact, it is useful to define an email notification strategy. A sample of such an email notification strategy is outlined here:

- Fast builds (unit/integration tests, runs in less than 5 minutes): notification is sent to the team lead and to developers having committed changes.
- Slow builds (acceptance test builds, run after the fast builds): notification is sent to team lead, testers, and developers having committed changes.
- Nightly builds (QA metrics, performance tests, and so on; only run if the other builds work): all team members—these provide a snapshot picture of project health before the daily status meeting.

Indeed, you should consider what notification strategy is appropriate for each build job on a case-by-case basis, rather than applying a blanket policy for all build jobs.

8.3. More Advanced Email Notification

By default, Jenkins email notification is a rather blunt tool. Notification messages are always sent to basically the same group of people. You can’t send messages to different people depending on what went wrong, or implement any sort of escalation policy. It would be useful, for example, to be able to notify the developers who committed changes the first time a build breaks, and send a different message to the team lead or the entire team if the build breaks a second time

The Email-ext plugin lets you define a more refined email notification strategy. This plugin adds an Editable Email Notification checkbox (see Figure 8.2, “Configuring advanced email notification”), which effectively replaces the standard Jenkins email notification. Here, you can define a default recipient list and fine-tune the contents of the email message, and also define a more precise notification

strategy with different messages and recipient lists for different events. Note that once you've installed and configured this plugin for your build job, you can deactivate the normal E-mail Notification configuration.

Figure 8.2. Configuring advanced email notification

This plugin has two related but distinct functions. Firstly, it lets you customize the email notification message. You can choose from a large number of predefined tokens to create your own customized message title and body. You include a token in your message template using the familiar dollar notation (e.g., `${BUILD_NUMBER}` or `$BUILD_NUMBER`). Some of the tokens accept parameters, which you can specify using a `name=value` format (e.g., `${BUILD_LOG, maxLines=100}` or `${ENV, var="PATH"}`). Among the more useful tokens are:

`${DEFAULT_SUBJECT}`

The default email subject configured in the Jenkins system configuration page

`${DEFAULT_CONTENT}`

The default email content configured in the Jenkins system configuration page

`${PROJECT_NAME}`

The project's name

`${BUILD_NUMBER}`

Current build number

`${BUILD_STATUS}`

Current build status (failing, success, etc.)

`${CAUSE}`

The cause of the build

`${BUILD_URL}`

A link to the corresponding build job page on Jenkins

`${FAILED_TESTS}`

Information about failing unit tests, if any have failed

`${CHANGES}`

Changes made since the last build

`${CHANGES_SINCE_LAST_SUCCESS}`

Changes made since the last successful build

You can get a full list of the available tokens, and the options for those that accept parameters, by clicking on the Help icon opposite the Context Token Reference label.

The Advanced button lets you define a more sophisticated notification strategy, based on the concept of triggers (see Figure 8.3, “Configuring email notification triggers”). Triggers determine when email notification messages should be sent out. The supported triggers include the following:

Failure

Any time the build fails.

Still Failing

Any successive build failures.

Unstable

Any time a build is unstable.

Still Unstable

Any successive unstable builds.

Success

Any successful build.

Fixed

When the build changes from Failure or Unstable to Successful.

Before Build

Sent before every build begins.

Figure 8.3. Configuring email notification triggers

You can set up as many (or as few) triggers as you like. The recipients list and message template can be customized for each trigger—for example, by using the Still Failing and Still Unstable triggers, you can set up a notification strategy that only notifies the developers who committed changes the first time a build job fails, but proceeds to notify the team leader if it fails a second time. You can choose to send the message only to the developers who have committed to this build (“Send to committers”), or to also include everyone who has committed since the last successful build. This helps ensure that everyone who may be involved in causing the build to break will be notified appropriately.

You can also customize the content of the message by clicking on the More Configuration option (as shown for the Still Failing trigger in Figure 8.3, “Configuring email notification triggers”). This way, you can customize different messages to be sent for different occasions.

The triggers interact intelligently. So, if you configure both the Failing and the Still Failing triggers, only the Still Failing trigger will be activated on the second build failure.

An example of such a customized message is illustrated in Figure 8.4, “Customized notification message”.

Figure 8.4. Customized notification message

Overall, however, as a notification strategy, email isn't without its faults. Some developers shut down their email clients at times to avoid being interrupted. In large organizations, the number of email messages arriving each day can be considerable, and build failure notifications can be hidden among a host of other less important messages. So, build failures may not always get the high-priority attention they require in a finely-tuned CI environment. In the following sections, we will look at some other notification strategies that can be used to raise team awareness of failed builds and encourage developers to get them fixed faster.

8.4. Claiming Builds

When a build does fail, it can be useful to know that someone has spotted the issue and is working on it. This avoids having more than one developer waste time by trying to fix the same problem separately.

The Claim plugin lets developers indicate that they've taken ownership of the broken build, and are attempting to fix it. You install this plugin in the usual way. Once installed, developers can claim a failed build as their own, and optionally, add a comment to explain the suspected cause of the build and what the developer intends to do about it. The claimed build will then be marked as such in the build history, so that fellow developers can avoid wasting time with unnecessary investigation.

To activate claiming for a build job, you need to tick the “Allow broken build claiming” option in the build job configuration page. From this point on, you'll be able to claim a broken build in the build details page (see Figure 8.5, “Claiming a failed build”). Claimed builds will display an icon in the build history indicating that they've been claimed. You can also make a build claim “sticky,” so that all subsequent build failures for this job will also be automatically claimed by this developer, until the issue is resolved.

Figure 8.5. Claiming a failed build

8.5. RSS Feeds

Jenkins also provides convenient RSS feeds for its build results, both for overall build results across all of your builds (or just the builds on a particular view), or build results for a specific build. RSS Feed icons are available at the bottom of build dashboards (see Figure 8.6, “RSS Feeds in Jenkins”) and at the bottom of the build history panel in the individual build jobs, giving you access to either all of the build results, or just the failing builds.

Figure 8.6. RSS Feeds in Jenkins

The URLs for RSS feeds are simple, and work for any Jenkins page displaying a set of build results. You just need to append `/rssAll` to get an RSS feed of all of the build results on a page, `/rssFailed` to only get the failing builds, or `/rssLatest` to get only the latest build results. But, the simplest way to obtain the URL is just to click on the RSS icon in the corresponding Jenkins screen.

There are an abundance of RSS readers out there, both commercial and open source, available for virtually every platform and device, so this can be a great choice to keep tabs on build results. Many common browsers (Firefox in particular) and email clients also support RSS feeds. Some readers have trouble with authentication, however, so if your Jenkins instance is secured, you may need to do a little extra configuration to see your build results.

RSS feeds can be a great information source on overall build results, letting you see the state of your builds at a glance without having to connect to the server. Nevertheless, most RSS Readers are by nature passive devices—you can consult the state of your builds, but the RSS reader will usually not be able to prompt you if a new build failure occurs.

8.6. Build Radiators

The concept of information radiators is commonly used in Agile circles. According to Agile guru Alistair Cockburn:

An Information radiator is a display posted in a place where people can see it as they work or walk by. It shows readers information they care about without having to ask anyone a question. This means more communication with fewer interruptions.

In the context of a CI server, an information radiator is a prominent device or display that allows team members and others to easily see if any builds are currently broken. It typically shows either a summary of all the current build results, or just the failing ones, and is displayed on a large, prominently located wall-mounted flat screen. This sort of specialized information radiator is often known as a build radiator.

When used well, build radiators are among the most effective of the passive notification strategies. They're very effective at ensuring that everyone is aware of failing builds. In addition, unlike some of the Extreme Feedback Devices that we discuss later on in this chapter, a build radiator can summarize many build jobs, including many failing build jobs, and so can still be effectively used in a multiteam context.

There are several build radiator solutions for Jenkins. One of the easiest to use is the Jenkins Radiator View plugin. This plugin adds a new type of job that you can create: the (see Figure 8.7, “Creating a build radiator view”).

Figure 8.7. Creating a build radiator view

Configuring the build radiator view is similar to configuring the more conventional list views—you just specify the build jobs you want included in the view, either by choosing them individually or by using a regular expression.

Since the build radiator view takes up the entire screen, modifying or deleting a build radiator is a bit tricky. In fact, the only way to open the view configuration screen is to append `/configure` to the view URL: so if your build radiator is called “build-radiator,” you can edit the view configuration by opening <http://my.hudson.server/view/build-radiator/configure>.

The build radiator view (see Figure 8.8, “Displaying a build radiator view”) displays a large red or yellow box for each failing or unstable build, with the build job name in prominent letters, as well as other details. You can configure the build radiator view to display passing builds as well as failing ones (they’ll be displayed in small green boxes). However, a good build radiator should really only display the failing builds, unless all the builds are passing.

Figure 8.8. Displaying a build radiator view

8.7. Instant Messaging

Instant Messaging (or IM) is widely used today as a fast, lightweight medium for both professional and personal communication. IM is, well, instant, which gives it an edge over email when it comes to fast notification. It’s also “push” rather than “pull”—when you receive a message, it will pop up on your screen and demand your attention. This makes it a little harder to ignore or put off than a simple email message.

Jenkins provides good support for notification via IM. The Instant Messaging plugin provides generic support for communicating with Jenkins using IM. Protocol-specific plugins can then be added for the various IM protocols such as Jabber and IRC.

8.7.1. IM Notification with Jabber

Many IM servers today are based on Jabber, an open source, XML-based instant messaging protocol. Jenkins provides good support for Jabber, so that developers can receive real-time notification of build failures. In addition, the plugin runs an IM bot that listens to the chat channels and lets developers run commands on the Jenkins server via chat messages.

Setting up IM support in Jenkins is straightforward. First, you need to install both the Jenkins instant-messaging plugin and the Jenkins Jabber notifier plugin using the standard plugin manager page and restart Jenkins (see Figure 8.9, “Installing the Jenkins IM plugins”).

Figure 8.9. Installing the Jenkins IM plugins

Once this is done, you need to configure your Instant Messaging server. Any Jabber server will do. You can use a public service like Google Chat, or set up your own IM server locally (the Java-based open

source chat server OpenFire¹ is a good choice). Using a public service for internal communications may be frowned upon by system administrators, and you may have difficulty getting through corporate fire walls. Setting up your own internal chat service, on the other hand, makes great sense for a development team or organization in general, as it provides another channel of communication that works well for technical questions or comments between developers. The following examples will be using a local OpenFire server, but the general approach will work for any Jabber-compatible server.

The first step involves creating a dedicated account on your Jabber server for Jenkins. This is just an ordinary chat account, but it needs to be distinct from your developer accounts (see Figure 8.10, “Jenkins needs its own dedicated IM user account”).

Figure 8.10. Jenkins needs its own dedicated IM user account

Once you've set up an IM account, you need to configure Jenkins to send IM notifications via this account. Go to the main configuration page and tick the Enable Jabber Notification checkbox (see Figure 8.11, “Setting up basic Jabber notification in Jenkins”). Here, you provide the Jabber ID and password for your IM account. Jenkins can usually figure out the IM server from the Jabber ID (if it's different, you can override this in the Advanced options). If you're using group chat rooms (another useful communication strategy for development teams), you can provide the name of these chat rooms here too. This way, Jenkins will be able to process instructions posted into the chat rooms as well as those received as direct messages.

Figure 8.11. Setting up basic Jabber notification in Jenkins

This is all you need for a basic setup. However, you may need to provide some extra details in the Advanced sector for details that are specific to your installation (see Figure 8.12, “Advanced Jabber configuration”). Here, you can specify the name and port of your Jabber server, if these can't be derived from the Jenkins Jabber ID. You can also provide a default suffix that can be applied to Jenkins user IDs to generate the corresponding Jabber IDs. Most importantly, if you've secured your Jenkins server, you'll need to provide a proper Jenkins username and password so that the IM bot can respond to instructions correctly.

Figure 8.12. Advanced Jabber configuration

Once this is configured, you need to set up a Jabber notification strategy for each of your build jobs. Open the build job configuration page and click on the Jabber Notification option.

¹ <http://www.igniterealtime.org/projects/openfire/index.jsp>

First of all, you define a recipient list for the messages. You can send messages to individuals (just use the corresponding Jabber ID, such as `joe@jabber.acme.com`) or to chat rooms that you've set up. For chat rooms, you normally need to add a “*” to the start of the chat room ID (e.g., “*gameoflife@conference.jabber.acme.org”). However, if the chat room ID contains “@conference.”, Jenkins will figure out that it's a chat room and append the “*” automatically. The chat room approach is more flexible, though you do have to trust developers to be connected permanently to the chat room for this strategy to be truly effective.

You also need to define a notification strategy to determine which build results will cause a message to be sent out. Options include:

all

Send a notification for every build.

failure

Only send notifications for failed or unstable builds.

failure and fixed

Send notifications for every failed or unstable build, and the first successful build following a failed or unstable one.

change

Send notification whenever the build outcome changes.

If you're using chat rooms, you can also ask Jenkins to send notifications to the chat rooms whenever a build starts (using the “Notify on build starts” option).

For SCM-triggered builds, Jenkins can also notify additional recipients, using the default suffix discussed earlier to build the Jabber ID from the SCM username. You can opt to notify:

SCM committers

All users having committed changes for the current build, and therefore suspected of breaking the build.

SCM culprits

SCM committers of all builds since the last successful one.

SCM fixers

SCM committers to the first successful build after a failed or unstable one.

Upstream committers

Also notifies committers to upstream builds as well as the current one. This works automatically for Maven build jobs, but needs fingerprinting to be activated for other build types.

At the time of writing, you can only have one notification strategy, so some of the advanced options you saw in Section 8.3, “More Advanced Email Notification” aren't yet possible with IM.

Developers will be notified via their favorite IM client (see Figure 8.13, “Jenkins Jabber messages in action”). Developers can also interact with the build server via the chat session, using a set of simple commands. Some examples of a few of the more useful commands are shown here:

- `!build game-of-life`—Start the game-of-life build immediately.
- `!build game-of-life 15m`—Start the game-of-life build in 15 minutes.
- `!comment game-of-life 207 'oops'`—Add a build description to a given build.
- `!status game-of-life`—display the status of the latest build for this build job.
- `!testresult game-of-life`—display the full test results for the latest build.
- `!health game-of-life`—display a more complete summary of the health status of the latest build.

You can get a full list of commands by sending the `!help` message to the Jenkins user.

Figure 8.13. Jenkins Jabber messages in action

8.7.2. IM Notification using IRC

Another popular form of IM is Internet Relay Chat, or IRC. IRC is traditionally focused on group discussions (though direct messaging is also supported), and is a very popular form of communication for developers, particularly in the open source world.

The Jenkins IRC plugin lets you interact with your Jenkins server via an IRC channel, both to receive notification messages and to issue commands to the server. Like the Jabber plugin, you also need to install the Instant Messaging plugin for this to work.

8.8. IRC Notification

Contributed by Juven Xu

Internet Relay Chat (or IRC) is a popular form of instant messaging, primarily designed for group communication via channels. For example, Jenkins has a channel set up on Freenode² so users and developers can discuss Jenkins related topics. You'll see many users ask questions and, most of the time, more experienced users will be prompt in providing useful answers.

Just like IM through Jabber, you can configure Jenkins to “push” notification through IRC. Some IRC clients such as xchat³ support alert configuration so that when the message arrives, it can blink the tray

² <http://jenkins-ci.org/content/chat>

³ <http://xchat.org/>

icon or make a beep sound. To set up IRC support on Jenkins, first you need to install the IRC plugin⁴ and the Instance Messaging plugin⁵. Simply go to the standard plugin manager, tick their checkboxes, and then restart Jenkins (see Figure 8.14, “Install the Jenkins IRC plugins”).

Figure 8.14. Install the Jenkins IRC plugins

Once it’s installed, you need to enable the IRC plugin and configure it to fit into your own environment. Basically, this involves providing the hostname and port of the IRC server you’re using, a dedicated IRC channel, and a nickname for the IRC plugin. It’s good practice to set up a dedicated channel for CI notification, so as people chat in other channels, they won’t be disturbed. You may also want to configure extra details in the Advanced sector. All of these are available in the Configure System page (see Figure 8.15, “Advanced IRC notification configuration”).

Figure 8.15. Advanced IRC notification configuration

In addition to the hostname, port, channel, and nickname we mentioned earlier, you can also configure your IRC server or NickServ password if your environment requires them. Command prefixes need to be configured if you want to interact with the server via IRC messages. This is basically the same as using Jabber (see Section 8.7, “Instant Messaging”). Finally, you may want to let the IRC plugin use the `/notice` command instead of the default `/msg` command. `/notice` is the same as `/msg` except that the message will be contained in dashes, which will prevent a response from most robots.

Once the global configuration is ready, you can enable IRC notification for each build job and set up a notification strategy. Open the build job configuration page, go to the Post-build Actions section, and click on the IRC Notification option. If you want to set up a notification strategy rather than using the default one, click the “Advanced...” button (see Figure 8.16, “Advanced build job IRC notification configuration”).

Figure 8.16. Advanced build job IRC notification configuration

Notification strategies (when to send notification messages, and to whom) are discussed in Section 8.7, “Instant Messaging”. Both the Jabber plugin and the IRC plugin depend on the Instant Messaging Plugin, so they share a number of common core features. Some options are specific to the IRC plugin, however. Here, for example, you can define a customized channel if you don’t like the global default. Finally, for

⁴ <http://wiki.jenkins-ci.org/display/JENKINS/IRC+Plugin>

⁵ <http://wiki.jenkins-ci.org/display/JENKINS/Instant+Messaging+Plugin>

a channel notification message, you can choose what information to send in the notification messages. Your options are build summary, SCM changes, and failed tests.

Once you save the configuration, you should be good to go. Based on what you've configured, this plugin will join the appropriate IRC channels and send notification messages for build jobs.

In Figure 8.17, “IRC notification messages in action”, for example, the IRC plugin joins the #ci-book channel on freenode. First, user juven committed a change with SCM message “feature x added” and the IRC plugin let everyone on the channel know that the build was successful. Then, juven committed another change for feature y, but this time the build failed. John noticed it and fixed the build error. The IRC plugin now happily said “Yippie, build fixed!” Note that some lines in this screen are highlighted. This is because I logged in as user “juven” and I configured my XChat IRC client to highlight messages containing my nickname.

Figure 8.17. IRC notification messages in action

8.9. Desktop Notifiers

The best push notification strategies integrate smoothly into the developer's daily work environment. This is why IM can be an effective strategy if developers are already in the habit of using IM for other work-related activities.

Desktop notification tools also fall into this category. These are tools that run locally on the developer machine, either as an independent application or widget, or as part of the developer's IDE.

If you're using Eclipse, the Jenkins Eclipse plugin⁶ displays a health icon at the bottom of the Eclipse window. If you click on the icon, you can see a detailed view of the Jenkins projects (see Figure 8.18, “Jenkins notifications in Eclipse”). In the Eclipse preferences, you provide the URL of your Jenkins server along with any required authentication details. The configuration is fairly simple, however, and you can only connect to a single Jenkins instance for a given Eclipse workspace.

Figure 8.18. Jenkins notifications in Eclipse

If you're using the NetBeans IDE, you already have integration with Hudson and Jenkins. Open the Services window and add the servers under Hudson Builders. (If you open a Maven project whose `ciManagement` specifies `hudson` or `jenkins` as the `system`, the corresponding server will be registered automatically.) This integration has various features beyond build notifications in the status line, such as Test Results window integration, build log and change log display, workspace browsing, and a job setup wizard.

⁶ <http://code.google.com/p/hudson-eclipse/>

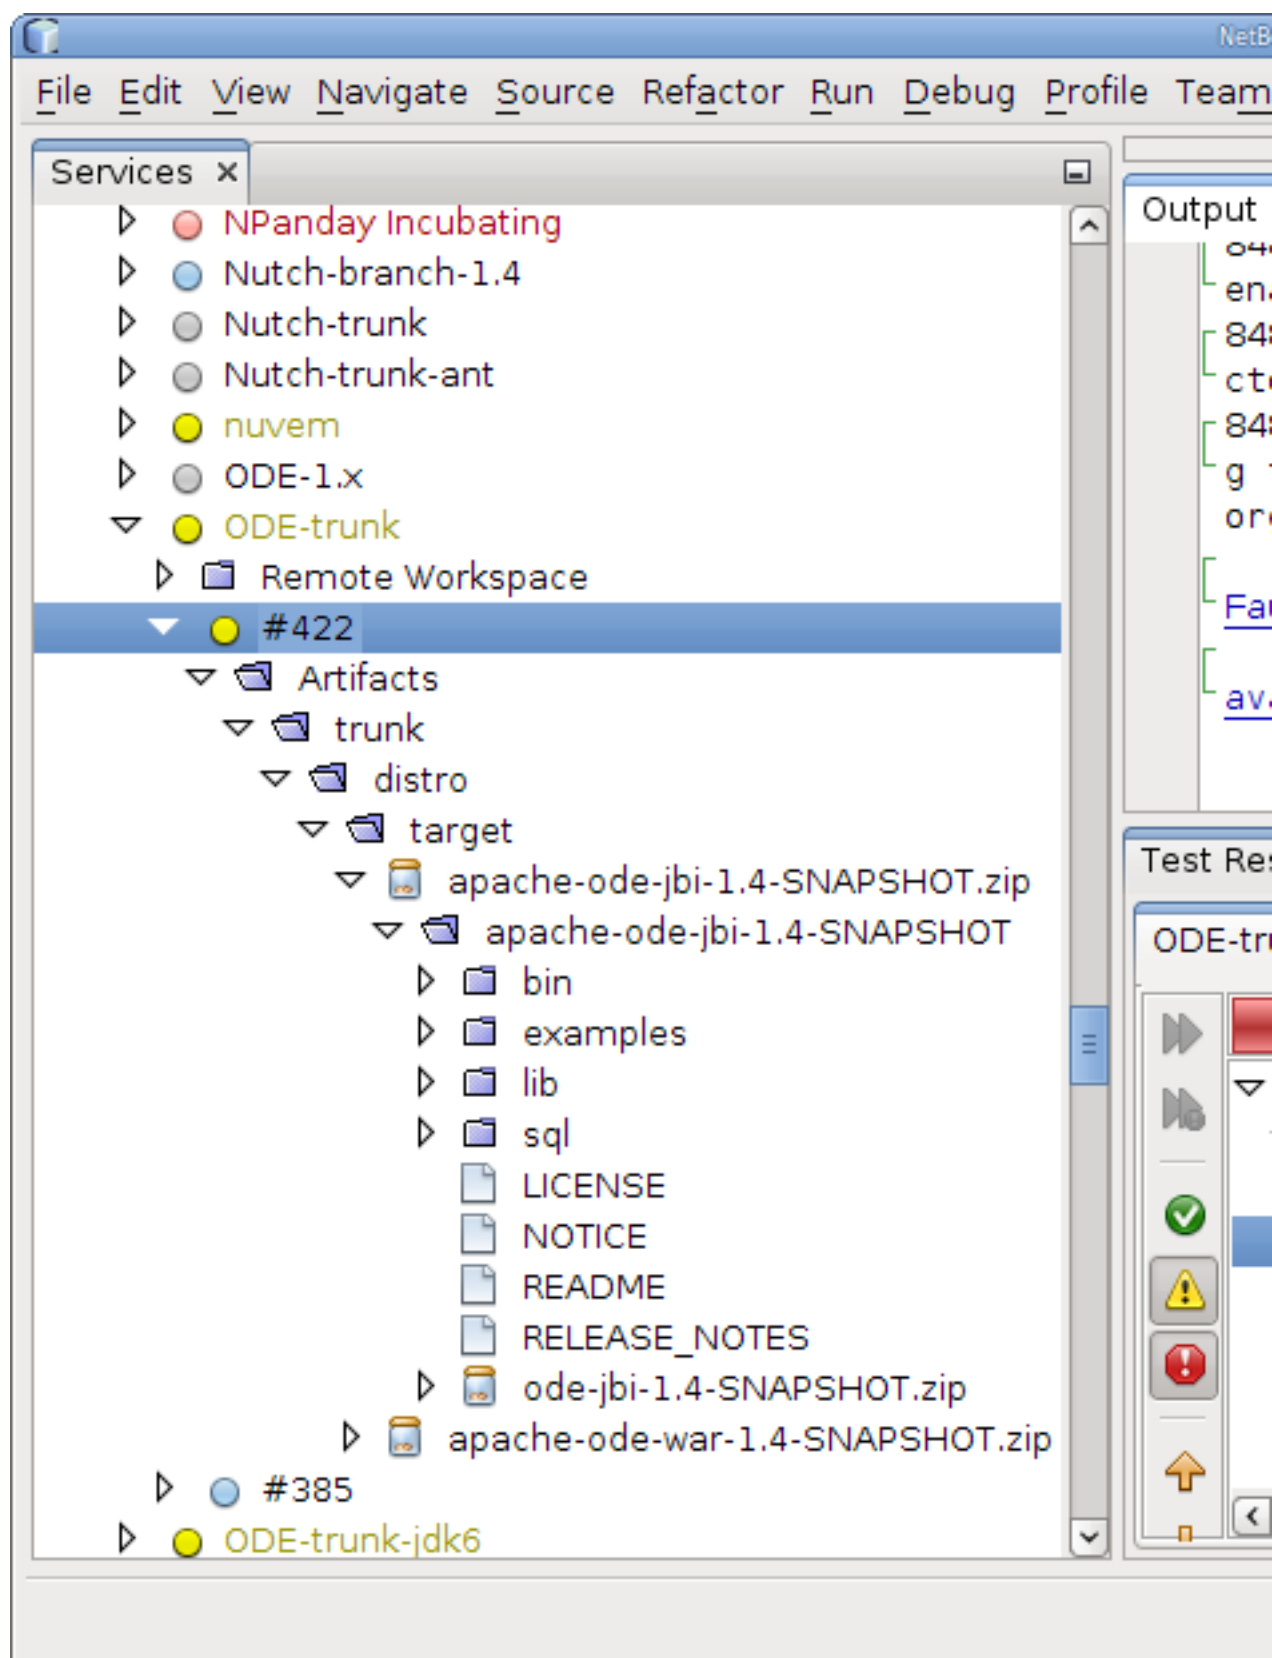


Figure 8.19. Jenkins connection in NetBeans

The Jenkins Tray Application plugin (see Figure 8.20, “Launching the Jenkins Tray Application”) lets you start up a small Java client application using Java Web Start from your Jenkins dashboard.

Figure 8.20. Launching the Jenkins Tray Application

This application sits in your system tray, lets you view the current state of your builds at a glance, and also brings up pop-up windows notifying you of new build failures (see Figure 8.21, “Running the Jenkins Tray Application”).

Figure 8.21. Running the Jenkins Tray Application

This is certainly a useful application, but it suffers from a few limitations. At the time of writing, the Jenkins Tray Application did not support accessing secured Jenkins servers. In addition, the developer needs to remember to restart it each morning. This may seem a minor issue, but in general, when it comes to notification strategies, the less you have to ask of your developers the better.

One of the best options for Jenkins desktop notification is to use a service like Notifo (see Section 8.10, “Notification via Notifo”), which provides both desktop and mobile clients. We'll see how this works in detail in the next section.

8.10. Notification via Notifo

Notifo⁷ is a fast and economical service to send real-time notifications to your smartphone or desktop. In the context of a Jenkins server, you can use it to set up free or low-cost real-time notifications for your Jenkins build results. Individual accounts (which you need to be able to receive notifications) are free. You need to set up a service account to send notification messages from your Jenkins server. This is where Notifo earns their keep, though at the time of writing a service account can send up to 10,000 notifications per month free of charge, which is usually plenty for an average Jenkins instance. One of the strong points of a real-time notification service like Notifo is that notification messages can be sent to the same users on different devices: smartphones and desk top clients, in particular.

Setting up Jenkins notifications with Notifo is relatively straightforward. First, go to the Notifo website and sign up to create an account. Each team member who wants to be notified will need their own Notifo account. They will also need to install the Notifo client on each device on which they need to receive notification messages. At the time of writing, Notifo clients were available for Windows and Mac OS X desktops, and iPhones, with support for Linux and other smartphones on the way.

Next, you need to set up a Notifo service account for your Jenkins server. You can do this with one of your developer accounts, or create a new account for the purpose. Log on to the Notifo website, and go to the My Services menu. Here, click on Create Service (see Figure 8.22, “Creating a Notifo service

⁷ <http://www.notifo.com>

for your Jenkins instance”), and fill in the fields. The most important is the Service Username, which needs to be unique. You can also specify the Site URL and the Default Notification URL to point to your Jenkins instance so that users can open the Jenkins console by clicking on the notification message.

Figure 8.22. Creating a Notifo service for your Jenkins instance

To receive notification messages from the Jenkins server, developers now need to subscribe to this service. You can then add developers to the list of subscribers in the service Subscribers page by sending them subscription requests. Once the service has been created and the users are all subscribed, you can configure your project to send out Notifo notifications (see Figure 8.23, “Configuring Notifo notifications in your Jenkins build job”). You need to provide the API username of the Jenkins service you set up, as well as the API Secret, both of which you can see in the Notifo Service Dashboard.

Figure 8.23. Configuring Notifo notifications in your Jenkins build job

Once this is set up, Jenkins will send almost real-time notifications of build failures to any Notifo clients the developer is running, whether it's on a desktop or on a mobile device (see Figure 8.24, “Receiving a Notifo notification on an iPhone”).

Figure 8.24. Receiving a Notifo notification on an iPhone

At the time of writing, sophisticated notification strategies are not supported—you just provide a list of Notifo usernames who need to be notified. Nevertheless, this remains a very effective notification tool for frontline developers.

8.11. Mobile Notification

If your Jenkins server is visible on the Internet (even if you've set up authentication on your Jenkins server), you can also monitor your builds via your iPhone or Android mobile device. The free Hudson Helper application (see Figure 8.25, “Using the Hudson Helper iPhone app”), for example, lets you list your current build jobs (either all of the build jobs on the server, or only the build jobs in a particular view). You can also view the details of a particular build job, including the current status, failing tests, and build time, and even start and stop builds.

Figure 8.25. Using the Hudson Helper iPhone app

For Android phones, you can also install the Hudson Mood widget which will also provide updates and alerts about build failures.

Note that these mobile applications rely on a data connection, so while they'll typically work well locally, you shouldn't rely on them if the developer in question is out of the country.

8.12. SMS Notification

These days SMS is another ubiquitous communication channel that has the added advantage of reaching people even when they're out of the office. For a build engineer, this can be a great way to monitor critical builds, even when developers or team leads are away from their desks.

SMS gateways⁸ are services that let you send SMS notifications via specially-formatted email addresses (for example, 123456789@mysmsgateway.com might send an SMS message to 123456789). Many mobile vendors provide this service, as do many third-party service providers. There's no built-in support for SMS Gateways in Jenkins, but the basic functionality of these gateways makes integration relatively easy: you simply add the special email addresses to the normal notification list. Alternatively, using the advanced email configuration, you can set up a separate rule containing only the SMS email addresses (see Figure 8.26, “Sending SMS notifications via an SMS Gateway Service”). Doing this makes it easier to fine-tune the message contents to suit an SMS message format.

Figure 8.26. Sending SMS notifications via an SMS Gateway Service

Once you've done this, your users will receive prompt notification of build results in the form of SMS messages (see Figure 8.27, “Receiving notification via SMS”). The main disadvantage of this approach is arguably that it isn't free, and requires the use of a third-party commercial service. That said, it's really the only notification technique capable of reaching developers when they're out of Internet range or who don't have a data-enabled smartphone. Indeed, this technique is popular among system administrators, and can be very useful for certain critical build jobs.

Figure 8.27. Receiving notification via SMS

8.13. Making Noise

If you have your Jenkins instance running on a machine that is physically located in proximity to the development team, you may also want to add sounds into the mix of notification strategies. This can be an effective strategy for small co-located teams, though it becomes trickier if the build server is set up on a virtual machine or elsewhere in the building.

There are two ways to integrate audio feedback into your build process in Jenkins: the Jenkins Sounds plugin and the Jenkins Speaks plugin. Both can be installed via the Plugin Manager page in the usual manner.

⁸ http://en.wikipedia.org/wiki/SMS_gateway

The Jenkins Sounds plugin is the most flexible of the two. It allows you to build a detailed notification strategy based on the latest build result and also (optionally) on the previous build result as well (see Figure 8.28, “Configuring Jenkins Sounds rules in a build job”). For example, you can configure Jenkins to play one sound the first time a build fails, a different sound if the build fails a second time, and yet another sound when the build is fixed.

Figure 8.28. Configuring Jenkins Sounds rules in a build job

To set this up, you need to tick the Jenkins Sounds checkbox in the Post-build Actions section of your build job configuration page. You can add as many sound configuration rules as you like. Adding a rule is simple enough. First, you need to choose which build result will trigger the sound. You also need to specify the previous build results for which this rule is applicable: Not Build (NB), Aborted (Ab), Failed (Fa), Unsuccessful (Un), or Successful (Su).

The Jenkins Sounds plugin includes a large list of pre-defined sounds, which usually offer plenty of choices for even the most discerning build administrator, but you can add your own to the list if you really want to. Sounds are stored as a ZIP or JAR file containing sound files in a flat directory structure (i.e., no subdirectories). The list of sounds shown by the plugin is simply the list of filenames, minus the extensions. The plugin supports AIFF, AU, and WAV files.

In the System Configuration page, you can give Jenkins a new sound archive file, using the `http://` notation if your sound archive file is available on a local web server, or the `file://` notation if it's available locally (see Figure 8.29, “Configuring Jenkins Sounds”). Once you've saved the configuration, you can test the sounds in your sound archive via the Test Sound button in the Advanced section.

Figure 8.29. Configuring Jenkins Sounds

The Jenkins Sounds plugin is an excellent choice if you want to complement your more conventional notification techniques. Short, recognizable sounds are a great way to grab a developer's attention and let the team know that something needs fixing. The developers will then be a bit more receptive when the more detailed notifications follow.

Another option is the Jenkins Speaks plugin. With this plugin, you can get Jenkins to broadcast a customized announcement (in a very robotic voice) when your build fails (see Figure 8.30, “Configuring Jenkins Speaks”). You can configure the exact message using Jelly. Jelly is an XML-based scripting language used widely in the lower levels of Jenkins.

Figure 8.30. Configuring Jenkins Speaks

The advantage of this approach lies in its precision: since you can use Jenkins variables in the Jelly script, you can get Jenkins to say just about anything you want about the state of the build. Here's a simple example:

```
<j:choose>
  <j:when test="${build.result!='SUCCESS'}">
    Your attention please. Project ${build.project.name} has failed
    <j:if test="${build.project.lastBuild.result!='SUCCESS'}"> again</j:if>
  </j:when>
  <j:otherwise><!-- Say nothing --></j:otherwise>
</j:choose>
```

If you leave this field blank, the plugin will use a default template that you can configure in the System Configuration page. In fact, it's usually a good idea to do this, and only to use a project-specific script if you really need to.

The disadvantage is that the robotic voice can be a little hard to understand. For this reason, it's a good idea to start your announcement with a generic phrase such as “Your attention please,” or to combine it with the Jenkins Sounds plugin, so that you have developers’ attention before the actual message is broadcast. Using hyphens in your project names (e.g., game-of-life rather than gameoflife) will also help the plugin know how to pronounce your project names.

Both these approaches are useful for small teams, but can be limited for larger ones, when the server isn't physically located in close proximity to the development team. Future versions may support playing sounds on a separate machine, but at the time of writing this feature was not available.

8.14. Extreme Feedback Devices

Many more imaginative notification tools and strategies exist, and there's plenty of scope for improvisation if you're willing to improvise with electronics a little. This includes devices such as Ambient Orbs, Lava Lamps, traffic lights, or other more exotic USB-controlled devices. The Build Radiator (see Section 8.6, “Build Radiators”) also falls into this category if you project it onto a big enough screen.

One device that integrates very nicely with Jenkins is the Nabaztag. The Nabaztag (see Figure 8.31, “A Nabaztag”) is a popular WiFi-enabled robotic rabbit that can flash colored lights, play music, or even speak. One advantage of the Nabaztag is that, since it works via WiFi, it isn't constrained to be located near the build server, and so will work even if your Jenkins instance is in a server room or on a virtual machine. As far as extreme feedback devices go, these little fellows are hard to beat.

Figure 8.31. A Nabaztag

And best of all, there's a Jenkins plugin available for the Nabaztag. Once you've installed the Nabaztag plugin and restarted Jenkins, it's easy to configure. In Jenkins's main Configuration page, go to the Global Nabaztag Settings section and enter the serial number and secret token for your electronic bunny

(see Figure 8.32, “Configuring your Nabaztag”). You can also provide some default information about how your build bunny should react to changes in build status (should it report on starting and successful builds, for example), what voice it should use, and what message it should say when a build fails, succeeds, is fixed, or fails again. Then, to activate Nabaztag notification for a particular build job, you need to tick the Nabaztag Publisher option in your build job configuration. Depending on your environment, for example, you may or may not want all of your builds to send notifications to your Nabaztag.

Figure 8.32. Configuring your Nabaztag

With the notable exception of the build radiator, many of these devices have similar limitations to the Jenkins Speaks and Jenkins Sounds plugins (see Section 8.13, “Making Noise”)—they're best suited for small, co-located teams, working on a limited number of projects. Nevertheless, when they work, they can be a useful addition to your general notification strategy.

8.15. Conclusion

Notification is a vital part of your overall CI strategy. After all, a failed build is of little use if there's no one listening. Nor is notification a one-size-fits-all affair. You need to think about your organization, and tailor your strategy to suite the local corporate culture and predominant tool set.

Indeed, it's important to define and implement a well thought-out notification strategy that suits your environment. Email, for example, is ubiquitous, so this will form the backbone of many notification strategies, but if you work in a larger team or with a busy technical lead, you may want to consider setting up an escalation strategy based on the advanced email options (see Section 8.3, “More Advanced Email Notification”). But, you should complement this with one of the more active strategies, such as instant messaging or a desktop notifier. If your team already uses a chat or IRC channel to communicate, try to integrate this into your notification strategy as well. And SMS notification is a great strategy for really critical build jobs.

You should also ensure that you have both passive and active (or pull and push) notification strategies. A prominent build radiator or an extreme feedback device, for example, sends a strong message to the team that fixing builds is a priority task, and can help install a more agile team culture.

Chapter 9. Code Quality

9.1. Introduction

Few would deny the importance of writing quality code. High quality code contains fewer bugs, and is easier to understand and maintain. However, the precise definitions of code quality can be more subjective, varying between organizations, teams, and even individuals within a team.

This is where coding standards come into play. Coding standards are rules, sometimes relatively arbitrary, that define the coding styles and conventions that are considered acceptable within a team or organization. In many cases, agreeing on a set of standards, and applying them, is more important than the standards themselves. Indeed, one of the most important aspects of quality code is that the code is easy to read and understand. If developers within a team all apply the same coding standards and practices, the code will be more readable, at least for members of that team. And if the standards are commonly used within the industry, the code will also be more readable for new developers arriving on the team.

Coding standards include both aesthetic aspects such as code layout and formatting, naming conventions, and so forth, as well as potentially bad practices such as missing curly brackets after a condition in Java. A consistent coding style lowers maintenance costs, makes code clearer and more readable, and makes it easier to work on code written by other team members.

Only an experienced developer can really judge code quality in all its aspects. That's the role of code reviews and, among other things, practices like pair programming. In particular, only a human eye can decide if a piece of code is truly well written, and if it actually does what the requirements ask of it. However, code quality metric tools can help a great deal. In fact, it's unrealistic to try to enforce coding standards without the use of such tools.

These tools analyze your application source code or byte code, and check whether the code respects certain rules. Code quality metrics can encompass many aspects of code quality, from coding standards and best practices right through to code coverage, with everything from compiler warnings to TODO comments in between. Certain metrics concentrate on measurable characteristics of your code base, such as the number of lines of code (NLOC), average code complexity, or the number of lines per class. Others focus on more sophisticated static analysis, or on looking for potential bugs or poor practices in your code.

There are a wide range of code quality reporting plugins available for Jenkins. Many are for Java static analysis tools, such as Checkstyle, PMD, FindBugs, Cobertura, and JDepend. Others, such as fxcop and NCover, are focused on .NET applications.

With all of these tools, you need to configure your build job to generate the code quality metrics data before Jenkins can produce any reports.

The notable exception to this rule is Sonar. Sonar can extract code quality metrics from any Maven project, with no additional configuration required in your Maven project. This is great when you have large numbers of existing Maven projects that you need to integrate into Jenkins, and you want to configure consistent code quality reporting across all of your projects.

In the rest of this chapter, we'll see how to set up code quality reporting in your Jenkins builds, and also how you can use it as an effective part of your build process.

9.2. Code Quality in Your Build Process

Before you look at how to report on code quality metrics in Jenkins, it can be useful to take a step back and look at the larger picture. Code Quality metrics are of limited value in isolation—they need to be part of a broader process improvement strategy.

The first level of code quality integration should be the IDE. Modern IDEs have great support for many code quality tools—Checkstyle, PMD, and FindBugs all have plugins for Eclipse, NetBeans, and IntelliJ, which provide rapid feedback for developers on code quality issues. This is a much faster and more efficient way to provide feedback for individual developers, and to teach developers about the organizational or project coding standards.

The second level is your build server. In addition to your normal unit and integration test build jobs, set up a dedicated code quality build, which runs after the normal build and test. The aim of this process is to produce project-wide code quality metrics, to keep tabs on how the project is doing as a whole, and to address any issues from a high level. The effectiveness of these reports can be increased by a weekly code quality review in which code quality issues and trends are discussed within the team.

It's important to run this job separately because code coverage analysis and static analysis tools can be quite slow to run. It's also important to keep any code coverage tests well away from builds, as the code coverage process produces instrumented code which should never be deployed to a repository for production use.

Code quality reporting is, by default, a relatively passive process. No one will know the state of the project if they don't seek out the information on the build server. While this is better than nothing, if you're serious about code quality, there's a better way. Rather than simply reporting on code quality, set up a dedicated code quality build, which runs after the normal build and test, and configure the build to fail if code quality metrics are not at an acceptable level. You can do this in Jenkins or in your build script, although one advantage of configuring this outside of your build script is that you can change code quality build failing criteria more easily without changing the project source code.

As a final word, remember that coding standards are guidelines and recommendations, not absolute rules. Use failing code quality builds and code quality reports as indicators of a possible area of improvement, not as measurements of absolute value.

9.3. Popular Java and Groovy Code Quality Analysis Tools

There are many open source tools that can help identify poor coding practices.

In the Java world, three static analysis tools have withstood the test of time, and are widely used in very complementary ways. Checkstyle excels at checking coding standards and conventions, coding practices, as well as other metrics such code complexity. PMD is a static analysis tool similar to Checkstyle, but is more focused on coding and design practices. FindBugs is an innovative tool issued from the ongoing research work of Bill Pugh and his team at the University of Maryland that focuses on identifying potentially dangerous and buggy code. And if you're working with Groovy or Grails, you can use CodeNarc, which checks Groovy coding practices and convention.

All of these tools can be easily integrated into your build process. In the following sections, we'll look at how to set up these tools to generate the XML reports that Jenkins can then use for its own reporting.

9.3.1. Checkstyle

Checkstyle¹ is a static analysis tool for Java. Originally designed to enforce a set of highly-configurable coding standards, Checkstyle now also checks for poor coding practices, as well as overly complex and duplicated code. Checkstyle is a versatile and flexible tool that should have its place in any Java-based code quality analysis strategy.

Checkstyle supports a very large number of rules, including ones relating to naming conventions, annotations, javadoc comments, class and method size, code complexity metrics, poor coding practices, and many others.

Duplicated code is another important code quality issue—duplicated or near-duplicated code is harder to maintain and to debug. Checkstyle provides some support for the detection of duplicated code, but more specialized tools such as CPD do a better job in this area.

One of the nice things about Checkstyle is how easy it is to configure. You can start off with the Sun coding conventions and tweak them to suit your needs, or start from scratch. Using the Eclipse plugin (or even directly in XML), you can pick and choose from several hundred different rules, and fine-tune the settings of the rules you do choose (see Figure 9.1, “It's easy to configure Checkstyle rules in Eclipse”). This is important, as different organizations, teams, and even projects have different requirements and preferences with regards to coding standards. It's better to have a precise set of rules that can be adhered to, rather than a broad set of rules that will be ignored. It's especially important where large legacy code bases are involved. In these cases, it's often better to start off with a more limited set of rules than to be overwhelmed with a large number of relatively minor formatting issues.

Figure 9.1. It's easy to configure Checkstyle rules in Eclipse

¹ <http://checkstyle.sourceforge.net>

Configuring Checkstyle in your build is usually straightforward. If you're using Ant, you need to download the checkstyle JAR file from the website² and make it available to Ant. You could place it in your Ant `lib` directory, but this would mean customizing the Ant installation on your build server (and any slave nodes), so it isn't a very portable solution. A better approach would be to place the Checkstyle JAR file in one of your project directories, or to use Ivy or the Maven Ant Task library to declare a dependency on Checkstyle. If you opt for keeping the Checkstyle JAR file in the project directories, you could declare the Checkstyle task as shown here:

```
<taskdef resource="checkstyletask.properties"
  classpath="lib/checkstyle-5.3-all.jar"/>
```

Then, to generate Checkstyle reports in an XML format that Jenkins can use, you could do the following:

```
<target name="checkstyle">
  <checkstyle config="src/main/config/company-checks.xml">
    <fileset dir="src/main/java" includes="**/*.java"/>
    <formatter type="plain"/>
    <formatter type="xml"/>
  </checkstyle>
</target>
```

Now, just invoke this task (e.g., `ant checkstyle`) to generate the Checkstyle reports.

In Maven 2, you could add something like the following to the `<reporting>` section:

```
<reporting>
  <plugins>
    <plugin>
      <groupId>org.apache.maven.plugins</groupId>
      <artifactId>maven-checkstyle-plugin</artifactId>
      <version>2.4</version>
      <configuration>
        <configLocation>
          src/main/config/company-checks.xml
        </configLocation>
      </configuration>
    </plugin>
  </plugins>
</reporting>
```

For a Maven 3 project, you need to add the plugin to the `<reportPlugins>` element of the `<configuration>` section of the `maven-site-plugin`:

```
<project>
  <properties>
    <sonar.url>http://buildserver.acme.org:9000</sonar.url>
  </properties>
  <build>
    ...
```

² <http://checkstyle.sourceforge.net>


```

<plugins>
...
<plugin>
  <groupId>org.apache.maven.plugins</groupId>
  <artifactId>maven-site-plugin</artifactId>
  <version>3.0-beta-2</version>
  <configuration>
    <reportPlugins>
      <plugin>
        <groupId>org.apache.maven.plugins</groupId>
        <artifactId>maven-checkstyle-plugin</artifactId>
        <version>2.4</version>
        <configuration>
          <configLocation>
            ${sonar.url}/rules_configuration/export/java/My_Rules/checkstyle.xml
          </configLocation>
        </configuration>
      </plugin>
    </reportPlugins>
  </configuration>
</plugin>
</plugins>
</build>
</project>

```

Now, running `mvn checkstyle:checkstyle` or `mvn site` will analyse your source code and generate XML reports that Jenkins can use.

Note that in the last example, you used a Checkstyle ruleset that we've uploaded to a Sonar server (defined by the `${sonar.url}` property). This strategy makes it easy to use the same set of Checkstyle rules for Eclipse, Maven, Jenkins, and Sonar.

Recent versions of Gradle also offer some integrated Checkstyle support. You can set up Checkstyle for your builds as shown here:

```

apply plugin: 'code-quality'

```

This will use the checkstyle ruleset in `config/checkstyle/checkstyle.xml` by default. You can override this with the `checkstyleConfigFileName` property: at the time of writing, however, you can't get the Gradle code quality plugin to obtain the Checkstyle rules from a URL.

You can generate the Checkstyle reports here by running `gradle checkstyleMain` or `gradle check`.

9.3.2. PMD/CPD

PMD³ is another popular static analysis tool. It focuses on potential coding problems such as unused or suboptimal code, code size and complexity, and good coding practices. Some typical rules include

³ <http://pmd.sourceforge.net>

“Empty If Statement,” “Broken Null Check,” “Avoid Deeply Nested If Statements,” “Switch Statements Should Have Default,” and “Logger Is Not Static Final.” There’s a fair amount of overlap with some of the Checkstyle rules, though PMD does have some more technical rules such as rules related to JSF and Android.

PMD also comes with CPD, a robust open source detector of duplicated and near-duplicated code.

PMD is a little less flexible than Checkstyle, though you can still pick and choose the rules you want to use in Eclipse, and then export them as an XML file (see Figure 9.2, “Configuring PMD rules in Eclipse”). You can then import this rule set into other Eclipse projects, into Sonar, or use it in your Ant or Maven builds.

Figure 9.2. Configuring PMD rules in Eclipse

PMD comes with an Ant task that you can use to generate the PMD and CPD reports. First, though, you need to define these tasks, as shown in the following example:

```
<path id="pmd.classpath">
  <pathelement location="org.apache.maven.model.Build@2bef09c0"/>
  <fileset dir="lib/pmd">
    <include name="*.jar"/>
  </fileset>
</path>

<taskdef name="pmd" classname="net.sourceforge.pmd.ant.PMDTask"
  classpathref="pmd.classpath"/>

<taskdef name="cpd" classname="net.sourceforge.pmd.cpd.CPDTask"
  classpathref="pmd.classpath"/>
```

Next, you can generate the PMD XML report by invoking the PMD task as illustrated here:

```
<target name="pmd">
  <taskdef name="pmd" classname="net.sourceforge.pmd.ant.PMDTask"
    classpathref="pmd.classpath"/>

  <pmd rulesetfiles="basic" shortFileNames="true">
    <formatter type="xml" toFile="target/pmd.xml" />
    <fileset dir="src/main/java" includes="**/*.java"/>
  </pmd>
</target>
```

And, to generate the CPD XML report, you could do something like this:

```
<target name="cpd">
  <cpd minimumTokenCount="100" format="xml" outputFile="/target/cpd.xml">
    <fileset dir="src/main/java" includes="**/*.java"/>
  </cpd>
</target>
```

```
</cpd>
</target>
```

You can place this XML ruleset in your project classpath (for example, in `src/main/resources` for a Maven project), or in a separate module (if you want to share the configuration between projects). An example of how to configure Maven 2 to generate PMD and CPD reports using an exported XML ruleset as shown here:

```
<reporting>
  <plugins>
    <plugin>
      <groupId>org.apache.maven.plugins</groupId>
      <artifactId>maven-pmd-plugin</artifactId>
      <version>2.5</version>
      <configuration>
        <!-- PMD options -->
        <targetJdk>1.6</targetJdk>
        <aggregate>true</aggregate>
        <format>xml</format>
        <rulesets>
          <ruleset>pmd-rules.xml</ruleset>
        </rulesets>

        <!-- CPD options -->
        <minimumTokens>20</minimumTokens>
        <ignoreIdentifiers>true</ignoreIdentifiers>
      </configuration>
    </plugin>
  </plugins>
</reporting>
```

If you're using Maven 3, place the plugin definition in the `<maven-site-plugin>` configuration section. This example also shows how to use a ruleset in another dependency (in this case the `pmd-rules.jar` file):

```
<project>
  ...
  <build>
    ...
    <plugins>
      ...
      <plugin>
        <groupId>org.apache.maven.plugins</groupId>
        <artifactId>maven-site-plugin</artifactId>
        <version>3.0-beta-2</version>
        <configuration>
          <reportPlugins>
            <plugin>
              <groupId>org.apache.maven.plugins</groupId>
              <artifactId>maven-pmd-plugin</artifactId>
              <version>2.5</version>
              <configuration>
                <!-- PMD options -->
```

```

        <targetJdk>1.6</targetJdk>
        <aggregate>true</aggregate>
        <format>xml</format>
        <rulesets>
            <ruleset>/pmd-rules.xml</ruleset>
        </rulesets>

        <!-- CPD options -->
        <minimumTokens>50</minimumTokens>
        <ignoreIdentifiers>true</ignoreIdentifiers>
    </configuration>
</plugin>
</reportPlugins>
</configuration>
<dependencies>
    <dependency>
        <groupId>com.wakaleo.code-quality</groupId>
        <artifactId>pmd-rules</artifactId>
        <version>1.0.1</version>
    </dependency>
</dependencies>
</plugin>
</plugins>
</build>
</project>

```

Now, you can run either `mvn site` or `mvn pmd:pmd pmd:cpd` to generate the PMD and CPD reports.

Unfortunately there's currently no built-in Gradle support for PMD or CPD, so you have to fall back on invoking the PMD Ant plugin directly, as shown here:

```

configurations {
    pmdConf
}

dependencies {
    pmdConf 'pmd:pmd:4.2.5'
}

task pmd << {
    println 'Running PMD static code analysis'
    ant {
        taskdef(name:'pmd', classname:'net.sourceforge.pmd.ant.PMDTask',
            classpath: configurations.pmdConf.asPath)

        taskdef(name:'cpd', classname:'net.sourceforge.pmd.cpd.CPDTask',
            classpath: configurations.pmdConf.asPath)

        pmd(shortFileNames:'true', failOnRuleViolation:'false',
            rulesetFiles:'conf/pmd-rules.xml') {
            formatter(type:'xml', toFile:'build/pmd.xml')
            fileset(dir: "src/main/java") {
                include(name: '**/*.java')
            }
        }
    }
}

```

```

        fileset(dir: "src/test/java") {
            include(name: '**/*.java')
        }
    }

    cpd(minimumTokenCount:'50', format: 'xml',
        ignoreIdentifiers: 'true',
        outputFile:'build/cpd.xml') {
        fileset(dir: "src/main/java") {
            include(name: '**/*.java')
        }
        fileset(dir: "src/test/java") {
            include(name: '**/*.java')
        }
    }
}
}
}

```

This configuration will use the PMD rule set in the `src/config` directory, and generate a PMD XML report called `pmd.xml` in the `build` directory. It will also run CPD and generate a CPD XML report called `cpd.xml` in the `build` directory.

9.3.3. FindBugs

FindBugs is a powerful code quality analysis tool that checks your application byte code for potential bugs, performance problems, or poor coding habits. FindBugs is the result of research carried out at the University of Maryland lead by Bill Pugh that studies byte code patterns coming from bugs in large real-world projects, such as the JDKs, Eclipse, and source code from Google applications. FindBugs can detect some fairly significant issues such as null pointer exceptions, infinite loops, and unintentionally accessing the internal state of an object. Unlike many other static analysis tools, FindBugs tends to find a smaller number of issues, but of those issues, a larger proportion will be important.

FindBugs is less configurable than the other tools we've seen, though in practice you generally don't need to fine-tune the rules as much as the other tools we've discussed. You can list the individual rules you want to apply, but you can't configure a shared XML file between your Maven builds and your IDE, for example.

FindBugs comes bundled with an Ant task. You can define the FindBugs task in Ant as shown below. FindBugs needs to refer to the FindBugs home directory, which is where the binary distribution of FindBugs has been unzipped. To make the build more portable, we store the FindBugs installation in our project directory structure, in the `tools/findbugs` directory:

```

<property name="findbugs.home" value="tools/findbugs" />

<taskdef name="findbugs" classname="edu.umd.cs.findbugs.anttask.FindBugsTask" >
  <classpath>
    <fileset dir="${findbugs.home}/lib" includes="**/*.jar"/>
  </classpath>
</taskdef>

```

Then, to run FindBugs, you could set up a “findbugs” target, as shown in the following example. Note that FindBugs runs against your application byte-code, not your source code, so you need to compile your source code first:

```
<target name="findbugs" depends="compile">
  <findbugs home="${findbugs.home}" output="xml" outputFile="target/findbugs.xml">
    <class location="${classes.dir}" />
    <auxClasspath refId="dependency.classpath" />
    <sourcePath path="src/main/java" />
  </findbugs>
</target>
```

If you're using Maven 2, you don't need to keep a local copy of the FindBugs installation. You just need to configure FindBugs in the reporting section as shown here:

```
<reporting>
  <plugins>
    <plugin>
      <groupId>org.codehaus.mojo</groupId>
      <artifactId>findbugs-maven-plugin</artifactId>
      <version>2.3.1</version>
      <configuration>
        <effort>Max</effort>
        <xmlOutput>true</xmlOutput>
      </configuration>
    </plugin>
  </plugins>
</reporting>
```

Or, for a Maven 3 project:

```
<project>
  ...
  <build>
    ...
    <plugins>
      ...
      <plugin>
        <groupId>org.apache.maven.plugins</groupId>
        <artifactId>maven-site-plugin</artifactId>
        <version>3.0-beta-2</version>
        <configuration>
          <reportPlugins>
            <plugin>
              <groupId>org.codehaus.mojo</groupId>
              <artifactId>findbugs-maven-plugin</artifactId>
              <version>2.3.1</version>
              <configuration>
                <effort>Max</effort>
                <xmlOutput>true</xmlOutput>
              </configuration>
            </plugin>
          </reportPlugins>
        </configuration>
      </plugin>
    </plugins>
  </build>
</project>
```

```

        </configuration>
    </plugin>
</plugins>
</build>
</project>

```

In both cases, you can generate the XML reports by running `mvn site` or `mvn findbugs:findbugs`. The XML reports will be placed in the `target` directory.

At the time of writing there's no special support for FindBugs in Gradle, so you need to invoke the FindBugs Ant plugin.

9.3.4. CodeNarc

CodeNarc is a static analysis tool for Groovy code, similar to PMD for Java. It checks Groovy source code for potential defects, poor coding practices and styles, overly complex code, and so on. Typical rules include “Constant If Expression,” “Empty Else Block,” “GString As Map Key,” and “Grails Stateless Service.”

For Ant or Maven-based projects, the CodeNarc Ant plugin is the simplest option (a Maven plugin is under development at the time of writing). A typical Ant configuration for use with Jenkins would look like this:

```

<taskdef name="codenarc" classname="org.codenarc.ant.CodeNarcTask"/>
<target name="runCodeNarc">
    <codenarc ruleSetFiles="rulesets/basic.xml,rulesets/imports.xml"
        maxPriority1Violations="0">

        <report type="xml">
            <option name="outputFile" value="reports/CodeNarc.xml" />
        </report>

        <fileset dir="src">
            <include name="**/*.groovy"/>
        </fileset>
    </codenarc>
</target>

```

You can integrate CodeNarc into a Grails project simply by installing the CodeNarc plugin:

```
$ grails install-plugin codenarc
```

This will configure CodeNarc to analyse the Groovy files in your Grails application code, as well as in the `src/groovy` and `test` directories.

Gradle 0.8 also provides support for CodeNarc in the code-quality plugin that you can configure in your builds as shown here:

```
apply plugin: 'code-quality'
```

This will use the CodeNarc configuration file in `config/codenarc/codenarc.xml` by default. You can override this with the `codeNarcConfigFileName` property.

You can generate the CodeNarc reports here by running `gradle codenarcMain` or, more simply, `gradle check`.

9.4. Reporting on Code Quality Issues with the Violations Plugin

One of the most useful code quality plugins for Jenkins is the Violations plugin. This plugin won't analyse your project source code (you need to configure your build to do that), but it does a great job of reporting on the code quality metrics generated for individual builds and trends over time. The plugin produces reports on code quality metrics coming from a large range of static analysis tools, including:

For Java

Checkstyle, CPD, PMD, FindBugs, and jcreport

For Groovy

codenarc

For JavaScript

jslint

For .Net

gendarme and stylecop

Installing the plugin is straightforward. Just go to the Plugin Manager screen and select the Jenkins Violations plugin. Once you've installed the plugin and restarted Jenkins, you'll be able to use the plugin for your projects.

The Violations plugin doesn't generate the code quality metrics data itself—you need to configure your build to do that, as shown in the previous section. An example of doing this for a Maven build job is illustrated in Figure 9.3, “Generating code quality reports in a Maven build”. Notice that here we're invoking the Maven plugin goals directly. We could also just run `mvn site`, but if we're only interested in the code quality metrics, and not the other elements of the Maven-generated site, calling the plugins directly will result in faster builds.

Figure 9.3. Generating code quality reports in a Maven build

Once you've set this up, you can configure the Violations plugin to generate reports and, if required, trigger notifications, based on the report results. Just go to the Post-build Actions and check the Report Violations checkbox. The details of the configuration vary depending on the project type. Let's look at freestyle build jobs first.

9.4.1. Working with Freestyle Build Jobs

Freestyle build jobs offer the most configuration flexibility, and are your only option for non-Java projects.

When you use the Violations plugin with a freestyle build job, you need to specify the paths to each of the XML reports generated by the static analysis tools you've used (see Figure 9.4, “Configuring the violations plugin for a freestyle project”). The plugin can produce several reports from the same tool, which is useful for Maven multimodule projects—just use a wildcard expression to identify the reports you want (for example, `**/target/checkstyle.xml`).

Figure 9.4. Configuring the violations plugin for a freestyle project

The Violations plugin will generate a graph tracking the number of each type of issue over time (see Figure 9.5, “Violations over time”). The graph displays a different-colored line for each type of violations you're tracking, as well as a summary of the latest results.

Figure 9.5. Violations over time

You can also click on this graph to drill down into a particular build. Here, you can see the number of issues raised for that particular build (see Figure 9.6, “Violations for a given build”), with various breakdowns by violation type, severity, and file.

Figure 9.6. Violations for a given build

Finally, you can drill down into a particular class, to display the detailed list of issues, along with where they appear in the source code.

The Violations plugin also allows for a more proactive management of code quality. You can use the results of the code quality analysis reports to influence the weather icon on the Jenkins dashboard. This weather icon is normally related to the number of failing builds in the previous five builds, but Jenkins can also take into account other factors, such as code quality results. Displaying a cloudy or stormy icon for a project on the dashboard is a much better way of raising awareness about code quality issues than simply relying on graphs and reports on the build job page.

To set this up, you need to go back to the Report Violations section in the Post-build Actions. The first three columns in Figure 9.4, “Configuring the violations plugin for a freestyle project” show a sunny icon, a stormy icon, and a yellow ball. The one with the sunny icon is the maximum number of violations tolerated in order to keep the sunny weather icon on the dashboard page. The second column, with the stormy weather icon, is the number of violations that will cause a stormy icon to appear on the dashboard. If you have a number of violations between these two extremes, you'll get one of the cloudy icons.

You can set different values for different tools. The exact thresholds will vary between teams, projects, and also between tools. For example, Checkstyle will typically raise a lot more issues than FindBugs or CPD, with PMD being somewhere in between. You need to adjust the values you use to reflect how these tools work on your code base and your expectations.

You can take this even further with the third column (the one with the yellow ball). This column lets you set a number of violations that will cause the build to be declared unstable. Remember, when a build becomes unstable Jenkins will send out notification messages, so this is an even more proactive strategy.

For example, in Figure 9.4, “Configuring the violations plugin for a freestyle project”, we’ve configured the minimum number of Checkstyle violations to 10, which means that the sunny weather icon will only appear if there are 10 or fewer Checkstyle violations. If there are more than 10, the weather will degrade progressively, until at the 200 violations mark, it will become stormy. And if there are 500 or more Checkstyle violations, the project will be flagged as unstable.

Now, look at the configuration for CPD, the duplicated code detector that comes with PMD. In this project, we’ve adopted a zero-tolerance policy for duplicated code, so the sunny icon value is set to zero. The stormy icon is set to 10, so if there are 10 or more copy-paste violations, the project weather indicator will appear as stormy. And if the project has 15 or more copy-paste violations, it will be declared unstable.

On the Dashboard page, this project will appear both with a stormy weather icon and as unstable, even though there are no test failures (see Figure 9.7, “Configuring the violations plugin for a freestyle project”). This particular build is unstable because there are 16 CPD violations. In addition, if you place your mouse over the weather icon, Jenkins will display some more details about how it calculated this particular status.

Figure 9.7. Configuring the violations plugin for a freestyle project

9.4.2. Working with Maven Build Jobs

Maven build jobs in Jenkins use the Maven conventions and information in the project `pom.xml` file to make configuration easier and more lightweight. When you use the Violations plugin with a Maven build job, Jenkins uses these conventions to reduce the amount of work you need to do to configure the plugin. You don’t need to tell Jenkins where to find the XML reports for many of the static analysis tools (for example, Checkstyle, PMD, FindBugs, and CPD), as Jenkins can figure this out based from the Maven conventions and plugin configurations (see Figure 9.8, “Configuring the violations plugin for a Maven project”). If you do need to override these conventions, you can choose the Pattern option in the “XML filename pattern” drop-down list, and enter a path as you do for freestyle build jobs.

Figure 9.8. Configuring the violations plugin for a Maven project

The Violations plugin works well with multimodule Maven projects, but at the time of writing it needs a little tweaking to obtain best results. Maven build jobs understand the structure of multimodule projects (see Figure 9.9, “Jenkins Maven build jobs understand Maven multimodule structures”). Furthermore, you can drill down into any module and get a detailed view of the build results for that build job.

Figure 9.9. Jenkins Maven build jobs understand Maven multimodule structures

This is a very useful feature, but it means you need to do a little extra work to get all of the benefits out of the Violations plugins for the individual modules. By default, the violations plugin will display an aggregated view of the code quality metrics like the one in Figure 9.5, “Violations over time”. You can also click on the violations graph, and view the detailed reports for each module.

However, for this to work correctly, you need to activate the Violations plugin individually for each module in addition to the main project. To do this, click on the module you want to configure in the Modules screen, and then click on the “Configure” menu. Here, you'll see a small subset of the usual configuration options (see Figure 9.10, “Activating the Violations plugin for an individual module”). Here, you just need to activate the Violations option, and configure the thresholds, if required. On the positive side, this means that you can define different threshold values for different modules.

Figure 9.10. Activating the Violations plugin for an individual module

Once you've done this, when you click on the violations aggregate graph on the project build job home page, Jenkins will list the individual violations graphs for each module.

9.5. Using the Checkstyle, PMD, and FindBugs Reports

You can also report individually on results from Checkstyle, PMD, and FindBugs. In addition to the Violations plugin, there are also Jenkins plugins that produce trend graphs and detailed reports for each of these tools individually. We'll look at how to do this for Checkstyle, but the same approach also applies for PMD and FindBugs. You can even use the Analysis Collector Plugin to display the combined results in a graph similar to the one produced by the Violations plugin.

You can install these plugins through the Plugin Manager in the usual way. The plugins in question are called, unsurprisingly, Checkstyle plugin, PMD plugin, and FindBugs plugin. All of these plugins use the Static Analysis Utilities plugin, which you need to install as well (see Figure 9.11, “Installing the Checkstyle and Static Analysis Utilities plugins”).

Figure 9.11. Installing the Checkstyle and Static Analysis Utilities plugins

Once you've installed these plugins, you can set up the reporting in your project configuration. Tick the “Publish Checkstyle analysis results” checkbox. In a freestyle build, you'll need to specify a path pattern to find the Checkstyle XML reports. In a Maven 2 build, Jenkins will figure out where to look for them by itself.

This will provide basic Checkstyle reporting, but as usual you can fine-tune things further by clicking on the Advanced button. In a Maven 2 build, you can configure the health threshold values (how many violations will cause the build to go from sunny to stormy), and also filter the priority violations you want to include in this calculation. For example, you may only want high priority issues to be taken into account for the weather icon status.

The freestyle builds have a few more options you can configure: in particular, you can cause the build to become unstable (yellow ball) or even to fail (red ball) if there are more than a given number of violations, or if there are more than a given number of new violations (see Figure 9.12, “Configuring the Checkstyle plugin”). So, in the configuration in the illustration, if there are more than 50 new checkstyle violations of any priority in a build, the build will be flagged as unstable. This certainly has its uses for Checkstyle, but it can also come in very handy with FindBugs, where high priority issues often represent dangerous and potentially show-stopping bugs.

Figure 9.12. Configuring the Checkstyle plugin

Now, when the build runs, Jenkins will now generate a trend graph and detailed reports for the Checkstyle violations (see Figure 9.13, “Displaying Checkstyle trends”). From here, you can drill down to see violations per priority, per category, per run type, per package, and so on.

Figure 9.13. Displaying Checkstyle trends

As we mentioned earlier, the same approach also works with the PMD plugin and the FindBugs plugin. These plugins are a great way to provide more focused reporting on the results of a particular tool, and they also give you more control over the impact that these violations will have on the build results.

9.6. Reporting on Code Complexity

Code complexity is another important aspect of code quality. Code complexity is measured in a number of ways, but one commonly used (and easy-to-understand) complexity metric is Cyclometric Complexity, which involves measuring the number of different paths through a method. Using this metric, complex code typically has large numbers of nested conditional statements and loops, which make the code harder to understand and debug.

There's also a code quality theory that correlates code complexity and code coverage, to give a general idea of how reliable a particular piece of code is. This is based on the (very understandable) idea that code that's both complex and poorly tested is more likely to contain bugs than simple, well-tested code.

The Coverage Complexity Scatter Plot plugin is designed to let you visualize this information in your Jenkins builds (see Figure 9.14, “A coverage/complexity scatter plot”). Dangerously complex and/or untested methods will appear high on the graph, and the more well-written and well-tested methods will appear lower down.

Figure 9.14. A coverage/complexity scatter plot

The scatter graph gives you a good overview of the state of your code in terms of complexity and test coverage, but you can also drill down to investigate further. If you click on any point in the graph, you can see the corresponding methods, with their test coverage and complexity (see Figure 9.15, “You can click on any point in the graph to investigate further”).

Figure 9.15. You can click on any point in the graph to investigate further

At the time of writing, this plugin requires Clover, so your build needs to have generated a Clover XML coverage report, and you need to have installed and configured the Clover Jenkins plugin (see Section 6.6.2, “Measuring Code Coverage with Clover”). However, support for Cobertura and other tools is planned.

9.7. Reporting on Open Tasks

When it comes to code quality, static analysis isn't the only tool you can use. Another indicator of the general health of your project can be found in the number of `FIXME`, `TODO`, `@deprecated`, and similar tags scattered through the source code. If there are a lot of these, this can be a sign that your code base has a lot of unfinished work, and is therefore not in a very finalized state.

The Jenkins Task Scanner plugin lets you keep track of these sorts of tags in your source code, and optionally flag a build with a bad weather icon on the dashboard if there are too many open tasks.

To set this up, you need to install both the Static Analysis Utilities plugin and the Task Scanner plugin. Once installed, you can activate the plugin in your project by checking the “Scan workspace for open tasks” checkbox in the Build Settings section of your project job configuration.

Configuring the Task Scanner plugin is pretty straightforward (see Figure 9.16, “Configuring the Task Scanner plugin is straightforward”). You simply enter the tags you want to track, with different priorities if you consider certain tags to be more important than others. By default, the plugin will scan all the Java source code in the project, but you can redefine this behavior by entering the Files to scan field. In Figure 9.16, “Configuring the Task Scanner plugin is straightforward”, for example, you also check XML and JSP files for tags.

Figure 9.16. Configuring the Task Scanner plugin is straightforward

The Advanced button gives you access to a few more sophisticated options. Probably the most useful are the Health thresholds, which let you define the maximum number of issues tolerated before the build can no longer be considered “sunny,” and the minimum number of issues required for “stormy weather” status.

The plugin generates a graph that shows tag trends by priority (see Figure 9.17, “The Open Tasks Trend graph”). If you click on the Open Tasks report, you can also see a breakdown of tasks by Maven module, package or file, or even list the open tasks.

Figure 9.17. The Open Tasks Trend graph

9.8. Integrating with Sonar

Sonar⁴ is a tool that centralizes a range of code quality metrics into a single website (see Figure 9.18, “Code quality reporting by Sonar”). It uses several Maven plugins (Checkstyle, PMD, FindBugs, Cobertura or Clover, and others) to analyse Maven projects and generate a comprehensive set of code quality metrics reports. Sonar reports on code coverage, rule compliance, and documentation, but also on more high-level metrics such as complexity, maintainability, and even technical debt. You can use plugins to extend Sonar's features and add support for other languages (such as support for CodeNarc for Groovy source code). The rules used by the various tools are managed and configured centrally on the Sonar website, and the Maven projects being analyzed don't require any particular configuration. This makes Sonar a great fit for working on Maven projects where you have limited control over the POM files.

Figure 9.18. Code quality reporting by Sonar

In one of the most common usages of Sonar, Sonar automatically runs a set of Maven code quality related plugins against your Maven project, and stores the results into a relational database. The Sonar server, which you run separately, then analyzes and displays the results as shown in Figure 9.18, “Code quality reporting by Sonar”.

Jenkins integrates well with Sonar. The Jenkins Sonar Plugin lets you define Sonar instances for all of your projects, and then activate Sonar in particular builds. You can run your Sonar server on your Jenkins instance, or on a different machine. The only constraint is that the Jenkins instance must have JDBC access to the Sonar database, as it injects code quality metrics directly into the database, without going through the Sonar website (see Figure 9.19, “Jenkins and Sonar”).

Figure 9.19. Jenkins and Sonar

⁴ <http://www.sonarsource.org>

Sonar also has an Ant bootstrap (with a Gradle bootstrap in the making at the time of writing) for non-Maven users.

You install the plugin in the usual way, via the Plugin Manager. Once installed, you configure the Jenkins Sonar plugin in the Configure System screen, in the Sonar section. This involves defining your Sonar instances—you can configure as many instances of Sonar as you need. The default configuration assumes that you're running a local instance of Sonar with the default embedded database. This is useful for testing purposes but not very scalable. For a production environment, you'll typically run Sonar on a real database such as MySQL or Postgres, and you'll need to configure the JDBC connection to the production Sonar database in Jenkins. You do this by clicking on the Advanced button and filling in the appropriate fields (see Figure 9.20, “Configuring Sonar in Jenkins”).

Figure 9.20. Configuring Sonar in Jenkins

The other thing you need to configure is when the Sonar build will kick off in a Sonar-enabled build job. You usually configure Sonar to run with one of the long-running Jenkins build jobs, such as the code quality metrics build. It isn't very useful to run the Sonar build more than once a day, as Sonar stores metrics in 24-hour slices. The default configuration will kick off a Sonar build in a Sonar-enabled build job whenever the job is triggered by a periodically scheduled build or by a manual build.

To activate Sonar in your build job with the system-wide configuration options, just check the Sonar option in the Post-build Actions (see Figure 9.21, “Configuring Sonar in a build job”). Sonar will run whenever your build is started by one of the trigger mechanisms defined above.

Figure 9.21. Configuring Sonar in a build job

You typically set up Sonar to run on a regular basis, for example every night or once a week. So, you can activate Sonar on your normal unit/integration test build job simply by adding a schedule (see Figure 9.22, “Scheduling Sonar builds”). This avoids duplicated configuration details between jobs. Or, if you already have a scheduled build job that runs with an appropriate frequency (such as a dedicated code quality metrics build), you can activate Sonar on this build job.

Figure 9.22. Scheduling Sonar builds

If you click on the Advanced button, you can specify other more sophisticated options, such as running your Sonar build on a separate branch, passing Maven additional command-line options (such as extra memory), or overriding the default trigger configuration.

By default, Sonar will run even if the normal build fails. This is usually what you want, as Sonar should record build and test failures as well as successful results. However, if required, you can deactivate this option too in the Advanced options.

9.9. Conclusion

Code quality is an important part of the build process, and Jenkins provides excellent support for the wide range of code quality-related tools out there. As a result, Jenkins should be a key part of your code quality strategy.

Chapter 10. Advanced Builds

10.1. Introduction

In this chapter, you'll look at some more advanced build job setups. We'll discuss parameterized builds, which allow Jenkins to prompt the user for additional parameters to pass into the build job, and multiconfiguration build jobs, which let you run a single build job through a large number of variations. We'll look at how to run build jobs in parallel, and wait for the outcome of one or more build jobs before continuing. And, you'll see how to implement build promotion strategies and build pipelines so that Jenkins can be used not only as a build server, but also as a deployment server.

10.2. Parameterized Build Jobs

Parameterized builds are a powerful concept that enable you to add another dimension to your build jobs.

The Parameterized Build plugin lets you configure parameters for your build job that can be either entered by the user when the build job is triggered, or (as you'll see later), from another build job.

For example, you might have a deployment build job, where you want to choose the target environment in a drop-down list when you start the build job. Or, you may want to specify the version of the application you want to deploy. Or, when running a build job involving web tests, you might want to specify the browser to run your Selenium or WebDriver tests in. You can even upload a file to be used by the build job.

Note that it's the job of the build script to analyze and process the parameter values correctly—Jenkins simply provides a user interface for users to enter values for the parameters, and passes these parameters to the build script.

10.2.1. Creating a Parameterized Build Job

You install the Parameterized Build plugin as usual, via the Plugin Manager screen. Once you've done this, configuring a parameterized build job is straightforward. Just tick the “This build is parameterized” option and click Add Parameter to add a new build job parameter (see Figure 10.1, “Creating a parameterized build job”). You can add parameters to any sort of build, and you can add as many parameters as you want for a given build job.

Figure 10.1. Creating a parameterized build job

To add a parameter to your build job, just pick the parameter type in the drop-down list. This will let you configure the details of your parameter (see Figure 10.2, “Adding a parameter to the build job”). You can choose from several different parameter types, such as Strings, Booleans, and drop-down lists. Depending on the type you choose, you'll have to enter slightly different configuration values, but the

basic process is identical. All parameter types, with the exception of the File parameter, have a name and a description, and most often a default value.

In Figure 10.3, “Adding a parameter to the build job”, for example, you're adding a parameter called `VERSION` to a deployment build job. The default value (`RELEASE`) will be initially displayed when Jenkins prompts the user for this parameter, so if the user doesn't change anything, this value will be used.

Figure 10.2. Adding a parameter to the build job

When the user starts a parameterized build job (parameterized build jobs are very often started manually), Jenkins will propose a page where the user can enter values for each of the build job's parameters (see Figure 10.3, “Adding a parameter to the build job”).

Figure 10.3. Adding a parameter to the build job

10.2.2. Adapting Your Builds to Work with Parameterized Build Scripts

Once you've added a parameter, you need to configure your build scripts to use it. Choosing the parameter name well is important here, as this is also the name of the variable that Jenkins will pass through as an environment variable when it runs the build job. To illustrate this, consider the very basic build job configuration in Figure 10.4, “Demonstrating a build parameter”, where you're simply echoing the build parameter back out to the console. Note that, to make the environment variables more portable across operating systems, it's good practice to put them all in upper case.

Figure 10.4. Demonstrating a build parameter

When you run this, you get a console output along the following lines:

```
Started by user anonymous
Building on master
[workspace] $ /bin/sh -xe /var/folders/y+/y+a+wZ-jG6WKHEm9KwnSvE+++TI/-Tmp-/
jenkins5862957776458050998.sh
+ echo Version=1.2.3
Version=1.2.3
Notifying upstream projects of job completion
Finished: SUCCESS
```

You can also use these environment variables from within your build scripts. For example, in an Ant or Maven build, you can use the special `env` property to access the current environment variables:

```
<target name="printversion">
```

```
<property environment="env" />
<echo message="${env.VERSION}" />
</target>
```

Another option is to pass the parameter into the build script as a property value. The following is a more involved example from a Maven POM file. In this example, Maven is configured to deploy a specific WAR file. You provide the version of the WAR file to be deployed in the `target.version` property, which is used in the dependency declaration, as shown below:

```
...
<dependencies>
  <dependency>
    <groupId>com.wakaleo.gameoflife</groupId>
    <artifactId>gameoflife-web</artifactId>
    <type>war</type>
    <version>${target.version}</version>
  </dependency>
</dependencies>
<properties>
  <target.version>RELEASE</target.version>
  ...
</properties>
```

When you invoke Maven, you pass in the parameter as one of the build properties (see Figure 10.5, “Adding a parameter to a Maven build job”). You can then use a tool like Cargo to do the actual deployment—Maven will download the requested version of the WAR file from the local Enterprise Repository Manager, and deploy it to an application server.

Figure 10.5. Adding a parameter to a Maven build job

That, in a nutshell, is how you can integrate build job parameters into your build. In addition to plain old String parameters, however, there are a few more sophisticated parameter types that you'll look at in the following paragraphs (see Figure 10.6, “Many different types of parameters are available”).

Figure 10.6. Many different types of parameters are available

10.2.3. More Advanced Parameter Types

Password parameters are, as you'd expect, very similar to String parameters, except that they're displayed as a password field.

There are many cases where you wish to present a limited set of parameter options. In a deployment build, you might want to let the user choose one of a number of target servers. Or, you may want to present a list of supported browsers for a suite of acceptance tests. Choice parameters let you define a set of values that will be displayed as a drop-down list (see Figure 10.7, “Configuring a Choice parameter”). You need to provide a list of possible values, one per line, starting with the default value.

Figure 10.7. Configuring a Choice parameter

Boolean parameters are, as you'd expect, parameters that take a value of `true` or `false`. They're presented as checkboxes.

Two more exotic parameter types, which behave a little differently to the others, are Run parameters and File parameters.

Run parameters let you select a particular run (or build) of a given build job (see Figure 10.8, “Configuring a Run parameter”). The user picks from a list of build run numbers. The URL of the corresponding build run is stored in the specified parameter.

Figure 10.8. Configuring a Run parameter

The URL (which will look something like `http://jenkins.myorg.com/job/game-of-life/197/`) can be used to obtain information or artifacts from that build run. For example, you could obtain the JAR or WAR file archived in a previous build and run further tests with this particular binary in a separate build job. For example, to access the WAR file of a previous build in a multimodule Maven project, the URL would look something like this:

```
http://buildserver/job/game-of-life/197/artifact/gameoflife-web/target/
gameoflife.war
```

So, using the parameter configured in Figure 10.8, “Configuring a Run parameter”, you could access this WAR file using the following expression:

```
${RELEASE_BUILD}gameoflife-web/target/gameoflife.war
```

File parameters let you upload a file into the build job workspace, so that it can then be used by the build script (see Figure 10.9, “Configuring a File parameter”). Jenkins will store the file into the specified location in the project workspace, where you can access it in your build scripts. You can use the `WORKSPACE` variable to refer to the current Jenkins workspace directory, so you could manipulate the file uploaded in Figure 10.9, “Configuring a File parameter” by using the expression `${WORKSPACE}/deploy/app.war`.

Figure 10.9. Configuring a File parameter

10.2.4. Building from a Subversion Tag

The parameterized trigger has special support for Subversion, allowing you to build against a specific Subversion tag. This is useful if you want to run a release build using a tag generated by a previous build

job. For example, an upstream build job may tag a particular revision. Alternatively, you might use the standard Maven release process (see Section 10.7.1, “Managing Maven Releases with the M2Release Plugin”) to generate a new release. In this case, a tag with the Maven release number will automatically be generated in Subversion.

This approach is useful for projects that need to be partially or entirely rebuilt before they can be deployed to a given platform. For example, you may need to run the Ant or Maven build using different properties or profiles for different platforms, so that platform-specific configuration files can be embedded in the deployed WAR or EAR files.

You can configure a Jenkins build to run against a selected tag by using the “List Subversion Tag” parameter type (see Figure 10.10, “Adding a parameter to build from a Subversion tag”). You just need to provide the Subversion repository URL pointing to the tags directory of your project.

Figure 10.10. Adding a parameter to build from a Subversion tag

When you run this build, Jenkins will propose a list of tags to choose from (see Figure 10.11, “Building from a Subversion tag”).

Figure 10.11. Building from a Subversion tag

10.2.5. Building from a Git Tag

Building from a Git tag isn't as simple as doing so from a Subversion tag, though you can still use a parameter to indicate which tag to use. Indeed, because of the very nature of Git, when Jenkins obtains a copy of the source code from Git, it clones the Git repository, including all of the tags. Once you have the latest version of the repository on your Jenkins server, you can then proceed to checkout a tagged version using `git checkout <tagname>`.

To set this up in Jenkins, you first need to add a String parameter to your build job (called `RELEASE` in this example—see Figure 10.12, “Configuring a parameter for a Git tag”). Unlike the Subversion support, there's no way to list the available Git tags in a drop-down list, so users will need to know the name of the tag they want to release.

Figure 10.12. Configuring a parameter for a Git tag

Once you've added this parameter, you need to checkout the corresponding tag after the repository has been cloned locally. So, if you have a freestyle build, the first build step would be a command-line call to Git to check out the tag referenced by the `RELEASE` parameter (see Figure 10.13, “Building from a Git tag”). Of course, a more portable way to do this would be to write a simple Ant or Groovy script to do the same thing in a more OS-neutral way.

Figure 10.13. Building from a Git tag

10.2.6. Starting a Parameterized Build Job Remotely

You can also start a parameterized build job remotely, by invoking the URL of the build job. The typical form of a parameterized build job URL is illustrated here:

```
http://jenkins.acme.org/job/myjob/buildWithParameters?PARAMETER=Value
```

So, in the example shown above, you could trigger a build like this:

```
http://jenkins.acme.org/job/parameterized-build/buildWithParameters?VERSION=1.2.3
```

When you use a URL to start a build job in this way, remember that the parameter names are case-sensitive, and that the values need to be escaped (just like any other HTTP parameter). And if you're using a Run parameter, you need to provide the name of the build job and the run number (e.g., `game-of-life#197`), not just the run number.

10.2.7. Parameterized Build Job History

Finally, it's indispensable to know what parameters were used to run a particular parameterized build. For example, in an automated deployment build job, it's useful to know exactly what version was actually deployed. Fortunately, Jenkins stores these values in the build history (see Figure 10.14, “Jenkins stores what parameter values were used for each build”), so you can always go back and take a look.

Figure 10.14. Jenkins stores what parameter values were used for each build

10.3. Parameterized Triggers

When you trigger another build job from within a parameterized build job, it's often useful to be able to pass the parameters of the current build job to the new one. Suppose, for example, that you have an application that needs to be tested against several different databases. As you've seen, you could do this by setting up a parameterized build job that accepts the target database as a parameter. You may want to kick off a series of builds, all of which need this parameter.

If you try to do this using the conventional “Build other projects” option in the Post-Build Actions section, it won't work. In fact, you can't trigger a parameterized build in this way.

However, you can do this using the Jenkins Parameterized Trigger plugin. This plugin lets you configure your build jobs to both trigger parameterized builds, and to pass arbitrary parameters to these builds.

Once you install this plugin, you'll find the option of “Triggering parameterized builds on other projects” in your build job configuration page (see Figure 10.16, “Adding a parameterized trigger to a build job”).

This lets you start another build job in a number of ways. In particular, it lets you kick off a subsequent build job, passing the current parameters to this new build job, which is impossible to do with a normal triggered build. The best way to see how this works is through an example.

In Figure 10.15, “A parameterized build job with DATABASE parameter” we have an initial build job. This build job takes a single parameter, `DATABASE`, which specifies the database to be used for the tests. As you've seen, the user will be prompted to enter this value whenever the build is started.

Figure 10.15. A parameterized build job with DATABASE parameter

Now, suppose you want to trigger a second build job to run more comprehensive integration tests once this first build job has finished. However, you need it to run the tests against the same database. You can do this by setting up a parameterized trigger to start this second build job (see Figure 10.16, “Adding a parameterized trigger to a build job”).

Figure 10.16. Adding a parameterized trigger to a build job

In this case, you're simply passing through the current build parameters. This second build job will automatically be started after the first one, with the `DATABASE` parameter value provided by the user. You can also fine-tune the triggering policy by telling Jenkins when the build should be triggered. Typically, you'd only trigger a downstream build after your build has completed successfully, but with the Parameterized Trigger plugin you can also configure builds to be triggered even if the build is unstable, only when the build fails, or ask for it to be triggered no matter what the outcome of the first build. You can even set up multiple triggers for the same build job.

Naturally, the build job that you trigger must be a parameterized build job (as illustrated in Figure 10.17, “The build job you trigger must also be a parameterized build job”), and you must pass through all of the parameters it requires.

Figure 10.17. The build job you trigger must also be a parameterized build job

This feature actually has much broader applications than simply passing through the current build parameters. You can also trigger a parameterized build job with an arbitrary set of parameters, or use a combination of parameters that were passed to the current build along with your own ones. Or, if you have a lot of parameters, you can load them from a properties file. In Figure 10.18, “Passing a predefined parameter to a parameterized build job”, you're passing both the current build parameters (the `DATABASE` variable in this case), and an additional parameter called `TARGET_PLATFORM`.

Figure 10.18. Passing a predefined parameter to a parameterized build job

10.4. Multiconfiguration Build Jobs

Multiconfiguration build jobs are an extremely powerful feature of Jenkins. A multiconfiguration build job can be thought of as a parameterized build job that can be automatically run with all the possible combinations of parameters that it can accept. They're particularly useful for testing your application using a single build job, but under a wide variety of conditions (browsers, databases, and so forth).

10.4.1. Setting Up a Multiconfiguration Build

To create a new multiconfiguration build job, simply choose this option on the New Job page (see Figure 10.19, “Creating a multiconfiguration build job”).

Figure 10.19. Creating a multiconfiguration build job

A multiconfiguration build job is just like any other build job, but with one very important additional element: the Configuration Matrix (see Figure 10.20, “Adding an axis to a multiconfiguration build”). This is where you define the different configurations that will be used to run your builds.

Figure 10.20. Adding an axis to a multiconfiguration build

You can define different axes of configuration options, including running the build job on different slaves or on different JDKs, or providing your own custom properties to the build. For example, in the build jobs discussed earlier, you might want to test your application for different databases and different operating systems. You could define one axis defining slave machines with the different operating systems you want your build to run on, and another axis defining all the possible database values. Jenkins will then run the build job for each possible database and each possible operating system.

Let's look at the types of axis you can define.

10.4.2. Configuring a Slave Axis

The first option is to configure your build to run simultaneously on different slave machines (see Chapter 11, Distributed Builds). Of course, the idea of having a set of slave machines is usually that you can run your build job on any of them. But, there might be cases where it makes sense to be a little more choosy. For example, you might want your tests to run on Windows, Mac OS X, and Linux. In this case, you create a new axis for your slave nodes, as shown in Figure 10.21, “Defining an axis of slave nodes”. You can choose the nodes you want to use in two ways: by label or by individual node. Using labels lets you identify categories of build nodes (for example, Windows machines), without tying the build to any one machine. This is a more flexible option, and makes it easier to expand your build capacity as required. Sometimes, however, you may really want to run a build on a specific machine. In this case, you can use the “Individual nodes” option, and choose the machine in this list.

Figure 10.21. Defining an axis of slave nodes

If you need more flexibility, you can also use a Label Expression, which lets you define which slave nodes should be used for builds on a particular axis using boolean expressions and logical operators to combine labels. For example, suppose you've defined labels for slave machines based on operating system (“windows”, “linux”) and installed databases (“oracle”, “mysql”, “db2”). To define an axis running tests only on Windows machines installed with MySQL, you could use an expression like `windows && mysql`.

We discuss working with slave nodes and distributed builds in more detail in Chapter 11, Distributed Builds.

10.4.3. Configuring a JDK Axis

If you're deploying your application to a broad client base where you have limited control over the target environment, you may need to test your application using different versions of Java. In cases like this it's useful to be able to set up a JDK axis in a multiconfiguration build. When you add a JDK axis, Jenkins will automatically propose the list of JDK versions that it knows about (see Figure 10.22, “Defining an axis of JDK versions”). If you need to use additional JDKs, just add them to your Jenkins configuration page.

Figure 10.22. Defining an axis of JDK versions

10.4.4. Custom Axis

The third type of axis lets you define different ways to run your build job, based on arbitrary variables that you define. For example, you might provide a list of databases you need to test against, or a list of browsers to use in your web tests. These are similar to parameters for a parameterized build job, except that you provide the complete list of possible values, and rather than prompting for you to enter a value, Jenkins will run the build with all of the values you provide (Figure 10.23, “Defining a user-defined axis”).

Figure 10.23. Defining a user-defined axis

10.4.5. Running a Multiconfiguration Build

Once you've set up the axes, you can run your multiconfiguration build just like any other. However, Jenkins will treat each combination of variables as a separate build job. Jenkins displays the aggregate results in a table, where all of the combinations are shown (see Figure 10.24, “Multiconfiguration build results”). If you click on any of the balls, Jenkins will take you to the detailed results for that particular build.

Figure 10.24. Multiconfiguration build results

By default, Jenkins will run the build jobs in parallel. However, there are times when this isn't a good idea. For example, many Java web applications use Selenium or WebDriver tests running against a local instance of Jetty that's automatically started by the build job. Build scripts like this need to be specially configured to be able to run in parallel on the same machine, to avoid port conflicts. Concurrent database access during tests can be another source of problems if concurrency isn't designed into the tests. If your builds are not designed to run in parallel, you can force Jenkins to run the tests sequentially by ticking the Run each configuration sequentially checkbox at the bottom of the Configuration Matrix section.

By default, Jenkins will run all possible combinations of the different axes. So, in the above example, you have three environments, two JDKs, and four databases. This results in a total of 24 builds. However, in some cases, it may not make sense (or be possible) to run certain combinations. For example, suppose you have a build job that runs automated web tests. If one axis contains the web browsers to be tested (Firefox, Internet Explorer, Chrome, etc.) and another the Operating Systems (Linux, Windows, Mac OS), it would make little sense to run Internet Explorer with Linux or Mac OS.

The Combination Filter option lets you set up rules about which combinations of variables are valid. This field is a Groovy boolean expression that uses the names of the variables you defined for each axis. The expression must evaluate to true for the build to execute. For example, suppose you have a build job running web tests in different browsers on different operating systems (see Figure 10.25, “Setting up a combination filter”). The tests need to run Firefox, Internet Explorer, and Chrome on Windows, Mac OS X, and Linux. However, Internet Explorer only runs on Windows, and Chrome does not run on Linux.

Figure 10.25. Setting up a combination filter

To set this up with a Combination Filter, you could use an expression like the following:

```
(browser=="firefox")  
|| (browser=="iexplorer" && os=="windows")  
|| (browser=="chrome" && os != "linux")
```

This would result in only the correct browser/operating system combinations being executed (see Figure 10.26, “Build results using a combination filter”). Executed builds are displayed in the usual colors, whereas skipped builds are shown in gray.

Figure 10.26. Build results using a combination filter

Another reason to use a build filter is that there are simply too many valid combinations to run in a reasonable time. In this case, the best solution may be to upscale your build server. The second-best

solution, on the other hand, might be to only run a subset of the combinations, possibly running the full set of combinations on a nightly basis. You can do this by using the special `index` variable. Including the expression `(index%2 == 0)`, for example, ensures that only one build job in two is actually executed.

You may also want certain builds to be executed before the others, as a sanity check. For example, you might want to run the default (and, theoretically, the most reliable) configuration for your application first, before continuing on to more exotic combinations. To do this, you can use the “Execute touchstone builds first” option. Here, you enter a filter value (like the one seen above) to define the first build or builds to be executed. You can also specify if the build should proceed only if these builds are successful, or even if they're unsuccessful. Once these builds have completed as expected, Jenkins will proceed with the other combinations.

10.5. Generating Your Maven Build Jobs Automatically

Contributed by Evgeny Goldin

As mentioned in the previous section, the number of build jobs that your Jenkins server will host can vary. As the number of build jobs grows, it becomes harder not only to view them in the Jenkins dashboard, but to configure them as well. Imagine what would it take to configure 20 to 50 Jenkins jobs one-by-one! In addition, many of those jobs may have common configuration elements, such as Maven goals or build memory settings, which result in duplicated configuration and higher maintenance overhead.

For example, if you decide to run `mvn clean install` instead of `mvn clean deploy` for your release jobs and switch to alternative deployment methods, such as those provided by Artifactory plugin¹, you'll have no choice but to visit all relevant jobs and update them manually.

Alternatively, you could take advantage of the fact that Jenkins is a simple and straightforward tool that keeps all of its definitions in text files. Indeed, you can update your jobs' `config.xml` file directly in the `.jenkins/jobs` directory. While this approach will work, it's still far from ideal as it involves quite a lot of manual picking and fragile replacements in Jenkins XML files.

There's a third way to achieve the nirvana of massive job updates: generate your configuration files automatically using some sort of definition file. The Maven Jenkins Plugin² does exactly that, generating `config.xml` files for all jobs using standard Maven definitions kept in a single `pom.xml` file.

10.5.1. Configuring a Job

When configuring a single job with the Maven Jenkins Plugin, you can define all the usual Jenkins configuration elements, such as Maven goals, POM location, repository URLs, e-mail addresses, number of days to keep the logs, and so on. The plugin tries to bring you as close to possible to Jenkins' usual way of configuring a job manually.

¹ <http://wiki.jenkins-ci.org/display/JENKINS/Artifactory+Plugin>

² <http://evgeny-goldin.com/wiki/Maven-jenkins-plugin>

Let's take a look at a Google Guice³ build job:

```
<job>
  <id>google-guice-trunk</id>
  <description>Building Google Guice trunk.</description>
  <descriptionTable>
    <row>
      <key>Project Page</key>
      <value>
        <a href="http://code.google.com/p/google-guice/">
          <b><code>code.google.com/p/google-guice</code></b>
        </a>
      </value>
      <escapeHTML>>false</escapeHTML>
      <bottom>>false</bottom>
    </row>
  </descriptionTable>
  <jdkName>jdk1.6.0</jdkName>
  <mavenName>apache-maven-3</mavenName>
  <mavenOpts>-Xmx256m -XX:MaxPermSize=128m</mavenOpts>
  <daysToKeep>5</daysToKeep>
  <useUpdate>>false</useUpdate>
  <mavenGoals>-e clean install</mavenGoals>
  <trigger>
    <type>timer</type>
    <expression>0 0 * * *</expression>
  </trigger>
  <repository>
    <remote>http://google-guice.googlecode.com/svn/trunk</remote>
  </repository>
  <mail>
    <recipients>jenkins@evgeny-goldin.org</recipients>
  </mail>
</job>
```

This job uses a number of standard configurations such as `<jdkName>`, `<mavenName>`, and `<mavenOpts>`. The code is checked out from a Subversion repository (defined in the `<repository>` element), and a cron `<trigger>` runs the job nightly at 00:00. Email notifications are sent to the people specified with the `<mail>` element. This configuration also adds a link back to the project's page in the description table that's generated automatically for each job.

The generated job is displayed in your Jenkins server as illustrated in Figure 10.27, "A job generated by the Maven Jenkins plugin".

Figure 10.27. A job generated by the Maven Jenkins plugin

Here's another job building the Jenkins master branch at GitHub:

³ <http://code.google.com/p/google-guice/>

```

<job>
  <id>jenkins-master</id>
  <jdkName>jdk1.6.0</jdkName>
  <numToKeep>5</numToKeep>
  <mavenName>apache-maven-3</mavenName>
  <trigger>
    <type>timer</type>
    <expression>0 1 * * *</expression>
  </trigger>
  <scmType>git</scmType>
  <repository>
    <remote>git://github.com/jenkinsci/jenkins.git</remote>
  </repository>
  <mail>
    <recipients>jenkins@evgeny-goldin.org</recipients>
    <sendForUnstable>>false</sendForUnstable>
  </mail>
</job>

```

This would generate the job shown in Figure 10.28, “jenkins-master job generated”.

Figure 10.28. jenkins-master job generated

The plugin’s documentation⁴ provides a detailed reference of all settings that can be configured.

10.5.2. Reusing Job Configuration with Inheritance

Being able to generate Jenkins jobs using centralized configuration, such as a Maven POM file, solves the problem of creating and updating many jobs at once. All you have to do now is modify the job definitions, re-run the plugin, and load definitions updated with Manage Jenkins#“Reload Configuration from Disk”. This approach also has the advantage of making it easy to store your job configurations in your version control system, which in turn makes it easier to keep track of changes made to the build configurations.

But, you still need to solve the problem of maintaining jobs that share a number of identical properties, such as Maven goals, email recipients, or code repository URLs. For that, the Maven Jenkins Plugin provides jobs inheritance, demonstrated in the following example:

```

<jobs>
  <job>
    <id>google-guice-inheritance-base</id>
    <abstract>true</abstract>
    <jdkName>jdk1.6.0</jdkName>
    <mavenName>apache-maven-3</mavenName>
    <daysToKeep>5</daysToKeep>
    <useUpdate>true</useUpdate>
    <mavenGoals>-B -e -U clean install</mavenGoals>
  </job>

```

⁴ <http://evgeny-goldin.com/wiki/Maven-jenkins-plugin#.3Cjob.3E>

```

    <mail><recipients>jenkins@evgeny-goldin.org</recipients></mail>
  </job>

  <job>
    <id>google-guice-inheritance-trunk</id>
    <parent>google-guice-inheritance-base</parent>
    <repository>
      <remote>http://google-guice.googlecode.com/svn/trunk/</remote>
    </repository>
  </job>

  <job>
    <id>google-guice-inheritance-3.0-rc3</id>
    <parent>google-guice-inheritance-base</parent>
    <repository>
      <remote>http://google-guice.googlecode.com/svn/tags/3.0-rc3/</remote>
    </repository>
  </job>

  <job>
    <id>google-guice-inheritance-2.0-maven</id>
    <parent>google-guice-inheritance-base</parent>
    <mavenName>apache-maven-2</mavenName>
    <repository>
      <remote>http://google-guice.googlecode.com/svn/branches/2.0-maven/
      </remote>
    </repository>
  </job>
</jobs>

```

In this configuration, `google-guice-inheritance-base` is an abstract parent job holding all common properties: JDK name, Maven name, days to keep the logs, SVN update policy, Maven goals, and mail recipients. The three following jobs are very short, merely specifying that they extend a `<parent>` job and add any missing configurations (repository URLs in this case). When generated, they inherit all of the properties from the parent job automatically.

Any inherited property can be overridden, as demonstrated in `google-guice-inheritance-2.0-maven` job where Maven 2 is used instead of Maven 3. If you want to “cancel” an inherited property, you’ll need to override it with an empty value.

Jobs inheritance is a very powerful concept that allows jobs to form hierarchical groups of any kind and for any purpose. You can group your CI, nightly, or release jobs this way, centralizing shared execution triggers, Maven goals, or mail recipients in parent jobs. This approach borrowed from an OOP world solves the problem of maintaining jobs sharing a number of identical properties.

10.5.3. Plugin Support

In addition to configuring a job and reusing its definitions, you can apply special support for a number of Jenkins plugins. Right now, a simplified usage of Parameterized Trigger and Artifactory plugins is provided, with support for other popular plugins planned for future versions.

Below is an example of invoking jobs with the Parameterized Trigger plugin. Using this option assumes you have this plugin installed already:

```
<job>
  <id>google-guice-inheritance-trunk</id>
  ...
  <invoke>
    <jobs>
      google-guice-inheritance-3.0-rc3,
      google-guice-inheritance-2.0-maven
    </jobs>
  </invoke>
</job>

<job>
  <id>google-guice-inheritance-3.0-rc3</id>
  ...
</job>

<job>
  <id>google-guice-inheritance-2.0-maven</id>
  ...
</job>
```

The `<invoke>` element lets you invoke other jobs each time the current job finishes successfully. You can create a pipeline of jobs this way, making sure each job in the pipeline invokes the following one. Note that if there is more than one Jenkins executor available at the moment of invocation, the specified jobs will start running in parallel. For serial execution you'll need to connect each upstream job to a downstream one with `<invoke>`.

By default, invocation happens only when the current job is stable. This can be modified, as shown in the following examples:

```
<invoke>
  <jobs>jobA, jobB, jobC</jobs>
  <always>true</always>
</invoke>

<invoke>
  <jobs>jobA, jobB, jobC</jobs>
  <unstable>true</unstable>
</invoke>

<invoke>
  <jobs>jobA, jobB, jobC</jobs>
  <stable>false</stable>
  <unstable>false</unstable>
  <failed>true</failed>
</invoke>
```

The first invocation in the example above always invokes the downstream jobs. It can be used for a pipeline of jobs that should always be executed even if some of them fail.

The second invocation in the example above invokes downstream jobs even if an upstream job is unstable: the invocation happens regardless of test results. It can be used for a pipeline of jobs that are less sensitive to tests and failures.

The third invocation in the example above invokes downstream jobs only when an upstream job fails but not when it's stable or unstable. You can find this configuration useful when a failing job needs to perform additional actions beyond traditional email notifications.

Artifactory⁵ is a general purpose binary repository that can be used as a Maven repository manager. The Jenkins Artifactory plugin⁶, shown in Figure 10.29, “Artifactory Jenkins plugin configuration”, provides a number of benefits for Jenkins build jobs. We've already reviewed some of them in Section 5.9.4, “Deploying to an Enterprise Repository Manager”, including an ability to deploy artifacts upon job completion or to send build environment info together with artifacts to make them easier to trace.

Figure 10.29. Artifactory Jenkins plugin configuration

You can also use the Artifactory Jenkins plugin in conjunction with the Maven Jenkins Plugin to deploy artifacts to Artifactory, as shown in the following example:

```
<job>
  ...
  <artifactory>
    <name>http://artifactory-server/</name>
    <deployArtifacts>true</deployArtifacts>
    <includeEnvVars>true</includeEnvVars>
    <evenIfUnstable>true</evenIfUnstable>
  </artifactory>
</job>
```

Default deployment credentials are specified when Jenkins is configured in the Manage Jenkins#Configure System screen. They can be also specified for each Jenkins job. The default Maven repositories are `libs-releases-local` and `libs-snapshots-local`. You can find more details in the plugin's documentation at <http://wiki.jenkins-ci.org/display/JENKINS/Artifactory+Plugin>.

10.5.4. Freestyle Jobs

In addition to Maven jobs, the Maven Jenkins Plugin allows you to configure Jenkins freestyle jobs. An example is shown here:

```
<job>
  <id>free-style</id>
  <jobType>free</jobType>
  <scmType>git</scmType>
  <repository>
    <remote>git://github.com/evgeny-goldin/maven-plugins-test.git</remote>
```

⁵ <http://jfrog.org>

⁶ <http://wiki.jenkins-ci.org/display/JENKINS/Artifactory+Plugin>


```

</repository>
<tasks>
  <maven>
    <mavenName>apache-maven-3</mavenName>
    <jvmOptions>-Xmx128m -XX:MaxPermSize=128m -ea</jvmOptions>
    <properties>plugins-version = 0.2.2</properties>
  </maven>
  <shell><command>pwd; ls -al; du -hs .</command></shell>
</tasks>
</job>

```

Freestyle jobs let you execute a shell or batch command, run Maven or Ant, and invoke other jobs. They provide a convenient run-time environment for system scripts or any other kind of activity not readily available with Jenkins or one of its plugins. Using this approach, you can generate freestyle build job configuration files in a similar way to the approach you've seen for Maven build jobs, which can help make your build environment more consistent and maintainable.

10.6. Coordinating Your Builds

Triggering downstream build jobs is easy enough. However, when setting up larger and more complicated build job setups, you sometimes would like builds to be able to run concurrently, or possibly wait for certain build jobs to finish before proceeding. In this section, you'll look at techniques and plugins that can help you do this.

10.6.1. Parallel Builds in Jenkins

Jenkins has built-in support for parallel builds—when a build job starts, Jenkins will assign it to the first available build node, so you can potentially have as many parallel builds running as you have build nodes available.

If you need to run slight variations of the same build job in parallel, multiconfiguration build jobs (see Section 10.4, “Multiconfiguration Build Jobs”) are an excellent option. They can come in handy as a way of accelerating your build process. A typical application of multiconfiguration build jobs in this context is to run integration tests in parallel. One strategy is to set up an integration test build job that can be run in different ways to execute different subsets of the integration tests. You could define separate Maven profiles, for example, or configure your build to use a command-line parameter to decide which tests to run. Once you've set up your build script in this way, it's easy to configure a multiconfiguration build job to run the subsets of your integration tests in parallel.

You can also get Jenkins to trigger several downstream builds in parallel, simply by listing them all in the “Build other projects” field (see Figure 10.30, “Triggering several other builds after a build job”). The subsequent build jobs will be executed in parallel as much as possible. However, as you'll see further on, this may not always be exactly what you need.

Figure 10.30. Triggering several other builds after a build job

10.6.2. Dependency Graphs

Before you investigate the finer points of parallel builds, it's useful to be able to visualize the relationships between your build jobs. The Dependency Graph View plugin analyzes your build jobs and displays a graph describing the upstream and downstream connections between your jobs. This plugin uses graphviz⁷, which you'll need to install on your server if you don't already have it.

This plugin adds a Dependency Graph icon in the main menu, which displays a graph showing the relationships between all the build jobs in your project (at the dashboard level), or all of the build jobs related to the current build job (when you're inside a particular project [see Figure 10.31, “A build job dependency graph”]). What's more, if you click on a build job in the graph, Jenkins will take you directly to the project page of that build job.

Figure 10.31. A build job dependency graph

10.6.3. Joins

When setting up more complicated build pipelines, you frequently come across situations where a build job can't proceed until a number of other build jobs have completed, but these upstream build jobs don't need to be executed sequentially. For example, in Figure 10.31, “A build job dependency graph”, imagine that the **phoenix-deploy-to-uat** build job actually requires three jobs to succeed before it can be executed: **phoenix-compatibility-tests**, **phoenix-load-tests**, and **phoenix-performance-tests**.

You can set this up by using the Joins plugin, which you'll need to install in the usual way via the Update center. Once installed, you configure a join in the build job that initiates the join process (in our example, this would be **phoenix-web-tests**). In our example, you need to modify the **phoenix-web-tests** build job so that it triggers the **phoenix-compatibility-tests**, **phoenix-load-tests**, and **phoenix-performance-tests** first, and then, if these three succeed, the **phoenix-deploy-to-uat** build job.

You do this by simply configuring the Join Trigger field with the name of the **phoenix-deploy-to-uat** build job (see Figure 10.32, “Configuring a join in the phoenix-web-tests build job”). The “Build other projects” field isn't modified, and still lists the build jobs to be triggered immediately after the current one. The Join Trigger field contains the build jobs to be built once all of the immediate downstream build jobs have finished.

Figure 10.32. Configuring a join in the phoenix-web-tests build job

As a result, you no longer need the original build trigger for the final build job, as it's now redundant.

⁷ <http://www.graphviz.org>

This new flow shows up nicely in the dependency graphs as illustrated in Figure 10.33, “A more complicated build job dependency graph”.

Figure 10.33. A more complicated build job dependency graph

10.6.4. Locks and Latches

In other situations, you might be able to run a series of builds in parallel to some degree, but certain build jobs can't be run in parallel because they access concurrent resources. Of course, well-designed build jobs should strive to be as independent as possible, but sometimes this can be difficult. For example, different build jobs may need to access the same test database or files on the hard disk, and doing so simultaneously could potentially compromise the results of the tests. Or, a performance build job may need exclusive access to the test server, in order to have consistent results each time.

The Locks and Latches plugin lets you get around this problem to some extent. This plugin lets you set up “locks” for certain resources, in a similar way to locks in multithreaded programming. Suppose, for example, in the build jobs depicted in Figure 10.33, “A more complicated build job dependency graph”, that the load tests and the performance tests run against a dedicated server, but only one build job can run against this server at any one time. Imagine furthermore that the performance tests for other projects also run against this server.

To avoid contention over the performance server, you could use the Locks and Latches plugin to set up a “lock” reserving access to this server to a single build job at a time. First, in the System Configuration page, you need to add a new lock in the Locks section (see Figure 10.34, “Adding a new lock”). This lock will then be available to all build jobs on the server.

Figure 10.34. Adding a new lock

Next, you need to set up each build job that will be using the contended resource. In the Build Environment section, you'll find a Locks field. Tick the checkbox and select the lock you just created (see Figure 10.35, “Configuring a build job to use a lock”). Once you do this for each of the build jobs that need to access the resource in question, only one of these build jobs will ever be able to run at a given time.

Figure 10.35. Configuring a build job to use a lock

10.7. Build Pipelines and Promotions

CI isn't just about automatically building and testing software, but can also help in the broader context of the software product development and release life cycle. In many organizations, the life of a particular

version of an application or product starts out in development. When it's deemed ready, it's passed on to a QA team for testing. If they consider the version acceptable, they pass it on to selected users for more testing in a User Acceptance Testing (UAT) environment. And if the users are happy, it's shipped out into production. Of course, there are almost as many variations on this as there are software development teams, but one common principle is that specific versions of your software are selected, according to certain quality-related criteria, to be “promoted” to the next stage of the life cycle. This is known as build promotion, and the broader process is known as a build pipeline. In this section, you'll look at how you can implement build pipelines using Jenkins.

10.7.1. Managing Maven Releases with the M2Release Plugin

An important part of any build pipeline is a well-defined release strategy. This involves, among other things, deciding how and when to cut a new release, and how to identify it with a unique label or version number. If you're working with Maven projects, using the Maven Release plugin to handle version numbers comes as a highly recommended practice.

Maven projects use well-defined and well-structured version numbers. A typical version number is made up of three digits (e.g., “1.0.1”). Developers work on SNAPSHOT versions (e.g., “1.0.1-SNAPSHOT”), which, as the name would indicate, are not designed to be definitive. The definitive releases (e.g., “1.0.1”) are built once and deployed to the local enterprise repository (or the central Maven repository for open source libraries), where they can be used in turn by other projects. The version numbers used in Maven artifacts are a critical part of Maven’s dependency management system, and it's strongly advised to stick to the Maven conventions.

The Maven Release plugin automates the process of updating Maven version numbers in your projects. In a nutshell, it verifies, builds and tests your application, bumps up the version numbers, updates your version control system with the appropriate tags, and deploys the released versions of your artifacts to your Maven repository. These are tedious tasks to do by hand, so the Maven Release plugin is an excellent way to automate things.

However, the Maven Release plugin can be fickle, too. Uncommitted or modified local files can cause the process to fail, for example. The process is also time-consuming and CPU intensive, especially for large projects: it builds the application and runs the entire set of unit and integration tests several times, checks out a fresh copy of the source code from the repository, and uploads many artifacts to the Enterprise repository. Indeed, this isn't the sort of thing you want running on a developer machine.

So, it makes good sense to run this process on your build server.

One way to do this is to set up a special manual build job to invoke the Maven Release plugin. However, the M2Release plugin proposes a simpler approach. Using this plugin, you can add the ability to build a Maven release version in an existing build job. This way you can avoid duplicating build jobs unnecessarily, making build job maintenance easier.

Once you've installed this plugin, you can define any build job to also propose a manual Maven Release step. You do this by ticking the “Maven release build” checkbox in the Build Environment section (see

Figure 10.36, “Configuring a Maven release using the M2Release plugin”). Here, you define the goals you want to execute to trigger the build (typically `release:prepare release:perform`).

Figure 10.36. Configuring a Maven release using the M2Release plugin

Once you've set this up, you can trigger a Maven release manually using a new menu option called “Perform Maven Release” (see Figure 10.37, “The Perform Maven Release menu option”).

Figure 10.37. The Perform Maven Release menu option

This will kick off a special build job using the goals you provided in the plugin configuration (see Figure 10.38, “Performing a Maven release in Jenkins”). Jenkins gives you the option to either use the default version numbers provided by Maven (for example, version 1.0.1-SNAPSHOT will be released as version 1.0.1, and the development version number bumped up to 1.0.2-SNAPSHOT), or to provide your own custom numbers. If you want to release a major version, for example, you might choose to manually specify 1.1.0 as the release version number and 1.1.1-SNAPSHOT as the next development version number.

If you have a multimodule Maven project, you can choose to provide a single version number configuration for all modules, or provide a different version number update for each module, something that's generally not recommended.

Figure 10.38. Performing a Maven release in Jenkins

Depending on your SCM configuration, you may also need to provide a valid SCM username and password to allow Maven to create tags in your source code repository.

The professional edition of the Nexus Enterprise Repository provides a feature called Staging Repositories, which is a way of deploying artifacts to a special staging area for further tests before releasing them officially. If you're using this feature, you need to fine-tune your build server configuration for best results.

Nexus Professional works by creating a new staging area for each unique IP address, deploy user, and HTTP User agent. A given Jenkins build machine will always have the same IP address and user. However, you'll typically want to have a separate staging area for each build. The trick, then, is to configure Maven to use a unique HTTP User-Agent for the deployment process. You can do this by configuring the `settings.xml` file on your build server to contain something along the following lines (the ID must match the ID for the release repository in the deployment section of your project):

```
<server>
  <id>nexus</id>
```

```
<username>my_login</username>
<password>my_password</password>
<configuration>
  <httpHeaders>
    <property>
      <name>User-Agent</name>
      <value>Maven m2Release (java:25.131-b11 ${env.BUILD_TAG }</value>
    </property>
  </httpHeaders>
</configuration>
</server>
```

10.7.2. Copying Artifacts

During a build process involving several build jobs, such as the one illustrated in Figure 10.33, “A more complicated build job dependency graph”, it can sometimes be useful to reuse artifacts produced by one build job in a subsequent build job. For example, you may want to run a series of web tests in parallel on separate machines, using local application servers for improved performance. In this case, it makes sense to retrieve the exact binary artifact that was produced in the previous build, rather than rebuilding it each time or, if you're using Maven, relying on a SNAPSHOT build deployed to your enterprise repository. Indeed, both these approaches may run the risk of inconsistent build results: if you use a SNAPSHOT from the enterprise repository, for example, you'll be using the latest SNAPSHOT build, which may not necessarily be the one built in the upstream build job.

The Copy Artifact plugin lets you copy artifacts from an upstream build and reuse them in your current build. Once you've installed this plugin and restarted Jenkins, you'll be able to add a new type of build step called “Copy artifacts from another project” to your freestyle build jobs (see Figure 10.39, “Adding a “Copy artifacts from another project” build step”).

Figure 10.39. Adding a “Copy artifacts from another project” build step

This new build step lets you copy artifacts from another project into the workspace of the current project. You can specify any other project, though most typically it will be one of the upstream build jobs. And of course you can specify, with a great deal of flexibility and precision, the exact artifacts that you want to copy.

You need to specify where to find the files you want in the other build job's workspace, and where Jenkins should put them in your current project's workspace. This can be a flexible regular expression (such as `**/*.war`, for any WAR file produced by the build job), or it can be much more precise (such as `gameoflife-web/target/gameoflife.war`). Note that by default, Jenkins will copy the directory structure along with the file you retrieve, so if the WAR file you're after is nested inside the `target` directory of the `gameoflife-web` module, Jenkins will place it inside the `gameoflife-web/target` directory in your current workspace. If this isn't to your tastes, you can tick the “Flatten directories” option to tell Jenkins to put all of the artifacts at the root of the directory you specify (or, by default, in your project workspace).

In many cases, you'll simply want to retrieve artifacts from the most recent successful build. However, sometimes you may want more precision. The “Which builds” field lets you specify where to look for artifacts in a number of other ways, including the latest saved build (builds which have been marked to “keep forever”), the latest successful build, or even a specific build number.

If you've installed the Build Promotion plugin (see Section 10.7.3, “Build Promotions”), you can also select the latest promoted artifact in a particular promotion process. To do this, choose “Specify by permalink”, then choose the appropriate build promotion process. This is an excellent way of ensuring a consistent and reliable build pipeline. For example, you can configure a build promotion process to trigger a build that copies a generated WAR file from the latest promoted build and then deploys it to a particular server. This ensures that you deploy precisely the right binary file, even if other builds have occurred since.

If you're copying artifacts from a multimodule Maven build job, Jenkins will, by default, copy all of the artifacts from that build. However, often times you're only interested in one specific artifact (such as the WAR file artifact in a web application, for example).

This plugin is particularly useful when you need to run functional or performance tests on your web application. It's often a useful strategy to place these tests in a separate project, and not as part of your main build process. This makes it easier to run these tests against different servers or run the subsets of the tests in parallel, all the while using the same binary artifact to deploy and test.

For example, imagine that you have a default build job called *gameoflife* that generates a WAR file, and you'd like to deploy this WAR file to a local application server and run a series of functional tests. Furthermore, you want to be able to do this in parallel on several distributed machines.

One way to do this would be to create a dedicated Maven project designed to run the functional tests against an arbitrary server. Then, you'd set up a build job to run these functional tests. This build job would use the Copy Artifact plugin to retrieve the latest WAR file (or even the latest promoted WAR file, for more precision), and deploy it to a local Tomcat instance using Cargo. This build job could then be set up as a configurable (“matrix”) build job, and run in parallel on several machines, possibly with extra configuration parameters to filter the tests run by each build. Each build run would then be using its own copy of the original WAR file. An example of a configuration like this is illustrated in Figure 10.40, “Running web tests against a copied WAR file”.

Figure 10.40. Running web tests against a copied WAR file

The Copy Artifact plugin isn't limited to fetching files from conventional build jobs. You can also copy artifacts from multiconfiguration build jobs (see Section 10.4, “Multiconfiguration Build Jobs”). Artifacts from each executed configuration will be copied into the current workspace, each in its own directory. Jenkins will build a directory structure using the axes that were used in the multiconfiguration build. For example, imagine you need to produce a highly-optimized version of your product for a number of different targeted databases and application servers. You could do this with a

multiconfiguration build job like the one illustrated in Figure 10.41, “Copying from a multiconfiguration build”.

Figure 10.41. Copying from a multiconfiguration build

The Copy Artifacts plugin can duplicate any and all of the artifacts produced by this build job. If you specify a multiconfiguration build as the source of your artifacts, the plugin will copy artifacts from all of the configurations into the workspace of the target build job, using a nested directory structure based on the multiconfiguration build axes. For example, if you define the target directory as `multi-config-artifacts`, Jenkins will copy artifacts into a number of subdirectories in the target directory, each with a name corresponding to the particular set of configuration parameters. So, using the build job illustrated in Figure 10.41, “Copying from a multiconfiguration build”, the JAR file customized for Tomcat and MySQL would be copied to the `$WORKSPACE/multi-config-artifacts/APP_SERVER/tomcat/DATABASE/mysql` directory.

10.7.3. Build Promotions

In the world of CI, not all builds are created equal. For example, you may want to deploy the latest version of your web application to a test server, but only after it has passed a number of automated functional and load tests. Or, you may want testers to be able to flag certain builds as being ready for UAT deployment, once they’ve completed their own testing.

The Promoted Builds plugin lets you identify specific builds that have met additional quality criteria, and to trigger actions on these builds. For example, you may build a web application in one build job, run a series of automated web tests in a subsequent build, and then deploy the WAR file generated to the UAT server for further manual testing.

Let’s see how this works in practice. In the project illustrated above, a default build job (**phoenix-default**) runs some unit and integration tests, and produces a WAR file. This WAR file is then reused for more extensive integration tests (in the **phoenix-integration-tests** build job) and then for a series of automated web tests (in the **phoenix-web-test** build job). If the build passes the automated web tests, you’d like to deploy the application to a functional testing environment where it can be tested by human testers. The deployment to this environment is implemented in the **phoenix-test-deploy** build job. Once the testers have validated a version, it can be promoted into UAT, and then into production. The full promotion strategy is illustrated in Figure 10.42, “Build jobs in the promotion process”.

Figure 10.42. Build jobs in the promotion process

This strategy is easy to implement using the Promoted Builds plugin. Once you’ve installed this in the usual way, you’ll find a new “Promote builds when” checkbox on the job configuration page. You use this option to set up build promotion processes. You define one or more build promotion

processes in the initial build job of process (**phoenix-default** in this example), as illustrated in Figure 10.43, “Configuring a build promotion process”. A build job may be the starting point of several build promotion processes, some automated, and some manual. In Figure 10.43, “Configuring a build promotion process”, for example, there's an automated build promotion process called **promote-to-test** and a manual one called **promote-to-uat**. Automated build promotion processes are triggered by the results of downstream build jobs. Manual promotion processes (indicated by ticking the ‘Only when manually approved’ checkbox) can only be triggered by user intervention.

Figure 10.43. Configuring a build promotion process

Let's look at configuring the automated **promote-to-test** build process.

The first thing you need to define is how this build promotion process will be triggered. Build promotion can be either automatic, based on the result of a downstream build job, or manually activated by a user. In Figure 10.43, “Configuring a build promotion process”, the build promotion for this build job will be automatically triggered when the automated web tests (executed by the **phoenix-web-tests** build job) are successful.

You can also have certain build jobs that can only be promoted manually, as illustrated in Figure 10.44, “Configuring a manual build promotion process”. Manual build promotion is used for cases where human intervention is needed to approve a build promotion. Deployment to UAT or production are common examples of this. Another example is where you want to temporarily suspend automatic build promotions for a short period, such as nearing a release.

Manual builds, as the name suggests, need to be manually approved to be executed. If the promotion process is to trigger a parameterized build job, you can also provide parameters that the approver will need to enter when approving. In some cases, it can also be useful to designate certain users who are allowed to activate the manual promotion. You can do this by specifying a list of users or groups in the Approvers list.

Figure 10.44. Configuring a manual build promotion process

Sometimes, it's useful to give some context to the person approving a promotion. When you set up a manual promotion process, you can also specify other conditions which must be met, in particular downstream (or upstream) build jobs which must have been built successfully (see Figure 10.45, “Viewing the details of a build promotion”). These will appear in the “Met Qualifications” (for the successful build jobs) and in “Unmet Qualifications” (for the build jobs that failed or haven't been executed yet).

Figure 10.45. Viewing the details of a build promotion

Next, you need to tell Jenkins what to do when the build is promoted. You do this by adding actions, just like in a freestyle build job. This makes build promotions extremely flexible, as you can add virtually any action available to a normal freestyle build job, including any additional steps made available by the plugins installed in your Jenkins instance. Common actions include invoking Maven or Ant scripts, deploying artifacts to a Maven repository, or triggering another build job.

One important thing to remember here is that you can't rely on files in the workspace when promoting your build. Indeed, by the time you promote the build, either automatically or manually, other build jobs may have deleted or rewritten the files you need to use. For this reason, it's unwise, for example, to deploy a WAR file directly from the workspace to an application server from within a build promotion process. A more robust solution is to trigger a separate build job and to use the Copy Artifacts plugin (see Section 10.7.2, “Copying Artifacts”) to retrieve precisely the right file. In this case, you'll be copying artifacts that you've configured Jenkins to conserve, rather than copying the files directly from the workspace.

For build promotion to work correctly, Jenkins needs to be able to precisely link downstream build jobs to upstream ones. The most accurate way to do this is by using fingerprints. In Jenkins, a fingerprint is the MD5 checksum of a file produced by or used in a build job. By matching fingerprints, Jenkins is able to identify all of the builds which use a particular file.

In the context of build promotion, a common strategy is to build your application once, and then to run tests against the generated binary files in a series of downstream build jobs. This approach works well with build promotion, but you need to ensure that Jenkins fingerprints the files that are shared or copied between build jobs. In the example shown in Figure 10.43, “Configuring a build promotion process”, for instance, you need to do two things (Figure 10.46, “Using fingerprints in the build promotion process”). First, you need to archive the generated WAR file so that it can be reused in the downstream project. Secondly, you need to record a fingerprint of the archived artifacts. You do this by ticking the “Record fingerprints of files to track usage” option, specifying the files you want to fingerprint. A useful shortcut is simply to fingerprint all archived files, since these are the files that will typically be retrieved and reused by the downstream build jobs.

Figure 10.46. Using fingerprints in the build promotion process

This is all you need to do to configure the initial build process. The next step is to configure the integration tests executed in the **phoenix-integration** build job. Here, you use the Copy Artifact plugin to retrieve the WAR file generated by the **phoenix-default** build job (see Figure 10.47, “Fetching the WAR file from the upstream build job”). Since this build job is triggered immediately after the **phoenix-default** build job, you can simply fetch the WAR file from the latest successful build.

Figure 10.47. Fetching the WAR file from the upstream build job

This isn't quite all you need to do for the integration tests, however. The **phoenix-integration** build job is followed by the **phoenix-web** build job, which executes the automated web tests. To ensure that the same WAR file is used at each stage of the build process, you need to retrieve it from the upstream **phoenix-integration** build job, and not from the original **phoenix-default** build job (which may have been executed again in the meantime). So, you also need to archive the WAR file in the **phoenix-integration** build job (see Figure 10.48, “Archiving the WAR file for use in the downstream job”).

Figure 10.48. Archiving the WAR file for use in the downstream job

In the **phoenix-web** build job, you then fetch the WAR file from the **phoenix-integration** build job, using a configuration very similar to the one shown above (see Figure 10.49, “Fetching the WAR file from the integration job”).

Figure 10.49. Fetching the WAR file from the integration job

For the build promotion process to work properly, there's one more important thing you need to configure in the **phoenix-web** build job. As we discussed earlier, Jenkins needs to be able to be sure that the WAR file used in these tests is the same one generated by the original build. You do this by activating fingerprinting on the WAR file you fetched from the **phoenix-integration** build job (which, remember, was originally built by the **phoenix-default** build job). Since you've copied this WAR file into the workspace, a configuration like the one in Figure 10.50, “You need to determine the fingerprint of the WAR file we use” will work just fine.

Figure 10.50. You need to determine the fingerprint of the WAR file we use

The final step is to configure the **phoenix-deploy-to-test** build job to retrieve the last promoted WAR file (rather than just the last successful one). To do this, you use the Copy Artifact plugin again, but this time you choose the “Specified by permalink” option. Here Jenkins will propose, among other things, the build promotion processes configured for the build job you're copying from. So, in Figure 10.51, “Fetching the latest promoted WAR file”, you're fetching the last promoted WAR file build by the **phoenix-default** job, which is precisely what you want.

Figure 10.51. Fetching the latest promoted WAR file

Your promotion process is now ready for action. When the automated web tests succeed for a particular build, the original build job will be promoted and the corresponding WAR file deployed to the test environment. Promoted builds are indicated by a star in the build history (see Figure 10.52, “Promoted

builds are indicated by a star in the build history”). By default, the stars are yellow, but you can configure the color of the star in the build promotion setup.

Figure 10.52. Promoted builds are indicated by a star in the build history

You can also use the “Promotion Status” menu entry (or click on the colored star in the build history) to view the details of a particular build promotion, and even to rerun a promotion manually (see Figure 10.45, “Viewing the details of a build promotion”). Any build promotion can be triggered manually by clicking on “Force promotion” (if this build job has never been promoted) or “Re-execute promotion” (if it has).

10.7.4. Aggregating Test Results

When distributing different types of tests across different build jobs, it's easy to lose a global vision about the overall test results. Test results are scattered among the various build jobs, without a central place to see the total number of executed and failing tests.

A good way to avoid this problem is to use the Aggregated Test Results feature of Jenkins. This will retrieve any test results recorded in the downstream jobs, and aggregate them in the upstream build job. You can configure this in the initial (upstream) build job by ticking the “Aggregate downstream test results” option (see Figure 10.53, “Reporting on aggregate test results”).

Figure 10.53. Reporting on aggregate test results

The aggregate test results can be seen in the build details page (see Figure 10.54, “Viewing aggregate test results”). Unfortunately, these aggregate test results don't appear in the overall test results, but you can display the full list of tests executed by clicking on the Aggregate Test Result link on the individual build page.

Figure 10.54. Viewing aggregate test results

For this to work correctly, you need to ensure that you've configured fingerprinting for the binary files you use at each stage. Jenkins will only aggregate downstream test results from builds containing an artifact with the same fingerprint.

10.7.5. Build Pipelines

The final plugin you'll be looking at in this section is the Build Pipeline plugin. The Build Pipelines plugin takes the idea of build promotion further, and helps you design and monitor deployment pipelines. A deployment pipeline is a way of orchestrating your build through a series of quality gates, with automated or manual approval processes at each stage, culminating with deployment into production.

The Build Pipeline plugin provides an alternative way to define downstream build jobs. A build pipeline, unlike conventional downstream dependencies, is considered to be a linear process, a series of build jobs executed in sequence.

To use this plugin, start by configuring the downstream build jobs for each build job in the pipeline, using the “Build other projects” field just as you'd normally do. The Build Pipelines plugin uses the standard upstream and downstream build configurations, and for automatic steps this is all you need to do. However, the Build Pipeline plugin also supports manual build steps, where a user has to manually approve the next step. For manual steps, you also need to configure In the **Post-build Actions** of your upstream build job: just tick the “Build Pipeline Plugin -> Specify Downstream Project” box, select the next step in your project, and tick the “Require manual build executor” option (see Figure 10.55, “Configuring a manual step in the build pipeline”).

Figure 10.55. Configuring a manual step in the build pipeline

Once you've set up your build process to your satisfaction, you can configure the build pipeline view. You can create this view just like any other view (see Figure 10.56, “Creating a Build Pipeline view”).

Figure 10.56. Creating a Build Pipeline view

There's a trick when it comes to configuring the view, however. At the time of writing, there's no menu option or button that lets you configure the view directly. In fact, you need to enter the URL manually. Fortunately, this isn't difficult: just add `/configure` to the end of the URL shown when you're displaying this view. For example, if you've named your view “phoenix-build-pipeline”, as shown here, the URL to configure this view would be `http://my_jenkins_server/view/phoenix-build-pipeline`. (see Figure 10.57, “Configuring a Build Pipeline view”).

Figure 10.57. Configuring a Build Pipeline view

The most important thing to configure in this screen is the initial job. This marks the starting point of your build pipeline. You can define multiple build pipeline views, each with a different starting job. You can also configure the maximum number of build sequences to appear on the screen at once.

Once you've configured the starting point, you can return to the view to see the current state of your build pipeline. Jenkins displays the successive related build jobs horizontally, using a specific color to indicate the outcome of each build (Figure 10.58, “A Build Pipeline in action”). There's a column for each build job in the pipeline. Whenever the initial build job kicks off, a new row appears on this page. As the build progresses through the successive build jobs in the pipeline, Jenkins will add a colored box in the successive columns, indicating the outcome of each stage. You can click on the box to drill down

into a particular build result for more details. Finally, if a manual execution is required, a button will be displayed where the user can trigger the job.

Figure 10.58. A Build Pipeline in action

This plugin is still relatively new, and doesn't integrate with all of the other plugins you've seen here. In particular, it's really designed for a linear build pipeline, and doesn't cope well with branches or parallel build jobs. Nevertheless, it does give an excellent global vision of a build pipeline.

10.8. Conclusion

CI build jobs are much more than simply the scheduled execution of build scripts. In this chapter you've reviewed a number of tools and techniques enabling you to go beyond your typical build jobs, combining them so that they can work together as part of a larger process. You've seen how parameterized and multiconfiguration build jobs add an element of flexibility to ordinary build jobs by allowing you to run the same build job with different sets of parameters. Other tools help coordinate and orchestrate groups of related build jobs. The Joins and Locks and Latches plugins helps you coordinate build jobs running in parallel. And the Build Promotions and Build Pipelines plugins, with the help of the Copy Artifacts plugin, make it relatively easy to design and configure complex build promotion strategies for your projects.

Chapter 11. Distributed Builds

11.1. Introduction

Arguably one of the more powerful features of Jenkins is its ability to dispatch build jobs across a large number of machines. It's quite easy to set up a farm of build servers, either to share the load across multiple machines, or to run build jobs in different environments. This is a very effective strategy which can potentially increase the capacity of your CI infrastructure dramatically.

Distributed builds are generally used either to absorb extra load, for example absorbing spikes in build activity by dynamically adding extra machines as required, or to run specialized build jobs in specific operating systems or environments. For example, you may need to run particular build jobs on a particular machine or operating system. Perhaps if you need to run web tests using Internet Explorer, you'll need to use a Windows machine. Or, one of your build jobs may be particularly resource-heavy, and need to be run on its own dedicated machine so as not to penalize your other build jobs.

Demand for build servers can also fluctuate over time. If you're working with product release cycles, you may need to run a much higher number of builds jobs towards the end of the cycle, for example, when more comprehensive functional and regression test suites may be run more frequently.

In this chapter, we'll discuss how to set up and manage a farm of build servers using Jenkins.

11.2. The Jenkins Distributed Build Architecture

Jenkins uses a master/slave architecture to manage distributed builds. Your main Jenkins server (the one you've been using so far) is the master. In a nutshell, the master's job is to handle scheduling build jobs, dispatching builds to the slaves for the actual execution, monitoring the slaves (possibly taking them online and offline as required), and recording and presenting the build results. Even in a distributed architecture, a master instance of Jenkins can also execute build jobs directly.

The job of the slaves is to do as they're told, which involves executing build jobs dispatched by the master. You can configure a project to always run on a particular slave machine, or a particular type of slave machine, or simply let Jenkins pick the next available slave.

A slave is a small Java executable that runs on a remote machine and listens for requests from the Jenkins master instance. Slaves can (and usually do) run on a variety of operating systems. The slave instance can be started in a number of different ways, depending on the operating system and network architecture. Once the slave instance is running, it communicates with the master instance over a TCP/IP connection. We'll look at different setups in the rest of this chapter.

11.3. Master/Slave Strategies in Jenkins

There are a number of different ways to configure a distributed build farm using Jenkins, depending on your operating systems and network architecture. In all cases, the fact that a build job is being run on a slave, and how that slave is managed, is transparent to the end-user: the build results and artifacts will always end up on the master server.

Creating a new Jenkins slave node is a straightforward process. First, go to the Manage Jenkins screen and click on Manage Nodes. This screen displays the list of slave agents (also known as “Nodes” in more politically correct terms), shown in Figure 11.1, “Managing slave nodes”. From here, you can set up new nodes by clicking on the New Node button. You can also configure some of the parameters related to your distributed build setup (see Section 11.5, “Node Monitoring”).

Figure 11.1. Managing slave nodes

There are several different strategies when it comes to managing Jenkins slave nodes, depending on your target operating systems and other architectural considerations. These strategies affect the way you configure your slave nodes, so you need to consider them separately. In the following sections, we'll look at the most frequently used ways to install and configure Jenkins slaves:

- Starting the slave agents from the master using SSH
- Starting the slave agents manually using Java Web Start
- Installing the slave agents as a Windows service
- Starting the slave agents directly from the command line on the slave machines

Each of these strategies has its uses, advantages, and inconveniences. Let's look at each in turn.

11.3.1. The Master Starts the Slave Agent Using SSH

If you're working in a Unix environment, the most convenient way to start a Jenkins slave is undoubtedly using SSH. Jenkins has its own built-in SSH client, and almost all Unix environments support SSH (usually running as `sshd`) out of the box.

To create a Unix-based slave, click on the New Node button as we mentioned above. This will prompt you to enter the name of your slave, and its type (see Figure 11.2, “Creating a new slave node”). At the time of writing, only “dumb slaves” are supported out of the box; “dumb” slaves are passive beasts that simply respond to build job requests from the master node. This is the most common way to set up a distributed build architecture, and the only option available in a default installation.

Figure 11.2. Creating a new slave node

In this screen, you simply need to provide a name for your slave. When you click on OK, Jenkins will let you provide more specific details about your slave machine (see Figure 11.3, “Creating a Unix slave node”).

Figure 11.3. Creating a Unix slave node

The name is simply a unique way of identifying your slave machine. It can be anything, but it may help if the name reminds you of the physical machine it's running on. It also helps if the name is file-system and URL-friendly. It will work with spaces, but you'll make life easier for yourself if you avoid them. So, “Slave-1” is better than “Slave 1”.

The description is also purely for human consumption, and can be used to indicate why you'd use this slave rather than another.

Like on the main Jenkins configuration screen, the number of executors lets you define how many concurrent build jobs this node can execute.

Every Jenkins slave node also needs a place that it can call home, or, more precisely, a dedicated directory on the slave machine that the slave agent can use to run build jobs. You define this directory in the Remote FS root field. You need to provide a local, OS-specific path, such as `/var/jenkins` for a Unix machine, or `C:\jenkins` on Windows. Nothing mission-critical is stored in this directory—everything important is transferred back to the master machine once the build is done. So, you usually don't need to be as concerned with backing up these directories as you should be on the master.

Labels are a particularly useful concept when your distributed build architecture begins to grow in size. You can define labels, or tags, for each build node, and then configure a build job to run only on a slave node with a particular label. Labels might relate to operating systems (Unix, Windows, MacOSx, etc.), environments (staging, UAT, development, etc.), or any criteria that you find useful. For example, you could configure your automated WebDriver/Selenium tests to run using Internet Explorer, but only on slave nodes with the “Windows” label.

The Usage field lets you configure how intensively Jenkins will use this slave. You have the choice of three options: use it as much as possible, reserve it for dedicated build jobs, or bring it online as required.

The first option, “Utilize this slave as much as possible”, tells Jenkins to use this slave freely as soon as it becomes available, for any build job that it can run. This is by far the most commonly used one, and is generally what you want.

There are times, however, when this second option comes in handy. In the project configuration, you can tie a build job to a specific node—this is useful when a particular task, such as automated deployment or a performance test suite, needs to be executed on a specific machine. In this case, the “Leave this machine for tied jobs only” option makes good sense. You can take this further by setting the maximum number of Executors to 1. In this case, not only will this slave be reserved for a particular type of job, but it will only ever be able to run one of these build jobs at any one time. This is a very useful configuration

for performance and load tests, where you need to reserve the machine so that it can execute its tests without interference.

The third option is “Take this slave on-line when in demand and off-line when idle” (see Figure 11.4, “Taking a slave off-line when idle”). As the name indicates, this option tells Jenkins to bring this slave online when demand is high, and to take it offline when demand subsides. This lets you keep some build slaves in reserve for periods of heavy use, without having to maintain a slave agent running on them permanently. When you choose this option, you also need to provide some extra details. The “In demand delay” indicates how many minutes jobs must have been waiting in the queue before this slave will be brought online. The Idle delay indicates how long the slave needs to be idle before Jenkins will take it off-line.

Figure 11.4. Taking a slave off-line when idle

The launch method is where you decide how Jenkins will start the node, as we mentioned earlier. For the configuration we're discussing here, you'd choose “Launch slave agents on Unix machines via SSH”. The Advanced button lets you enter the additional details that Jenkins needs to connect to the Unix slave machine: a host name, a login and password, and a port number. You can also provide a path to the SSH private key file on the master machine (e.g., `id_dsa` or `id_rsa`) to use for “password-less” Public/Private Key authentication.

You can also configure when Jenkins starts and stops the slave. By default, Jenkins simply keeps the slave running and uses it whenever required (the “Keep this slave on-line as much as possible” option). If Jenkins notices that the slave has gone offline (for example due to a server reboot), it will attempt to restart the slave if it can. Alternatively, Jenkins can be more conservative with your system resources and take the slave offline when Jenkins doesn't need it. To do this, simply choose the “Take this slave on-line when in demand and off-line when idle” option. This is useful if you have regular spikes and lulls of build activity, as an unused slave can be taken offline to conserve system resources for other tasks, and brought back online when required.

Jenkins also needs to know where it can find the build tools it needs for your build jobs on the slave machines. This includes JDKs as well as build tools such as Maven, Ant, and Gradle. If you've configured your build tools to be automatically installed, you'll usually have no extra configuration to do for your slave machines; Jenkins will download and install the tools as required. On the other hand, if your build tools are installed locally on the slave machine, you'll need to tell Jenkins where it can find them. You do this by ticking the Tool Locations checkbox, and providing the local paths for each of the tools you'll need for your build jobs (see Figure 11.5, “Configuring tool locations”).

Figure 11.5. Configuring tool locations

You can also specify environment variables. These will be passed into your build jobs, and can be a good way to allow your build jobs to behave differently depending on where they're being executed.

Once you've done this, your new slave node will appear in the list of computers on the Jenkins Nodes page (see Figure 11.6, “Your new slave node in action”).

Figure 11.6. Your new slave node in action

11.3.2. Starting the Slave Agent Manually Using Java Web Start

Another option is to start a slave agent from the slave machine itself using Java Web Start. This approach is useful if the server can't connect to the slave, for example if the slave machine is running on the other side of a firewall. It works no matter what operating system your slave is running. However, it's more commonly used for Windows slaves. It does suffer a few major drawbacks: the slave node can't be started, or restarted, automatically by Jenkins. So, if the slave goes down, the master instance can't restart the slave.

When you do this on a Windows machine, you need to start the Jenkins slave manually at least once. This involves opening a browser on the machine, opening the slave node page on the Jenkins master, and launching the slave using a very visible JNLP icon. However, once you have launched the slave, you can install it as a Windows service.

There are also times in a Unix environment when you need to do this from the command line. You may need to do this because of firewalls or other networking issues, or because SSH isn't available in your environment.

Let's step through both these processes.

The first thing you need to do in all cases is create a new slave. As with any other slave node, you do this by clicking on the New Node menu entry in the Nodes screen. When entering the details concerning your slave node, make sure you choose “Launch slave agents via JNLP” in the Launch Method field (see Figure 11.7, “Creating a slave node for JNLP”). Also, remember that if this is to be a Windows slave node, the Remote FS root needs to be a Windows path (such as `C:\jenkins-slave`). This directory doesn't have to exist: Jenkins will create it automatically if it's missing.

Figure 11.7. Creating a slave node for JNLP

Once you've saved this configuration, next log in to the slave machine and open the Slave node screen in a browser, as shown in Figure 11.8, “Launching a slave via Java Web Start”. You'll see a large orange Launch button—if you click on this button, you should be able to start a slave agent directly from within your browser.

Figure 11.8. Launching a slave via Java Web Start

If all goes well, this will open up a small window indicating that your slave agent is now running (see Figure 11.9, “The Jenkins slave agent in action”).

Figure 11.9. The Jenkins slave agent in action

Browsers are fickle, however, and Java Web Start isn't always easy to use. This approach usually works best with Firefox, although you must have the JRE installed beforehand to make Firefox Java-aware. Using JNLP with Internet Explorer requires some (considerable) fiddling to associate *.jnlp files with the Java Web Start executable, which is the file **javaws** that you'll find in the Java `bin` directory. In fact, it's probably easier just to start it from the command line as discussed below.

A more reliable, albeit low-level, approach is to start the slave from the command line. To do this, simply invoke the **javaws** executable from a command window as shown here:

```
C:> javaws http://build.myorg.com/jenkins/computer/windows-slave-1/slave-agent.jnlp
```

The exact command that you need to execute, including the correct URL, is conveniently displayed in the Jenkins slave node window just below the JNLP launch button (see Figure 11.8, “Launching a slave via Java Web Start”).

If security is activated on your Jenkins server, Jenkins will communicate with the slave on a specific nonstandard port. If, for some reason, this port is inaccessible, the slave node will fail to start and will display an error message similar to the one shown in Figure 11.10, “The Jenkins slave failing to connect to the master”.

Figure 11.10. The Jenkins slave failing to connect to the master

This is usually a sign that a firewall is blocking a port. By default, Jenkins picks a random port to use for communicating with its slaves. However, if you need to have a specific port that your firewall has authorized, you can force Jenkins to use a fixed port in the System configuration screen by selecting Fixed in the “TCP port for JNLP slave agents” option, as shown in Figure 11.11, “Configuring the Jenkins slave port”.

Figure 11.11. Configuring the Jenkins slave port

11.3.3. Installing a Jenkins Slave as a Windows Service

Once you have the slave up and running on your Windows machine, you can save yourself the bother of having to restart it manually each time your machine reboots by installing it as a Windows service. To do this, select the “Install as Windows Service” menu option in the File menu of the slave agent window (see Figure 11.12, “Installing the Jenkins slave as a Windows service”).

Figure 11.12. Installing the Jenkins slave as a Windows service

Once this is done, your Jenkins slave node will start automatically whenever the machine starts up, and can be administered just like any other Windows service (see Figure 11.13, “Managing the Jenkins Windows service”).

Figure 11.13. Managing the Jenkins Windows service

11.3.4. Starting the Slave Node in Headless Mode

You can also start a slave agent in headless mode, directly from the command line. This is useful if you don’t have a user interface available, for example when you’re starting a JNLP slave node on a Unix machine. If you’re working with Unix machines, it’s generally easier and more flexible just to use an SSH connection, but sometimes there are network or architecture constraints that prevent you from using SSH. In cases like this, it’s still possible to run a slave node from the command line.

To start the slave node this way, you need to use Jenkins’ `slave.jar` file. You can find this in `JENKINS_HOME/war/WEB-INF/slave.jar`. Once you have located this file and copied it onto the Windows slave machine, you can run it as follows:

```
java -jar slave.jar \  
-jnlpUrl http://build.myorg.com/jenkins/computer/windows-slave-1/slave-agent.jnlp
```

If your Jenkins server requires authentication, just pass in the `-auth username:password` option:

```
java -jar slave.jar \  
-jnlpUrl http://build.myorg.com/jenkins/computer/windows-slave-1/slave-agent.jnlp  
-auth scott:tiger
```

Once you’ve started the slave agent, be sure to install it as a Windows service, as discussed in the previous section.

11.3.5. Starting a Windows Slave as a Remote Service

Jenkins can also manage a remote Windows slave as a Windows service, using the Windows Management Instrumentation (WMI) service which is installed out of the box on Windows 2000 and later (see Figure 11.14, “Letting Jenkins control a Windows slave as a Windows service”). When you choose this option, you just need to provide a Windows username and password. The name of the node must be the hostname of the slave machine.

This is certainly convenient, as it doesn’t require you to physically connect to the Windows machine to set it up. However, it does have limitations—in particular, you can’t run any applications requiring

a graphical interface, so you can't use a slave set up this way for web testing, for example. In practice this can be a little tricky to set up, as you may need to configure the Windows firewall to open the appropriate services and ports. If you run into trouble, make sure that your network configuration allows TCP connections to ports 135, 139, and 445, and UDP connections to ports 137 and 138 (see <https://wiki.jenkins-ci.org/display/JENKINS/Windows+slaves+fail+to+start+via+DCOM> for more details).

Figure 11.14. Letting Jenkins control a Windows slave as a Windows service

11.4. Associating a Build Job with a Slave or Group of Slaves

In the previous section, you saw how to assign labels to your slave nodes. This is a convenient way to group your slave nodes according to characteristics such as operating system, target environment, database type, or any other criteria that are relevant to your build process. A common application of this practice is to run OS-specific functional tests on dedicated slave nodes, or to reserve a particular machine exclusively to performance tests.

Once you've assigned labels to your slave nodes, you also need to tell Jenkins where it can run the build jobs. By default, Jenkins will simply use the first available slave node, which usually results in the best overall turn-around time. If you need to tie a build job to a particular machine or group of machines, you need to tick the “Restrict where this project can be run” checkbox in the build configuration page (see Figure 11.15, “Running a build job on a particular slave node”). Next, enter the name of the machine or a label identifying a group of machines into the Label Expression field. Jenkins will provide a dynamic dropdown showing the available machine names and labels as you type.

Figure 11.15. Running a build job on a particular slave node

This field also accepts boolean expressions, allowing you to define more complicated constraints about where your build job should run. How to use these expressions is best illustrated by an example. Suppose you have a build farm with Windows and Linux slave nodes (identified by the labels “windows” and “linux”), distributed over three sites (“sydney”, “sanfrancisco”, and “london”). Your application also needs to be tested against several different databases (“oracle”, “db2”, “mysql”, and “postgres”). You also use labels to distinguish slave nodes used to deploy to different environments (test, uat, and production).

The simplest use of label expressions is to determine where a build job can or can't be executed. If your web tests require Internet Explorer, for example, you'll need them to run on a Windows machine. You could express this by simply including the corresponding label:

```
windows
```

Alternatively, you might want to run tests against Firefox, but only on Linux machines. You could exclude Windows machines from the range of candidate build nodes by using the `!` negation operator:

```
!windows
```

You can also use the **and** (`&&`) and **or** (`||`) operators to combine expressions. For example, suppose the Postgres database is only tested for Linux. You could tell Jenkins to run a particular build job only on Linux machines installed with postgres using the following expression:

```
linux && postgres
```

Or, you could specify that a particular build job is only to be run on a UAT environment in Sydney or London:

```
uat && (sydney || london)
```

If your machine names contain spaces, you'll need to enclose them in double quotes:

```
"Windows 7" || "Windows XP"
```

There are also two more advanced logical operators that you may find useful. The **implies** operator (`=>`) lets you define a logical constraint of the form “if A is true, then B must also be true.” For example, suppose you have a build job that can run on any Linux distribution, but if it's executed on a Windows box, it must be Windows 7. You could express this constraint as follows:

```
windows => "Windows 7"
```

The other logical operator is the **if-and-only-if** (`<=>`) operator. This operation lets you define stronger constraints of the form “If A is true, then B must be true, but if A is false, then B must be false.” For example, suppose that Windows 7 tests are only to be run in a UAT environment. You could express this as shown here:

```
"Windows 7" <=> uat
```

11.5. Node Monitoring

Jenkins doesn't just dispatch build jobs to slave agents and hope for the best: it pro-actively monitors your slave machines, and will take a node offline if it decides that the node is incapable of safely performing a build. You can fine-tune exactly what Jenkins monitors in the Manage Nodes screen (see Figure 11.16, “Jenkins proactively monitors your build agents”). Jenkins monitors the slave agents in several different ways. It monitors the response time: an overly slow response time can indicate either a network problem or that the slave machine is down. It also monitors the amount of disk space, temporary directory space, and swap space available to the Jenkins user on the slave machine, since build jobs can be notoriously disk-hungry. It also keeps tabs on the system clocks, since if the clocks aren't correctly synchronized, odd errors can sometimes happen. If any of these criteria is not up to scratch, Jenkins will automatically take the server offline.

Figure 11.16. Jenkins proactively monitors your build agents

11.6. Cloud Computing

Cloud computing involves using hardware resources on the Internet as an extension and/or replacement of your local computing architecture. Cloud computing is expanding into many areas of the enterprise, including email and document sharing (Gmail and Google Apps are particularly well-known examples, but there are many others), off-site data storage (such as Amazon S3), as well as more technical services such as source code repositories (such as GitHub, Bitbucket, etc.), and many others.

Of course, externalized hardware architecture solutions have been around for a long time. The main thing that distinguishes cloud computing with more traditional services is the speed and flexibility with which a service can be brought up and brought down when it's no longer needed. In a cloud computing environment, a new machine can be running and available within seconds.

However, cloud computing in the context of CI is not always as simple as it might seem. For any cloud-based approach to work, some of your internal resources may need to be available to the outside world. This can include opening access to your version control system, your test databases, and to any other resources that your builds and tests require. All these aspects need to be considered carefully when choosing a cloud-based CI architecture, and may limit your options if certain resources simply can't be accessed from the Internet. Nevertheless, cloud-based CI has the potential of providing huge benefits when it comes to scalability.

In the following sections, we'll look at how to use the Amazon EC2 cloud computing services to set up a cloud-based build farm.

11.6.1. Using Amazon EC2

In addition to selling books, Amazon is one of the more well-known providers of cloud computing services. If you're willing to pay for the service, Amazon can provide build machines that can be either used permanently as part of your build farm, or brought online as required when your existing build machines become overloaded. This is an excellent and reasonably cost-efficient way to absorb extra build load on an as-needed basis, and without the headache of extra physical machines to maintain.

If you want the flexibility of a cloud-based CI architecture, but don't want to externalize your hardware, another option is to set up a Eucalyptus cloud. Eucalyptus is an open source tool that enables you to create a local private cloud on existing hardware. Eucalyptus uses an API that's compatible with Amazon EC2 and S3, and works well with Jenkins.

11.6.1.1. Setting up your Amazon EC2 build farm

Amazon EC2 is probably the most popular and well-known commercial cloud computing service. To use this service, you'll need to create an EC2 account with Amazon if you don't already have one. The process required to do this is well documented on the Amazon website, so we'll not dwell on it here.

Once you've created your account, you'll be able to create the virtual machines and machine images that will make up your EC2-based build farm.

When using Amazon EC2, you create virtual machines, called instances, using the Amazon Web Services (AWS) Management Console (see Figure 11.17, “You manage your EC2 instances using the Amazon AWS Management Console”). This web page is where you manage your running instances and create new ones. You create these instances from predefined images, called Amazon Machine Images (AMIs). There are many AMI images, both from Amazon and in the public domain, that you can use as a starting point, covering most of the popular operating systems. Once you've created a new instance, you can connect to it using either SSH (for Unix machines) or Windows Remote Desktop Connection (for Windows machines), to configure it further.

Figure 11.17. You manage your EC2 instances using the Amazon AWS Management Console

To set up a build farm, you'll also need to configure your have one, just go to the Key Pairs menu in the Security build server to be able to access your EC2 instances. In particular, you'll need to install the Amazon EC2 API tools, set up the appropriate private/public keys, and allow SSH connections from your server or network to your Amazon instances. Again, the details of how to do this are well documented for all the major operating systems on the EC2 website.

You can use Amazon EC2 instances in two ways—either create slave machines on Amazon EC2 and use them as remote machines, or have Jenkins create them for you dynamically on demand. Or, you can have a combination of the two. Both approaches have their uses, and we'll discuss each of them in the following sections.

11.6.1.2. Using EC2 instances as part of your build farm

Creating a new EC2 instance is as simple as choosing the base image you want to use. You'll just need to provide some details about the instance, such as its size and capacity, and the key pair you want to use to access the machine. Amazon will then create a new running virtual machine based on this image. Once you've set it up, an EC2 instance is essentially a machine like any other. It's easy and convenient to set up permanent or semipermanent EC2 machines as part of your build infrastructure. You may even opt to use an EC2 image as your master server.

Setting up an existing EC2 instance as a Jenkins slave is little different to setting up any other remote slave. If you're setting up a Unix or Linux EC2 slave, you'll need to refer to the key pair (see Figure 11.18, “Configuring an Amazon EC2 slave”) that you used to create the EC2 instance on the AWS Management console. Depending on the flavor of Linux you're using, you may also need to provide a username. Most distributions connect as root, but some, such as Ubuntu, need a different user name.

Figure 11.18. Configuring an Amazon EC2 slave

11.6.1.3. Using dynamic instances

The second approach involves creating new Amazon EC2 machines dynamically, when they're required. Setting up dedicated instances is not difficult, but it doesn't scale well. A better approach is to let Jenkins create new instances as required. To do this, you'll need to install the Jenkins Amazon EC2 plugin. This plugin lets your Jenkins instance start slaves on the EC2 cloud on demand, and then kill them off when they're no longer needed. The plugin works both with Amazon EC2 and the Ubuntu Enterprise Cloud. We'll be focusing on Amazon EC2 here. Note that at the time of writing the Amazon EC2 Plugin only supports managing Linux EC2 images.

Once you've installed the plugin and restarted Jenkins, go to the main Jenkins configuration screen and click on Add a New Cloud (see Figure 11.19, “Configuring an Amazon EC2 slave”). Choose Amazon EC2. You'll need to provide your Amazon Access Key ID and Secret Access Key so that Jenkins can communicate with your Amazon EC2 account. You can access these in the Key Pairs screen of your EC2 dashboard.

Figure 11.19. Configuring an Amazon EC2 slave

You'll also need to provide your RSA private key. If you don't have one, just go to the Key Pairs menu in the Security Credentials screen and create one. This will create a new key pair for you and download the private key. Keep the private key in a safe place (you will need it if you want to connect to your EC2 instances via SSH).

In the advanced options, you can use the Instance Cap field to limit the number of EC2 instances that Jenkins will launch. This limit refers to the total number of active EC2 instances, not just the ones that Jenkins is currently running. This is useful as a safety measure, as you pay for the time your EC2 instances spend running.

Once you've configured your overall EC2 configuration, you need to define the machines you'll work with. You do this by specifying the Amazon Mirror Image (AMI) identifier of the server image you'd like to start. Amazon provides some starter images, and many more are available from the community.

The predefined Amazon and public AMI images are useful starting points for your permanent virtual machines, but for the purposes of implementing a dynamic EC2-based cloud, you need to define your own AMI with the essential tools (Java, build tools, SCM configuration, and so forth) preinstalled. Fortunately, this is a simple process: just start off with a generic AMI (preferably one compatible with the Jenkins EC2 plugin), and install everything your builds need. Make sure you use an EBS image. This way, changes you make to your server instance persist on an EBS volume so that you don't lose them when the server shuts down. Then, create a new image by selecting the Create Image option in the Instances screen on the EC2 management console (see Figure 11.20, “Creating a new Amazon EC2 image”). Make sure SSH is open from your build server's IP address in the default security group on Amazon EC2. If you don't do this, Jenkins will time out when it tries to start up a new slave node.

Once you've prepared your image, you'll be able to use it for your EC2 configuration.

Figure 11.20. Creating a new Amazon EC2 image

Now, Jenkins will automatically create a new EC2 instance using this image when it needs to, and delete (or “terminate,” in Amazon terms) the instance once it's no longer needed. Alternatively, you can bring a new EC2 slave online manually from the Nodes screen using the Provision via EC2 button (see Figure 11.21, “Bringing an Amazon EC2 slave online manually”). This is a useful way to test your configuration.

Figure 11.21. Bringing an Amazon EC2 slave online manually

11.7. Using the CloudBees DEV@cloud Service

Another option you might consider is running your Jenkins instance using a dedicated cloud-based Jenkins architecture, such as the DEV@cloud service offered by CloudBees. CloudBees provides Jenkins as a service as well as various development services (like Sonar) around Jenkins. Using a dedicated Jenkins-specific service, there's no need to install (or manage) Jenkins masters or slaves on your machines. A master instance is automatically configured for you, and when you submit a job to be built, CloudBees provisions a slave for you and takes it back when the job is done.

How does this approach compare with the Amazon EC2-based architecture we discussed in the previous section? The main advantage of this approach is that there's much less work involved in managing your CI architecture instances. Using the Amazon EC2 infrastructure means you don't need to worry about hardware, but you still need to configure and manage your server images yourself. The CloudBees DEV@cloud architecture is more of a high-level, CI-centric service, which provides not only a Jenkins server but also other related tools such as SVN or Git repositories, user management, and Sonar. In addition, the pricing model (pay by the minute) is arguably better suited to a cloud-based CI architecture than the pay-by-the-hour approach used by Amazon.

Amazon EC2-based services are often, though not always, used in a “hybrid cloud” environment where you're offloading your jobs to the cloud, but a bulk of your builds remain in-house. The CloudBees DEV@cloud service is a public cloud solution where the whole build is happening on the cloud (though CloudBees does also offer a similar solution running on a private cloud).

Creating a CloudBees DEV@cloud account is straightforward, and you can use a free one to experiment with the service (note that the free CloudBees service only has a limited set of plugins available; you'll need to sign up for the professional version to use the full plugin range). To sign up for CloudBees, go to the signup page¹. You'll need to enter some relevant information such as a user name, email information, and an account name. Once signed up, you will have access to both DEV@cloud and RUN@cloud (essentially the entire CloudBees platform) services.

¹ <https://grandcentral.cloudbees.com/account/signup>

At this point, you'll have to subscribe to the DEV@cloud service. For our purposes, you can get away with simply choosing the “free” option. You'll have to wait for a few minutes as CloudBees provisions a Jenkins master for you. The next step is to validate your account (this helps CloudBees prevent dummy accounts from running spurious jobs on the service). Click on the validation link, and enter your phone number. An automated incoming phone call will give your pin; enter the pin on the form. Once this is done, you can start running builds.

Your first port of call when you connect will be the management console (called GrandCentral). Click on the “Take me to Jenkins” button to go to your brand new Jenkins master instance.

From here, your interaction with the DEV@cloud platform is exactly like with a standalone Jenkins instance. When you create a new build job, just point to your existing source code repository and hit build. DEV@cloud will provision a slave for you and kick off a build (it may take a minute or two for the slave to be provisioned).

11.8. Conclusion

In CI, distributed builds are the key to a truly scalable architecture. Whether you need to be able to add extra build capacity at the drop of a hat, or your build patterns are subject to periodic spikes in demand, a distributed build architecture is an excellent way to absorb extra load. Distributed builds are also a great way to delegate specialized tasks, such as OS-specific web testing, to certain dedicated machines.

Once you start down the path of distributed builds, cloud-based distributed build farms are a very logical extension. Putting your build servers in the cloud makes it easier and more convenient to scale your build infrastructure when required, as much as required.

Chapter 12. Automated Deployment and Continuous Delivery

12.1. Introduction

CI shouldn't stop once your application compiles correctly. Nor should it stop once you can run a set of automated tests or automatically check and audit the code for potential quality issues. The next logical step, once you've achieved all of these, is to extend your build automation process to the deployment phase. This practice is globally known as Automated Deployment or Continuous Deployment.

In its most advanced form, Continuous Deployment is the process whereby any code change, subject to automated tests and other appropriate verifications, is immediately deployed into production. The aim is to reduce cycle time and reduce the time and effort involved in the deployment process. This, in turn, helps development teams reduce the time it takes to deliver individual features or bug fixes, and, as a consequence, significantly increase their throughput. Reducing or eliminating the periods of intense activity leading up to a traditional release and deployment also frees up time and resources for process improvement and innovation. This approach is comparable to the philosophy of continual improvement promoted by lean processes such as Kanban.

Systematically deploying the latest code into production isn't always suitable, however, no matter how good your automated tests are. Many organizations aren't well prepared for new versions appearing unannounced every week; users might need to be trained, products may need to be marketed, and so forth. A more conservative variation on this theme, often seen in larger organizations, is to have the entire deployment process automated but to trigger the actual deployment manually in a one-click process. This is known as Continuous Delivery, and it has all the advantages of Continuous Deployment without the disadvantages. Variations on Continuous Delivery may also involve automatically deploying code to certain environments (such as test and QA) while using a manual one-click deployment for the other environments (such as UAT and Production). The most important distinguishing characteristic of Continuous Delivery is that any and every successful build that has passed all the relevant automated tests and quality gates can potentially be deployed into production via a fully automated one-click process and be in the hands of the end-user within minutes. However, the process isn't automatic: it's the business, rather than IT, that decides the best time to deliver the latest changes.

Both Continuous Deployment and Continuous Delivery are rightly considered to represent a very high level of maturity in terms of build automation and SDLC practices. These techniques cannot exist without an extremely solid set of automated tests. Nor can they exist without a CI environment and a robust build pipeline—indeed it typically represents the final stage and goal of the build pipeline. However, considering the significant advantages that Continuous Deployment/Delivery can bring to an organization, it's a worthy goal. During the remainder of this chapter, we'll use the general term of “Continuous Deployment” to refer to both Continuous Deployment and Continuous Delivery. Indeed,

Continuous Delivery can be viewed as Continuous Deployment with the final step (deployment into production) being a manual one dictated by the business rather than the development team.

12.2. Implementing Automated and Continuous Deployment

In its most elementary form, Automated Deployment can be as simple as writing your own scripts to deploy your application to a particular server. The main advantage of a scripted solution is simplicity and ease of configuration. However, a simple scripted approach may run into limits if you need to perform more advanced deployment activities, such as installing software on a machine or rebooting the server. For more advanced scenarios, you may need to use a more sophisticated deployment/configuration management solution such as Puppet or Chef.

12.2.1. The Deployment Script

An essential part of any Automated Deployment initiative is a scriptable deployment process. While this may seem obvious, there are still many organizations where deployment remains a cumbersome, complicated, and labor-intensive process, including manual file copying, manual script execution, handwritten deployment notes, and so forth. The good news is that, in general, it doesn't have to be this way. With a little work, it's usually possible to write a script of some sort to automate most, if not all, of the process.

The complexity of a deployment script varies enormously from application to application. For a simple website, a deployment script may be as simple as resyncing a directory on the target server. Many Java application servers have Ant or Maven plugins that can be used to deploy applications. For a more complicated infrastructure, deployment may involve deploying several applications and services across multiple clustered servers in a precisely coordinated manner. Most deployment processes tend to fall somewhere between these extremes.

12.2.2. Database Updates

Deploying your app to the application server is often only one part of the puzzle. Databases, relational or otherwise, almost always play a central role in any application architecture. Of course, ideally, your database would be perfect from the start, but this is rarely the case in the real world. Indeed, when you update your application, you'll generally also need to update one or more databases as well.

Database updates are usually more difficult to manage smoothly than application updates, as both the structure and the contents of the database may be impacted. However, managing database updates is a critical part of both the development and the deployment process, and deserves some reflection and planning.

Some application frameworks, such as Ruby on Rails and Hibernate, can manage structural database changes automatically to some extent. Using these frameworks, you can typically specify if you want to create a new database schema from scratch at each update, or whether you want to update the database

schema while conserving the existing data. While this sounds useful in theory, in fact it's very limited for anything other than noncritical development environments. In particular, these tools don't handle data migration well. For example, if you rename a column in your database, the update process will simply create a new column: it won't copy the data from the old column into the new column, nor will it remove the old column from the updated table.

Fortunately, this isn't the only approach you can use. Another tool that attempts to tackle the thorny problem of database updates is Liquibase¹. Liquibase is an open source tool that can help manage and organize upgrade paths between versions of a database using a high-level approach.

Liquibase works by keeping a record of database updates applied in a table in the database, so that it's easy to bring any target database to the correct state for a given version of the application. As a result, you don't need to worry about running the same update script twice—Liquibase will only apply the update scripts that haven't already been applied to your database. Liquibase is also capable of rolling back changes, at least for certain types of changes. However, since this will not work for every change (for example, data in a deleted table cannot be restored), it's best not to place too much faith in this particular feature.

In Liquibase, you keep track of database changes as a set of “change sets,” each of which represents the database update in a database-neutral XML format. Change sets can represent any changes you would make in a database, from adding and deleting tables, to creating or updating columns, indexes, and foreign keys:

```
<databaseChangeLog
xmlns="http://www.liquibase.org/xml/ns/dbchangelog/1.6"
xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
xsi:schemaLocation="http://www.liquibase.org/xml/ns/dbchangelog/1.6
http://www.liquibase.org/xml/ns/dbchangelog/dbchangelog-1.6.xsd">
  <changeSet id="1" author="john">
    <createTable tableName="department">
      <column name="id" type="int">
        <constraints primaryKey="true" nullable="false"/>
      </column>
      <column name="name" type="varchar(50)">
        <constraints nullable="false"/>
      </column>
      <column name="active" type="boolean" defaultValue="1"/>
    </createTable>
  </changeSet>
</databaseChangeLog>
```

Change sets can also reflect modifications to existing tables. For example, the following change set represents the renaming of a column:

```
<changeSet id="1" author="bob">
  <renameColumn tableName="person" oldColumnName="fname" newColumnName="firstName"/>
</changeSet>
```

¹ <http://www.liquibase.org/>

Since this representation records the semantic nature of the change, Liquibase is capable of handling both the schema updates and data migration associated with this change correctly.

Liquibase can also handle updates to the contents of your database, as well as to its structure. For example, the following change set inserts a new row of data into a table:

```
<changeSet id="326" author="simon">
  <insert tableName="country">
    <column name="id" valueNumeric="1"/>
    <column name="code" value="AL"/>
    <column name="name" value="Albania"/>
  </addColumn>
</changeSet>
```

Each changeset has an ID and an author, which makes it easier to keep track of who made a particular change and reduces the risk of conflict. Developers can test their change sets on their own database schema, and then commit the change sets to version control when ready. The next obvious step is to configure a Jenkins build to run the Liquibase updates against the appropriate database automatically before any integration tests or application deployment is done, usually as part of the ordinary project build script.

Liquibase integrates well into the build process—it can be executed from the command line, or integrated into an Ant or Maven build script. Using Maven, for example, you can configure the Maven Liquibase Plugin as shown here:

```
<project>
  <build>
    <plugins>
      <plugin>
        <groupId>org.liquibase</groupId>
        <artifactId>liquibase-plugin</artifactId>
        <version>1.9.3.0</version>
        <configuration>
          <propertyFileWillOverride>true</propertyFileWillOverride>
          <propertyFile>src/main/resources/liquibase.properties</propertyFile>
        </configuration>
      </plugin>
    </plugins>
  </build>
  ...
</project>
```

Using Liquibase with Maven this way, you could update a given target database to the current schema using this plugin:

```
$ mvn liquibase:update
```

The default database connection details are specified in the `src/main/resources/liquibase.properties` file, and might look something like this:

```
changeLogFile = changelog.xml
```



```
driver = com.mysql.jdbc.Driver
url = jdbc:mysql://localhost/ebank
username = scott
password = tiger
verbose = true
dropFirst = false
```

However, you can override any of these properties from the command line, which makes it easy to set up a Jenkins build to update different databases.

Other similar commands let you generate an SQL script (if you need to submit it to your local DBA for approval, for example), or rollback to a previous version of the schema.

This is of course just one example of a possible approach. Other teams prefer to manually maintain a series of SQL update scripts, or write their own in-house solutions. The important thing is to have a solution that you can use reliably and reproducibly to update different databases to the correct state when deploying your applications.

12.2.3. Smoke Tests

Any serious automated deployment needs to be followed up by a series of automated smoke tests. A subset of the automated acceptance tests can be a good candidate for smoke tests. Smoke tests should be unobtrusive and relatively fast. They should be safe to run in a production environment, which may restrict the number of modifications the test cases can do in the system.

12.2.4. Rolling Back Changes

Another important aspect to consider when setting up Automated Deployment is how to back out if something goes wrong, particularly if you're thinking of implementing Continuous Deployment. Indeed, it's critical to be able to roll back to the previous version if required.

How to do this depends a lot on your application. While it is relatively straight-forward to redeploy a previous version of an application using Jenkins (we'll look at a technique to do this further on in this chapter), the application is often not the only player in the game. In particular, you'll need to consider how to restore your database to a previous state.

We saw how it's possible to use Liquibase to manage database updates, and, of course, many other strategies are possible. However, rolling back a database version presents its own challenges. Liquibase, for example, lets you revert some, but not all, changes to the database structure. However, data lost (in dropped tables, for example) can't be recovered using Liquibase alone.

The most reliable way to revert your database to a previous state is probably to make a backup of the database just before the upgrade, and use this backup to restore the database to its previous state. One effective strategy is to automate this process in Jenkins in the deployment build job, and then to save both the database backup and the deployable binary file as artifacts. This way, you can easily restore the database using the saved backup and then redeploy the application using the saved binary. We'll look at an example of this strategy in action further on in this chapter.

12.3. Deploying to an Application Server

Jenkins provides plugins to help you deploy your application to a number of commonly-used application servers. The Deploy plugin lets you deploy to Tomcat, JBoss, and GlassFish. The Deploy WebSphere plugin tries to cater to the particularities of the IBM WebSphere Application Server.

For other application servers, you'll typically have to integrate the deployment process into your build scripts, or resort to custom scripts to deploy your application. For other languages, too, your deployment process will vary, but it will often involve some use of shell scripting. For example, for a Ruby on Rails application, you may use a tool like Capistrano or Chef, or simply a shell script. For a PHP application, an FTP or SCP file transfer may suffice.

Let's first look at some strategies for deploying your Java applications to an application server.

This is known as a hot-deploy, where the application is deployed onto a running server. This is generally a fast and efficient way of getting your application online. However, depending on your application and on your application server, this approach has been known to result in memory leaks or resource locking issues—older versions of Tomcat, for example, were particularly well-known for this. If you run into this sort of issue, you may have to force the application to restart after each deployment, or possibly schedule a nightly restart of the application server on your test machine.

12.3.1. Deploying a Java Application

In this section we'll look at an example of how to deploy your Java web or JEE application to an application server such as Tomcat, JBoss, or GlassFish.

One of the fundamental principles of automated deployment is to reuse your binaries. It's inefficient, and potentially unreliable, to rebuild your application during the deployment process. Indeed, imagine that you run a series of unit and integration tests against a particular version of your application before deploying it to a test environment for further testing. If you rebuild the binary before deploying it to the test environment, the source code may have changed since the original revision, which means you may not know exactly what you're deploying.

A more efficient process is to reuse the binaries generated by a previous build. For example, you may configure a build job to run unit and integration tests before generating a deployable binary file (typically a WAR or EAR file). You can do this very effectively using the Copy Artifact plugin (see Section 10.7.2, “Copying Artifacts”). This plugin lets you copy an artifact from another build job workspace into the current build job workspace. This, when combined with a normal build trigger or with the Build Promotion plugin, lets you deploy precisely the binary file that you built and tested in the previous phase.

This approach does put some constraints on the way you build your application. In particular, any environment-specific configuration must be externalized to the application; JDBC connections or other such configuration details shouldn't be defined in configuration files embedded in your WAR file, for example, but rather be defined using JONI or in an externalized properties file. If this isn't the case, you

may need to build from a given SCM revision, as discussed for Subversion in Section 10.2.4, “Building from a Subversion Tag”.

12.3.1.1. Using the Deploy plugin

If you're deploying to a Tomcat, JBoss, or GlassFish server, the most useful tool at your disposition will probably be the Deploy plugin. This plugin makes it relatively straightforward to integrate deployment to these platforms into your Jenkins build process. If you're deploying to IBM Websphere, you can use the Websphere Deploy plugin to similar ends.

Let's see how this plugin works in action, using the simple automated build and deployment pipeline illustrated in Figure 12.1, “A simple automated deployment pipeline”.

Figure 12.1. A simple automated deployment pipeline

Here, the default build (**gameoflife-default**) runs the unit and integration tests, and builds a deployable binary in the form of a WAR file. The metrics build (**gameoflife-metrics**) runs additional checks regarding coding standards and code coverage. If both these builds are successful, the application will be automatically deployed to the test environment by the **gameoflife-deploy-to-test** build job.

In the **gameoflife-deploy-to-test** build job, you use the Copy Artifact plugin to retrieve the WAR file generated in the **gameoflife-default** build job and copies it into the current build job's workspace (see Figure 12.2, “Copying the binary artifact to be deployed”).

Figure 12.2. Copying the binary artifact to be deployed

Next, you use the Deploy plugin to deploy the WAR file to the test server. Of course, it's generally possible, and not too difficult, to write a hand-rolled deployment script to get your application on to your application server. In some cases, this may be your only option. However, if a Jenkins plugin exists for your application server, using it can simplify things considerably. If you're deploying to Tomcat, JBoss, or GlassFish, the Deploy plugin may work for you. This plugin uses Cargo to connect to your application server and deploy (or redeploy) your application. Just select the target server type, and specify the server's URL along with the username and password of a user with deployment rights (see Figure 12.3, “Deploying to Tomcat using the Deploy Plugin”).

Figure 12.3. Deploying to Tomcat using the Deploy Plugin

This is known as a hot-deploy, where the application is deployed onto a running server. This is generally a fast and efficient way of getting your application online, and should be the preferred solution because

of its speed convenience. However, depending on your application and on your application server, this approach has been known to result in memory leaks or resource locking issues—older versions of Tomcat, for example, were particularly well-known for this. If you run into this sort of issue, you may have to force the application to restart after each deployment, or possibly schedule a nightly restart of the application server on your test machine.

12.3.1.2. Redeploying a specific version

When you deploy your application automatically or continually, it becomes of critical importance to precisely identify the version of the application currently deployed. There are several ways you can do this, which vary essentially depending on the role Jenkins plays in the build/deployment architecture.

Some teams use Jenkins as the central place of truth, where artifacts are both built and stored for future reference. If you store your deployable artifacts on Jenkins, then it may make perfect sense to deploy your artifacts directly from your Jenkins instance. This isn't hard to do: in the next section we'll look at how to do this using a combination of the Copy Artifacts, Deploy, and Parameterized Trigger plugins.

Alternatively, if you're using an enterprise repository such as Nexus or Artifactory to store your artifacts, then this repository should act as the central point of reference: Jenkins should build and deploy artifacts to your central repository, and then deploy them from there. This is typically the case if you're using Maven as your build tool, but teams using tools like Gradle or Ivy may also use this approach. Repository managers such as Nexus and Artifactory, particularly their commercial editions, make this strategy easier to implement by providing features such as build promotion and staging repositories that help manage the release state of your artifacts.

Let's look at how you might implement each of these strategies using Jenkins.

12.3.1.3. Deploying a version from a previous Jenkins build

Redeploying a previously-deployed artifact in Jenkins is relatively straightforward. In Section 12.3.1.1, “Using the Deploy plugin”, you saw how to use the Copy Artifacts and Deploy plugins to deploy a WAR file built by a previous build job to an application server. What you need to do now is to let the user specify the version to be deployed, rather than just deploying the latest build.

We can do this using the Parameterized Trigger plugin (see Section 10.2, “Parameterized Build Jobs”). First, you add a parameter to the build job, using the special “Build selector for Copy Artifact” parameter type (see Figure 12.4, “Adding a “Build selector for Copy Artifact” parameter”).

Figure 12.4. Adding a “Build selector for Copy Artifact” parameter

This adds a new parameter to your build job (see Figure 12.5, “Configuring a build selector parameter”). Here you need to enter a name and a short description. The name you provide will be used as an environment variable passed to the subsequent build steps.

Figure 12.5. Configuring a build selector parameter

The build selector parameter type lets you pick a previous build in a number of ways, including the latest successful build, the upstream build that triggered this build job, or a specific build. All of these options will be available to the user when he or she triggers a build. The Default Selector lets you specify which of these options will be proposed by default.

When the user selects a particular build job, the build number will also be stored in the environment variables for use in the build steps. The environment variable is called `COPYARTIFACT_BUILD_NUMBER_MY_BUILD_JOB`, where `MY_BUILD_JOB` is the name of the original build job (in upper case and with characters other than A–Z converted to underscores). For example, if you copy an artifact from build number 4 of the `gameoflife-default` project, the `COPYARTIFACT_BUILD_NUMBER_GAMEOFLIFE_DEFAULT` environment variable would be set to 4.

The second part of the configuration is to tell Jenkins what to fetch, and from which build job. In the Build section of our project configuration, you add a “Copy artifacts from another project” step. Here you specify the project where the artifact was built and archived (`gameoflife-default` in our example). You also need to make Jenkins use the build specified in the parameter we defined earlier. You do this by choosing “Specified by a build parameter” in the “Which build” option, and providing the variable name you specified earlier in the build selector name field (see Figure 12.6, “Specify where to find the artifacts to be deployed”). Then, just configure the artifacts to copy as you did in the previous example.

Figure 12.6. Specify where to find the artifacts to be deployed

Finally, you deploy the copied artifact using the Deploy plugin, as illustrated in Figure 12.3, “Deploying to Tomcat using the Deploy Plugin”.

So, let’s see how this build works in practice. When you kick off a build manually, Jenkins will propose a list of options letting you select the build to redeploy (see Figure 12.7, “Choosing the build to redeploy”).

Figure 12.7. Choosing the build to redeploy

Most of these options are fairly self-explanatory.

The “latest successful build” is the most recent build excluding any failing builds. So, this option will typically just redeploy the latest version. If you use this option, you’ll probably want to select the “Stable builds only” checkbox, which will exclude any unstable builds as well.

If you’ve opted to discard old builds, you’ll be able to flag certain build jobs to be kept forever (see Section 5.3.1, “General Options”). In this case, you can choose to deploy the “Latest saved build”.

A sensible option for an automated build job at the end of a build pipeline is “Upstream build that triggered this job”. This way, you can be sure that you're deploying the artifact that was generated by (or promoted through) the previous build job, even if other builds have happened since. It's worth noting that although this sort of parameterized build job is often used to manual deploy a specific artifact, it can also be effectively used as part of an automated build process. If it's not triggered manually, it will simply use whatever value you define in the “default selector” field.

You can also choose the “Specified by permalink” option (see Figure 12.8, “Using the “Specified by permalink” option”). This lets you choose from a number of shortcut values, such as the last build, the last stable build, the last successful build, and so on.

Figure 12.8. Using the “Specified by permalink” option

However, if you want to redeploy a particular version of your application, a more useful option is “Specific build” (see Figure 12.9, “Using a specific build”). This option lets you provide a specific build number to be deployed. This is the most flexible way to redeploy an application—you'll just need to know the number of the build you need to redeploy, but this usually isn't too hard to find by looking at the build history of the original build job.

Figure 12.9. Using a specific build

This is a convenient way to deploy or to redeploy artifacts from previous Jenkins build jobs. However, in some cases you may prefer to use an artifact stored in an enterprise repository like Nexus or Artifactory. We'll look at an example of how to do this in the next section.

12.3.1.4. Deploying a version from a Maven repository

Many organizations use an enterprise repository manager such as Nexus and Artifactory to store and share binary artifacts such as JAR files. This strategy is commonly used with Maven, but also with other build tools such as Ant (with Ivy or the Maven Ant Tasks) and Gradle. Using this approach in a CI environment, both snapshot and release dependencies are built on your Jenkins server, and then deployed to your repository manager (see Figure 12.10, “Using a Maven Enterprise Repository”). Whenever a developer commits source code changes to the version control system, Jenkins will pick up the changes and build new snapshot versions of the corresponding artifacts. Jenkins then deploys these snapshot artifacts to the local enterprise repository manager, where they can be made available to other developers on the team or on other teams within the organization. We discussed how to get Jenkins to automatically deploy Maven artifacts to an enterprise repository in Figure 12.10, “Using a Maven Enterprise Repository”. A similar approach can also be done using Gradle or Ivy.

Figure 12.10. Using a Maven Enterprise Repository

Maven conventions use a well-defined system of version numbers, distinguishing between SNAPSHOT and RELEASE versions. SNAPSHOT versions are considered to be potentially unstable builds of the latest code base, whereas RELEASE versions are official releases having undergone a more formal release process. Typically, SNAPSHOT artifacts are reserved for use within a development team, whereas RELEASE versions are considered ready for further testing.

A similar approach can be used for deployable artifacts such as WAR or EAR files—they're built and tested on the CI server, then automatically deployed to the enterprise repository, often as part of a build pipeline involving automated tests and quality checks (see Section 10.7, “Build Pipelines and Promotions”). SNAPSHOT versions are typically deployed to a test server for automated and/or manual testing, in order to decide whether a version is ready to be officially released.

The exact strategy used to decide when a release version is to be created, and how it's deployed, varies greatly from one organization. For example, some teams prefer a formal release at the end of each iteration or sprint, with a well-defined version number and corresponding set of release notes that's distributed to QA teams for further testing. When a particular version gets the go-ahead from QA, it can then be deployed into production. Others, using a more lean approach, prefer to cut a new release whenever a new feature or bug fix is ready to be deployed. If a team is particularly confident in their automated tests and code quality checks, it may even be possible to automate this process completely, generating and releasing a new version either periodically (say every night) or whenever new changes are committed.

There are many ways to implement this sort of strategy. In the rest of this section, we'll see how to do it using a conventional multimodule Maven project. Our sample project is a web application called **gameoflife**, consisting of three modules: **gameoflife-core**, **gameoflife-services** and **gameoflife-web**. The **gameoflife-web** module produces a WAR file that includes JAR files from the other two modules. It's this WAR file that you want to deploy:

```
tuatara:gameoflife johnsmart$ ls -l
total 32
drwxr-xr-x  16 johnsmart  staff   544 16 May 09:58 gameoflife-core
drwxr-xr-x   8 johnsmart  staff   272  4 May 18:12 gameoflife-deploy
drwxr-xr-x   8 johnsmart  staff   272 16 May 09:58 gameoflife-services
drwxr-xr-x  15 johnsmart  staff   510 16 May 09:58 gameoflife-web
-rw-r--r--@  1 johnsmart  staff 12182  4 May 18:07 pom.xml
```

Earlier on in this chapter you saw how to use the Deploy plugin to deploy a WAR file generated by the current build job to an application server. What you want to do now is to deploy an arbitrary version of the WAR file to an application server.

In Section 10.7.1, “Managing Maven Releases with the M2Release Plugin”, we discussed how to configure Jenkins to invoke the Maven Release Plugin to generate a formal release version of an application. The first step of the deployment process starts here, so we'll assume that this has been configured and that a few releases have already been deployed to our enterprise repository manager.

The next step involves creating a dedicated project to manage the deployment process. This project will be a standard Maven project.

The first thing you need to do is to set up a dedicated deployment project. In its simplest form, this project will simply fetch the requested version of your WAR file from your enterprise repository to be deployed by Jenkins. In the following `pom.xml` file, you use the `maven-war-plugin` to fetch a specified version of the **gameoflife-web** WAR file from our enterprise repository. The version you want is specified in the `target.version` property:

```
<project xmlns="http://maven.apache.org/POM/4.0.0" xmlns:xsi="http://www.w3.org/2001/
  XMLSchema-instance"
  xsi:schemaLocation="http://maven.apache.org/POM/4.0.0 http://
maven.apache.org/maven-v4_0_0.xsd">
  <modelVersion>4.0.0</modelVersion>
  <groupId>com.wakaleo.gameoflife</groupId>
  <artifactId>gameoflife-deploy-with-jenkins</artifactId>
  <version>0.0.1-SNAPSHOT</version>
  <packaging>war</packaging>
  <dependencies>
    <dependency>
      <groupId>com.wakaleo.gameoflife</groupId>
      <artifactId>gameoflife-web</artifactId>
      <type>war</type>
      <version>${target.version}</version>
    </dependency>
  </dependencies>
  <build>
    <plugins>
      <plugin>
        <artifactId>maven-war-plugin</artifactId>
        <configuration>
          <warName>gameoflife</warName>
          <overlays>
            <overlay>
              <groupId>com.wakaleo.gameoflife</groupId>
              <artifactId>gameoflife-web</artifactId>
            </overlay>
          </overlays>
        </configuration>
      </plugin>
    </plugins>
  </build>
  <properties>
    <target.version>RELEASE</target.version>
  </properties>
</project>
```

Next, you configure a Jenkins build job to invoke this `pom.xml` file using a property value provided by the user (see Figure 12.11, “Deploying an artifact from a Maven repository”). Note that we’ve set the default value to `RELEASE` so that, by default, the most recent release version will be deployed. Otherwise, the user can provide the version number of the version to be deployed or redeployed.

Figure 12.11. Deploying an artifact from a Maven repository

The rest of this build job simply checks out the deployment project and invokes the `mvn package` goal, and then deploys the WAR file using the Deploy plugin (see Figure 12.12, “Preparing the WAR to be deployed”). The `target.version` property will be automatically passed into the build job and used to deploy the correct version.

Figure 12.12. Preparing the WAR to be deployed

Similar techniques can be used for other project types. If you're deploying to an application server that isn't supported by the Deploy plugin, you also have the option of writing a custom script in whatever language is most convenient, and getting Jenkins to pass the requested version number as a parameter as described above.

12.3.2. Deploying Scripting-based Applications Like Ruby and PHP

Deploying projects using scripting languages such as PHP and Ruby is generally simpler than deploying Java applications, though the issues related to database updates are similar. Indeed, very often these deployments essentially involve copying files onto a remote server. To obtain the files in the first place, you have the choice of either copying them from another build job's workspace using the Copy Artifacts option, or checking the source code out directly from the source code repository, if necessary using a specific revision or tag as described for Subversion in Section 10.2.4, “Building from a Subversion Tag” and for Git in Section 10.2.5, “Building from a Git Tag”. Then, once you have the source code in your Jenkins workspace, you simply need to deploy it onto the target server.

A useful tool for this sort of deployment is the Publish Over series of plugins for Jenkins (Publish Over FTP, Publish Over SSH, and Publish Over CIFS). These plugins provide a consistent and flexible way to deploy your application artifacts to other servers over a number of protocols, including CIFS (for Windows shared drives), FTP, and SSH/SFTP.

The configuration for each of these plugins is similar. Once you've installed the plugins, you need to set up the host configurations, which are managed centrally in the main configuration screen. You can create as many host configurations as you like—they'll appear in a drop-down list in the job configuration page.

Configuration of the hosts is fairly self-explanatory (see Figure 12.13, “Configuring a remote host”). The name is the name that will appear in the drop-down list in the build job configurations. You can configure authentication using a username and password for FTP, or either an SSH key or a username and password for SSH. You also need to provide an existing directory on the remote server that will act at the root directory for this configuration. In the Advanced options, you can also configure the SSH port and timeout options.

Figure 12.13. Configuring a remote host

Once you've configured your hosts, you can set up your build jobs to deploy artifacts to these hosts. You can do this either as a build step (see Figure 12.14, “Deploying files to a remote host in the build

section”) or as a post-build action (see Figure 12.15, “Deploying files to a remote host in the post-build actions”). In both cases, the options are similar.

Figure 12.14. Deploying files to a remote host in the build section

First of all, you select the target host from the list of hosts you configured in the previous section. Next, you configure the files you want to transfer. You do this by defining one or more “Transfer sets.” A Transfer set is a set of files (defined by an Ant fileset expression) that you deploy to a specified directory on the remote server. You can also provide a prefix to be removed—this lets you strip off unnecessary directories that you don't want to appear on the server (such as the `target/site` directory path in the example). You can add as many transfer sets as you need to get the files you want onto the remote server. The plugin also provides options to execute commands on the remote server once the transfer is complete (“Exec command”) or to exclude certain files or flatten the directories.

Figure 12.15. Deploying files to a remote host in the post-build actions

12.4. Conclusion

Automated Deployment, and in its most advanced form, Continuous Deployment or Continuous Delivery, can be considered the culminating point of a modern CI infrastructure.

In this chapter we've reviewed several Automated Deployment techniques, mostly centered around Java-based deployments. However, the general principles discussed here apply for any technology. Indeed, the actual deployment process in many other technologies, in particular scripting languages such as Ruby and PHP, is considerably simpler than when using Java, and essentially involve copying files onto the production server. Ruby also benefits from tools such as Heroku and Capistrano to facilitate the task.

There are several important aspects you need to consider when setting up an Automated Deployment. First of all, Automated Deployment is the end-point of your CI architecture: you need to define a build pipeline to take your build from the initial compilation and unit tests, through more comprehensive functional and automated acceptance tests and code quality checks, culminating in deployment to one or more platforms. The degree of confidence you can have in your build pipeline depends largely on the degree of confidence you have in your tests. Or, in other words, the less reliable and comprehensive your tests, the earlier in the build process you'll have to fall back to manual testing and human intervention.

Finally, if at all possible, it's important to build your deployable artifact once and once only, and then reuse it in subsequent steps for functional tests and deployment to different platforms.

Chapter 13. Maintaining Jenkins

13.1. Introduction

In this chapter, we'll be discussing a few tips and tricks that you might find useful when maintaining a large Jenkins instance. We'll look at things like how to limit, and keep tabs on, disk usage, how to give Jenkins enough memory, and how to archive build jobs or migrate them from one server to another. Some of these topics are discussed elsewhere in the book, but here we'll be looking at things from the point of view of the system administrator.

13.2. Monitoring Disk Space

Build History takes disk space. In addition, Jenkins analyzes the build records when it loads a project configuration, so a build job with a thousand archived builds is going to take a lot longer to load than one with only fifty. If you have a large Jenkins server with tens or hundreds of build jobs, adjust your expectations accordingly.

Probably the simplest way to keep a cap on disk usage is to limit the number of builds a project maintains in its history. You can configure this by ticking the Discard Old Builds checkbox at the top of the project configuration page (see Figure 13.1, “Discarding old builds”). If you tell Jenkins to only keep the last 20 builds, it will start discarding (and deleting) older build jobs once it reaches this number. You can limit them by number (i.e., no more than 20 builds) or by date (i.e., builds no older than 30 days). It does this intelligently, though: if there has ever been a successful build, Jenkins will always keep at least the latest successful build as part of its build history, so you'll never lose your last successful build.

Figure 13.1. Discarding old builds

The problem with discarding old builds is that you lose the build history at the same time. Jenkins uses the build records to produce graphs of test results and build metrics. If you limit the number of builds to be kept to twenty, for example, Jenkins will only display graphs containing the last twenty data points, which can be a bit limited. This sort of information can be very useful to the developers, but it's often good to be able to see how the project metrics are doing throughout the whole life of the project, not just over the last week or two.

Fortunately, Jenkins has a work-around that can keep both developers and system administrators happy. In general, the items that take up the most disk space are the build artifacts: JAR files, WAR files, and so on. The build history itself is mostly XML log files, which don't take up too much space. If you click on the “Advanced...” button, Jenkins will let you discard the artifacts, but not the build data. In Figure 13.2, “Discarding old builds—advanced options”, for example, you're configured Jenkins to keep artifacts for a maximum of 7 days. This is a great option if you need to put a cap on disk usage, but still want to provide a full scope of build metrics for the development teams.

Figure 13.2. Discarding old builds—advanced options

Don't hesitate to be ruthless, keeping the maximum number of builds with artifacts quite low. Remember, Jenkins will always keep the last stable and the last successful builds, no matter what you tell it, so you'll always have at least one working artifact (unless of course the project has yet to successfully build). Jenkins also lets you mark an individual build as "Keep this log forever", to exclude certain important builds from being discarded automatically.

13.2.1. Using the Disk Usage Plugin

One of the most useful tools in the Jenkins administrator's tool box is the Disk Usage plugin. This plugin records and reports on the amount of disk space used by your projects. It lets you isolate and fix projects that are using too much disk space.

You install the Disk Usage plugin in the usual way, from the Plugin Manager screen. Once you've installed the plugin and restarted Jenkins, the Disk Usage plugin will record the amount of disk space used by each project. It will also add a Disk Usage link on the Manage Jenkins screen, which you can use to display the overall disk usage for your projects (see Figure 13.3, "Viewing disk usage").

Figure 13.3. Viewing disk usage

This list is sorted by overall disk usage, so the projects using the most disk space are at the top. The list provides two values for each project—the Builds column indicates the total amount of space used by all of the project's build history, whereas the Workspace column is the amount of space used to build the project. For ongoing projects, the Workspace value tends to be relatively stable (a project needs what it needs to build correctly), whereas the Builds column will increase over time, sometimes at a dramatic rate, unless you do something about it. You can keep the space needed by a project's history under control by limiting the number of builds being kept for a project, and by being careful about what artifacts are being stored.

To get an idea of how fast the disk space is being used up, you can also display the amount of disk space used in each project over time. To do this, you need to activate the plugin in the System Configuration screen (see Figure 13.4, "Displaying disk usage for a project").

Figure 13.4. Displaying disk usage for a project

This will record and display how much space your projects are using over time. The Disk Usage plugin displays a graph of disk usage over time (see Figure 13.5, "Displaying project disk usage over time"), which can give you a great view of how fast your project is filling up the disk, or, on the contrary, if the disk usage is stable over time.

Figure 13.5. Displaying project disk usage over time

13.2.2. Disk Usage and the Jenkins Maven Project Type

If you're using the Jenkins Maven build jobs, there are some additional details you should know about. In Jenkins, Maven build jobs will automatically archive your build artifacts by default. This may not be what you intend.

The problem is that these SNAPSHOT artifacts take up space—a lot of it. In an active project, Jenkins might be running several builds per hour, so permanently storing the generated JAR files for each build can be very costly. The problem is accentuated if you have multimodule projects, as Jenkins will archive the artifacts generated for each module.

In fact, if you need to archive your Maven SNAPSHOT artifacts, it's probably a better idea to deploy them directly to your local Maven repository manager. Nexus Pro, for example, can be configured to do this and Artifactory can be configured to delete old snapshot artifacts.

Fortunately, you can configure Jenkins to this, go to the “Build” section of your build job configuration screen and click on the Advanced button. This will display some extra fields, as shown in Figure 13.6, “Maven build jobs—advanced options”.

Figure 13.6. Maven build jobs—advanced options

If you tick the “Disable automatic artifact archiving” checkbox here, Jenkins will refrain from storing the JAR files your project build generates. This is a good way of making your friendly system administrator happy.

Note that sometimes you do need to store the Maven artifacts. For example, they often come in handy when implementing a build pipeline (see Section 10.7, “Build Pipelines and Promotions”). In this case, you can always choose to archive the artifacts you need manually, and then use the “Discard old builds” option to refine how long you keep them for.

13.3. Monitoring the Server Load

Jenkins provides build-in monitoring of server activity. On the Manage Jenkins screen, click on the Load Statistics icon. This will display a graph of the server load over time for the master node (see Figure 13.7, “Jenkins Load Statistics”). This graph keeps track of three metrics: the total number of executors, the number of busy executors, and queue length.

The **total number of executors** (the blue line) includes the executors on the master and on the slave nodes. This can vary when slaves are brought on and offline, and can be a useful indicator of how well dynamic provisioning of slave nodes is working.

The **number of busy executors** (the red line) indicates how many executors are executing builds. You should make sure you have enough spare capacity here to absorb spikes in build jobs. If all of executors are permanently occupied running build jobs, you should add more executors and/or slave nodes.

The **queue length** (the gray line) is the number of build jobs awaiting executing. Build jobs are queued when all executors are occupied. This metric doesn't include jobs that are waiting for an upstream build job to finish, so it gives a reasonable idea of when your server could benefit from extra capacity.

Figure 13.7. Jenkins Load Statistics

You can get a similar graph for slave nodes, using the Load Statistics icon in the slave node details page.

Another option is to install the Monitoring plugin. This plugin uses JavaMelody to produce comprehensive HTML reports about the state of your build server, including CPU and system load, average response time, and memory usage (see Figure 13.8, “The Jenkins Monitoring plugin”). Once you've installed this plugin, you can access the JavaMelody graphs from the Manage Jenkins screen, using the “Monitoring of Jenkins/Jenkins master” or “Jenkins/Jenkins nodes” menu entries.

Figure 13.8. The Jenkins Monitoring plugin

13.4. Backing Up Your Configuration

Backing up your data is a universally recommended practice, and your Jenkins server should be no exception. Fortunately, backing up Jenkins is relatively easy. In this section, you'll learn a few ways to do this.

13.4.1. Fundamentals of Jenkins Backups

In the simplest of configurations, all you need to do is to periodically back up your `JENKINS_HOME` directory. This contains all of your build job configurations, your slave node configurations, and your build history. This will work fine while Jenkins is running—there's no need to shut down your server while doing your backup.

The downside of this approach is that the `JENKINS_HOME` directory can contain a very large amount of data (see Section 3.13, “What’s in the Jenkins Home Directory”). If this becomes an issue, you can save a little by not backing up the following directories, which contain data that can be easily recreated on-the-fly by Jenkins:

```
$JENKINS_HOME/war
    The exploded WAR file
```

```
$JENKINS_HOME/cache
Downloaded tools

$JENKINS_HOME/tools
Extracted tools
```

You can also be selective about what you back up in your build job data. The `$JENKINS_HOME/jobs` directory contains job configuration, build history, and archived files for each of your build jobs. The structure of a build job directory is illustrated in Figure 13.9, “The builds directory”.

Figure 13.9. The builds directory

To understand how to optimize your Jenkins backups, you need to understand how the build job directories are organized. Within the jobs directory there's a subdirectory for each build job. This subdirectory contains two subdirectories of its own: `builds` and `workspace`. There's no need to backup the `workspace` directory, as it will simply be restored with a clean checkout if Jenkins finds it missing.

The `builds` directory, on the other hand, needs more attention. This directory contains the history of your build results and previously-generated artifacts, with a time-stamped directory for each previous build. If you're not interested in restoring build history or past artifacts, you don't need to store this directory. If you are, read on! In each of these directories, you'll find the build history (stored in the form of XML files such as JUnit test results) and archived artifacts. Jenkins uses the XML and text files to produce the graphs it displays on the build job dashboard, so if these are important to you, you should store these files. The `archive` directory contains binary files that were generated and stored by previous builds. These binaries may or may not be important to you, but they can take up a lot of space, so if you exclude them from your backups, you may be able to save a considerable amount of space.

Just as it's wise to make frequent backups, it's also wise to test your backup procedure. With Jenkins, this is easy to do. Jenkins home directories are totally portable, so all you need to do to test your backup is to extract your backup into a temporary directory and run an instance of Jenkins against it. For example, imagine we've extracted our backup into a temporary directory called `/tmp/jenkins-backup`. To test this backup, first set the `JENKINS_HOME` directory to this temporary directory:

```
$ export JENKINS_HOME=/tmp/jenkins-backup
```

Then, simply start Jenkins on a different port and see if it works:

```
$ java -jar jenkins.war --httpPort=8888
```

You can now view Jenkins running on this port and make sure that your backup worked correctly.

13.4.2. Using the Backup Plugin

The approach described in the previous section is simple enough to integrate into your normal backup procedures, but you may prefer something more Jenkins-specific. The Backup plugin (see Figure 13.10,

“The Jenkins Backup Manager Plugin”) provides a simple user interface that you can use to back up and restore your Jenkins configurations and data.

Figure 13.10. The Jenkins Backup Manager Plugin

This plugin lets you configure and run backups of both your build job configurations and your build history. The Setup screen gives you a large degree of control over exactly what you want backed up (see Figure 13.11, “Configuring the Jenkins Backup Manager”). You can opt to only back up the XML configuration files, or back up both the configuration files and the build history. You can also choose to backup (or not to backup) the automatically-generated Maven artifacts (in many build processes, these will be available on your local enterprise repository manager). You can also back up the job workspaces (typically unnecessary, as we discussed above) and any generated fingerprints.

Figure 13.11. Configuring the Jenkins Backup Manager

You can trigger a backup manually from the Backup Manager screen (which you can access from the Manage Jenkins screen). The backup takes some time, and will shut down Jenkins during the process (unless you deactivate this option in the backup configuration).

At the time of writing, there's no way to schedule this operation from within Jenkins, but you can start the backup operation externally by invoking the corresponding URL (e.g., `http://localhost:8080/backup/backup` if your Jenkins instance is running locally on port 8080). In a Unix environment, for example, this would typically be scheduled as a cron job using a tool like `wget` or `curl` to start the backup.

13.4.3. More Lightweight Automated Backups

If all you want to back up is your build job configuration, the Backup Manager plugin might be considered overkill. Another option is to use the Thin Backup plugin, which lets you schedule full and incremental backups of your configuration files. Because they don't save your build history or artifacts, these backups are very fast, and there's no need to shut down the server to do them.

Like the Backup plugin, this plugin adds an icon to the Jenkins System Configuration page. From here, you can configure and schedule your configuration backups, force an immediate backup, or restore your configuration files to a previous state. Configuration is straightforward (see Figure 13.12, “Configuring the Thin Backup plugin”), and simply involves scheduling full and incremental backups using a cron job syntax, and providing a directory in which to store the backups.

Figure 13.12. Configuring the Thin Backup plugin

To restore a previous configuration, just go to the Restore page and choose the date of the configuration you wish to reinstate (see Figure 13.13, “Restoring a previous configuration”). Once the configuration

has been restored to the previous state, you need to reload the Jenkins configuration from disk or restart Jenkins.

Figure 13.13. Restoring a previous configuration

13.5. Archiving Build Jobs

Another way to address disk space issues is to delete or archive projects that are no longer active. Archiving a project allows you to easily restore it later if you need to consult the project data or artifacts. Archiving a project is simple: just move the build project directory out of the job directory. Of course, typically, you would compress it into a ZIP file or a tarball first.

In the following example, you want to archive the `tweeter-default` project. So, first you go to the Jenkins `jobs` directory and create a tarball (compressed archive) of the `tweeter-default` build job directory:

```
$ cd $JENKINS_HOME/jobs
$ ls
gameoflife-default      tweeter-default
$ tar czf tweeter-default.tgz tweeter-default
$ ls
gameoflife-default      tweeter-default      tweeter-default.tgz
```

As long as the project you want to archive isn't running, you can now safely delete the project directory and move the archive into storage:

```
$ rm -Rf tweeter-default
$ mv tweeter-default.tgz /data/archives/jenkins
```

Once you've done this, you can simply reload the configuration from the disk in the Manage Jenkins screen (see Figure 13.14, “Reloading the configuration from disk”). The archived project will promptly disappear from your dashboard.

Figure 13.14. Reloading the configuration from disk

On a Windows machine, you can do exactly the same thing by creating a ZIP file of the project directory.

13.6. Migrating Build Jobs

There are times when you need to move or copy Jenkins build jobs from one Jenkins instance to another, without copying the entire Jenkins configuration. For example, you might be migrating your build jobs to a Jenkins instance on a brand new box, with system configuration details that vary from the original machine. Or, you might be restoring an old build job that you've archived.

As you've seen, Jenkins stores all of the data it needs for a project in a subdirectory of the `jobs` directory in your Jenkins home directory. This subdirectory is easy to identify—it has the same name as your project. Incidentally, this is one reason why your project names really shouldn't contain spaces, particularly if Jenkins is running under Unix or Linux—it makes maintenance and admin tasks a lot easier if the project names are also well-behaved Unix filenames.

You can copy or move build jobs between instances of projects simply enough by copying or moving the build job directories to the new Jenkins instance. The project job directory is self-contained—it contains both the full project configuration and all the build history. It's even safe enough to copy build job directories to a running Jenkins instance, though if you're also deleting them from the original server, you should shut this one down first. You don't even need to restart the new Jenkins instance to see the results of your import—just go to the Manage Jenkins screen and click on Reload Configuration From Disk. This will load the new jobs and make them immediately visible on the Jenkins dashboard.

There are a few gotchas, however. If you're migrating your jobs to a brand new Jenkins configuration, remember to install, or migrate, the plugins from your original server. The plugins can be found in the `plugins` directory, so you can simply copy everything from this directory to the corresponding directory in your new instance.

Of course, you might be migrating the build jobs to a new instance precisely because the plugin configuration on the original box is a mess. Some Jenkins plugins can be a bit buggy sometimes, and you may want to move to a clean installation with a well-known, well-defined set of vetted plugins. In this case, you may need to rework some of your project configurations once they've been imported.

The reason for this is straightforward. When you use a plugin in a project, the project's `config.xml` file will be updated with plugin-specific configuration fields. If, for some reason, you need to migrate projects selectively to a Jenkins installation without these plugins installed, Jenkins will no longer understand these parts of the project configuration. The same thing can also happen if the plugin versions are very different on the machines, and the data format used by the plugin has changed.

If you're migrating jobs to a Jenkins instance with a different configuration, it also pays to keep an eye on the system logs. Invalid plugin configurations will usually let you know through warnings or exceptions. While not always fatal, these error messages often mean that the plugin won't work as expected, or at all.

Jenkins provides some useful features to help you migrate your project configurations. If Jenkins finds data that it thinks is out of date or invalid, it will tell you so. On the Manage Jenkins screen, you'll get a message like the one in Figure 13.15, “Jenkins will inform you if your data isn't compatible with the current version”.

Figure 13.15. Jenkins will inform you if your data isn't compatible with the current version

From here, you can choose to either leave the configuration as it is (just in case you roll back to a previous version of your Jenkins instance, for example), or let Jenkins discard the fields it can't read. If you choose this option, Jenkins will bring up a screen containing more details about the error, and

can even help tidy up your project configuration files if you wish (see Figure 13.16, “Managing out-of-date build jobs data”).

Figure 13.16. Managing out-of-date build jobs data

This screen gives you more details about the project containing the dodgy data, as well as the exact error message. This gives you several options. If you're sure that you no longer need the plugin that originally created the data, you can safely remove the redundant fields by clicking on the Discard Unreadable Data button. Alternatively, you may decide that the fields belong to a useful plugin that hasn't yet been installed on the new Jenkins instance. In this case, install the plugin and all should be well. Finally, you can always choose to leave the redundant data and live with the error message, at least until you're sure that you won't need to migrate the job back to the old server some day.

However, Jenkins doesn't always detect all of the errors or inconsistencies—it still pays to keep one eye on the system logs when you migrate your build jobs. For example, the following is a real example from a Jenkins log file showing what can happen during the migration process:

```
Mar 16, 2010 2:05:06 PM hudson.util.CopyOnWriteList$ConverterImpl unmarshal
WARNING: Failed to resolve class
com.thoughtworks.xstream.mapper.CannotResolveClassException: hudson.plugins.cigame.GamePublisher : hudson.plugins.cigame.GamePublisher
    at com.thoughtworks.xstream.mapper.DefaultMapper.realClass(DefaultMapper.java:68)
    at com.thoughtworks.xstream.mapper.MapperWrapper.realClass(MapperWrapper.java:38)
    at com.thoughtworks.xstream.mapper.DynamicProxyMapper.realClass(DynamicProxyMapper.java:71)
    at com.thoughtworks.xstream.mapper.MapperWrapper.realClass(MapperWrapper.java:38)
```

The error is essentially telling us that it can't find a class called `hudson.plugins.cigame.GamePublisher`. In fact, the target installation is missing the CI Game plugin. In this case (as sometimes happens), no warning messages were appearing on the Manage Jenkins page, so Jenkins was unable to correct the configuration files itself.

The simplest solution in this case would be to install the CI Game plugin on the target server. But what if you don't want to install this plugin? We could leave the configuration files alone, but this might mask more significant errors later on—it would be better to tidy them up.

In that case, you need to inspect and update the project configuration files by hand. On this Unix box, I just used `grep` to find all the configuration files containing a reference to “cigame”:

```
$ cd $JENKINS_HOME/jobs
$ grep cigame */config.xml
project-a/config.xml:    <hudson.plugins.cigame.GamePublisher/>
project-b/config.xml:    <hudson.plugins.cigame.GamePublisher/>
project-c/config.xml:    <hudson.plugins.cigame.GamePublisher/>
```

In these config.xml files, I found the reference to the CI Game plugin in the <publishers> section (which is where the configuration for the reporting plugins generally goes):

```
<maven2-moduleset>
...
<publishers>
  <hudson.plugins.cigame.GamePublisher/>
  <hudson.plugins.claim.ClaimPublisher/>
</publishers>
...
</maven2-moduleset>
```

To fix the issue, all I have to do is to remove the offending line:

```
<maven2-moduleset>
...
<publishers>
  <hudson.plugins.claim.ClaimPublisher/>
</publishers>

...
</maven2-moduleset>
```

The exact location of the plugin configuration data will vary depending on the plugin, but in general the config.xml files are quite readable, and updating them by hand isn't too hard.

So, all in all, migrating build jobs between Jenkins instances isn't all that hard—you just need to know a couple of tricks for the corner cases. If you know where to look Jenkins provides some nice tools to make the process smoother.

13.7. Conclusion

In this chapter, you looked at a number of considerations that you should be aware of if you're responsible for maintaining a Jenkins server, including how to monitor disk and server usage, how to back up your build jobs and Jenkins configuration files, and also how to migrate build jobs and upgrade build data safely.

Appendix A. Automating Your Unit and Integration Tests

A.1. Automating Your Tests with Maven

Maven is a popular open source build tool in the Java world that makes use of practices such as declarative dependencies, standard directories and build life cycles, and convention over configuration to encourage clean, maintainable, high level build scripts. Test automation is strongly supported in Maven. Maven projects use a standard directory structure: it will automatically look for unit tests in a directory called (by default) `src/test/java`. There's little else to configure: just add a dependency to the test framework (or frameworks) your tests are using, and Maven will automatically look for and execute the JUnit, TestNG, or even Plain Old Java Objects (POJO) tests contained in this directory structure.

In Maven, you run your unit tests by invoking the `test` life cycle phase, as shown here:

```
$ mvn test
[INFO] Scanning for projects...
[INFO] -----
[INFO] Building Tweeter domain model
[INFO]      task-segment: [test]
[INFO] -----
...
-----
T E S T S
-----
Running com.wakaleo.training.tweeter.domain.TagTest
Tests run: 13, Failures: 0, Errors: 0, Skipped: 0, Time elapsed: 0.093 sec
Running com.wakaleo.training.tweeter.domain.TweeterTest
Tests run: 3, Failures: 0, Errors: 0, Skipped: 0, Time elapsed: 0.021 sec
Running com.wakaleo.training.tweeter.domain.TweeterUserTest
Tests run: 4, Failures: 0, Errors: 0, Skipped: 0, Time elapsed: 0.055 sec
Running com.wakaleo.training.tweeter.domain.TweetFeeRangeTest
Tests run: 10, Failures: 0, Errors: 0, Skipped: 0, Time elapsed: 0.051 sec
Running com.wakaleo.training.tweeter.domain.HamcrestTest
Tests run: 8, Failures: 0, Errors: 0, Skipped: 0, Time elapsed: 0.023 sec

Results :

Tests run: 38, Failures: 0, Errors: 0, Skipped: 0
```

In addition to executing your tests, and failing the build if any of the tests fail, Maven will produce a set of test reports (again, by default) in the `target/surefire-reports` directory, in both XML and text formats. For our CI purposes, it's the XML files that interest us, as Jenkins is able to understand and analyze these files for its CI reporting:

```
$ ls target/surefire-reports/*.xml
```

```
target/surefire-reports/TEST-com.wakaleo.training.tweeter.domain.HamcrestTest.xml
target/surefire-reports/TEST-com.wakaleo.training.tweeter.domain.TagTest.xml
target/surefire-reports/TEST-com.wakaleo.training.tweeter.domain.TweetFeeRangeTest.xml
target/surefire-reports/TEST-com.wakaleo.training.tweeter.domain.TweeterTest.xml
target/surefire-reports/TEST-com.wakaleo.training.tweeter.domain.TweeterUserTest.xml
```

Maven defines two distinct testing phases: unit tests and integration tests. Unit tests should be fast and lightweight, providing a large amount of test feedback in as little time as possible. Integration tests are slower and more cumbersome, and often require the application to be built and deployed to a server (even an embedded one) to carry out more complete tests. Both these sorts of tests are important. For a well-designed CI environment, it is important to be able to distinguish between them. The build should ensure that all of the unit tests are run initially—if a unit test fails, developers should be notified very quickly. Only if all of the unit tests pass is it worthwhile undertaking the slower and more heavyweight integration tests.

In Maven, integration tests are executed during the **integration-test** life cycle phase, which you can invoke by running `mvn integration-test` or (more simply) `mvn verify`. During this phase, it's easy to configure Maven to start up your web application on an embedded Jetty web server, or to package and deploy your application to a test server, for example. Your integration tests can then be executed against the running application. The tricky part, however, is telling Maven how to distinguish between your unit tests and your integration tests, so that they will only be executed when a running version of the application is available.

There are several ways to do this, but at the time of writing there's no official standard approach used across all Maven projects. One simple strategy is to use naming conventions: all integration tests might end in “IntegrationTest”, or be placed in a particular package. The following class uses one such convention:

```
public class AccountIntegrationTest {

    @Test
    public void cashWithdrawalShouldDeductSumFromBalance() throws Exception {
        Account account = new Account();
        account.makeDeposit(100);
        account.makeCashWithdraw(60);
        assertThat(account.getBalance(), is(40));
    }
}
```

In Maven, tests are configured via the **maven-surefire-plugin** plugin. To ensure that Maven only runs these tests during the **integration-test** phase, you can configure this plugin as shown here:

```
<project>
  ...
  <build>
    <plugins>
      <plugin>
        <artifactId>maven-surefire-plugin</artifactId>
        <configuration>
          <skip>true</skip>❶
```

```

</configuration>
<executions>
  <execution>❶
    <id>unit-tests</id>
    <phase>test</phase>
    <goals>
      <goal>test</goal>
    </goals>
    <configuration>
      <skip>false</skip>
      <excludes>
        <exclude>**/*IntegrationTest.java</exclude>
      </excludes>
    </configuration>
  </execution>
  <execution>❷
    <id>integration-tests</id>
    <phase>integration-test</phase>
    <goals>
      <goal>test</goal>
    </goals>
    <configuration>
      <skip>false</skip>
      <includes>
        <include>**/*IntegrationTest.java</include>
      </includes>
    </configuration>
  </execution>
</executions>
</plugin>
...

```

- ❶ Skip all tests by default—this deactivates the default Maven test configuration.
- ❷ During the unit test phase, run the tests but exclude the integration tests.
- ❸ During the integration test phase, run the tests but only include the integration tests.

This will ensure that the integration tests are skipped during the unit test phase, and only executed during the integration test phase.

If you don't want to put unwanted constraints on the names of your test classes, you can use package names instead. In the project illustrated in Figure A.1, “A project containing freely-named test classes”, all of the functional tests have been placed in a package called **webtests**. There's no constraint on the names of the tests, but you're using Page Objects to model your application user interface, so you also make sure that no classes in the **pages** package (underneath the **webtests** package) are treated as tests.

Figure A.1. A project containing freely-named test classes

In Maven, you could do this with the following configuration:

```
<plugin>
```

```

<artifactId>maven-surefire-plugin</artifactId>
<configuration>
  <skip>true</skip>
</configuration>
<executions>
  <execution>
    <id>unit-tests</id>
    <phase>test</phase>
    <goals>
      <goal>test</goal>
    </goals>
    <configuration>
      <skip>>false</skip>
      <excludes>
        <exclude>**/webtests/*.java</exclude>
      </excludes>
    </configuration>
  </execution>
  <execution>
    <id>integration-tests</id>
    <phase>integration-test</phase>
    <goals>
      <goal>test</goal>
    </goals>
    <configuration>
      <skip>>false</skip>
      <includes>
        <include>**/webtests/*.java</include>
      </includes>
      <excludes>
        <exclude>**/pages/*.java</exclude>
      </excludes>
    </configuration>
  </execution>
</executions>
</plugin>

```

TestNG currently has more flexible support for test groups than JUnit. If you're using TestNG, you can identify your integration tests using TestNG Groups. In TestNG, test classes or test methods can be tagged using the `groups` attribute of the `@Test` annotation, as shown here:

```

@Test(groups = { "integration-test" })
public void cashWithdrawalShouldDeductSumFromBalance() throws Exception {
    Account account = new Account();
    account.makeDeposit(100);
    account.makeCashWithdraw(60);
    assertThat(account.getBalance(), is(40));
}

```

Using Maven, you could ensure that these tests were only run during the integration test phase using the following configuration:

```

<project>
  ...

```



```

<build>
  <plugins>
    <plugin>
      <artifactId>maven-surefire-plugin</artifactId>
      <configuration>
        <skip>true</skip>
      </configuration>
      <executions>
        <execution>
          <id>unit-tests</id>
          <phase>test</phase>
          <goals>
            <goal>test</goal>
          </goals>
          <configuration>
            <skip>false</skip>
            <excludedGroups>integration-tests</excludedGroups>❶
          </configuration>
        </execution>
        <execution>
          <id>integration-tests</id>
          <phase>integration-test</phase>
          <goals>
            <goal>test</goal>
          </goals>
          <configuration>
            <skip>false</skip>
            <groups>integration-tests</groups>❷
          </configuration>
        </execution>
      </executions>
    </plugin>
    ...

```

- ❶ Don't run the integration-tests group during the test phase.
- ❷ Run only the tests in the integration-tests group during the integration-test phase.

It often makes good sense to run your tests in parallel where possible, as it can speed up your tests significantly (see Section 6.9, “Help! My Tests Are Too Slow!”). Parallel tests are particularly intensive when doing lots of IO, disk, or network access (such as web tests). These are precisely the sort of tests you want to speed up.

TestNG provides good support for parallel tests. For instance, using TestNG, you could configure your test methods to run in parallel on ten concurrent threads like this:

```

<plugin>
  <groupId>org.apache.maven.plugins</groupId>
  <artifactId>maven-surefire-plugin</artifactId>
  <version>2.5</version>
  <configuration>
    <parallel>methods</parallel>
    <threadCount>10</threadCount>

```

```
</configuration>
</plugin>
```

As of JUnit 4.7, you can also run your JUnit tests in parallel using a similar configuration. In fact, the configuration shown above will work for JUnit 4.7 onwards.

You can also set the `<parallel>` configuration item to `classes` instead of `methods`, which will try to run the test classes in parallel, rather than running each method serially. This might be slower or faster, depending on the number of test classes you have, but might be safer for some test cases not designed with concurrency in mind.

Your mileage will vary, so you should experiment with the numbers to get the best results.

A.2. Automating Your Tests with Ant

Setting up automated testing in Ant is also relatively easy, though it requires a bit more plumbing than with Maven. In particular, Ant does not come packaged with the JUnit libraries or Ant tasks out of the box, so you have to install them somewhere yourself. The most portable approach is to use a Dependency Management tool such as Ivy, or to place the corresponding JAR files in a directory within your project structure.

To run your tests in Ant, you call the `<junit>` task. A typical Jenkins-friendly configuration is shown in this example:

```
<property name="build.dir" value="target" />
<property name="java.classes" value="${build.dir}/classes" />
<property name="test.classes" value="${build.dir}/test-classes" />
<property name="test.reports" value="${build.dir}/test-reports" />
<property name="lib" value="${build.dir}/lib" />

<path id="test.classpath">❶
  <pathelement location="/home/my pc/Documents/jenkinsbook/hudsonbook-content/tools/junit/*.jar" />
  <pathelement location="${java.classes}" />
  <pathelement location="${lib}" />
</path>

<target name="test" depends="test-compile">
  <junit haltonfailure="no" failureproperty="failed">❷
    <classpath>❸
      <path refid="test.classpath" />
      <pathelement location="${test.classes}" />
    </classpath>
    <formatter type="xml" />❹
    <batchtest fork="yes" forkmode="perBatch"❺ todir="${test.reports}">
      <fileset dir="${test.src}">❻
        <include name="**/*Test*.java" />
      </fileset>
    </batchtest>
  </junit>
```

```
<fail message="TEST FAILURE" if="failed" />❶
</target>
```

- ❶ You need to set up a classpath containing the `junit` and `junit-ant` JAR files, as well as the application classes and any other dependencies the application needs to compile and run.
- ❷ The tests themselves are run here. The `haltonfailure` option is used to make the build fail immediately if any tests fail. In a CI environment, this isn't exactly what you want, as you need to get the results for any subsequent tests as well. So, you set this value to `no` and use the `failureproperty` option to force the build to fail once all of the tests have finished.
- ❸ The classpath needs to contain the JUnit libraries, your application classes and their dependencies, and your compiled test classes.
- ❹ The JUnit Ant task can produce both text and XML reports, but for Jenkins you only need the XML ones.
- ❺ The `fork` option runs your tests in a separate JVM. This is generally a good idea, as doing so can avoid classloader issues related to conflicts with Ant's own libraries. However, the default behaviour of the JUnit Ant task is to create a new JVM for each test, which slows down the tests significantly. The `perBatch` option is better, as it only creates one new JVM for each batch of tests.
- ❻ Define the tests you want to run in a `fileset` element. This provides a great deal of flexibility, and makes it easy to define other targets for different subsets of tests (integration, web, and so on).
- ❼ Force the build to fail after the tests have finished, if any of them failed.

If you prefer TestNG, Ant is of course well supported here as well. Using TestNG with the previous example, you could do something like this:

```
<property name="build.dir" value="target" />
<property name="java.classes" value="${build.dir}/classes" />
<property name="test.classes" value="${build.dir}/test-classes" />
<property name="test.reports" value="${build.dir}/test-reports" />
<property name="lib" value="${build.dir}/lib" />

<path id="test.classpath">
  <pathelement location="${java.classes}" />
  <pathelement location="${lib}" />
</path>

<taskdef resource="testngtasks" classpath="lib/testng.jar"/>

<target name="test" depends="test-compile">
  <testng classpathref="test.classpath"
    outputDir="${testng.report.dir}"
    haltonfailure="no"
    failureproperty="failed">
    <classfileset dir="${test.classes}">
      <include name="**/*Test*.class" />
    </classfileset>
  </testng>
  <fail message="TEST FAILURE" if="failed" />
</target>
```

TestNG is a very flexible testing library, and the TestNG task has many more options than this. For example, to only run tests defined as part of the “integration-test” group that you saw earlier, you could do this:

```
<target name="integration-test" depends="test-compile">
  <testng classpathref="test.classpath"
    groups="integration-test"
    outputDir="${testng.report.dir}"
    haltonfailure="no"
    failureproperty="failed">
    <classfileset dir="${test.classes}">
      <include name="**/*Test*.class" />
    </classfileset>
  </testng>
  <fail message="TEST FAILURE" if="failed" />
</target>
```

Or, to run your tests in parallel, using four concurrent threads, you could do this:

```
<target name="integration-test" depends="test-compile">
  <testng classpathref="test.classpath"
    parallel="true"
    threadCount=4
    outputDir="${testng.report.dir}"
    haltonfailure="no"
    failureproperty="failed">
    <classfileset dir="${test.classes}">
      <include name="**/*Test*.class" />
    </classfileset>
  </testng>
  <fail message="TEST FAILURE" if="failed" />
</target>
```

Index

A

- acceptance tests, automated, 6, 97, 111-113
- Acceptance-Test Driven Development, 6
- active (push) notifications, 139
- Active Directory, Microsoft, as security realm, 125
- administrator
 - for Jenkins internal user database, 122
 - for matrix-based security, 129
- aggregate test results, 206-206
- Amazon EC2 cloud computing service, 218-221
- Amazon EC2 plugin, 220
- Amazon Machine Image (AMI), 219
- Amazon Web Services (AWS), 219
- AMI (Amazon Machine Image), 219
- analysis (see code coverage metrics; code quality metrics; tests)
- Ant, 53-54
 - automating tests, 252-254
 - code coverage metrics with Cobertura, 106-107
 - code quality metrics
 - with Checkstyle, 162
 - with CodeNarc, 169
 - with FindBugs, 167
 - with PMD and CPD, 164
 - configuring, 53-54
 - environment variables, accessing from, 78
 - in freestyle build steps, 75-76
 - installing, 53
- ANT_OPTS environment variable, 40
- application server
 - automated deployment to, 228-236
 - Java applications, 228-235
 - scripting-based applications, 235-236
 - deploying Jenkins to, 13, 39-40
 - upgrading Jenkins in, 46
- archives of binary artifacts, 19
 - deploying to Enterprise Repository Manager, 87-89

- disabling, 86
 - in freestyle build jobs, 81-83
- archiving build jobs, 243-243
- Artifactory
 - Enterprise Repository Manager, 88, 89
 - Jenkins support for, 4
- Artifactory plugin, 194
- artifacts (see binary artifacts)
- Atlassian Crowd, as security realm, 126-127
- Audit Trail plugin, 135-135
- auditing user actions, 134-137
- authorization, 121, 121
 - (see also security)
 - matrix-based security, 128-132
 - no restrictions on, 121-122
 - project-based security, 132-133
 - role-based security, 133-134
- automated deployment, 223-227
 - to application server, 228-236
 - database updates with, 224-227
 - deployment script for, 224
 - rolling back changes in, 227
 - smoke tests for, 227
- automated nightly builds, 5
- automated tests (see tests)
- AWS (Amazon Web Services), 219

B

- Backup plugin, 241
- backups, 45, 240-243
- batch scripts, 54, 76-77
- BDD (Behavior-Driven Development), 97
- BDD (Behaviour Driven Development), 112
- binary artifacts
 - archiving, 19
 - deploying to Enterprise Repository Manager, 87-89
 - disabling, 86
 - in freestyle build jobs, 81-83
 - reusing in build pipeline, 200-202
- Boolean parameters, 182
- build agents
 - configuring for multiple JDK versions, 51

- monitoring, 217
- build history
 - in builds directory, 44-45
 - details regarding, 22-23
 - disk usage of, 237-239
 - number of builds to keep, 58
 - parameterized, 184
 - permissions for, 131
 - results summary for, 19, 22
- build jobs, 57, 57
 - (see also freestyle build jobs; Maven build jobs)
 - archiving, 243-243
 - binary artifacts from (see binary artifacts)
 - code coverage metrics in (see code coverage metrics)
 - code quality metrics in (see code quality metrics)
 - copying, 58
 - creating, 16-19, 57-58
 - delaying start of, 49
 - dependencies between, 196
 - distributed across build servers, 209-209
 - associating slave nodes to jobs, 216-217
 - cloud-based build farm for, 218-221
 - creating slave nodes, 210
 - master/slave architecture for, 209-216
 - monitoring slave nodes, 217
 - starting slave nodes, 210-216
 - external, monitoring, 57
 - failed
 - claiming, 143
 - details regarding, 100-101
 - example of, 20-23
 - indicator for, 20, 22
 - notifications for, 139, 142
 - global properties for, 50-50
 - history of (see build history)
 - Javadocs generation in, 24-25
 - joins in, 196-197
 - locking resources for, 197-197
 - migrating, 243-246
 - multiconfiguration, 186-189
 - combination filter for, 188
 - configuration matrix for, 188
 - creating, 186-186
 - custom axis for, 187
 - JDK axis for, 187
 - running, 187-189
 - slave axis for, 186-187
 - naming, 17
 - parameterized, 179-184
 - build scripts for, 180-181
 - creating, 179
 - history of, 184
 - run against a Git tag, 183-183
 - run against a Subversion tag, 182-183
 - starting remotely, 184-184
 - types of parameters, 179, 181-182
 - reports resulting from (see reporting)
 - run numbers for, as parameters, 182
 - running in parallel, 195-197
 - scheduling (see build triggers)
 - source code location for, 17
 - status of, while running, 19
 - steps in, adding, 18-19, 24, 26
 - success of, indicator for, 20
 - tests in (see tests)
 - triggering manually, 19, 20, 73
 - types of, 17, 57
 - unstable build from, 100
 - criteria for, 82, 109, 172, 174
 - indicator for, 27
 - notifications for, 84, 139, 142
 - triggering another build after, 70, 84
- Build Pipeline plugin, 207
- build pipelines, 197-208
 - aggregating test results for, 206-206
 - deployment pipelines from, 206-208
 - Maven version numbers for, 198-199
 - promotions in, 198, 202-206
 - reusing artifacts in, 200-202
- Build Promotion plugin, 228
- build radiators, 144-145
- build scripts (see scripts)
- build server, 5
 - CPU requirements for, 31

- installing Jenkins on, 30-32
- memory requirements for, 30
- monitoring load of, 239-240
- multiple, running builds on (see distributed builds)
- upgrading, 119
- virtual machine for, 31, 119
- build tools, configuring, 52-54
- build triggers
 - configuring, 17-18
 - for freestyle build jobs, 69-73
 - manual, 19, 20, 73
 - parameterized, 184-185
 - polling SCM for version control changes, 71
 - at regular intervals, 70-71
 - remotely triggering from version control system, 72-73
 - when another build has finished, 70
- BUILD_ID environment variable, 77
- BUILD_NUMBER environment variable, 77
- BUILD_TAG environment variable, 77
- BUILD_URL environment variable, 78
- builds directory, 44-45

C

- CAS (Central Authentication Service), 127
- Checkstyle, 161-163, 173
- Checkstyle plugin, 173
- CI (Continuous Integration), 1-2, 5-6
- claiming failed builds, 143
- cloud computing, for builds, 119, 218-221
- CloudBees (sponsor), xxvii
- Clover, 110-111
- Clover plugin, 111
- Cobertura, 25-28, 103-110
 - with Ant, 106-107
 - configuring in build jobs, 108-109
 - with Maven, 103-105
 - reports from, 110-110
- Cobertura plugin, 108
- code complexity, 174-175
- code coverage metrics, 6, 102-111
 - with Clover, 110-111
 - with Cobertura, 25-28, 103-110
 - software for, 103
- code examples, using, xxviii
- code quality metrics, 6, 159-160
 - in build jobs, 160
 - with Checkstyle, 161-163, 173
- code complexity, 174-175
- with CodeNarc, 169-170
- with CPD, 164-167
- with FindBugs, 167-169, 173
- with IDE, 160
- open tasks, 175-176
- plugins for, 159
- with PMD, 163-167, 173
- software for, 160, 161
- with Sonar, 160, 176-177
- with Violations plugin, 170-173

- code reviews, 159
- CodeNarc, 169-170
- coding standards, 159
- commit messages, excluding from triggering build jobs, 62
- config.xml file, 44
- configuration, 13-16, 47-49
 - Ant, 53-54
 - build tools, 52-54
 - Configure System screen, 47, 49-50
 - CVS, 54
 - email server, 54-55
 - Git, 15-16
 - global properties, 50-50
 - JDK, 15, 51-52
 - Load Statistics screen, 48
 - Manage Nodes screen, 49
 - Manage Plugins screen, 48
 - Maven, 14-15, 52-53
 - notifications, 15
 - Prepare for Shut down screen, 49
 - proxy, 55-56
 - quiet period before build starts, 49
 - Reload Configuration from Disk screen, 48
 - Script Console, 48
 - Subversion, 54

- System Information screen, 48
- System Log screen, 48
- system message on home page, 49
- version control systems, 54-54
- Configure System screen, 47, 49-50
- contact information for this book, xxix
- continuous delivery, 2
- continuous deployment, 2, 6, 223-227
 - to application server, 228-236
 - database updates with, 224-227
 - deployment script for, 224
 - rolling back changes in, 227
 - smoke tests for, 227
- Continuous Integration (see CI)
- contributors for this book, xxv
- conventions used in this book, xxiv
- Copy Artifact plugin, 200, 228, 230
- Coverage Complexity Scatter Plot plugin, 175
- CPD, 164-167
- CppUnit, 98
- CPUs, build server requirements for, 31, 31
- cron jobs (see external jobs)
- Crowd, Atlassian, as security realm, 126
- CVS
 - configuring, 54
 - delaying build jobs, 49, 59
 - Jenkins supporting, 15
 - polling with, 72
- CVS_BRANCH environment variable, 78

D

- database
 - rolling back changes to, 227
 - updating with automated deployment, 224-227
 - user database, 122, 122-124
- Dependency Graph View plugin, 196
- Deploy plugin, 228, 229-230, 230
- Deploy Websphere plugin, 228, 229
- deployment (see automated deployment; continuous deployment)
- deployment pipelines, 206-208
- deployment script, 224
- desktop notifiers, 150-153

- disk space
 - for build directory, 44
 - monitoring, 237-239
- Disk Usage plugin, 238-238
- distributed builds, 31, 209-209
 - with cloud-based build farm, 218-221
 - master/slave architecture for, 209-216
 - slave nodes for
 - associating with build jobs, 216-217
 - creating, 210
 - installing as Windows service, 214-215
 - monitoring, 217
 - starting as remote service, 215
 - starting in headless mode, 215
 - starting using SSH, 210-213
 - starting with Java Web Start, 213-214
- DocLinks plugin, 113
- documentation (see Javadocs)

E

- Eclipse
 - code quality metrics with Checkstyle, 161
 - code quality metrics with PMD, 164
 - desktop notifiers with, 150
- Eclipse plugin, 150
- email notifications, 15, 123, 139-143, 143
 - (see also notifications)
- email server, configuring, 54-55
- Email-ext plugin, 140-143
- Enterprise Repository Manager, 87-89
- environment variables, 77
 - (see also specific environment variables)
 - build parameters as, 180
 - using in build steps, 77-79
- Eucalyptus cloud, 218
- EXECUTOR_NUMBER environment variable, 77
- Extended Read Permission plugin, 133
- external jobs, monitoring, 57

F

- File parameters, 182

- FindBugs, 167-169, 173
- FindBugs plugin, 173
- fingerprints, 204, 206
- fingerprints directory, 43
- fonts used in this book, xxiv
- freestyle build jobs, 17, 57, 58-60
 - archiving binary artifacts, 81-83
 - blocking for upstream projects, 59
 - build history for, number of builds to keep, 58
 - build steps in, 73-80
 - Ant build scripts, 75-76
 - batch scripts, 76-77
 - environment variables in, 77-79
 - Groovy scripts, 79-80
 - Maven build steps, 18, 74-75
 - shell scripts, 76-77
 - build triggers for, 69-73
 - code quality metrics in, with Violations, 171-172
 - creating, 17-19
 - delaying start of, 59
 - description of, for project home page, 58
 - disabling, 59
 - failed, 100
 - generating automatically, 194
- Git used with, 63-69
 - branches to build, 64, 66
 - build triggers, 67-68
 - checking out to local branch, 65
 - cleaning after checkout, 66
 - commit author, including in changelog, 66
 - excluding regions from triggering, 64
 - excluding users from triggering, 65
 - Git executable, specifying, 66
 - merging before build, 65
 - post-build merging and pushing actions, 68-69
 - pruning branches before build, 66
 - recursively update submodules, 66
 - repository address, 64
 - source code browsers for, 67
 - SSH keys, 63
 - workspace location, overriding, 65
 - workspace, wiping out before build, 66
- Gradle projects in, 91-93
- Grails projects in, 90-91
- naming, 58
- NAnt build scripts in, 94
- .NET projects in, 93-94
- notifications sent after, 83-84
- post-build actions, 80-84, 99
- reporting on test results, 81-81, 99
- Ruby and Ruby on Rails projects in, 94-95
- running, 84
- starting other build jobs in, 84
- Subversion used with, 60-62
 - excluding commit messages from triggering, 62
 - excluding regions from triggering, 62
 - excluding users from triggering, 62
 - source code browsers for, 61
 - workspace for, overriding, 60
- functional (regression) tests, 97, 98
 - number of, 119
 - performance of, 119
 - running in parallel, 120

G

- Game of Life example application, 16-27
- Gerrit Trigger plugin, 67
- Git, 9
 - branches to build, 64, 66
 - build triggers, 67-68
 - checking out to local branch, 65
 - cleaning after checkout, 66
 - commit author, including in changelog, 66
 - excluding regions from triggering builds, 64
 - excluding users from triggering builds, 65
 - with freestyle build jobs, 63-69
 - installing, 10
 - merging before build, 65
 - post-build merging and pushing actions, 68-69
 - pruning branches before build, 66
 - recursively update submodules, 66
 - repository address, 64
 - source code browsers for, 67

- SSH keys, 63
- tags, building against, 183-183
- workspace location, overriding, 65
- workspace, wiping out before build, 66
- Git plugin, 15-16, 63-64
- GitHub plugin, 69
- GitHub project, 3
- GitHub repository, 9, 67, 69
 - account for, setting up, 11
 - cloning a local copy of, 11
 - forking, 11-12
- GlassFish application server, deploying Java applications to, 228-235
- GlassFish Servlet container, 126
- global properties, 50-50
- Gmail, configuring, 55
- Goldin, Evgeny (contributor), xxv
- Gradle
 - builds in, running with Jenkins, 91-93
 - code quality metrics
 - with Checkstyle, 163
 - with CodeNarc, 169
 - Jenkins support for, 4
- Grails
 - builds in, running with Jenkins, 90-91
 - code quality metrics with CodeNarc, 169
- Groeschke, Rene (contributor), xxvi
- Groovy scripts
 - authentication script, 127-128
 - code quality metrics with CodeNarc, 169-170
 - environment variables in, 79
 - running in build jobs, 79-80
 - running on Script Console, 48
- groups
 - Active Directory, 125, 126
 - Atlassian Crowd, 127
 - LDAP, 124
 - Unix, 126

H

- headless mode, starting slave nodes in, 215
- Hibernate, database updates with, 224
- home directory for Jenkins, 32-33, 42-45, 49
- home page, 13, 49
- hot-deploy, 228, 229
- HTML Publisher plugin, 112-113
- HTTP proxy server, 55
- Hudson, xxiii, 2, 3, 4
 - (see also Jenkins)
- HUDSON_HOME environment variable, 32
- HUDSON_URL environment variable, 77

I

- IDE, code quality metrics with, 160
- IM (see instant messaging)
- information radiators, 144-145
- installation
 - Ant, 53
 - Git, 10
 - JDK, 51
 - Jenkins, 29-30
 - from binary distribution, 30
 - on build server, 30-32
 - on CentOS, 34-34
 - on Debian, 33-34
 - on Fedora, 34-34
 - with Java Web start, 12-13
 - on Linux, 30
 - on OpenSUSE, 35-35
 - on Red Hat, 34-34
 - on SUSE, 35-35
 - on Ubuntu, 33-34
 - on Unix, 30
 - from WAR file, 13, 29
 - on Windows, 29, 30
 - as Windows service, 40-42
- JRE, 10
- Maven, 14-15, 52
- plugins, 25, 25
 - (see also specific plugins)
- upgrading, 45-46
- instant messaging (IM), 145-148
 - IRC for, 148-150
 - Jabber protocol for, 145-148
- Instant Messaging plugin, 145
- integration tests, 97, 98

- number of, 119
 - performance of, 119
- IRC (Internet Relay Chat), 148-150
- IRC plugin, 148, 149
- J**
- Jabber Notifier plugin, 145
- Jabber protocol, 145-148
- Java applications
 - deploying from Maven repository, 232-235
 - deploying to application server, 228-235
 - redeploying a specific version, 230
 - redeploying from previous build, 230-232
 - test reports from, 99
- Java Development Kit (see JDK)
- Java Runtime Environment (JRE), installing, 10
- Java version installed, checking, 29
- Java Web Start
 - installing and starting Jenkins using, 12-13
 - starting slave nodes using, 213-214
- JAVA_ARGS parameter, 34
- JAVA_HOME environment variable, 51, 77
- JAVA_OPTS environment variable, 40
- Javadocs, 24-25
- JBoss application server, deploying Java applications to, 228-235
- JDK (Java Development Kit), 9
 - configuring, 15
 - configuring multiple versions of, 51-52
 - installing, 51
 - requirements for, 29
 - versions of, for multiconfiguration build jobs, 187
- JEE applications (see Java applications)
- Jenkins, 2-4
 - community for, 4
 - configuring (see configuration)
 - CVS supported by, 15
 - dedicated user for, 31
 - environment, requirements for, 9-12
 - help icons in, 14
 - history of, xxiii, 3-4
 - home directory for, 32-33, 42-45, 49
 - home page for, 13, 49
 - installing (see installation)
 - Java requirements for, 29
 - maintenance of, 237-246
 - archiving build jobs, 243-243
 - backups, 45, 240-243
 - migrating build jobs, 243-246
 - monitoring disk space, 237-239
 - monitoring server load, 239-240
 - memory requirements for, 30, 40-40
 - as Open Source project, 4
 - port running on, 30, 30
 - rapid release cycle of, 4
 - reasons to use, 4
 - running
 - on Apache server, 38-39
 - from application server, 13
 - on application server, 39-40
 - from command line, 13, 30
 - from Java Web Start, 12
 - as stand-alone application, 35-38
 - stopping, 13
 - upgrading, 45-46
 - version control systems supported by, 15, 17, 60
- Jenkins Console, 13
- Jenkins M2 Extra Steps plugin, 90
- JENKINS_HOME environment variable, 32-33, 49
- JENKINS_JAVA_CMD parameter, 34
- JENKINS_JAVA_OPTIONS parameter, 34
- JENKINS_PORT parameter, 34
- JMeter, 113-118
- JOB_NAME environment variable, 77
- JOB_URL environment variable, 77
- JobConfigHistory plugin, 135-137
- jobs directory, 43-45
- joins, in build jobs, 196-197
- JRE (Java Runtime Environment), installing, 10
- JUnit reports, 98
 - for acceptance tests, 112
 - configuring in freestyle build job, 19, 99
 - format for, 18

K

Kawaguchi, Kohsuke (developer of Hudson), 3

L

LDAP repository, as security realm, 124-125

LDAP/Active Directory, 4

lightweight backups, 242

Linux, 33

- (see also specific Linux platforms)

- upgrading Jenkins on, 45

Liquibase, 225-227

Load Statistics screen, 48

locking resources for build jobs, 197-197

Locks and Latches plugin, 197

LTS (Long-Term Support) releases, 3

M

M2Eclipse, 4

mail server, configuring, 54-55

maintenance, 237-246

- archiving build jobs, 243-243

- backups, 45, 240-243

- migrating build jobs, 243-246

- monitoring disk space, 237-239

- monitoring server load, 239-240

Manage Jenkins screen, 14, 47-49

Manage Nodes screen, 49

Manage Plugins screen, 48

master/slave architecture for distributed builds, 209-216

matrix build jobs (see multiconfiguration build jobs)

matrix-based security, 128-132

Maven, 9

- automating tests, 247-252

- build steps in freestyle build jobs, 18, 74-75

- Cobertura with, 103-105

- code quality metrics

 - with Checkstyle, 162

 - with CodeNarc, 169

 - with FindBugs, 168

 - with PMD and CPD, 165

- configuring, 14-15, 52-53

- environment variables in, 78

- Hudson support for, 4

- installing, 14-15, 52

- SNAPSHOT dependencies, 74, 85-85

- SNAPSHOT versions, 62

- version numbers for, 198-199

Maven build jobs, 17, 57, 84-90

- archiving binary artifacts, disabling, 86

- build steps in, 85, 90

- code quality metrics in, with Violations, 172-173

- creating, 85

- deploying artifacts to Enterprise Repository Manager, 87-89

- disk usage of, 239-239

- generating automatically, 189-195

 - Artifactory plugin with, 194

 - configuring, 189-191

 - inheritance of configuration, 191-192

 - Parameterized Trigger plugin with, 193

- incremental builds, 86

- modules for, managing, 89

- Post-build Actions, 87

- private repository for, 86

- reporting on test results, 99

- running modules in parallel, 86

- test results of, 100

Maven Jenkins plugin, 189, 194

Maven Release plugin, 198

MAVEN_OPTS environment variable, 40

McCullough, Matthew (contributor), xxv

memory, requirements for, 30, 40-40

metrics (see reporting)

Microsoft Active Directory, as security realm, 125-126

migrating build jobs, 243-246

mobile devices, notifications to, 153

MSBuild plugin, 93

MSTest plugin, 93

multiconfiguration build jobs, 57, 186-189

- combination filter for, 188

- configuration matrix for, 188

- creating, 186-186
- custom axis for, 187
- JDK axis for, 187
- running, 187-189
- slave axis for, 186-187

N

- Nabaztag plugin, 156
- NAnt build scripts, 94
- NAnt plugin, 94
- .NET projects, 93-94
- NetBeans, 150
- Nexus
 - Enterprise Repository Manager, 89, 199
 - Hudson support for, 4
- nightly builds (see automated nightly builds)
- NODE_LABELS environment variable, 77
- NODE_NAME environment variable, 77
- notifications, 139
 - active (push) notifications, 139
 - build radiators, 144-145
 - configuring, 15
 - desktop notifiers, 150-153
 - email, 123, 139-143
 - from freestyle build job, 83-84
 - instant messaging, 145-148
 - to mobile devices, 153
 - using Nabaztag, 156
 - passive (pull), 139
 - RSS feeds, 143-144
 - to smartphones, 152-154
 - SMS messages, 154-154
 - sounds in, 154
 - spoken, 155
- Notifo, 152-153
- NTLM proxy authentication, 55
- NUnit, 98

O

- Odd-e (sponsor), xxvii
- open tasks, reporting on, 175-176

P

- P environment variable, 127
- parameterized build jobs, 179-184, 186
 - (see also multiconfiguration build jobs)
 - build scripts for, 180-181
 - creating, 179
 - history of, 184
 - run against a Git tag, 183-183
 - run against a Subversion tag, 182-183
 - starting remotely, 184-184
 - types of parameters, 179, 181-182
- Parameterized Build plugin, 179
- Parameterized Trigger plugin, 184, 193, 230
- parameterized triggers, 184-185
- passive (pull) notifications, 139
- Password parameters, 181
- performance
 - of code coverage analysis, 103
 - of application, 113-118
 - of tests, 101-101, 119-120
- permissions (see authorization)
- PHP applications, deploying to application server, 235-236
- PHPUnit, 98
- pipelines (see build pipelines)
- plugins
 - Active Directory, 125
 - Amazon EC2, 220
 - architecture of, Jenkins compared to Hudson, 5
 - Artifactory, 88, 194
 - Audit Trail, 135-135
 - Backup, 241
 - Build Pipeline, 207
 - Build Promotion, 228
 - CAS, 127
 - Checkstyle, 173
 - Clover, 111
 - Cobertura, 108
 - Copy Artifact, 200, 228, 230
 - Coverage Complexity Scatter Plot, 175
 - Crowd, for Atlassian Crowd, 126
 - Dependency Graph View, 196

- Deploy, 228, 229-230, 230
- Deploy Websphere, 228, 229
- Disk Usage, 238-238
- DocLinks, 113
- Eclipse, 150
- Email-ext, 140-143
- Extended Read Permission, 133
- FindBugs, 173
- Gerrit Trigger, 67
- Git, 15-16, 63-64
- GitHub, 69
- HTML Publisher, 112-113
- installing, 25-26
- Instant Messaging, 145
- IRC, 148, 149
- Jabber Notifier, 145
- Jenkins M2 Extra Steps, 90
- JobConfigHistory, 135-137
- Locks and Latches, 197
- managing, 48
- Maven Jenkins, 189, 194
- Maven Release, 198
- MSBuild, 93
- MSTest, 93
- Nabaztag, 156
- NAnt, 94
- Parameterized Build, 179
- Parameterized Trigger, 184, 193, 230
- PMD, 173
- Promoted Builds, 202
- Publish Over, 235
- Role Strategy, 133
- Script Security Realm, 127-128
- SFEE, 127
- Sounds, 155
- Speaks, 155
- Task Scanner, 175
- Thin Backup, 242
- Tray Application plugin, 152-152
- upgrading, 46
- Violations, 170-173
- xUnit, 99
- plugins directory, 43

- PMD, 163-167, 173
- PMD plugin, 173
- Prepare for Shutdown screen, 49
- project-based security, 132-133
- project-level permissions, in role-based security, 134
- Promoted Builds plugin, 202
- promotions, 198, 202-206
- properties
 - build parameters as, 181
 - global, 50-50
- proxy, configuring, 55-56
- Publish Over plugins, 235

Q

- quiet period before build starts, 49, 59

R

- radiators, information, 144-145
- regression tests (see functional (regression) tests)
- Reload Configuration from Disk screen, 48
- remote service, starting slave nodes as, 215
- reporting
 - acceptance test results, 112-113
 - code coverage metrics, 6, 25-28
 - from Clover, 111
 - from Cobertura, 110-110
 - code quality metrics, 6
 - with Checkstyle, 173
 - with FindBugs, 173
 - open tasks, 175-176
 - with PMD, 173
 - Violations plugin for, 170-173
- Javadocs API documentation, 24-25
- performance test results, 117-118
- test results, 22-23
 - aggregating, 206-206
 - configuring, 98-99
 - displaying, 99-101
 - JUnit reports, 18, 81-81
 - in RSS feeds, 143-144
- Role Strategy plugin, 133

- role-based security, 133-134
- RSS feeds, of build results, 143-144
- Ruby applications, 94-95, 235-236
- Ruby on Rails projects, 94-95, 224
- Run parameters, 182

S

- SCM (Source Code Management), 17, 60-69
 - (see also version control systems)
- Script Console screen, 48
- Script Security Realm plugin, 127-128
- scripting-based applications, deploying to application server, 235-236
- scripts, 75
 - (see also Ant; Maven)
 - batch scripts, 54, 76-77
 - custom authentication scripts, 127
 - deployment script, 224
 - Groovy scripts, 79-80
 - languages supported, 80
 - parameterized, 180-181
 - shell scripts, 54, 76-77
- security, 121-122
 - authorization, 121
 - matrix-based security, 128-132
 - no restrictions on, 121-122
 - project-based security, 132-133
 - role-based security, 133-134
 - enabling, 121
- security realms, 121
 - Atlassian Crowd, 126-127
 - CAS, 127
 - customizing, 127-128
 - enabling sign-ups, 122
 - enabling user sign-ups, 123
 - Jenkins internal user database, 122, 122-124
 - LDAP repository, 124-125
 - Microsoft Active Directory, 125-126
 - Servlet container, 126
 - SFEE, 127
 - Unix users and groups, 126
- Servlet container
 - as security realm, 126
 - running Jenkins stand-alone using, 35
- SFEE (Source Forge Enterprise Edition), 127
- shell scripts, 54, 76-77
- slave machines
 - for distributed builds, 31, 209-216
 - for multiconfiguration build jobs, 186-187
- smartphones, notifications to, 152-153
- smoke tests, 227
- SMS messages, notifications using, 154-154
- SNAPSHOT dependencies, 74, 85-85
- SNAPSHOT versions, 62
- Sonar
 - code quality metrics with, 160, 176-177
 - frequency of builds, 70
- Sonatype tools, 4, 4
- Sounds plugin, 155
- sounds, in notifications, 154
- source code browsers
 - with Git, 67
 - with Subversion, 61
- Source Code Management (see SCM; version control systems)
- Source Forge Enterprise Edition (see SFEE)
- Speaks plugin, 155
- sponsors for this book, xxvi
- SSH keys, 11, 63
- SSH, starting slave node using, 210-213
- stand-alone application
 - running Jenkins as, 35-38
 - upgrading Jenkins as, 45
- start page (see home page)
- String parameters, 179
- Subversion
 - configuring, 54
 - excluding commit messages from triggering builds, 62
 - excluding regions from triggering builds, 62
 - excluding users from triggering builds, 62
 - with freestyle build jobs, 60-62
 - Jenkins supporting, 15
 - source code browsers for, 61
 - tags, building against, 182-183
- SVN_REVISION environment variable, 78

System Information screen, 48
System Log screen, 48

T

Task Scanner plugin, 175
TDD (Test Driven Development), 97
Test Result Trend graph, 23
Test-Driven development, 6
Test::Unit, 98
TestNG, 98, 99, 102
tests
 acceptance tests, 6, 97, 111-113
 automating, 5, 6, 97-98
 with Ant, 252-254
 with Maven, 247-252
 in freestyle build jobs, 99
 functional (regression) tests, 97, 98
 ignoring, 101-102
 integration tests, 97, 98
 in Maven build jobs, 99
 performance of, 101-101, 119-120
 performance tests, 113-118
 reports from, 22-23
 aggregating, 206-206
 configuring, 98-99
 displaying, 99-101
 JUnit reports, 18, 81-81
 smoke tests, 227
 Test-Driven development, 6
 unit tests, 97, 98
 web tests, 98, 98
Thin Backup plugin, 242
Tomcat application server
 deploying Java applications to, 228-235
 deploying Jenkins using, 13
Tomcat Servlet container, 126
Tray Application plugin, 152-152

U

U environment variable, 127
Ubuntu Enterprise Cloud, 220
unit tests, 97, 98

Unix, 30
 (see also specific Unix platforms)
 users and groups, as security realm, 126
unstable builds, 100
 criteria for, 82, 109, 172, 174
 indicator for, 27
 notifications for, 84, 139, 142
 triggering another build job after, 70, 84
updates directory, 43
upgrades, 45-46
user database, 122, 122, 122-124
 (see also security, security realms)
userContent directory, 43
users
 administrator
 for Jenkins internal user database, 122
 for matrix-based security, 129
 auditing actions of, 134-137
 authorization for (see authorization)
 claiming failed builds, 143
 excluding from triggering builds, 62, 65
 for Jenkins, on build server, 31
users directory, 43

V

version control systems, 9, 54
 (see also CVS; Git; Subversion)
 configuring, 17, 54-54
 polling for changes to trigger build, 71
 remotely triggering builds from, 72-73
 supported by Jenkins, 54-54, 60
version numbers, Maven, 198-199
Violations plugin, 170-173
virtual machine, for build server, 31, 119
Visual Studio MSBuild, 93-94

W

Wakaleo Consulting (sponsor), xxvi
war directory, 43
WAR file, installing Jenkins from, 13
web tests, 98, 98
WebSphere Application Server, 228, 229

Windows

- installation package for Jenkins, 29

Windows services

- installing Jenkins as, 40-42

- installing slave node as, 214-215

- starting slave nodes as, 215

WMI (Windows Management Instrumentation),
215

workspace directory, 44

WORKSPACE environment variable, 77

X

XML format for test reports (see JUnit reports)

Xu, Juven (contributor), xxvi

xUnit, 98, 99

xUnit plugin, 99

Colophon

The animal on the cover of Jenkins: The Definitive Guide is an ornate chorus frog (*Pseudacris ornata*). These small amphibians, only 1–1.5 inches long, can be found on the coastal plains of North America from North Carolina to central Florida and eastern Louisiana. They prefer areas of shallow water without dense vegetation, such as ponds, roadside ditches, and flooded meadows.

The coloration of ornate chorus frogs varies depending on locale, and individuals can be predominantly black, white, brown, red, green, or some variation thereof. All specimens, though, display a dark stripe or collection of spots running from the nostril to the shoulder through the eye, and most have various other spots or stripes as well. The species breeds from November to March, and the calls of males can be heard from in or near areas of shallow water.

Ornate chorus frogs also owe their name to the sound of their mating call: *Pseudacris* comes from the ancient Greek for “false locust.” The name was assigned in 1836 by American naturalist John Edwards Holbrook after he observed that the rapid shrill sound resembled that made by the infamous insect.

The cover image is from Cassell’s Natural History. The cover font is Adobe ITC Garamond. The text font is Linotype Birka; the heading font is Adobe Myriad Condensed; and the code font is LucasFont’s TheSansMonoCondensed.
